

ngspice-46 — Full Change Report

Every modification and its reason (original → current)

Ngspice + OpenVAF Enhancements project

July 2026

Contents

ngspice-46 — Full Change Report	3
1. The OSDI ABI, version 0.4 → 0.7	3
2. The OSDI runtime — <code>src/osdi/</code>	4
3. Analyses — <code>src/spicelib/analysis/</code>	6
4. Frontend — <code>src/frontend/</code>	7
5. Parser and device registry — <code>src/spicelib/</code>	9
6. Maths — <code>src/math/</code>	10
7. Build system	12
8. Per-enhancement index	12

ngspice-46 — Full Change Report

Every modification applied to ngspice-46 in this project, with its reason. The baseline is the pristine ngspice-46 source tree (as vendored in the project's `original/` snapshot, which already contains OpenVAF-Reloaded's stock OSDI support); the current state is the tree committed in this repository under `ngspice-46/`. Each entry links the enhancement write-up in [enhancements_doc/](#) that carries the full engineering detail and the verifying example suite. The companion document [openvaf_changes_full-report.md](#) covers the compiler side.

Maintenance note: this report is updated whenever an enhancement touches ngspice sources. If you are reading it alongside a newer tree, the per-enhancement index at the end tells you the last change it covers.

Scope summary. Two new source files and 51 modified ones carry all the functional changes (~2,100 substantive diff lines). A further ~100 `Makefile.in` files and `config.h.in` differ only because the auto-tools were regenerated with a current libtool ([E-77](#)); they contain no hand-written changes and are not listed individually. Everything below is behavior, interface, or diagnostics.

1. The OSDI ABI, version 0.4 → 0.7

`src/osdi/osdi.h` is the authoritative C contract between ngspice and the `.osdi` objects `openvaf-r` produces. The project extended it four times — twice additively (new exported symbols old loaders simply never look up), twice with genuine version bumps (layout changes):

Change	Kind	Enhancement	Why
OSDI_ABSDELAY_COUNTS / OSDI_ABSDELAY_INFOS exports	additive symbols	E-1	<code>absdelay()</code> needs simulator-side waveform history and synthetic delay-equation rows; the descriptors tell ngspice which nodes/offsets each delay slot uses
OSDI_LAST_CROSSING_COUNTS / OSDI_LAST_CROSSING_INFOS exports	additive symbols	E-6	<code>last_crossing()</code> needs waveform history and crossing interpolation — same additive pattern as <code>absdelay</code>
<code>Osdinode</code> . <code>nodeset</code> field	ABI 0.5 (node stride change)	E-45	Verilog-A net initializers (<code>electrical a = 5.0;</code>) become solver initial-guess nodesets; every node record carries the hint
<code>num_ac_stim_src</code> / <code>ac_stim_sources</code> / <code>load_ac_stim</code> descriptor tail	ABI 0.6 (descriptor grows)	E-51	full <code>ac_stim()</code> support: a Verilog-A module can <i>be</i> the AC stimulus, so the descriptor must enumerate its AC right-hand-side sources
<code>load_noise</code> destination stride 1 → 2 (signed (<code>flat</code> , <code>react</code>) power pairs)	ABI 0.7	E-54	correct noise physics: operating-point-dependent factors and <code>ddt()</code> -shaped ($j\omega$) noise need a complex amplitude per source, $re + j\omega \cdot im$, so coherent same-named sources (LRM 4.6.4) sum with correct sign and frequency shape

Two additive flag conventions ride on the existing `eval()` interface (not version bumps):

- `EVAL_FLAG_IS_FINAL_STEP` (bit 1 $\ll$ 21) in the `eval input` flags — `@(final_step)` firing ([E-53](#));

- *eval return* flags for *\$finish/\$stop* requests and *\$discontinuity*-driven step rejection ([E-55](#)).

Pairing consequence: this ngspice requires OSDI ≥ 0.7 and rejects older objects with a clear version message; *.osdi* files from older compilers must be recompiled ([E-54](#)).

2. The OSDI runtime — **src/osdi/**

osdi.h — **PARAM_KIND_*** unsigned shifts

The parameter-kind flag macros were **PARAM_KIND_MASK** ($3 \ll 30$) and **PARAM_KIND_OPVAR** ($2 \ll 30$) — both shift into the sign bit of a signed int, which is undefined behavior (a UBSan build flags it on every OSDI load, at *osdiinit.c*:41). The shifts are now unsigned ($3u \ll 30$, $2u \ll 30$, and the two zero/one companions for consistency). The bit patterns are unchanged ($0xC0000000 / 0x80000000$), so the OSDI ABI is byte-identical, and they now match the u32 constants *openvaf-r* emits in *metadata/osdi_0_4.rs*. Found by the 2026-07 ASan/UBSan pass.

osdiaccept.c — new file ([E-55](#), [E-1](#), [E-6](#))

The accepted-timepoint hook (**DEVaccept**) OSDI previously lacked. It advances the *absdelay/last_crossing* waveform histories only at *accepted* points (so rejected timesteps don't corrupt the delay lines) and reads the per-attempt latched *\$finish/\$stop* request flags at the one boundary where honoring them is safe. Why: acting on *\$stop* mid-Newton-iteration was treated as a step failure and ground the timestep into a rejection loop; *\$finish* was ignored outright.

osdidefs.h (104 diff lines)

- **OsdiExtraInstData** grows the per-instance fields behind every runtime feature: *dt/temp* offsets with given-flags (instance temperature), *eval_flags*, *absdelay* waveform-history rows and pre-allocated KLU/SPARSE matrix-pointer sets for the delay-equation rows (re-pointed on every DC↔AC transition, mirroring the regular Jacobian handling), and the *last_crossing* history ([E-1](#), [E-6](#), [E-55](#)).
- **ALIGN** macro renamed **OSDI_ALIGN**: the macOS SDK's *<sys/param.h>* defines an unrelated **ALIGN**, producing a redefinition warning in every OSDI translation unit ([E-77](#)).

osdiitf.h / **osdiext.h**

The registry-entry structure gains the *absdelay/last_crossing* descriptor pointers, and **OSDIpendingRequests()** is exported so the analyses can ask "did any device request *\$finish/\$stop* at this accepted point?" ([E-1](#), [E-6](#), [E-55](#)).

osdiregistry.c (68 diff lines)

- Loads the *absdelay/last_crossing* descriptor arrays from each *.osdi* and threads them into the registry entries ([E-1](#), [E-6](#)).
- Requires OSDI ≥ 0.7 ([E-54](#)).
- Descriptor-count hardening ([E-228](#)): *load_object_file* strided/indexed the descriptor arrays by the counts the *.osdi* itself declares (**OSDI_NUM_DESCRIPTOR**s, and per-descriptor *num_params/num_opvars/num_terminals/...*) with no bounds, so a corrupt, truncated, or ABI-mismatched object that still *dlopens* but declares an out-of-range count read past the array (the inner loop dereferences each *param_opvar* entry's *name[]* pointers) → **SIGSEGV** during *pre_osdi*. Two generous ceilings (**OSDI_MAX_DESCRIPTOR**s 4096, **OSDI_MAX_DESC_COUNT** $1 \ll 20$, far above any real model) now reject such a file with a clean diagnostic — the outer count before it sizes the registry / strides the array, and each descriptor's own count fields in the loop before they

index anything (a single load-time choke point protecting all downstream consumers). Targets the egregious/out-of-range count (the truncated / byte-corrupted / ABI-drift signature — an ABI mismatch reads a pointer's bits as a count → huge); a count that lies within the plausible range is an internally inconsistent binary indistinguishable from a valid one and is out of scope, as is a hostile object (a `.osdi` is native code whose loading executes it).

osdiinit.c (8 diff lines)

- Registers the DEVaccept hook (E-55).
- Publishes the synthetic instance-temperature parameter under both spellings, `dt` and the conventional `dtemp` every built-in device uses; `n1 a b mod dtemp=10` used to fail with *"unknown parameter"* while the underlying plumbing was exact (E-80).

osdiloadd.c (441 diff lines — the largest OSDI change)

- Stamps the `absdelay` synthetic rows (`y_synth`, `z`) for DC, AC and transient, driving the delay from the recorded history (E-1).
- Evaluates `last_crossing` from the recorded waveform with linear interpolation between the bracketing timepoints (E-6).
- Passes the final-step flag through to `eval()` (E-53).
- Latches eval-return flags per timepoint attempt (`point_eval_flags`, OR-ed across Newton iterations, reset per attempt) so `$finish/$stop` under solution-dependent conditions survive to the accepted-point boundary (E-55).
- Folds the stride-2 signed noise-power pairs (`fac · |fac| semantics`) when transferring `load_noise` results (E-54, E-42).

osdisetup.c (178 diff lines)

- Creates the synthetic delay-equation unknowns and matrix entries for `absdelay` (E-1).
- Applies the OSDI nodesets (`Osdinode.nodeset`) as solver initial guesses (E-45).
- A `$fatal/$finish` raised during setup (models rejecting their configuration by design — HiSIM's `$port_connected` guards) used to surface as ngspice's baffling *"impossible error — can't occur"*; it now reads *"a Verilog-A device rejected its configuration during setup"* (E-56).
- An internal OSDI node that appears in no Jacobian entry — a net structurally decoupled from the matrix, e.g. an explicit ground `gnd` reference whose branch contribution $V(p, gnd) <+ \dots$ drops the `gnd` column — used to be allocated its own solver row. That all-zero row/column is benign under Sparse (it decouples to $V=0$) but makes the KLU matrix structurally singular, so KLU returned a wrong DC answer (groundcontrib: $v(p)=0$ not 1.5; hierbranch: branch currents 0). Such nodes are now tied to ground (node 0) at setup, matching how OpenVAF already treats them and fixing both solvers (E-116).

osdiacld.c (89 diff lines)

- AC loading gains the `ac_stim` right-hand-side injection: the partitioned AC-stimulus source array is loaded through `load_ac_stim` and added to the AC RHS with exact magnitude and phase, `$mfactor` applied linearly (deterministic signals scale with multiplicity) (E-51, E-26).
- Complex ($j\omega$ -shaped) noise coupling entries participate in the AC matrix (E-54).

osdinoise.c (93 diff lines)

- Per-instance grouping of same-named noise sources: coherent summation of complex amplitudes $(a + j\omega \cdot b) \cdot T$ before squaring, per LRM 4.6.4 — same-named sources correlate (including exact anti-phase cancellation), differently-named ones stay independent (E-42).
- Reads the stride-2 signed pairs of ABI 0.7 (E-54).

osditrunc.c (34 diff lines)

- Honors `$discontinuity(n)`'s next-step clamp: a negative `bound_step` sentinel from the model limits the step after the event (E-24).
- Honors the step-rejection return flag: when an event fired inside the step, `OSDIt trunc` requests `delta/8` (with a `20·CKTdelmin` termination floor) so the integrator bisection onto the discontinuity instead of extrapolating across it (E-55).

osdi/Makefile.am

Adds `osdiaccept.c` to the build (E-55).

3. Analyses — **src/spicelib/analysis/**

dctran.c (46 diff lines)

- Advances OSDI delay/crossing histories via the accept path and skips history corruption on rejected steps (E-1, E-6).
- Issues the dedicated `@(final_step)` evaluation at the converged end of a successful transient (E-53).
- Honors accepted-point `$finish` (ends the analysis cleanly, firing `final_step` at the requesting point), `$stop` (pauses resumably), and aborts with a clear error on `$fatal`'s `E_PANIC` instead of retrying it as nonconvergence (E-55).

dcop.c (9) and **acan.c** (10)

- Fire `@(final_step)` at their successful end (an operating point fires both step events — a single point is first *and* last) (E-53).
- An AC job's operating-point phase carries the analysis name bit (only the name — adding the reactive `CALC_*` bits would wrongly enable integration at an op), so `analysis("ac")` holds through the whole AC run per LRM 4.6.1 and `@(initial_step("ac"))` fires in AC (E-53).

dctrcurv.c (150) + **include/ngspice/trcvdefs.h** (13)

- Generic `.dc @inst[param]` sweeps (a new sweep code): the DC sweep previously hardcoded `Vsource/Isource/Resistor/temperature`, so sweeping any other device parameter was a fatal error. The generic sweep resolves `@inst[param]` through the device's own parameter tables, refreshes per point via `DEVtemperature` (for OSDI exactly the `alter` path), nests with other sweep variables, and restores the original value afterwards (E-62).
- Fires `final_step` at the last sweep point; honors a `$finish` raised mid-sweep (E-53, E-55).

noisean.c (37)

- A singular AC matrix during noise analysis crashed ngspice with a SIGABRT: the code ignored `NIacIter`'s return and the adjoint solve asserted on the unfactored matrix. It now aborts cleanly ("AC solution failed at ... Hz") (E-56).
- Honors deferred `$finish/$stop` raised at the noise operating point and fires `final_step` on success (E-55, E-53).

span.c (31) + **cktspdum.c** (14)

- 1-port S-parameter analyses enabled: the "we need at least two ports" error was over-strict (a 1-port is a plain reflection measurement; the matrix machinery is N-general) — and it was hiding

the cadjoint crash fixed in `dense.c` (E-64).

- **Rbase** published into the `sp` plot from port 1's `z0`, making the Touchstone writers work with no manual `let Rbase = 50` incantation (E-64).
- Stock NaN fixed in **donoise** noise-parameter extraction: the uncorrelated noise conductance G_u is analytically zero for fully-correlated single-source topologies, and floating-point rounding could land `sqrt(Ycor2 + Gu/Rn)`'s argument at $-10^{-18} \rightarrow \text{NaN}$ for the noise figure. Clamped to the physical range ≥ 0 — found by parity testing against an OSDI resistor that produced the exact textbook NF where the built-in returned NaN (E-63).

cktdisto.c (16)

- `.disto` silently reported zero distortion for OSDI devices — the distortion kernel needs higher-order Taylor coefficients the OSDI ABI cannot provide (first derivatives only), and such devices were skipped without a word. A prominent warning now names each affected OSDI device type (E-62).

4. Frontend — **src/frontend/**

plotting/pyplot.c (new) + **com_pyplot.c** (new) + **plotit.c, commands.c** (E-94)

A new interactive command, **pyplot**, plots simulated vectors with matplotlib — a Python counterpart to `gnuplot`. `ft_pyplot()` (`plotting/pyplot.c`) mirrors `ft_gnuplot()`: it writes a `<file>.data` table and a `<file>.py` matplotlib script and `system()`s `python3` (synchronously for a PNG, backgrounded for an interactive window). A new "pyplot" device arm in `plotit()` routes to it, so it accepts the same plot expressions as `plot/gnuplot`; the `com_pyplot()` wrapper mirrors `com_gnuplot()`, and `pyplot` is registered in both command tables. Two set variables tune it: `pyplot_terminal=png` renders headless (Agg) to a PNG, and `pyplot_python` picks the interpreter (default `python3`). Purely additive (E-94). The output file base name is optional (E-95): `com_pyplot()` treats the first word as a file name only when it is not a plot expression (no `(`, and not a vector by `vec_get()`), otherwise it defaults to `pyplot` — so `pyplot v(out)` plots directly (the command's minimum argument count was lowered from 2 to 1). `ft_pyplot()` also grew stacked subplots (`set pyplot_subplots=N` $\rightarrow$ `plt.subplots(nrows, 1, sharex=True)`, `N` traces per panel; `vs` stays the x-axis, so it is not a panel separator) and matplotlib style sheets (`set pyplot_style=<name>`, `dark` $\rightarrow$ `dark_background`, guarded) (E-98). The headless `pyplot_terminal` was extended from PNG-only to the vector export formats `svg` and `pdf` (`set pyplot_terminal=svg/pdf` $\rightarrow$ `fig.savefig(<file>.<fmt>)`, matplotlib picking the writer from the extension), and a figure size control was added (`set pyplot_figsize="W,H"` $\rightarrow$ `plt.subplots(..., figsize=(W, H))`; quote the value so ngspice keeps the comma) (E-99). A **-eye** flag turns `pyplot` into a one-line eye-diagram renderer over the E-207 eye analysis: `com_pyplot()` detects the `-eye` marker, runs `com_eye()` on the trailing tokens (which folds the waveform and leaves `eye_wave/eye_t` + the scalar metrics in a fresh current eye plot), then a new **ft_pyplot_eye()** reads those back and writes a `hist2d`-based persistence eye script (LogNorm, turbo, empty cells masked; threshold + eye-centre lines, eye-height/width double-arrows, a metric title), launched by the same `system(python ...)` mechanism as `ft_pyplot()` and honouring the same `pyplot_terminal/pyplot_python/pyplot_backend/pyplot_figsize/pyplot_style` settings. `ft_pyplot()` and the `gnuplot/asciiplot` back-ends are untouched (E-208).

Four further groups of additions. Appearance controls — seven set variables (`pyplot_grid`, `pyplot_legend`, `pyplot_markers`, `pyplot_axhline`, `pyplot_axvline`, `pyplot_dpi`, `pyplot_transparent`) threaded into `ft_pyplot()`, all additive so a control-free plot is byte-identical to before (E-296). A **-fft** mode — `ft_pyplot()`'s `bool hist` became an `int mode` (LINE/HIST/FFT); `-fft` reuses the time-domain data table and the emitted script resamples onto a uniform grid (`np.interp`, because transient data is adaptively sampled) before `np.fft.rfft`, scaled by `2/sum(window)` so a tone reads back its amplitude (`pyplot_fft_window/_db/_points/_logf`) (E-297). Complex-aware AC views — a new `ft_`

`pyplot_ac()` renderer for `-bode` (mag dB / phase deg vs log-f, stacked, `np.unwrap`), `-nyquist` (imag vs real) and `-polar` (mag at phase), emitting a `<vec-index> <freq> <re> <im>` table that KEEPS the imaginary part an ordinary `pyplot` discards (E-298). And finishing touches — cross-run overlay (`pyplot tran1.v(out) tran2.v(out)`) now sizes the data table by the longest scale so a finer second run is not truncated, plus a `pyplot_cursor` hover crosshair (matplotlib's built-in `Cursor`, window-only) (E-299). A `pyplot_mplcursors` variable then selects the `mplcursors` package (data cursors that snap to a trace) as the cursor backend, turning the cursor on by itself, with a try/except fallback to the built-in `Cursor` if the package is absent where the emitted script runs (E-300). The cursor gating was then consolidated so `pyplot_cursor` is the single master switch (off by default, the only enable); `pyplot_mplcursors` only selects the backend when the cursor is on (E-301).

postcoms.c (414 diff lines) + **postcoms.h** + **commands.c** (16)

The Touchstone I/O subsystem (E-64, E-72):

- **wrsnp** (with `wrs2p` dispatching to the same handler): Touchstone v1 for any port count — the classic 2-port S11 S21 S12 S22 column order preserved byte-identically, $N \geq 3$ in the spec's row-major layout with at most four complex pairs per line;
- writer options `wrsnp <file> [ri|ma|db] [s|y|z] [hz|khz|mhz|ghz]`, any combination in any order: MA/DB formats, Y/Z parameters (normalized to Rbase, $Y \cdot R / Z/R$, per the v1 spec), frequency units — the option line reflects every choice;
- **rdsnp**: reads any Touchstone v1 file into a new plot with a Hz frequency scale and complex vectors matching the `.sp` plot's conventions (MA/DB converted back, Y/Z de-normalized, the 2-port column order handled, Rbase published so imports round-trip) — measured data diffs against simulation in one `let` expression.

Why: E-63 showed the S-parameter *analysis* was exact but its results could not leave ngspice in the industry-standard format (`wrs2p` demanded a never-created Rbase vector), and nothing could bring VNA data in.

outitf.c (16)

- Integer operating-point variables recorded per point: `getSpecial` masked with `IF_REAL` only, so an integer opvar saved with `.save @n1[flag]` produced garbage in transient vectors (E-32).
- The plot-memory guard's message printed `%Id` — a Windows-only `printf` length modifier that is undefined behavior elsewhere and rendered the message as the numberless "memory required (Id Bytes)". Now `%zu`: real byte counts on every platform (E-77).

resource.c (27)

`ft_ckptspace()` — the "approaching max data size" warning (`RSS > 95%` of `RSS + free RAM`) — now honors **set no_mem_check** (the same opt-out the fatal output-size estimate in `outitf.c` respects; one variable governs both memory guards) and warns once per excursion above the threshold, re-arming below 90%, instead of repeating on every check through a long analysis (E-81).

display.c (1)

`#undef NODEV` before the local macro definition — the macOS SDK's `<sys/param.h>` defines an unrelated `NODEV` (E-77).

5. Parser and device registry — `src/spicelib/`

`parser/ingtok.c` (10)

A `.model` card override silently dropped its first parameter when the type name and the opening parenthesis had no space between them (`modname devtype(p1=v1 ...)`): `INPgetNetTok`'s scan did not treat `(` as a token boundary, so the type token swallowed the paren and the first parameter name. Found while verifying event-function counters whose model-card overrides mysteriously kept defaults (E-8).

`parser/inpgmod.c` (5)

`find_model_parameter()` dereferenced `*(device->numModelParms)`, which is `NULL` for built-ins that take no model cards (`VCVS`, `CCCS`, ...) — any *referenced* `.model m vcvs()` card segfaulted, with no OSDI involved at all (an ordinary MOS instance sufficed). This was the true root of the long-documented “module named like a built-in crashes” gotcha. One `NULL` guard; both shapes now produce clean, located errors (E-76).

`parser/numparse.c` — out-of-range int test

`ft_numparse` decides whether a parsed number is integer-valued with the round-trip (`double`)(`int`)`val == val`. Casting a double that lies outside `int` range to `int` is undefined behavior, reached by any control-language literal `>= 2^31` (e.g. `let x = 3e9`) — a UBSan build flags it at `numparse.c:166`. The value is now range-checked (`val >= -2147483648.0 && val < 2147483648.0`) before the cast; a value outside `int` range is, by definition, not `int`-representable, so the answer is unchanged for every in-range input (boundary-tested at `INT_MIN`, `INT_MAX`, and $\pm(2^{31})$). Found by the 2026-07 ASan/UBSan pass.

`snp2va.c` — `lstsq_real` heap overflow on a tiny Touchstone file

`lstsq_real` (the Householder-QR least-squares primitive behind `pre_snp`'s vector fit) documented itself as requiring an overdetermined system ($m \geq n$), but its back-substitution read `b[i]` and `A[i*n+i]` for every unknown `i` up to `n-1`, past the `m`-row buffers when $m < n$. `pre_snp` on a Touchstone file with very few frequency points (≤ 3) drives the fit underdetermined — the stacked system has fewer rows than poles, and the buffers shrink to a `malloc(0)` — so the back-substitution read off the end: a heap-buffer-overflow (ASan: READ of size 8 ... 7 bytes after a 1-byte region), reachable from an ordinary coarse `.snp` measurement. The back-substitution now leaves any unknown with no constraining row ($i \geq m$) at zero instead of reading past the matrix; for the normal $m \geq n$ path the guard never fires and the fit is bit-identical (verified: a 21-point file fits to 2 poles, rms 1.77e-05, before and after). Found by the 2026-07 ASan/UBSan pass (heavy-deck sweep, via a degenerate `pre_snp` input).

`snp2va.c` — unbounded filename-derived port count → heap corruption (E-227)

`parse_touchstone` infers the port count `N` from the file extension `.snp` (`N = atoi(dot+2)`) with no upper bound. A filename such as `.s2147483647p` (`N = INT_MAX`) slips past the parse — the record stride `rec = 1 + 2*N*N` and the per-record inner-loop bound `N*N` share the same wrapped `int`, so `N*N` overflows to a small/zero value and the token layout stays self-consistent — but `N` is stored in `out->N` and then used to size the downstream vector-fit allocations (`malloc(nf*N*N*sizeof(cplx))`, `malloc(N*N*sizeof(cplx))`). The overflowing allocation and the writes that follow corrupt the heap (the program typically aborts later on an unrelated `malloc`, so it presents as flaky). The brute-force fallback already caps port inference at 512; the filename path now honors the same bound — `if (N > 512) N = 0`; falls back to inferring `N` from the data. Valid files (`.s2p`, ...) are far below the cap and unaffected. Found by the 2026-07 Touchstone-parser fuzzing pass (`rdsnp` was clean over 4,500 iters; `pre_snp` hit this on 103/3,600). Verify: `examples/snpfuzz_examples/verify_snpfuzz.py`.

devices/dev.c (51)

- Duplicate device-type registration warned and skipped: `osdi_add_device()` appended every descriptor unconditionally, and the model-card lookup scans front-to-back — so a module name duplicated across two loaded libraries silently resolved to the *first* registration (loading an updated model library gave the stale device with no hint), and a module named like a built-in corrupted model creation. Registration now checks the table case-insensitively and skips duplicates loudly (E-76).
- **pre_osdi** path dedup: loading an already-loaded file notes it and skips (E-76, E-81).
- **pre_osdi -f** force reload (E-229): the three dedup gates (path, `dlopen` handle, device-name) plus `dlopen`'s own path cache meant a recompiled `.osdi` never took effect without restarting ngspice. A `-f/-force` flag (parsed in `com_osdi`, `frontend/com_dl.c`) now force-reloads: `load_osdi(path, force)` stages a byte-for-byte copy of the recompiled file under a fresh unique path and loads *that* (a new path is always re-read, sidestepping the cache; no `dlclose`, which on macOS does not reliably unmap and where overwriting a mapped file faults the process), removes the staged copy once mapped, and `osdi_add_device(..., replace)` swaps the registered device to the new descriptor in place at its existing table index. The old SPICEdev + mapping are left resident so a circuit still built against the prior model stays valid. Without `-f` the re-load is still skipped (now with a hint pointing at `-f`).

osdi/osdiparam.c (E-93)

Setting a fixed (**localparam**) OSDI parameter from the netlist used to be swallowed silently: `openvaf` exports every `localparam` as a parameter, its "given" flag is forced false, and the written value is overwritten by the parameter-init default — so `.model ... N=8` on a frozen structural width parameter changed nothing, with no feedback. `OSDIparam/OSDImparam` now test the new `PARAM_FLAG_FIXED` descriptor flag (set by `openvaf`, a free bit of the parameter `flags` field) and emit *"parameter 'N' is a fixed (localparam) value and cannot be set from the netlist; ignored"* instead of silently dropping it (E-93).

6. Maths — src/maths/

dense/dense.c (11 substantive lines; the file also carries a whitespace/line-ending normalization from editing)

cadjoint() had no 1×1 base case: its cofactor loop allocated 0×0 and then negative-sized minors, killing ngspice with `malloc: can't allocate -8 bytes` — the crash that had been hiding behind the "at least two ports" guard in `span.c`. The adjugate of `[a]` is `[1]` (the empty minor's determinant is 1 by convention), making `cinverse` of a 1×1 equal $1/a$; a $100\ \Omega$ load on a $50\ \Omega$ port now writes a proper `.s1p` with $S_{11} = 1/3$ exactly (E-64).

sparse/spfactor.c (6)

The Markowitz-product overflow guards compare a `double` against `LARGEST_LONG_INTEGER (= LONG_MAX)`, which is not exactly representable as a double; the conversion is now explicit — identical semantics, intentional and warning-free (E-77).

KLU/klu_multiply.c (16)

`SMPmultiply`'s KLU path passed `NULL` internal $\leftrightarrow$ external ordering maps to `KLU_matrix_vector_multiply`, which dereferenced them unconditionally (`&NULL[n] = 0x14` for `n = 5`) and segfaulted — so `.option linesearch` crashed under `.option klu`. The line search is the only caller of `SMPmultiply` under KLU, so this path had never been exercised. A `NULL` map now means the identity ordering (`ext = i + 1`, since the KLU CSR is 0-based and the RHS/Solution vectors are 1-based), making the KLU

matrix-vector product — and hence the [E-111](#) line-search residual merit — correct under KLU, with a merit sequence numerically identical to Sparse 1.3 ([E-112](#)).

KLU/klusmp.c (SMPcaSolve)

SMPcaSolve is the complex adjoint (transposed) solve used by noise and S-parameters. Its Sparse branch calls spSolveTransposed, but the KLU branch called the *non-transposed* klu_z_solve — silently wrong for any asymmetric matrix (every circuit with a transistor or controlled source), which is why .noise was disabled under KLU. It now calls klu_z_tsolve (transposed), so KLU noise matches Sparse exactly (including OSDI device models). This is the core of [E-113](#), which also removes the KLU refusal guards in noisean.c / pzan.c (single-ended pole-zero runs under KLU). Balanced-output pole-zero was Sparse-only at first but [E-172](#) closed that too (a KLU SMPcAddCol branch plus a union-pattern reservation in CKTpzSetup), so no analysis is Sparse-only under KLU any more.

spicelib/analysis/cktsens.c

Sensitivity builds an auxiliary perturbation matrix delta_Y (holding $\partial Y / \partial p$) via SMPnewMatrix, which allocates a plain Sparse 1.3 matrix (delta_Y->CKTkluMODE = 0, no SMPkluMatrix). It is only *multiplied* against the solution (SMPmultiply → spMultiply), never factored, and DEVbindCSC always binds into ckt->CKTmatrix, not delta_Y — so it is correctly Sparse in every case. But two KLU-only setup blocks were gated on the main matrix's CKTkluMODE, so under .option klu they ran and dereferenced the NULL delta_Y->SMPkluMatrix — a segfault on every DC/AC .sens deck. The blocks now gate on delta_Y->CKTkluMODE (its own flag, 0), keeping delta_Y Sparse under KLU exactly as it already is under the default solver, while the main Y matrix stays KLU. DC and AC sensitivity now match Sparse exactly ([E-114](#)).

spicelib/analysis/distoan.c

Distortion is a complex analysis (Volterra series solved at the harmonic / intermodulation frequencies via NIdIter → SMPcSolve), but distoan.c had no KLU code at all: under .option klu the KLU matrix stayed in *real* mode, the complex solves ran against an unconverted matrix, and every harmonic came back zero — a silent wrong answer. The fix mirrors acan.c: before the solve loop it converts the matrix to complex (DEVbindCSCComplex per device + KLUmatrixIsComplex), and on exit converts back to real (DEVbindCSCComplexToReal) so a later real analysis in the same session is unaffected. KLU distortion now matches Sparse bit-for-bit (single-tone, two-tone intermodulation, and OSDI models) ([E-115](#)).

misc/equality.c (AlmostEqualUlp)

AlmostEqualUlp is ngspice's ULP-distance float comparison — the primitive behind breakpoint matching, .measure triggers, transient truncation, PSS, and source timing (18 files). It reinterprets each double as an int64_t, maps the sign, and compared them with intDiff = labs(aInt - bInt). That signed subtraction overflows **int64_t** when the two keys straddle zero, and labs(INT64_MIN) is itself undefined — surfaced by a UBSan build. It is not merely theoretical UB: AlmostEqualUlp(-2.0, +2.0) returned TRUE (declaring -2 and +2 "almost equal"), because those two keys differ by exactly INT64_MIN, whose labs stays negative and slips under maxUlp. The difference is now taken in **uint64_t** as larger-minus-smaller, where the wraparound is well defined and the true magnitude of any two-int64_t difference always fits; the comparison against maxUlp is unsigned. Behavior is identical wherever the old code did not overflow, and correct (FALSE) where it did. Found by the 2026-07 ASan/UBSan pass.

7. Build system

xspice/icm/makedefs.in (4)

The XSPICE codemodel (.cm) link rule hardcoded the pre-macOS-10.3 idiom `-bundle -flat_namespace -undefined suppress`, deprecated loudly by the modern linker on all 14 codemodel links (CI macOS builds included). Now `-bundle -undefined dynamic_lookup` — the modern spelling for plugins whose host symbols resolve at load time (E-77).

Autotools regeneration (the ~100 **Makefile.in** files + **config.h.in**)

Regenerated by `./autogen.sh` with libtool 2.5.4 during the zero-warning work — a build tree whose configure predates libtool 2.4.7 selects the deprecated darwin `-undefined suppress` flag whenever `MACOSX_DEPLOYMENT_TARGET` is unset. No hand-written content (E-77).

8. Per-enhancement index

Every enhancement that touched ngspice, oldest first:

E-1 *osdi.h, osdidefs.h, osdiitf.h, osdiregistry.c, osdiload.c, osdisetup.c, dctran.c (+accept path)*

`absdelay()`: synthetic delay rows + waveform history

E-6 *osdi.h, osdidefs.h, osdiitf.h, osdiregistry.c, osdiload.c*

`last_crossing()`: history + interpolated crossings

E-8 *parser/ingptok.c*

`.model` card first-parameter drop fix

E-24 *osditrunc.c*

`\$discontinuity` next-step clamp

E-25 *osdi callbacks (get_simparams)*

expose `analysis_name` + simulator `simparams`

E-32 *outitf.c*

integer `opvars` recorded per point

E-42 *osdinoise.c*

coherent same-named noise grouping (LRM 4.6.4)

E-45 *osdi.h (ABI 0.5), osdisetup.c*

Verilog-A nodesets applied to the solver

E-51 *osdi.h (ABI 0.6), osdiacld.c*

full `ac_stim` AC-RHS injection

E-53 *dctran.c, dcop.c, dctrcurv.c, acan.c, noisear.c, osdiload.c*

`@(final_step)` + analysis-name phase bits

E-54 *osdi.h (ABI 0.7), osdinoise.c, osdiload.c, osdiacld.c, osdiregistry.c*

node-free complex noise factors, stride-2 pairs

- E-55 *osdiaccept.c (new), osdiloadd.c, osditrunc.c, osdiinit.c, dctrans.c, dctrunc.c, osdiitf.h, Makefile.am*
`\$finish/\$stop/\$fatal` honored; `\$discontinuity` step rejection
- E-56 *noisean.c, osdisetup.c*
noise SIGABRT fix; setup-rejection diagnostics
- E-62 *dctrunc.c, trcvdefs.h, cktdisto.c*
generic `.dc @inst[param]` sweeps; `.disto` warning
- E-63 *span.c*
donoise NaN clamp (stock defect, found by parity)
- E-64 *span.c, cktspdum.c, dense.c, postcoms.c/.h, commands.c*
Touchstone export, auto-Rbase, 1-port `.sp`, `cadjoint 1x1`
- E-72 *postcoms.c/.h, commands.c*
Touchstone round 2: MA/DB/Y/Z/units + `rdsnp` reader
- E-76 *dev.c, inpgmod.c*
duplicate-registration warning, load dedup, stock `.model` segfault fix
- E-77 *osdidefs.h, display.c, outitf.c, spfactor.c, makedefs.in (+autotools regen)*
zero-warning build, `33 → 0`
- E-80 *osdiinit.c*
`dtemp` instance-parameter alias
- E-81 *resource.c, dev.c*
memory-warning opt-out + once-per-excursion; reload-note hint
- E-93 *osdi.h, osdiparam.c*
warn (not silently ignore) when a netlist sets a `PARAM_FLAG_FIXED` (localparam) OSDI parameter
- E-94 *plotting/pyplot.c (new), com_pyplot.c (new), plotit.c, commands.c*
new `pyplot` command — plot vectors with `matplotlib` (a Python `gnuplot`)
- E-95 *com_pyplot.c, commands.c*
make the `pyplot` output file name optional (defaults to `pyplot`)
- E-98 *plotting/pyplot.c*
`pyplot` stacked subplots (`pyplot_subplots=N`) + `matplotlib` style sheets (`pyplot_style`)
- E-99 *plotting/pyplot.c*
`pyplot` vector export formats (`pyplot_terminal=svg/pdf`) + figure size (`pyplot_figsize`)
- E-110 *optdefs.h, tsdefs.h, cktsopt.c, cktntask.c*
`.option errpreset=conservative|moderate|liberal` — one knob for a coordinated tolerance/robustness set; explicit options override regardless of order (`moderate` = historical defaults)
- E-111 *optdefs.h, tsdefs.h, cktdefs.h, cktsopt.c, cktjob.c, cktntask.c, cktdest.c, niiter.c*
`.option linesearch` — globalized (damped) Newton via Armijo backtracking on a new KCL-residual merit $\|F\| = \|G \cdot x - b\|$; result-neutral, off by default

E-112 *maths/KLU/klu_multiply.c*

KLU support for `.option linesearch` — `SMPmultiply`'s KLU path passed NULL ordering maps that `klu_matrix_vector_multiply` dereferenced (SIGSEGV); NULL now means identity ordering. Line search verified merit-identical KLU vs Sparse

E-113 *maths/KLU/klusmp.c, spicelib/analysis/noisean.c, pzan.c*

KLU support for noise + single-ended pole-zero — `SMPcaSolve`'s adjoint KLU branch used the non-transposed solve (silently wrong noise on asymmetric matrices); now `klu_z_solve`. Guards removed; balanced-output pz keeps a targeted guard

E-114 *spicelib/analysis/cktsens.c*

KLU support for sensitivity — the auxiliary perturbation matrix `delta_Y` is Sparse, but two KLU setup blocks gated on the *main* matrix's flag dereferenced its NULL `SMPkluMatrix` (segfault on every DC/AC `.sens`); now gated on `delta_Y`'s own flag. DC/AC `.sens` match Sparse exactly

E-115 *spicelib/analysis/distoan.c*

KLU support for distortion — `.disto` is a complex analysis but `distoan.c` had no KLU code, so under KLU the matrix stayed real and every harmonic came back zero (silent wrong answer); now converts real↔complex around the solve loop like `acan.c`. Matches Sparse bit-for-bit

E-116 *osdi/osdisetup.c*

KLU wrong-DC fix — an OSDI internal node with no Jacobian entry (a decoupled ground reference) was given an all-zero solver row that made the KLU matrix structurally singular; now tied to ground. Fixes the `groundcontrib` and `hierbranch` KLU_XFAILs

E-117 *configure.ac (+configure), spicelib/analysis/dcpss.c*

Periodic steady state (PSS) productionized — built by default (was experimental `--enable-pss`, so `.pss` was unimplemented in shipped builds); ~230 lines of shooting-loop trace routed through `set ngdebug` (232 → 31 lines by default); fail-fast KLU guard (PSS hangs under KLU) directing `.pss` to `.option sparse`

E-118 *spicelib/analysis/dcpss.c*

PSS runs under KLU — the shooting loop hung under KLU (`klu_refactor`'s reused pivots inflated the truncation error into a ~20M-step timestep explosion); forcing a full re-factor (NISHOULDREORDER) each PSS step makes KLU converge to the same result as Sparse. E-117's Sparse-only guard removed

E-119 *pssdefs.h, spicelib/analysis/dcpss.c*

Retain the PSS periodic operating point — PSS sampled the node voltages per period for its DFT then freed them, and never captured device states; now the voltages and `CKTstate0` states are captured per sample and retained on the `PSSan` job (with frequency + dims) as the substrate for the periodic small-signal suite (PAC/pnoise/PXF). Self-checks that the retained samples reproduce the fundamental

E-120 *spicelib/analysis/dcpss.c*

Periodic small-signal Jacobian harmonics — walk the retained operating point, re-linearize each sample (`CKTload MODEINITSMSIG`) and stamp `G+jC` (`CKTacLoad ω=1`), read the osc-node diagonal → `g(t)`, `c(t)`, DFT to harmonics. The blocks the PAC conversion matrix is built from. (Fixed a complex-mode-not-set bug where `C(t)` accumulated across samples.) Verified: RC gives $G=1/R1$, $C=C1$, no harmonics

E-121 *spicelib/analysis/dcpss.c*

PAC conversion-matrix engine — extend E-120 from the osc diagonal to every matrix nonzero, complex-DFT to harmonics G_k , C_k , assemble the $(2M+1)N$ harmonic conversion matrix $H_{\{nm\}} = G_{\{n-m\}} + j\omega_m \cdot C_{\{n-m\}}$, inject a unit current at the osc node in sideband 0 and solve by dense complex LU (`pss_solve`), report the sideband conversion gains. The numerical heart of PAC/pnoise/PXF. Verified: linear RC gives sideband-0 = AC driving-point $|Z|(f_0/2) = 303.3 \Omega$ with the ± 1 sidebands at floating-point zero (no spurious conversion)

E-122 *include/ngspice/pssdefs.h, spicelib/analysis/psssetp.c, spicelib/parser/inp2dot.c, spicelib/analysis/dcpss.c*

`.pac` command (periodic AC) — the user-facing analysis on the E-121 engine. `.pac Fguess StabTime OscNode Points Harmonics SC_iter Steady_coeff <dec|oct|lin> Npts Fstart Fstop` runs PSS then sweeps the input frequency, solving the conversion matrix at each point and emitting the 0-th-sideband node responses as a complex PAC Analysis plot (`print/plot/wrdata`). Reuses the PSS analysis via a `PSSdoPAC` flag; engine factored into `pac_extract_harmonics` (once) + `pac_solve_at` (per freq). Verified: swept sideband-0 $|b(f)| ==$ analytic AC driving-point $|Z(f)|$ to $1.6e-7$ across 10 kHz–1 MHz

E-123 *include/ngspice/pssdefs.h, spicelib/analysis/psssetp.c, spicelib/parser/inp2dot.c, spicelib/analysis/dcpss.c*

finish `.pac` — (1) source-referenced stimulus: capture the netlist AC-source RHS from `CKTacLoad` as the sideband-0 stimulus B_0 (a periodic-AC transfer/conversion gain), unit-current-at-osc-node fallback; (2) multi-sideband output: optional trailing `maxsideband Ksb` emits every conversion sideband $f_{in} + k \cdot f_0$ ($k = -Ksb \dots Ksb$) as its own vector — sideband 0 keeps the node name, others `<node>_usb<k>/<node>_lsb<k>` via `IFnewUid`. Verified: RC with `V1 AC 1` gives sideband-0 = low-pass transfer $1/\sqrt{1+(2\pi fRC)^2}$ ($0.998 \rightarrow 0.157$) with `b_usb1/b_lsb1 ~0` (no conversion)

E-124 *include/ngspice/pssdefs.h, spicelib/analysis/psssetp.c, spicelib/parser/inp2dot.c, spicelib/analysis/dcpss.c*

`.pnoise` command (periodic noise) — fold each device's noise through the conversion matrix. `pac_solve_adjoint` solves $H^T \Psi = e_{\{out, 0\}}$; `pnoise_sweep` loads the sideband- k adjoint into `CKTrhs/CKTirhs` and calls the existing device noise routines (`NevalSrc`, `OSDI load_noise`) once per sideband, accumulating $\Sigma_k S \cdot |\Delta \Psi_k|^2$ — reuses every device noise model via a minimal local `NOISEAN` context. `.pnoise <pss> OutNode InSrc <dec|oct|lin> Np Fstart Fstop`; emits `onoise_spectrum/inoise_spectrum`. Verified: linear RC reduces exactly to `.noise` ($4kTR/(1+(2\pi fRC)^2)$), matches the `.noise` reference to every digit

E-125 *include/ngspice/pssdefs.h, spicelib/analysis/psssetp.c, spicelib/parser/inp2dot.c, spicelib/analysis/dcpss.c*

`.pxf` command (periodic transfer function) — the adjoint of PAC, completing the `PSS` $\rightarrow$ `PAC` $\rightarrow$ `Pnoise` $\rightarrow$ `PXF` suite. `pxf_sweep` solves $H^T \Psi = e_{\{out, 0\}}$ per frequency and dots each sideband block with the AC-source pattern $B_0 \rightarrow x_{f_k} = \Sigma_j \Psi_k(j) \cdot B_0(j)$, the input $\rightarrow$ output transfer at each sideband. `.pxf <pss> OutNode <dec|oct|lin> Np Fstart Fstop [maxsb]`; emits `xf + xf_usb<k>/xf_lsb<k>`. Verified: by the identity $(H^{-1}B)_{out} = (H^{-T}e_{out})^T B$ the sideband-0 transfer is bit-identical to the PAC response ($0.998 \rightarrow 0.157$ low-pass), conversion sidebands $\sim 2e-16$

E-126 *include/ngspice/pssdefs.h, spicelib/analysis/psssetp.c, spicelib/parser/inp2dot.c, spicelib/analysis/dcpss.c*

cyclostationary noise — `.pnoise ... cyclo` evaluates each device's noise PSD $S(t)$ at every PSS sample's bias (`CKTload` per sample) and folds it through the time-domain adjoint transfer $A_s(j) = \Sigma_k \Psi_k(j) e^{j2\pi k s/P}$, averaging: `onoise(f) = (1/P) \Sigma_s S(t_s) \cdot |\Delta A_s|^2`. Reuses the device noise routines. Verified: (1) reduces exactly to `.noise` on the linear RC (bias-independent thermal $\rightarrow$ `Parseval`); (2) a flicker resistor carrying the RC current gives `onoise.f = R1^2 \cdot KF \cdot \langle I^2 \rangle = 4.88e-10` (uses the period-average $\langle I^2 \rangle$), matched to 5 digits

E-127 *include/ngspice/{optdefs,tskdefs,cktdefs}.h, spicelib/analysis/{cktsopt,cktnntask,cktdojob,cktop,cktdest}.c, maths/ni/niiter.c*

pseudo-transient continuation — `.option ptcont` embeds $f(x)=0$ in a backward-Euler pseudo-transient $f(x)+Gps \cdot (x-x_{prev})=0$ and marches $d\tau$ small $\rightarrow$ large ($Gps \rightarrow 0$). Gps diagonal via the `gmin` path; the $Gps \cdot x_{prev}$ RHS coupling (added in `NIiter`) makes each step a stable-trajectory move, not a static `gmin` step. New `pseudo_transient` in the `CKTop` cascade, off by default. Verified under `KLU+Sparse`: result-neutral on normal circuits; on a stiff no-limiting exp, reaches the correct DC 0.837922 V where plain Newton lands on a spurious 70.5 V (21/21)

E-128 *include/ngspice/{optdefs,tskdefs,cktdefs}.h, spicelib/analysis/{cktsopt,cktnntask,cktdojob}.c, spicelib/analysis/dctran.c*

LTE-based dynamic integration-order control — `.option dynorder` selects the Gear order per step from the local-truncation-error limit instead of the stock `1 $\leftrightarrow$ 2` toggle, so orders 3–6 (already coded in `NIcomCof` but never used) are actually exercised. The selector compares the raw LTE-limited step (uncapped `CKTtrunc`) at the current order and its ± 1 neighbours and moves with hysteresis ($1.2\times$), a settling hold after each change (let the BDF divided differences rebuild), and an order-dependent step-growth cap ($2\times$ at order $\leq 3 \rightarrow 1.3\times$ at 6). Off by default; bounded by `maxord`; inert at the default `maxord=2`. Stiff-transient guards (post-breakpoint order hold + leaky-bucket rejection-rate order drop) keep it from collapsing the step on a violent slew. Verified under `KLU+Sparse`: RC decay 3–5 $\times$ fewer steps at matched accuracy with monotone error; smooth RLC ringdown 8.9 $\times$ fewer steps AND more accurate (0.13 % vs 0.34 %); nonlinear rectifier matches stock to 5 sig figs; the transistor-level $\mu A741 \pm 5\text{ V}$ slew completes under `dynorder` and matches fixed Gear-2

E-129 *frontend/outitf.c*

sweep progress bar — the throttled `Reference value` status line (redrawn in place every 0.25 s during a sweep) now carries a live progress bar + percentage, e.g. `Reference value : 5.91926e-04 [=====] 74%`. `outp_progress_frac()` computes the 0–1 fraction per analysis: transient from $(CKTtime-TSTART)/(TSTOP-TSTART)$, AC/noise from the frequency's linear/log position in the band, DC from the accepted-point count over the nested step-count product; analyses with no span (`op, ...`) keep the plain line. Fixed-width (no stale chars), std-out status-line only (never in the `rawfile/wrdata`), same `!ft_norefprint && !cp_background` gating. Verified 22/22: printed % matches the analytic sweep fraction to 0.5 % across `tran/AC/DC/noise`; `op` shows no bar

E-130 *frontend/com_optimize.{c,h}, commands.c, options.c, include/ngspice/ftext.h, spicelib/analysis/cktdojob.c, frontend/outitf.c*

built-in Nelder-Mead optimizer — `optimize -param <name> <init> <lo> <hi> [-param ...] -analysis <cmd> -minimize <expr> [-maxiter N] [-tol T] [-verbose]` varies device/alter parameters, re-runs an analysis, and minimizes a scalar objective expression via a derivative-free downhill simplex in normalized [0,1] space (scale-invariant across orders-of-magnitude params). Sub-commands dispatched synchronously through `cp_coms` (not the deferring re-entrant `cp_evloop`); the hundreds of inner analyses are silenced via a new `ft_optimizing` flag gating the banner/row-count/E-129 progress bar at source (`-verbose` to show). Verified 9/9 against analytic optima: DC divider $R1 \rightarrow 2333.3\ \Omega$, AC low-pass $R1 \rightarrow 2756.6\ \Omega$, 2-D compound objective $R1=3k/R2=2k$ exactly

E-131 *frontend/com_checkpoint.{c,h}, commands.c, com_commands.h, Makefile.am, spicelib/analysis/dctran.c, include/ngspice/cktdefs.h*

transient checkpoint / restart — new `savestate <file> / loadstate <file>` commands serialize the full transient integration state (solution vector `CKTrhsOld`, device state history `CKTstates[]`, time/step/order/mode, pending breakpoints) to a binary file and resume it, including

in a fresh process — so a long run survives a crash, splits across sessions, or moves between machines (stock ngspice could only resume in memory). `DCtran()` gains a `CKTcheckpoint`-gated branch that opens a fresh output plot (no live plot to 666-relink across a reload), initializes the XSPICE breakpoint markers, and fixes up the breakpoint list before jumping into the stepping loop; the rhs length is keyed off `SMPmatSize+1` (not `CKTmaxEqNum`). A stored signature rejects a mismatched circuit; Sparse-solver only (KLU rejected with a clear message). Verified 19/19: resumed waveform bit-identical to an uninterrupted run for RC-step/pulse/diode built-ins and $\sim 2e-7$ for a compiled OSDI diode, across a separate process

E-132 *spicelib/analysis/dcpss.c, spicelib/parser/inp2dot.c, spicelib/analysis/psssetp.c, include/ngspice/pssdefs.h*

periodic S-parameters (`.psp`) — small-signal scattering parameters around a PSS operating point, including conversion between the input frequency and its sidebands $f_{in+k} \cdot f_0$ (mixers / switched circuits, where a static-DC `.sp` cannot see the conversion). Sits on the PSS→conversion-matrix suite (E-117–126): after PSS, `psp_sweep` excites each RF port (`portnum/z0`, the `.sp` framework) by driving its branch source ($V=1$, like `.sp`'s `VSRCspupdate`) through the shared $(2M+1)N$ conversion matrix, reads per-sideband port waves in the same Kurosawa power-wave convention, and forms $S^{(k)} = B^{(k)} \cdot A^{-1}$ (dense-complex `cinverse/cmultiply`). `pac_solve_at`'s matrix assembly factored into a reusable `pac_build_matrix`; `PSSdoPSP` flag + `psp` PSS param + `dot_psp` card. Because $S = B \cdot A^{-1}$ is excitation-basis-invariant, sideband 0 reduces exactly to `.sp` for a time-invariant network. Runs under both linear solvers (the conversion matrix is a standalone dense LU; PSS runs under both since E-118). Verified 8/8: sideband-0 matches `.sp` to $\sim 1e-16$ for 1/2/3-port resistive + reactive networks (magnitude and phase) incl. OSDI Verilog-A devices (G + reactive `ddt` stamps), conversion sidebands correctly ~ 0

E-133 *frontend/com_qpss.{c,h}, commands.c, com_commands.h, Makefile.am*

quasi-periodic (two-tone) steady state (`qpss`) — `qpss <expr> <f1> <f2> [periods] [max-order]` computes the two-tone steady-state spectrum: every mixing product $k_1 \cdot f_1 + k_2 \cdot f_2$ including third-order intermodulation (IM3 at $2f_1 - f_2$ / $2f_2 - f_1$) that single-tone AC/PSS cannot show. For commensurate tones (common beat $fb = \gcd(f_1, f_2)$) it runs an ordinary transient over a few beat periods to reach steady state, then evaluates the Fourier coefficient directly at each exact intermod frequency (a direct DFT, exact — no FFT-bin rounding) and labels it by the 2-D index (k_1, k_2) . Deliberately not a slow beat-frequency shooting PSS; a front-end command driving a transient, so solver-independent and works with built-in and OSDI devices. Verified 11/11: analytic fundamentals/IM3/3f of a cubic, no even-order products, the 3:1 IP3 slope law (fund $\times 2$, IM3 $\times 8$ per $2 \times$ drive), beat-frequency derivation, and an OSDI cubic matching the built-in

E-134 *spicelib/analysis/dcpss.c, frontend/com_hb.{c,h}, commands.c, com_commands.h, Makefile.am, include/ngspice/cktdefs.h*

Harmonic Balance (`hb <f0> <K>`) — solves the periodic steady state in the frequency domain by Newton (each node a truncated Fourier series), instead of integrating in time; the real analysis behind ngspice's unimplemented `WITH_HB` stub. Residual $F_k = I_{R,k}(V) + [dq/dt]_k - I_{s,k} = 0$ with the E-121 $(2K+1)N$ conversion matrix as the exact Jacobian. Per iteration: inverse-DFT $V \rightarrow v(t_s)$; drive DC+AC device loads at those voltages for the resistive current $I_R = G \cdot v - b$ (companion `b` saved before `acLoad`, so it's the ACTUAL current not the tangent $G \cdot v$), and $G(t)/C(t)$; DFT; dense complex Newton (`pss_solve`). Nonlinear reactive with NO charge extraction: $dq/dt = C(v) \cdot v'$, so the reactive current is the conversion matrix's $j\omega C$ term applied to V . Solver-independent (KLU + Sparse): the dense complex Newton is HB's own; the sparse solver only reads $G(t)/C(t)$ off the device matrix, so `hb_extract` carries the same `#ifdef KLU` complex-CSC binding (`DEVbindCSCComplex`) as the PAC extraction, and `com_hb` copies the task's KLU mode before building so a bare `hb` honours `.option klu` — verified bit-identical under both. Built-in + OSDI. Verified 8/8 vs transient/fourier, quadratic convergence: nonlinear R (analytic 3rd harmonic), R+C, a real diode rectifier (junction limiting), OSDI varactor $Q(v)$ 2nd harmonic,

KLU==Sparse parity

E-135 *spicelib/analysis/dcpss.c, examples/hb_examples/verify_hb.py*

HB source-stepping continuation — makes E-134 Harmonic Balance robust on strongly-driven circuits where a cold full-strength Newton diverges ($|F| \rightarrow 1e69$). HBAnalyze's Newton loop is wrapped in an adaptive homotopy: every independent source scaled by λ : $0 \rightarrow 1$, each level warm-started from the last converged V . The residual becomes $F = I_R + I_C - \lambda \cdot I_S$; on level success V is checkpointed and $d\lambda$ grows ($\times 1.7$), on failure (Newton exhausted, residual non-finite, or singular Jacobian) V is restored and $d\lambda$ halves; $d\lambda < 1e-5 \rightarrow$ reported as no reachable steady state. First level is full strength ($d\lambda=1$) so easy circuits converge at $\lambda=1$ immediately, bit-identical to the plain solve. Automatic (no new syntax); set `hb_verbose` shows the λ ramp, summary reports iterations + continuation steps. Verified 9/9: a strongly-driven diode rectifier (5 V into 20 Ω , sharp $I_S=1e-14$) diverges cold but converges in 3 continuation steps, matching transient fourier to $<0.1\%$ for $DC/f_0/2f_0$; the 8 existing HB checks stay bit-identical

E-136 *spicelib/analysis/dcpss.c, frontend/com_qpss.c, frontend/commands.c, include/ngspice/cktdefs.h, examples/qpss_examples/verify_qpss_hb.py*

two-tone Harmonic Balance (`qpss <expr> <f1> <f2> hb [K1] [K2]`) — the TRUE, incommensurate-capable quasi-periodic steady state, a frequency-domain HB engine alongside the E-133 transient qpss (which is unchanged, commensurate-only). Each node is a 2-D Fourier series $v(t)=\sum\sum V_{\{k1,k2\}} e^{j(k1\omega_1+k2\omega_2)t}$; QPSShb samples devices on a 2-D phase grid (θ_1, θ_2) — time never appears, so incommensurate tones (no beat period) just work — and 2-D DFTs to the conversion matrix $H_{\{(n),(m)\}}=G_{\{n-m\}}+j\omega_m \cdot C_{\{n-m\}}$ (`qp_build_matrix`), Newton-solved by `pss_solve` with the E-135 source stepping. Sources captured by an oversampled least-squares APFT ($\Gamma^H \Gamma$) $I_S = \Gamma^H b$ (a square Vandermonde is catastrophically ill-conditioned past a few harmonics). Reactive needs NO charge extraction ($dq/dt=C(v)v'$); the converged operating point is retained (`qpss_hb_saved`) for qpac (E-137). Solver-independent (KLU + Sparse) — same `#ifdef KLU` complex-CSC binding as PAC/HB. Built-in + OSDI. Verified 7/7: analytic cubic $|IM3(2,-1)|/|3rd(3,0)|=3$, even products ~ 0 , 3:1 IP3 slope, incommensurate 2 tones (E-133 cannot), HB==transient (commensurate), KLU==Sparse; E-133 `verify_qpss.py` stays 11/11

E-137 *spicelib/analysis/dcpss.c, frontend/com_qpac.{c,h}, commands.c, com_commands.h, Makefile.am, include/ngspice/cktdefs.h, examples/qpss_examples/verify_qpac.py*

two-tone small-signal QPAC (`qpac <f_in>`) — completes the quasi-periodic gap: the two-tone analogue of PAC. Run after `qpss ... hb`, QPACAnalyze injects a small signal at `f_in` around the retained QPSS operating point (`qpss_hb_saved`) and reports the response at every sideband $f_{in}+k_1 \cdot f_1+k_2 \cdot f_2$, mixing it through the SAME 2-D conversion matrix $H_{\{(n),(m)\}}=G_{\{n-m\}}+j\omega_m \cdot C_{\{n-m\}}$ (`qp_build_matrix` at `f_in`) that the QPSS Newton used as its Jacobian. Stimulus placed in the $(0,0)$ sideband — a netlist AC-source RHS B_0 (captured bias-independent at the op-point) or a unit-current fallback — one dense `pss_solve`. Adds no device evaluation $\rightarrow$ solver-independent (KLU + Sparse) by construction. Verified 7/7: reduce-to-AC (pump $\rightarrow 0$ 2 direct $(0,0) =$ plain .ac response = R , sidebands vanish), v^2 -pump conversion ratio $| (1,1) | / | (2,0) | = 2$, equal-tone symmetry, clean no-op-point error, KLU==Sparse; E-136 (7/7) + E-133 (11/11) unaffected

E-138 *spicelib/analysis/dcpss.c, frontend/com_qpnoise.{c,h}, commands.c, com_commands.h, Makefile.am, include/ngspice/cktdefs.h, examples/qpss_examples/verify_qpnoise.py*

two-tone small-signal QPnoise (`qpnoise <output_node> <f_in>`) — quasi-periodic noise, the two-tone analogue of pnoise (E-124), on the retained `qpss ... hb` operating point. Reports output + input-referred noise density at `f_in`, folding every device's noise over all sidebands $f_{in}+k_1 \cdot f_1+k_2 \cdot f_2$ (mixer/PA noise conversion invisible to a static .noise). QPnoiseAnalyze biases the devices at the QPSS operating point, then `qp_solve_adjoint` transposes the 2-D conver-

sion matrix (qp_build_matrix) and solves $H^T \Psi = e_{\text{out}, (0,0)}$ — one adjoint gives the transimpedance from every (node, sideband) to the output; each device's DEVnoise computes $S \cdot |\Psi|^2$ (reading the transfer from CKTrhs/CKTirhs) and the sum over all Nh harmonics is the folded onoise; input-referred via the QPAC gain. Reads only retained data → solver-independent (KLU + Sparse) by construction. With no pump the conversion matrix is block-diagonal so only (0,0) survives → reduces to .noise. Verified 6/6: reduce-to-noise (pump→0 → onoise == plain .noise == 4kTR exactly), conversion active under pump, inoise=onoise/gain², clean no-op-point error, KLU==Sparse; E-133 (11/11) + E-136 (7/7) + E-137 (7/7) unaffected

E-139 *spicelib/analysis/dcpss.c, frontend/com_qpnoise.c, frontend/commands.c, include/ngspice/cktdefs.h, examples/qpss_examples/verify_qpnoise.py*

cyclostationary QPnoise (qpnoise <output_node> <f_in> cyclo) — upgrades E-138 to a time-varying device PSD. Under a two-tone pump a device's bias swings over the period (a diode's shot noise $2qI_D(t)$ spikes when it conducts), so QPnoiseAnalyze gains a cyclo branch using the identity $\text{onoise} = (1/P) \cdot \sum_s S(t_s) \cdot |A_s|^2$, where $A_s(j) = \sum_{\{(k1,k2)\}} \Psi_{\{(k1,k2)\}}(j) \cdot e^{j2\pi(k1 s1/P1 + k2 s2/P2)}$ is the inverse 2-D DFT of the adjoint transfers (the time-domain transimpedance at phase sample s). Each sample re-biases the devices at the retained op-point (qp_synth, P1/P2 now stored in qp_harm) with per-sample junction settling (the E-134 MODEINITFLOAT walk — else a limited diode reports a stale bias and the "cyclostationary" result collapses onto the stationary one), evaluates $S(t_s)$, folds through A_s , and averages. By Parseval reduces to the stationary sum (and .noise) when S is constant. Verified 10/10 (6 E-138 + 4 cyclo): cyclo reduce-to-noise, Parseval (bias-independent thermal PSD → cyclo==stationary even under pump), a hard-pumped diode where cyclo differs from stationary by ~8× (switching-mixer cyclostationary enhancement, visible only with the settling), cyclo KLU==Sparse; E-133/136/137 unaffected

E-140 *spicelib/analysis/dcpss.c, frontend/com_hbosc.{c,h}, commands.c, com_commands.h, Makefile.am, include/ngspice/cktdefs.h, examples/phasenoise_examples/verify_phasenoise.py*

oscillator phase noise (hbosc <oscnode> <K> [fguess] [tstab] + phasenoise <fstart> <fstop> [pts]) — closes the periodic/phase-noise gap with the phase-noise piece. HBOSC-analyze is an autonomous harmonic balance: an oscillator has no source, so $F(V)=I_R+[dq/dt]=0$ is solved for the harmonics AND the unknown oscillation frequency ω_0 . The conversion matrix $dF/dV=H$ is singular (right null space = phase mode $u_k=j k V_k$, since the oscillator's phase is free), so Newton runs on the bordered system $[H, \partial F/\partial \omega_0; u^*, 0]$ (nonsingular by bordering, $\partial F/\partial \omega_0=I_C/\omega_0$, gauge row u^*), transient-seeded (hbosc runs a short transient from the deck's .ic, reads amplitude + zero-crossing frequency), converging quadratically (LC osc: 4 iters to $|F|=3e-12$; inductors fit the $G+j\omega C$ matrix via branch-current MNA). PhaseNoiseAnalyze builds the adjoint of H at OFFSET df with the unit at the carrier sideband ($m=1$) and folds device noise (DEVnoise); as $df \rightarrow 0$ the limit-cycle matrix goes singular through the phase mode so the transimpedance diverges as $1/df \rightarrow$ folded noise $1/df^2$, normalized to carrier power $2|V1|^2 \rightarrow L(df)$ in dBc/Hz. Verified 8/8 on an LC oscillator: autonomous HB converges to f_0 +describing-function amplitude, $L(df)$ has the -20 dB/dec skirt near carrier flattening into the noise floor at a physical level (≈ -147 dBc/Hz @1kHz), thermal scaling $L \propto T$ ($2 \times T \rightarrow +3$ dB exactly, pinning the absolute), clean no-op-point error, KLU==Sparse; pnoise (9/9) + qpnoise (10/10) unaffected

E-141 *spicelib/analysis/dcpss.c, frontend/com_qpxf.{c,h}, commands.c, com_commands.h, Makefile.am, include/ngspice/cktdefs.h, examples/qpss_examples/verify_qpxf.py*

two-tone small-signal QPXF (qpxf <output_node> <f_in>) — the quasi-periodic transfer function, the ADJOINT of QPAC (E-137), completing the two-tone small-signal suite (QPSS→QPAC→QPnoise→QPXF $\equiv$ PSS→PAC→pnoise→PXF). QPXFAnalyze runs one adjoint solve $H^T \Psi = e_{\text{out}, (0,0)}$ (qp_solve_adjoint, reused from QPnoise E-138) and dots each sideband block of Ψ with the netlist AC-source pattern B0 (the QPAC stimulus, captured at the op-point) → the transfer $H_{\{(k1,k2)\}} = \sum_j \Psi_{\{(k1,k2),j\}} \cdot B0_j$ from an input at every

sideband $f_{in} + k_1 \cdot f_1 + k_2 \cdot f_2$ to the output, all from ONE adjoint. By the reciprocity identity ($H^{-1}B$)_{out}=($H^T e_{out}$)^TB the sideband-(0,0) transfer is bit-identical to the QPAC response (the PXF↔PAC cross-check, E-125). Reads only retained data → solver-independent (KLU + Sparse). Verified 6/6: reciprocity (QPXF(0,0)==QPAC(0,0) bit-identical), conversion-sideband reciprocity, reduce-to-XF (pump→0 ? (0,0)=plain transfer=R, sidebands vanish), clean no-op-point error, KLU==Sparse; QPSS 11/11 + QPSS-HB 7/7 + QPAC 7/7 + QPnoise 10/10 unaffected

E-142 *spicelib/analysis/dcpss.c, frontend/com_qpac.{c,h}, com_qpnoise.c, com_qpxf.c, commands.c, include/ngspice/cktdefs.h, examples/qps_sweep/verify_qps_sweep.py*

input-frequency sweep for the two-tone small-signal analyses. qpac/qpnoise/qpxf gain a <dec|oct|lin> N fstart fstop sweep of f_{in} that emits a plottable ngspice plot (conversion gain / noise figure / image-rejection curves vs frequency), matching how .ac/.pnoise/.pxf sweep. QPACsweep/QPnoiseSweep/QPXFswep step f_{in} and reuse the exact single-frequency solve per point: qpac records per-node |(0,0) response|, qpnoise records onoise_spectrum/inoise_spectrum, qpxf records xf (in-band (0,0)) + xf_conv (total conversion $\sqrt{\sum |H_{\{sb \neq 0\}}|^2}$). The analysis fills plain arrays; the front-end builds the plot with a frequency scale via the nutmeg vector API (plot_alloc/dvec_alloc/vec_new) — the correct layer for a front-end command (the analysis OUTp framework dereferences a live analysis job's JOBtype, which a front-end command lacks → was crashing). Shared helpers qp_steptype/qp_sweep_maxpts/qp_emit_plot. Verified 5/5: swept value at 0.3G == single-frequency qpac/qpxf/qpnoise (machine precision), reactive roll-off vs f_{in} , correct point count; single-frequency QPAC 7/7 + QPnoise 10/10 + QPXF 6/6 unaffected

E-143 *frontend/com_optimize.c, frontend/commands.c, examples/optimize_examples/verify_optimize.py, examples/optimize_examples/optimize_ls_demo.cir*

gradient least-squares curve fitting for **optimize**. The built-in optimizer (E-130, derivative-free Nelder-Mead on a scalar -minimize) gains a least-squares mode: one or more weighted -target <expr> <value> [<weight>] measurements, optionally spread over several -analysis stages (each -analysis opens a stage; following -targets attach to it, so a single fit can combine e.g. a DC operating point AND an AC response), are fitted with Levenberg-Marquardt — a forward/backward finite-difference Jacobian J, normal equations $A=J^T J$, $g=J^T r$, damped solve $(A+\lambda \cdot \text{diag}(A)) \delta = -g$ via a small partial-pivot Gauss solver, standard λ up/down loop. Exploiting the sum-of-squares structure it converges in far fewer analysis runs than the simplex (27 vs 67 evals on a two-target RC fit). -method nm|lm overrides the auto choice (LS→lm, scalar→nm); the scalar -minimize/Nelder-Mead path is unchanged; -minimize and -target are mutually exclusive. Parameters are still applied in place with alter (no re-source) and the search runs in normalized [0,1] space. All changes in com_optimize.c (+ one help string); no new files, no ABI change. Verified 23/23 (E-130 checks [1]–[5] + E-143 [6]–[10]): 1-param/3-target/3-stage RC magnitude fit, 2-param DC+AC multi-analysis fit (R1=3k,R2=2k), LM-fewer-evals-than-NM, OSDI/Verilog-A diode extraction (recover both is→1e-14 and n→1.2 from two measured I-V points by weighted least squares), and input validation (lm-without-target, minimize+target, multi-analysis+scalar all rejected)

E-144 *frontend/com_optimize.c, frontend/inp.c, frontend/commands.c, examples/optimize_examples/verify_optimize.py, examples/optimize_examples/optimize_dparam_demo.cir*

optimize tunes symbolic .param values via a new knob kind -dparam. -param remains an alter target (device/instance, changed in place); -dparam names a netlist .param symbol (e.g. .param w=1u,R1={500*k}) which — being expanded at parse time — is changed with alterparam <name>=<value> + a reset that re-sources the deck (re-evaluating .param expressions and re-stamping device values). Per parameter a kind (OPT_ALTER / OPT_DECKPARAM) is tracked; when any -dparam is present opt_eval applies ALL deck params first, re-sources ONCE, THEN applies the in-place alter params (reset rebuilds from the deck and would otherwise wipe an earlier alter — this ordering is what lets the two kinds mix; verified on a joint .param+device fit). No

–dparam the E-130/-143 fast path is untouched (no reset, no perf hit). The two re-source banners (Reset re-loads circuit ... and Circuit: ... in inp.c) are gated by the existing ft_optimizing flag so hundreds of inner re-sources are silent (only the final leave-at-optimum run shows); ft_optimizing re-asserted after each reset. Also fixed 5 pre-existing –wsign-conversion warnings in inp.c's *ng_script_with_params handler (argc/size unsigned). Closes the last E-143 follow-up. No new files, no ABI change. Verified 31/31 (E-130/-143 [1]–[10] + E-144 [11]–[15]): scalar .param fit (rtop→2333.3), mixed –dparam+–param joint fit (rtop=3k,R2=2k), .param in expression (k→6), least-squares –dparam (LM), quiet re-source (≤1 banner); AC + OSDI-device-value .param verified manually

E-145 *frontend/com_optimize.c, frontend/commands.c, examples/optimize_examples/verify_optimize.py, examples/optimize_examples/optresm.va, examples/optimize_examples/optimize_mparam_demo.cir*

optimize tunes .model-card parameters via a third knob kind –mparam. –param remains an alter instance target and –dparam a symbolic .param (re-source); –mparam <@model[param]> names a .model-card parameter (e.g. @dmod[is], @rmod[r]) — which is NOT alter-reachable and NOT .dc-sweepable (only sources/resistors/instance params are) — and is changed in place with altermod <name>=<value>, no re-source (as cheap per eval as –param, unlike –dparam). A new kind value OPT_MODELPARAM; in opt_eval the deck params are re-sourced first (E-144), then the in-place knobs: alter for instances, altermod for models. @model[param] is passed straight to altermod (which accepts that accessor form), mirroring how –param @m1[w] emits alter @m1[w]=... Circuits with only –param/–mparam keep the E-130/-143 fast path (no reset). All three knob kinds mix in one run. Change confined to com_optimize.c (+1 help line); no new source files, no inp.c change, no ABI change. Verified 39/39 (E-130/-143/-144 [1]–[15] + E-145 [16]–[20]): OSDI model-param fit (@rmod[r]→3k), built-in diode model-param fit (@dmod[is]→1.22e-14), determined model+instance joint fit (r=3k,R2=2k), –mparam does 0 re-sources (in-place fast path), and all three kinds (–dparam+–mparam+–param) coexist and converge

E-146 *frontend/com_sweep.c (new), frontend/com_sweep.h (new), frontend/commands.c, frontend/com_commands.h, frontend/inp.c, frontend/Makefile.am, examples/sweep_examples/**

universal **sweep** command + .sweep card. Sweeps ANY circuit knob, generalizing .dc (which steps only sources/resistors/instance params). sw_kind auto-detects the knob: @X[y] with ft_sim→findModel(ckt,X) non-NULL ? model param (altermod), else instance (alter); a bare name with nupa_get_param(name,&found) found ? .param (alterparam+reset), else device (alter) — so model params and .params, impossible with .dc, are now sweepable. sweep <knob> (<start> <stop> <step> | lin|dec|oct N a b | list ...) [–analysis <cmd>] [–output <name>=<expr> ...]: sets the knob, runs the inner analysis (default op) at each point, evaluates each output (last value; default = all node voltages like .dc), and emits a summary plot named sweep with the knob values as the scale (nutmeg vector API, front-end layer per E-142). The per-point analysis plots are retained (tran1,tran2,...). .sweep card: inp.c collects a top-level .sweep line into the post-parse control list as a sweep command; a static sweep_active re-entrancy guard stops a .param re-source (reset re-runs the .sweep card) from recursing. New source files registered in commands.c/com_commands.h/Makefile.am. Verified 11/11 (verify_sweep.py): sweep R1 reproduces built-in .dc R1 bit-for-bit, model-param (@rmod[r]) + .param (rtop) sweeps vs analytic divider, auto-detection routing, .sweep card == command form (no recursion), AC named-output vs analytic |H(1kHz)|, transient settled value, lin/list/step point counts, multi-output

E-149 *maths/misc/randnumb.c, include/ngspice/randnumb.h, frontend/numparam/xpressn.c, frontend/numparam/spicenum.c, frontend/commands.c, examples/lhs_examples/**

Latin-Hypercube Monte Carlo sampling (mcsample), closing the low-discrepancy-sampling gap vs commercial tools. A stratified sampler in randnumb.c holds the mode/N/sample-index/dimension-counter/seed; each random dimension's stratum permutation (Fisher-Yates over 0..N–1) + per-sample jitter are generated lazily from a self-contained splitmix64 seeded

by (seed, dimension) — reproducible and order-independent. `mc_sample_uniform()` = `(perm[d][i]+jitter[d][i])/N`; `mc_sample_gauss()` maps it through `inv_normal_cdf` (Acklam probit, rel err < 1.15e-9) so Gaussian tails stratify too. Sample boundary: `spicenum.c`'s `nupa_signal(NUPADECKCOPY)` edge (once per reset re-source, guarded by the existing `firstsignalS` latch) calls `mc_sample_advance()` (step sample, rewind dimension) before that pass's `.param` draws. Draw hook (`xpressn.c`): `agauss/gauss` take one `mc_sample_gauss()` and `aunif/unif/limit` one `2*mc_sample_uniform()-1` when LHS is active, else the unchanged `gauss1()/drand()` (LHS-off is byte-identical to before). Command `com_mcsample` (`randnumb.c`, registered in `commands.c` both tables) parses `lhs <N> [seed <s>] / random / off`. Front-end only, so solver-independent. Verified 5/5 under both solvers (`verify_lhs.py`): stratification (each of 48 strata hit once vs random's ~32/48), two-param multi-dimension, ~130x lower Var(sample-mean) at N=32 (both converge to the same mean), seed reproducibility, analytic mean/sigma correctness

E-150 *frontend/com_sweep.c, frontend/commands.c, maths/misc/randnumb.c, include/ngspice/randnumb.h, examples/_setup.py, examples/highsigma_examples/**

high-sigma rare-event estimation (`highsigma`), the last statistical ROW versus a commercial simulator -- 4-6 sigma failure probabilities (SRAM / standard-cell yield) that plain MC cannot reach (1e-7..1e-9 needs 1e7..1e9 runs). Scaled-sigma importance sampling: a new `MC_MODE_SSS` in the E-149 sampler draws each Gaussian `.param` from the lambda-inflated normal ($z = \text{lambda} * \text{gauss1}()$) and accumulates that draw's log likelihood-ratio $\log(\text{lambda}) - (z^2/2)(1 - 1/\text{lambda}^2)$ into `sss_logw`; `mc_sample_weight() = exp(sss_logw)` reweights the sample so the estimator is unbiased for the nominal probability. Uniform `.params` are bounded so SSS leaves them uninflated (weight 1). Direction-free -- no gradient / sensitivity / most-probable-failure-point. `mc_sss_config(N, lambda, seed)` seeds the global PRNG for reproducibility; `mc_sss_off()` reverts. Command `com_highsigma` lives in `com_sweep.c` (reuses its `sw_run_cmd` reset/analysis runner and `sw_eval_expr`), registered in `commands.c`: `highsigma <N> [-scale <lambda>] [-seed <s>] [-analysis <cmd>] -metric <expr> [-max <hi>] [-min <lo>]` loops N samples (reset -> analysis -> eval metric -> spec test -> read weight), reports P(fail), relative error, equivalent one-sided sigma-to-fail ($-\Phi^{-1}(P)$) and failure count, and leaves them in the `highsigma_*` vectors/vars. The spec is -max/-min values (not a >/< inside the expr, which a control command reads as an I/O redirect). Front-end only, solver-independent; the verify is heavy (thousands of re-sources) so `examples/_setup.py` marks `highsigma` `SPARSE_ONLY`. Verified vs analytic $\Phi(-\beta)$ (`verify_highsigma.py`): $\beta=4 \rightarrow$ sigma 4.000, P 3.16e-5 (true 3.17e-5); deep tail $\beta=5$ (P 2.87e-7) recovered from 6000 runs where plain MC expects ~0.0017 failures; two-sided spec ~doubles P; seed reproducibility exact; two independent Gaussians combine as $N(\sqrt{s_1^2+s_2^2})$

E-151 *maths/misc/randnumb.c, include/ngspice/randnumb.h, frontend/numparam/xpressn.c, frontend/com_sweep.c, frontend/commands.c, examples/_setup.py, examples/yield_examples/**

process/mismatch correlations + packaged Monte Carlo yield -- closes the last two partial (warn) statistical rows vs a commercial simulator. (1) CORRELATIONS: `mc_corr_config(k, matrix)` Cholesky-factors a $k \times k$ correlation matrix (row-major, unit diagonal; rejects a non-positive-definite one) into lower-triangular `corr_L`; `mc_corr_component(i)` lazily draws the k underlying z 's once per sample THROUGH `mc_sample_gauss()` (so LHS stratification / SSS weighting apply), computes $y = L*z$, caches it, and returns `y[i]`. The cache is reset in `mc_sample_advance()` which now runs its reset in EVERY mode (so correlation works under plain MC, not just LHS/SSS). New `.param` function `mvnorm(i)` (`frontend/numparam/xpressn.c`: added to `fmathS`, `XFU_MVNORM` enum in `fmathS` order, and the dispatch) returns the i -th correlated standard normal; with no matrix registered it degrades to an independent draw. `com_mccorr(mccorr <k> <m11>. <mkk> | off)` in `randnumb.c`. A reusable `mc_lhs_config(N, seed)` was factored out of `com_mcsample`. (2) YIELD: `com_montecarlo` in `com_sweep.c` (reuses `sw_run_cmd` / `sw_eval_expr`), registered in `commands.c`: `montecarlo <N> [-lhs] [-seed <s>] [-analysis <cmd>] (-spec <metric> [-max <hi>] [-min <lo>])...` runs N samples (reset -> analysis -> eval each

spec), counts a sample PASS only if every spec is within limits, and reports the yield with a Wilson 95% score interval and per-spec violation counts, leaving `montecarlo_yield/montecarlo_npass/montecarlo_n`. `-lhs` uses `mc_lhs_config` for a lower-variance estimate; corners are the ordinary `.lib` selection. Front-end only, solver-independent; heavy deck so `examples/_setup.py` marks `yield SPARSE_ONLY`. Verified (`verify_yield.py`, Sparse): `mccorr rho=+0.7/-0.6` -> empirical corr 0.711/-0.614 with correct means/sigmas; `mvnorm` without `mccorr` -> corr ~0; non-PD matrix rejected; single two-sided spec yield 0.8660 (true $P(|Z|<1.5)=0.8664$); two independent specs multiply (0.7508 vs 0.7506); `-lhs` unbiased and ~800x lower yield-estimate variance; positive correlation raises the joint yield (0.75 -> 0.82). Demo: a matched divider yields ~100% process-correlated ($\rho=0.9$) vs ~74% independent

E-152 *include/ngspice/optdefs.h, include/ngspice/tskdefs.h, include/ngspice/cktdefs.h, include/ngspice/smpdefs.h, spicelib/analysis/cktsopt.c, spicelib/analysis/cktntask.c, spicelib/analysis/cktdojob.c, maths/ni/niinit.c, maths/KLU/klusmp.c, examples/klu_tuning_examples/**

KLU matrix reordering + scaling controls. The KLU direct solver ran on the compiled-in `klu_defaults` (AMD ordering, max row scaling, BTF on) with no way to change them, and its one exposed knob `klu_memgrow_factor` was broken (`task->TSKkluMemGrowFactor = (val->rValue == 1.2)` stored a boolean, not the value). E-152 exposes three new `.options --klu_ordering=amd|colamd, klu_scale=none|sum|max, klu_btf=on|off` -- as IF_STRING options (friendly names, mapped to KLU's integer codes 0/1, 0/1/2, 1/0 in the `cktsopt.c` handlers) and fixes the memgrow assignment. Plumbing mirrors the existing memgrow path exactly: new `OPT_KLU_ORDERING/SCALE/BTF` enums (`optdefs.h`); `TSKklu*/CKTklu*` int fields on the task (`tskdefs.h`), circuit (`cktdefs.h`) and `SMPmatrix` (`smpdefs.h`) structs; defaults 0/2/1 that MATCH `klu_defaults` set in `cktntask.c` (so behaviour is unchanged unless the user overrides) plus the task copy; `task->ckt` in `cktdojob.c`; `ckt->CKTmatrix` in `niinit.c`; and in `klusmp.c` `SMPnewMatrix`, right after `klu_defaults()`, `Common->ordering/scale/btf = Matrix->CKTkluOrdering/Scale/BTF`. KLU-only; changes only HOW the matrix factors, never the physical solution; the numerics are untouched. Verified (`verify_klu_tuning.py`, KLU-only) on a resistor grid: every ordering/scaling/BTF setting gives the physically identical solution (max relative spread $2.8e-14$); AMD vs COLAMD and `scale=max` vs `scale=none` change the factorization arithmetic so the full-precision result differs in its last digits (rel diff $\sim 1.2e-14 / 2.8e-14$ -- a deterministic, nonzero proof the knobs reach KLU); invalid values warn; a plain `.option klu` equals `amd/max/btf-on` bit-for-bit; a wide-dynamic-range network solves correctly under `none/sum/max`

E-153 *maths/ni/niiter.c, maths/KLU/klusmp.c, include/ngspice/smpdefs.h, include/ngspice/optdefs.h, include/ngspice/tskdefs.h, include/ngspice/cktdefs.h, spicelib/analysis/cktsopt.c, spicelib/analysis/cktntask.c, spicelib/analysis/cktdojob.c, examples/trustregion_examples/**

Levenberg-Marquardt trust-region Newton (`.option trustregion`, off by default), completing the damped/trust-region-Newton row alongside the E-111 line search. In the Newton loop (`niiter.c`): the E-111 KCL-residual merit is reused (gated on `linesearch||trustregion`); when a dimensionless damping $\lambda > 0$ is in effect, $\mu = \lambda * ||\text{diag}(J)||$ is added to the effective diagonal `gmin` (`trGmin`, passed to `SMPreorder/SMPluFac`) AND `mux_k` to the RHS (the E-127 pseudo-transient coupling with $x_{\text{prev}}=x_k$), so the solve yields the exact damped step $x_{k+1}=x_k-(J+\mu I)^{-1} F(x_k)$; a post-solve accept/reject block re-loads at the trial point, computes $||F(x_{\text{new}})||$, and either accepts (`relax lambda=0.25`) or rejects (`restore x_k, grow lambda, force another iteration`), reusing the E-111 state-save/limiting-reset machinery; a convergence guard forbids converging while $\lambda > 0$. The $||\text{diag}(J)||$ scale (new `SMPdiagNorm` in `klusmp.c`, for both the KLU and Sparse matrices) makes λ dimensionless (Marquardt / scale-invariant). Unlike the line search (which only shortens a fixed Newton direction) large μ re-aims the step toward steepest descent, regularizing an ill-conditioned Jacobian. Result-neutral: the fixed point is $F=0$ for any μ and $\lambda > 0$ on success, so it converges to the same operating point -- verified BIT-IDENTICAL to plain Newton (diode, BJT, resistor divider, plus transient) under both solvers. Option plumbing mirrors `linesearch` (`OPT_TRUSTREGION`, `IF_FLAG`, `TSK/CKT` fields + `CKTTrLambda`). Solver-

independent. HONEST FINDING: instrumented step-rejection counting showed ZERO rejections on every circuit tried (diode strings, behavioral exponentials, cubics, negative-resistance oscillators) -- ngspice globalizes Newton at the DEVICE level (per-device junction limiting: limexp/pnjlim/fetlim, 30 families) which damps the controlling voltages before the residual is computed, so a residual-increasing overshoot never reaches the solver-level merit; the trust-region is thus a correct, safe solver-level regularization that stays inert on typical circuits

E-154 *frontend/com_envelope.c (new), frontend/com_envelope.h (new), frontend/commands.c, frontend/com_commands.h, frontend/Makefile.am, spicelib/analysis/envelope.c (new), spicelib/analysis/Makefile.am, include/ngspice/cktdefs.h*

Envelope Following (envelope <node> <fc> <tstop> [nppp] [m] [maxm] [reltol] [settle]), the last missing analysis in the RF/periodic-steady-state suite. For a carrier-driven circuit whose amplitude/phase envelope varies slowly over many carrier periods, it samples the state once per carrier period $T=1/f_c$ and integrates the slow drift, jumping M periods at a time. The exact per-period map is $X_{n+1}=\phi(X_n)$, ϕ = one-period DAE integration; the envelope obeys $dX/dn \sim \phi(X)-X$. The naive forward-Euler jump $X_{n+M}=X_n+M(\phi(X_n)-X_n)$ is UNSTABLE on high- Q circuits (the one-period monodromy has unit-circle eigenvalues, so $I+M(\phi-I)$ amplifies $\rightarrow$ blow-up; this is what shelved an earlier attempt). The new engine (spicelib/analysis/envelope.c, EFanalysis) uses the IMPLICIT backward-Euler jump $X_{n+M}=X_n+M(\phi(X_{n+M})-X_{n+M})$, solved by Newton $[(1+M)I - M*\phi] dY = -G$ with $\phi=d\phi/dY$ the one-period monodromy (finite-differenced, one extra period-integration per state, dense $N \times N$ solve capped for modest circuits). A-stable, so it tracks a resonator's envelope without diverging; its fixed point $\phi(X)=X$ is the true steady state. Step size M is chosen by step-doubling local-truncation-error control. The one-period map reuses the transient primitives (NIcomCof + NIiter per fixed substep, the dctrans state rotation) on a fixed grid of $nppp$ points in TRAPEZOIDAL mode (backward-Euler damps a high- Q resonance); ϕ is SELF-STARTING (a frozen-history variant plateaued at a false low fixed point) with a sub-divided backward-Euler first sub-step to bound the restart damping. The com_envelope.c front-end command runs a short settling transient, calls EFanalysis, and emits an envelope plot ($_amp = 2|V_1|$, $_dc$, $_re$, $_im$ vs time; nutmeg vector API). Solver-independent. Verified against a full .tran under both linear solvers: EF amplitude tracks the transient across a $Q \sim 3160$ tank ring-up to $<3\%$ and to the steady state ($<0.5\%$), stays bounded over 3000 carrier periods (26 envelope samples), and tracks a $Q \sim 316$ tank to $\sim 1.6\%$. Completes the RF suite

E-155 *frontend/com_reduce.c (new), frontend/com_reduce.h (new), frontend/commands.c, frontend/com_commands.h, frontend/Makefile.am, spicelib/analysis/rcreduce.c (new), spicelib/analysis/Makefile.am, include/ngspice/cktdefs.h, examples/reduce_examples/**

RC network reduction (reduce <fmax> [factor f] [file fname] [name subckt] [keep node ...]), moving the post-layout "RC reduction / model-order reduction" gap to done. Post-layout extraction produces enormous linear parasitic R/C networks (thousands-millions of interior nodes); reduce collapses the R/C network to a small ELECTRICALLY EQUIVALENT one preserving the port behaviour over DC..fmax, written as an ordinary .subckt of R's and C's. Method (spicelib/analysis/rcreduce.c, CKTreduceRC): TICER (Time-Constant Equilibration Reduction) -- Schur-complement (Gaussian) elimination of interior nodes of $Y(s)=G+sC$, kept first order in s so the reduced network is REALIZABLE as R's and C's (no model-order-reduction black box, no passive-synthesis step). Per elimination of node n , for each neighbour pair (a,b) : $G[a,b] -= G[a,n]G[n,b]/G[n,n]$; $C[a,b] -= (G[a,n]C[n,b]+C[a,n]G[n,b])/G[n,n] - G[a,n]G[n,b]C[n,n]/G[n,n]^2$. The conductance update is the exact Schur complement so DC is preserved EXACTLY; the capacitance is matched to first order. A node is eliminated only when its self time-constant frequency $f_n=G_n/(2\pi C_n) > factor*fmax$ (quasi-static in the band of interest); factor trades reduction against in-band accuracy, monotonically. PORTS (kept nodes) are auto-detected as every node touched by a device that is NOT a resistor/capacitor (sources, transistors, OSDI Verilog-A devices), plus ground and user keep nodes, via the generic GENnode/terminal-count device interface -- so built-

in and OSDI devices alike. The `com_reduce.c` front-end command enumerates the built-in resistor/capacitor instances (`CKTtypelook`), builds dense G/C over the RC nodes, runs TICER, and emits the reduced subckt (`CKTnodName` for node names). Solver-independent (reads devices + own dense solve). First cut is dense (node cap ~2500); sparse TICER is the follow-up. Verified under both linear solvers: a huge factor eliminates nothing and the emitted subckt reproduces the full AC response BIT-for-BIT; a moderate factor cuts nodes (e.g. 25->6, 4x) with <0.25 dB in-band error and DC exact; the accuracy/reduction tradeoff is monotone in factor; an OSDI device on a node auto-marks it a kept port

E-156 *spicelib/analysis/rcreduce.c, frontend/com_reduce.c, include/ngspice/cktdefs.h, examples/reduce_examples/**

Sparse RC reduction -- makes the E-155 reduce command scale to real post-layout size. The dense NN TICER engine (capped ~2500 nodes) is replaced by a SPARSE one: the network is stored as per-node adjacency lists (`RCadj`: a growable edge list of {neighbour, g, c}), and interior nodes are eliminated in a MINIMUM-DEGREE order (a lazy binary heap keyed on current degree, like sparse LU) so fill-in stays tiny -- a degree-2 chain node merges two series elements with ZERO fill. A FILL GUARD (`maxdeg`, default 12) declines to eliminate a node once its degree grows past the threshold, so a dense mesh core (whose boundary Schur complement is inherently dense) is left intact instead of densifying the whole network (measured: unguarded, one mesh elimination created 3308 fill edges; guarded, 9). The TICER Schur-complement update, DC-exact conductance, frequency criterion (eliminate only when $f_n = G_n / (2\pi C_n) > factor_{fmax}$), and automatic port detection are unchanged from E-155; in the edge representation the diagonal is implicit so eliminating a node just adds Schur edges among its neighbours. Node cap raised 2500 -> 5,000,000; a 65,017-node network reduces in ~4 s. New `maxdeg` command option (`com_reduce.c`) and `CKTreduceRC` parameter (`cktdefs.h`). Also fixes a terminal-ordering pitfall: the reduced .subckt's terminals are emitted in internal-index order (not the order keep nodes were typed), so instantiating with the wrong order silently swaps ports -- the command now prints the correct `x1 <ports...> <name>` instantiation line. Verified under both linear solvers: identity reduction bit-exact (0.00 dB), a moderate factor gives ~4x fewer nodes at <0.25 dB in-band with DC exact, monotone accuracy/reduction tradeoff, OSDI auto-port, and an 8001-node network (past the old dense cap) reduces successfully

E-157 *frontend/com_aging.c (new), frontend/com_aging.h (new), frontend/commands.c, frontend/com_commands.h, frontend/Makefile.am, examples/aging_examples/**

Device aging (`aging <t_target> [rate <opvar>] [param <ageparam>] [dynamic <tstop> [tstep]] [verbose]`), adding the reliability "stress -> degrade -> re-simulate (fresh/aged)" flow (HCI/NBTI/TDDDB) and moving those reliability gap rows to done. The command finds every aging-capable device in the loaded circuit, computes how much it has degraded after `t_target` seconds of operation, writes that back into the device, and RE-STAMPS the circuit so any analysis run afterwards (`op/dc/tran/ac`) sees the aged devices; it runs the fresh stress simulation first, leaving it as the current plot (a "fresh" baseline). The design is MODEL-AGNOSTIC: a device opts in purely by exposing, in its Verilog-A/OSDI model, a degradation-RATE operating-point variable (default `agerate`, in dose-units/second) and a per-instance AGE parameter (default `age`, (`*type=\"instance\"*`)). The engine's only job is to integrate the rate into a dose and feed it back; the model owns the physics that maps age to a parameter shift (a sublinear power law, an Arrhenius factor, ...). Participants are detected by scanning each device TYPE's instance-parameter table (`DEVices[t]->DEVpublic.instanceParms`) for the two keywords, so ordinary resistors/sources are silently skipped and probing never errors. Two modes: STATIC (default) reads the rate at the DC operating point and scales by the lifetime, `age = agerate(op)t_target`; DYNAMIC (`dynamic <tstop>`) runs a transient over one representative window, trapezoidally (time-weighted) integrates the rate, and extrapolates, `age = (INT agerate dt / tstop)t_target` -- capturing duty cycle. Because age is per-instance, devices sharing a .model but at different bias age independently. Implementation (`frontend/com_aging.c`) reuses the `com_optimize` synchronous command-dispatch pattern to run `op/tran`, the nutmeg expression engine to read `@inst[agerate]`, and alter `@inst[age]=...`

to write back; console chatter suppressed via `ft_optimizing` unless `verbose`. Solver-independent. Verified under both linear solvers with a demo square-law NMOS carrying an NBTI hook (`agemos.va`): enumeration picks exactly the ageable devices, the dose is exactly `ratet_target`, the threshold shift matches the analytic $\Delta V_{th} = dv_{th_ref}(\text{age}/\text{age_ref})^{0.25}$ law to 5 sig figs, aged drain current drops, degradation is monotone in lifetime, a near-threshold device loses a larger fraction of its current than a hard-driven one for the same shift, a 30%-duty pulsed gate ages at 0.30x the DC rate (dynamic), and a sub-threshold device accrues zero dose

E-158 *frontend/com_emir.c (new), frontend/com_emir.h (new), frontend/commands.c, frontend/com_commands.h, frontend/Makefile.am, examples/emir_examples/**

Power-grid EMIR (`emir [rail V] [thresh frac] [thick m] [jmax A/m2] [n exp] [tref s] [top k] [verbose]`), adding electromigration + IR-drop reliability sign-off and moving the last reliability gap row to done. After a DC solve of the power-distribution network the command reports two metrics. IR-DROP: for every node, the drop below the ideal rail ($\text{rail} - V(\text{node})$); the rail defaults to the highest node voltage (the supply pad) unless given; reports the worst node and every node past a threshold (default 10% of the rail). ELECTROMIGRATION: for every wire-segment resistor, the current DENSITY $J = |I|/(w\text{thickness})$, ranked; the worst-density segment, every segment past the current-density limit J_{max} , and a Black's-equation relative lifetime $\text{MTTF}/\text{ref} = (J_{\text{max}}/J)^n$ (Black: $\text{MTTF} \sim J^{-n}$; a segment exactly at J_{max} has $\text{MTTF} = \text{the reference lifetime } t_{\text{ref}}$, needing no process-specific prefactor). The physics the analysis exists to expose: EM is set by current DENSITY, not current -- a fat trunk carrying huge current can be safe while a thin wire at a fraction of that current voids first, which is why real grids taper wire width with carried current to hold J roughly constant. Implementation (`frontend/com_emir.c`) runs a fresh op (`com_optimize` synchronous-dispatch pattern), walks the current plot's node-voltage vectors (`plot_cur->pl_dvecs` filtered to `SV_VOLTAGE`) for IR-drop, and enumerates the resistor instances (the device type named "Resistor") reading each segment's current and width via the nutmeg expression engine (`@Rk[i]`, `@Rk[w]`) -- so it needs no device-struct access and works for any resistor; segments without a width are skipped and counted. Both tables are qsort-ranked (IR by drop, EM by density) and truncated to `top`. Solver-independent (reads a DC solution + per-resistor currents). Verified under both linear solvers on a 3-segment tapered ladder off a 1 V rail: worst IR drop 0.30 V (30%) at the far tap matching the exact ladder solve and linear in load current; $J = I/(w\text{thick})$ exact; the worst-density segment is the narrow low-current wire, not the high-current wide trunk; the Black MTTF ratio between two segments equals $(J_2/J_1)^n$; with J_{max} between the trunk and leaf density exactly the under-sized segments fail; rail auto-detect equals an explicit rail; and an OSDI (Verilog-A) current load at a tap is handled identically to a current source

E-159 *(no ngspice source change) examples/compactmodels_examples/**

Real production compact-model bring-up -- a validation/example enhancement that exercises the EXISTING toolchain on actual CMC (Compact Model Coalition) Verilog-A models, so no ngspice or openvaf-r source change. BSIM4 (the ~12,600-line industry-standard bulk MOSFET) is compiled through openvaf-r to OSDI (~6 s) and validated against ngspice's BUILT-IN BSIM4 -- the same model, so a rigorous self-check: the OSDI BSIM4.8 and native BSIM4.8.3 agree to ~2% on the I_d - V_{ds} output family and ~4.5% on the I_d - V_{gs} transfer curve (the residual is the point-release version gap). EKV (EPFL compact MOSFET, which ngspice has NO built-in for) is brought up purely through OSDI, giving textbook saturation. A general finding surfaced: OSDI keeps every internal node STATIC (no dynamic node collapsing), so BSIM4's internal drain/source nodes (`di/si`) -- which the default `rdsmod=0` expects the simulator to collapse onto the external `d/s` -- are left floating and the device conducts ZERO current; enabling the external series-resistance nodes with `rdsmod=1` (the model near-shorts them, 1e3 mho, when no S/D geometry is given) connects them and the device works. Models are compiled in place from the bundled OpenVAF integration-test sources (licenses stay with them). Verified under both linear solvers (6 checks)

E-160 *(no ngspice source change) examples/cmcsweep_examples/*, examples/_setup.py*

CMC compact-model coverage sweep -- the comprehensive follow-up to E-159, again a validation/example enhancement with no ngspice/openvaf-r source change (only `_setup.py` gains `cmcsweep` in its `SPARSE_ONLY` and `REGRESSION_EXCLUDE` sets so the slow 19-model compile runs single-solver and off the fast sweep). It drives EVERY real CMC (Compact Model Coalition) Verilog-A model bundled with OpenVAF through `openvaf-r -> OSDI -> ngspice` and reports a coverage matrix, the same idea as the E-84 LRM sweep but for production compact models. Result: 19 of 19 models COMPILE and 19 of 19 LOAD, spanning every device class -- MOSFET (BSIM3, BSIM4, BSIM6, BSIMBULK, EKV, HiSIM2, HiSIMSOTB, MVSG_CMC), FinFET (BSIMCMG, BSIMIMG), SOI (BSIMSOI, HiSIMHV), GaN HEMT (ASMHEMT), surface-potential (PSP102, PSP103), bipolar/SiGe-HBT (HICUML2, MEXTRAM), and diode (DIODE, DIODE_CMC). Deeper conduction+physics validation for a representative model per class: OSDI BSIM4 and BSIM3 match ngspice's BUILT-IN BSIM4/BSIM3 to ~2%/~7%; EKV gives monotonic saturating I-V; HICUML2 shows bipolar current gain $\beta \sim 100$ on a Gummel plot; DIODE_CMC turns on exponentially ($\times 16000$ over 0.2-1.0 V). Two findings: (1) the E-159 internal-node/`rdsmod=1` issue is BSIM4-SPECIFIC, not universal -- BSIM3 handles the zero-resistance case fine and conducts out of the box (its apparent zero at $V_{gs}=1$ V is just a high ~1.7 V default threshold); (2) HiSIMHV (6 terminals d,g,s,b,sub,temp) needs `cosubnode=1` to enable the substrate node, and with the default `cosubnode=0` its `$fatal` correctly guards the node-count mismatch (confirming `$fatal` works through OSDI). Models compiled in place from the bundled integration-test sources (licenses stay put). Verified (7 checks); the compile/load tiers are solver-independent

E-161 *(no ngspice source change) examples/dynmodels_examples/**

Dynamic (AC/RF) compact-model validation -- extends the E-159/160 DC validation into the dynamic domain, again a validation/example enhancement with no ngspice/openvaf-r source change. Where E-159/160 exercised DC I-V, this exercises the DYNAMIC behavior, which flows through a different code path: OSDI's REACTIVE (charge) Jacobian stamping and ngspice's `.ac` analysis, on the real production models (compiled in place from the bundled CMC sources). BSIM4 C-V: the gate capacitance $C_{gg}(V_{gs})$ is extracted from `.ac` as $\text{Im}(I_{\text{gate}})/\omega$ and swept over gate bias; it rises from ~40 fF in subthreshold (overlap+fringe) to ~129 fF in inversion (approaching $C_{ox}WL$) -- the textbook MOSFET C-V -- and the OSDI-compiled BSIM4 matches ngspice's BUILT-IN BSIM4 to under 1% at every bias (a tighter check than the ~2% DC I-V match, since C_{gg} is dominated by the version-independent oxide term). Cutoff frequency f_T (the frequency where AC current gain $|h_{21}| = |I_{\text{out}}/I_{\text{in}}|$ falls to 1): OSDI BSIM4 $f_T \sim 3.5$ GHz matches the built-in to ~1%; HICUML2 (SiGe HBT, which ngspice has no built-in for) has zero default transit time ($t_0=0$) giving infinite f_T , so a realistic dynamic parameter set ($t_0=10$ ps, 1 fF junction caps) is supplied and the resulting f_T lands right at the transit-time limit $1/(2\pi t_0) \sim 15.9$ GHz and rises with collector current -- textbook charge-control behavior. Verified under BOTH the Sparse and KLU solvers (`.ac` is supported by both), 4 checks

E-162 *frontend/inp.c, examples/hb_examples/verify_hb.py*

`.hb` dot-card for harmonic balance -- gives the E-134 `hb` command netlist dot-card parity with the PSS family (`.pss/.pac/.pnoise/.pxf/.psp/.sp` are dot-cards, while the HB family `hb/qpss/hbosc` were control-block commands only). A top-level `.hb <f0> <K> [points] [maxiter]` card now runs single-tone harmonic balance straight from the deck (no `.control` block needed). Rather than duplicating the HB machinery as a separate analysis JOB, `.hb` reuses the deck->control mechanism Enhancement-146 introduced for `.sweep`: during deck loading (`frontend/inp.c`) a top-level `.hb ...` line is stripped of its leading `.` and appended to the post-parse control list, so it executes as `hb ...` (the `com_hb` command) once the circuit is built -- inheriting all of the E-134 engine (argument parsing, the E-121 conversion-matrix Jacobian, E-135 source-stepping continuation, solver independence). A boundary check (next char after `.hb` is whitespace or end-of-line) keeps it from swallowing any future `.hb*` card. Before this, `.hb` produced unimplemented dot command `' .hb'` -- the only `.hb` handler in ngspice was a disabled `#ifdef WITH_HB` upstream stub (`dot_hb`) that merely redirected to the PSS analysis, not the

E-134 engine. Verified (examples/hb_examples/verify_hb.py, +2 checks): the .hb dot-card run in plain batch mode (no .control) converges and produces a harmonic-balance spectrum bit-for-bit identical to the hb command form on a junction-limited diode rectifier; it threads its optional [points]/[maxiter] arguments and coexists with a following .control block (deck order preserved); both Sparse and KLU give identical results. (Like .sweep, a bare command-style dot-card with no other analysis card prints a benign "no simulations run" batch notice even though HB ran.)

E-163 *frontend/inp.c, examples/qpss_examples/verify_qpss.py, examples/phasenoise_examples/verify_phasenoise.py*

.qpss / .hbosc / .phasenoise dot-cards -- completes the harmonic-balance family's netlist dot-card parity begun by E-162's .hb. Three more top-level command-style cards now run their analyses straight from the deck: .qpss <expr> <f1> <f2> [periods] [maxorder] (two-tone quasi-periodic steady state, E-133/136; add the hb keyword for the frequency-domain form), .hbosc <oscnode> <K> [fguess] [tstab] (autonomous harmonic balance for an oscillator, E-140; the deck needs a .ic to start the oscillation), and .phasenoise <fstart> <fstop> [points] (oscillator phase noise, E-140; run after .hbosc). Each reuses the same one-branch deck->control mechanism as .hb (E-162) / .sweep (E-146): in frontend/inp.c a top-level .qpss/.hbosc/.phasenoise line is stripped of its leading . and appended to the post-parse control list, executing as the corresponding command once the circuit is built -- inheriting the full E-133/136/140 engines and their solver independence. Because the bridge preserves deck order, .hbosc (which finds the oscillator PSS + frequency) and .phasenoise (which reads it) compose, so a complete oscillator phase-noise run needs no .control block. A per-card boundary check (the char after the name is whitespace or end-of-line) keeps the cards distinct -- in particular .hbosc is NOT matched by the .hb branch, whose own boundary check rejects the trailing osc. Verified: the .qpss dot-card produces a two-tone spectrum bit-for-bit identical to the qpss command (examples/qpss_examples/verify_qpss.py); .hbosc + .phasenoise in a deck reproduce the command form's oscillation frequency and phase-noise spectrum exactly with order preserved (examples/phasenoise_examples/verify_phasenoise.py); both under Sparse and KLU. No regression to .hb/.sweep.

E-164 *(no ngspice source change) examples/rfpa_examples/**

Large-signal RF characterization of a real transistor -- a validation/example enhancement (no ngspice/openvaf-r source change) that closes the production-model validation loop (DC E-159, coverage E-160, small-signal E-161, now large-signal). It drives a common-emitter amplifier built from the bundled HICUM/L2 SiGe HBT (compiled in place, with the E-161 dynamic params $t_0=10\text{ps} + 1\text{fF caps}$) hard and extracts the large-signal RF figures of merit via TRANSIENT + Fourier/FFT: AM-AM gain compression (the power gain expands 17.8->20.7 dB then compresses -- the exponential-transconductance signature of a bipolar -- defining P1dB), harmonic distortion (the third harmonic follows the textbook 3:1 slope, $\text{HD3} \sim A^3$), and two-tone third-order intermodulation (the IM3 products $2f_1-f_2 / 2f_2-f_1$ at $\sim -30\text{ dBc}$). IIP3, extracted from the single tone as $A/\sqrt{3*\text{HD3}}$ (the two-tone IM3 is exactly 3x the single-tone HD3 for a third-order nonlinearity), is constant at $\sim 0.134\text{ V}$ across drive level, confirming the third-order model. A real finding: the frequency-domain harmonic-balance engines (hb/qpss, E-134/136, which just gained netlist dot-cards in E-162/163) do NOT converge on this stiff, many-internal-node production model in an amplifier configuration -- hb returns error 103 at any drive level (even 1 mV) and with extra iterations/source-stepping, and the two-tone qpss is prohibitively slow (>2 min per point). Transient+FFT integrates reliably and is the robust route for large-signal RF on a full production compact model; HB/QPSS remain the tool of choice for lighter, well-conditioned circuits. Verified under both linear solvers (4 checks): transient small-signal gain matches .ac (7.744 vs 7.742), HD3 3:1 slope (62.9 vs 64), IIP3 constant (0.134 V, spread <2%), and gain compresses at high drive (7.74->5.97).

E-165 *(no ngspice source change) examples/modelnoise_examples/**

Production compact-model noise validation -- a validation/example enhancement (no ngspice/

openvaf-r source change) exercising the last untested small-signal path, OSDI's .noise stamping of the models' own white_noise / flicker_noise sources, on production models compiled in place from the OpenVAF integration-test sources. BSIM4 (MOSFET): the output-noise spectral density $S_v(f)$ of a common-source amplifier is compared to ngspice's BUILT-IN BSIM4 over the full band and matches to $\sim 1.5\%$ everywhere, covering both the low-frequency $1/f$ FLICKER region (amplitude density $\sqrt{S_v} \sim 1/\sqrt{f}$) and the flat high-frequency THERMAL floor -- a stringent check since the flicker coefficients, thermal-noise model, and their bias dependence all have to line up. HICUM/L2 (SiGe HBT): no ngspice built-in, so validated against physics -- its default noise is WHITE (shot noise on the junction currents + thermal noise on the parasitic resistances, no flicker by default), a flat spectrum in contrast to the MOSFET's $1/f$ rise; with a small source resistance so the intrinsic device noise dominates, the output-noise floor tracks the collector SHOT-noise line $\sqrt{2qI_c R_C^2}$ across two decades of bias current (within 6% at high bias where shot dominates) and scales as $\sqrt{I_c}$. Verified under BOTH the Sparse and KLU solvers (.noise works under both since E-113 fixed the KLU adjoint solve), 5 checks: BSIM4 spectrum matches built-in $<4\%$, BSIM4 $1/f$ flicker slope ($S_v(1\text{Hz})/S_v(10\text{Hz}) \sim \sqrt{10}$), BSIM4 flat thermal floor, HICUM white noise, HICUM shot-noise tracking + $\sqrt{I_c}$ scaling. Completes the production-model validation loop (DC E-159, coverage E-160, small-signal E-161, large-signal E-164, noise E-165).

E-169 *frontend/syntaxhl.c (new), include/ngspice/syntaxhl.h (new), frontend/commands.c, frontend/Makefile.am, main.c, examples/syntaxhl_examples/**

Interactive command-line syntax highlighting -- the interactive prompt now colors the command line AS IT IS TYPED. The command word is shown GREEN the moment it is a recognized command, RED when it can no longer become one, and left the terminal default while it is still a valid PREFIX being typed (so `pl0` is neutral, `plot` green, `plt` red); numbers are yellow, "quoted strings" magenta, and -option flags cyan. The command word is classified exactly as the interpreter resolves it -- the same case-insensitive walk of the live command table `cp_coms` that `do_command()` uses, plus the control keywords (`if/while/foreach/begin/end/...`) -- so the color always agrees with whether the word will actually run. Implementation: a pure engine `cp_highlight_line()` (`frontend/syntaxhl.c`) tokenizes a line and wraps each token in ANSI SGR color escapes; live coloring overrides GNU readline's `rl_redisplay_function` (installed by `cp_synhl_init()` from `main.c` right after `rl_outstream` is set). The replacement `redisplay` redraws the colorized line only when the prompt plus line fit on one terminal row and otherwise defers to readline's own `rl_redisplay()`, so a wrapped line is never corrupted; the cursor is repositioned with a plain `CR + CUF` (cursor-forward) so mid-line editing is unaffected. Because that manual cursor positioning bypasses readline's internal position tracking, readline's own `accept-line` stops emitting the closing newline, so the Enter key is bound to a wrapper (`cp_synhl_accept`) that emits the newline itself when coloring is active before running the normal `accept-line` -- otherwise a command's output would begin on the input line (when coloring is off the wrapper is a no-op and readline handles the newline as usual). Gated three ways: it is on by default only at an interactive TTY (`isatty(rl_outstream)`), is disabled by `set no_syntax_highlight` and by the `NO_COLOR` environment convention, and never fires in batch (`ngspice -b`) or any piped/redirected session -- so simulation output is byte-for-byte unchanged (confirmed by the full example regression). Because as-you-type coloring needs the terminal in raw mode it requires a readline-enabled build, which the shipped binaries are (CI builds with `--with-readline=yes` and the committed binaries `link libreadline`). A new `synhl <command line>` command prints the colorized form of a line non-interactively -- useful for previewing and, since it needs no terminal, what makes the coloring engine regression-testable in batch. Verified (`examples/syntaxhl_examples/verify_syntaxhl.py`, 11 checks): the engine via `synhl` in batch (valid command green, unknown red, valid prefix neutral, numbers/strings/options colored, case-insensitive) and the live as-you-type behavior driven through a pseudo-terminal (green/red/neutral on the fly, `NO_COLOR` and a non-tty suppress all color); the live layer auto-skips on a build without readline. Solver-independent (front-end only).

E-170 *frontend/syntaxhl.c, include/ngspice/syntaxhl.h, examples/syntaxhl_examples/**

Semantic syntax highlighting -- extends E-169 from lexical to SEMANTIC coloring of the interactive line: it now checks whether the SIGNALS and EXPRESSIONS typed actually exist and parse. (1) A signal reference `v(node)/i(source)/@device[param]` is looked up read-only in the current plot (`vec_fromplot`, no evaluation, no vector creation); it keeps the default color if it exists and is drawn RED if it does not. (2) Inside an otherwise-valid expression only the invalid signal reddens -- `print v(a)+v(zzz)` colors just `v(zzz)`, including signals nested in functions (`sqrt(v(a))-v(bad)`). (3) A genuinely malformed expression argument of an expression command (`plot/print/gnuplot/asciiplot`) reddens as a whole, tested with the real parser `ft_getpnames_from_string(FALSE)` with its tree freed. (4) Error output is drawn RED at an interactive terminal: `cp_err` is wrapped in a colorizing stream (`funopen` on BSD/macOS, `fopencookie` on Linux) that prepends the red SGR to each write; `cp_curerr` is pointed at the wrapper too so `ngspice's` per-command stream reset `cp_ioreset()` keeps it in place instead of restoring and closing it. A half-typed expression is NOT treated as an error and stays neutral -- a signal whose parenthesis has not closed (`v(bP)`), an expression with unbalanced parens (`v(a)+v(b)` or one ending in an operator (`v(a)+`). KEY FIX: the expression parser's error handler `PError` writes straight to `stderr`, bypassing the `cp_err` muting, so an early version spewed syntax error in line segment ... onto the prompt while typing; `expr_pares()` now also mutes `fd 2 (dup2 /dev/null)` across the parse and restores it. Signal-validity is scoped to the unambiguous `v()/i()/@` forms (bare words could be plot keywords like `vs`, functions like `sqrt`, or options `-> false reds`). Same three-way gating as E-169 (interactive TTY, set `no_syntax_highlight`, `NO_COLOR`); the parser + vector lookups run per keystroke but are read-only and do NOT corrupt command execution (a subsequent `let z=v(a)+v(b)` still computes correctly). Verified (`verify_syntaxhl.py`, 19 checks total): the E-169 lexical checks plus a semantic layer driven through a pseudo-terminal after simulating a small circuit -- valid signal not red, invalid signal red, invalid-in-expression reddens only the signal, malformed expression red as a whole, half-typed stays neutral with no parser-error spam, error output red, `NO_COLOR` suppresses it. Solver-independent (front-end only).

E-171 *maths/KLU/klusmp.c, spicelib/analysis/pzan.c, examples/klupz_examples/**

KLU/Sparse solver audit -- KLU pole-zero fixed (complex determinant + pivot tolerance) -- a deep audit of the KLU<->SMP bridge found that pole-zero analysis under 'option klu' SILENTLY PRODUCED GARBAGE for any circuit with complex poles or zeros: a series RLC reported four bogus real poles (-100, -10, -0.89, 0) instead of its conjugate pair -5000 +- j999987.5. Real-axis roots came out right (the mixed formula is accidentally correct for real pivots), which is why the existing single-real-pole RC regression never caught it. DEFECT 1 (`spDeterminant_KLU`): the complex determinant built each pivot as the mixed quantity $1/(U_x R_s), U_z R_s$ and took its complex reciprocal -- KLU's `Udiag` stores the ACTUAL pivots (its solve divides by them) while Sparse's `Diag` stores RECIPROCAL pivots (its solve multiplies), and the adaptation kept Sparse's reciprocal dance for the real part only. The real branch was worse: its product loop NEVER RAN (the loop index was left at N by the preceding permutation scan, so `det` was always +/-1.0), it divided instead of multiplying, and it never wrote `*piDeterminant` (`SMPcDProd` consumed an uninitialized stack value). Both branches computed the permutation sign as `#non-fixed-points/2`, wrong for any cycle longer than 2 (a 3-cycle is EVEN but counted odd). Rewritten: $\det(A) = \text{sign}(P) * \text{sign}(Q) * \text{prod}(U_{\text{diag}}[k] R_s[k])$ with parity computed exactly by cycle decomposition (new `PermutationParity` helper); the KLU determinant now matches Sparse to ~14 digits at every Muller trial point. DEFECT 2 (`SMPcReorder/SMPpreorder`): pole-zero calls `SMPcReorder` with `PivRel = 0.0`; Sparse's `spOrderAndFactor` SANITIZES a non-positive threshold to its stored default, but the KLU branch assigned it straight to `Common->tol`, making KLU's partial pivoting accept an EXACTLY-ZERO diagonal as pivot (`|diag| >= 0max` always holds). At PZ's `s=0` trial an inductor branch has a 0.0 diagonal (`-sL`), so the factorization returned `KLU_SINGULAR` and PZ recorded a spurious root at the origin -- and the poisoned deflated Muller search never expanded past `|s| ~ 10`, yielding the garbage above. Both reorder entry points now mirror Sparse's sanitization (in-range `PivRel` still takes effect; `DC/AC/tran` pass the 1e-3 default and are byte-identical). With both root causes fixed, the E-113-era guard in `pzan.c` ('pole-zero finite-zero computation is not supported with option KLU' -- added

when the zero search was seen to go spuriously singular, i.e. defect 2) is REMOVED; the finite-zero search now works under KLU and a genuine E_SHORT reports the same 'input shorted' diagnostic as Sparse. The balanced/differential-output PZ guard remains (SMPcAddCol genuinely has no KLU branch). Verified (examples/klupz_examples/verify_klupz.py, 6 checks): full pole/zero root-set parity between the two solvers plus analytic anchors on a series RLC (conjugate pair), RC lowpass (old-good case, no regression), lead network (finite real zero), RC highpass (origin zero), RLC bandstop (PURELY IMAGINARY zeros $+j1e6$, the hardest case), and RLC bandpass (the finite-zero search formerly guarded off) -- all six identical across solvers. Full 134-example regression clean. Residual scope (solver-independent, ~1990 Muller-on-determinant fragility): a twin-T's deflated conjugate zero pair can still stall under KLU (4 of 6 roots found) while on the bandpass it is SPARSE that hits its iteration limit (KLU converges cleanly).

E-172 *maths/KLU/klusmp.c, spicelib/analysis/cktpzset.c, spicelib/analysis/pzan.c, examples/_setup.py, examples/klupz_examples/**

KLU balanced pole-zero + full partial pivoting -- closes E-171's two scope items; no analysis card remains Sparse-only under KLU. (1) BALANCED / DIFFERENTIAL-OUTPUT pole-zero ('pz in 0 a b vol', non-ground output reference) was guarded off as unsupported: the algorithm's change of variables needs the per-trial column fold $\text{col}(b) += \text{col}(a)$ (SMPcAddCol), which had no KLU branch -- and KLU's CSC sparsity pattern is FIXED at COO->CSC conversion time, so the fold's fill-in cannot be created on the fly the way Sparse's spcCreateElement does. Fixed in two halves: CKTpzSetup reserves the UNION pattern before the conversion (every row of the solution column is also registered in the balance column via SMPmakeElt; duplicate (row,col) COO entries fold into a single CSC slot by the conversion's group labeling, so the reservation is safe), and SMPcAddCol gains a KLU branch that merge-walks the two columns' sorted row lists adding the complex values in place (with a defensive error if an addend row is missing, impossible after the reservation). The two pzan.c balanced-output guards are removed. (2) FULL PARTIAL PIVOTING for PZ factorizations: validating a fully-differential configuration exposed spurious far-field roots at $|s| \sim 1e19$. $.1e21$ (including a positive-real pole). An exact-rational (fraction-arithmetic) determinant of the very matrix KLU factored proved the factorization itself was numerically rotten at extreme $|s|$: at $s=4.52e19$ the exact determinant was $+4.36e11$ while the pivot product gave $-2.05e12$ -- wrong sign AND magnitude. Root cause is architectural: KLU's ordering is fixed at `klu_analyze` time (pattern-only AMD/BTF) while Sparse re-runs value-aware Markowitz ordering every PZ trial; PZ sweeps $|s|$ across ~20 decades, and with the relaxed default `tol=0.001` the fixed ordering accepted catastrophically-cancelling pivots. KLU's only value-adaptive lever is the within-block partial pivoting, so the E-171 sanitization fallback is upgraded from 'keep current tolerance' to FULL partial pivoting (`tol=1.0`) whenever the caller passes an out-of-range `PivRel` -- only pole-zero does (0.0); explicit in-range tolerances (DC/AC/tran $1e-3$) are honored unchanged. This single change eliminated the far-field junk AND cured the twin-T conjugate-zero-pair stall recorded in E-171's scope (what looked like vintage-Muller noise-floor fragility under KLU was largely this factorization inaccuracy; KLU now finds all 6 twin-T roots with notch-zero real parts $\sim 1e-12$, an order CLOSER to the exact 0 than Sparse's $\sim 1e-11$). Verified (verify_klupz.py, 6 -> 9 checks): the E-171 six plus the differential RC bridge (balanced output, formerly 'not supported'), a balanced output with a complex pole pair, and the twin-T notch -- all root sets identical across solvers. The dual-solver harness's balanced-pz KLU-skip detection is removed from examples/_setup.py (nothing remains Sparse-only). Full example regression 134/134.

E-173 *spicelib/analysis/cktpzeig.c (new), maths/dense/eig.c (new), maths/KLU/klusmp.c, spicelib/analysis/pzan.c, cktsopt.c, cktnntask.c, cktdojob.c, include/ngspice/{smpdefs,optdefs,tskdefs,cktddefs}.h, two Makefile.am, examples/pzeig_examples/**

Eigenvalue-based pole-zero ('.options pzeig') -- a modern alternative root finder for .pz. The vintage driver hunts roots of $\det(G+sC)$ with a Muller iteration on determinant values and is famously fragile (iteration limits -- 'giving up after 231 trials' on a plain RLC bandpass -- noise-floor stalls, chaotic search paths; E-171/172 fixed the KLU-side defects but the ~1990 algorithm itself stays

delicate under BOTH solvers). The new opt-in method computes every pole and zero at once as a dense eigenvalue problem, reusing the existing CKTpzSetup/CKTpzLoad machinery so it applies unchanged to the poles and zeros configurations, to balanced/differential outputs, and to both linear solvers. Steps: (1) the PZ-configured MNA matrix is AFFINE in s , $A(s) = G + sC$, so two loads ($s=0$, $s=1$) extract the dense pencil via a new solver-agnostic `SMPdenseExtractReal` (KLU: walk $A_p/A_i/A_x$ Complex; Sparse: walk `FirstInCol` through the `Int->Ext` translation maps); a third load at an arbitrary point VERIFIES affinity so a non-polynomial device falls back with a clear message instead of wrong roots. (2) C is structurally singular, so the pencil is solved by SHIFT-INVERT linearization: factor $(G + \sigma C)$ ONCE at a non-root shift (tried from a small ladder of shifts, with the circuit's own sparse solver), form the dense $M = (G + \sigma C)^{-1}C$ by n sparse solves; from $(G + sC)v=0 \Leftrightarrow Mv = \mu v$ with $\mu = 1/(\sigma - s)$, every finite root is $s = \sigma - 1/\mu$ and the pencil's infinite eigenvalues land harmlessly at $\mu = 0$ (dropped by a noise-floor threshold $64n\epsilon s^*||M||$). (3) eigenvalues of the real dense M from a new self-contained eigensolver `maths/dense/eig.c` -- the classical balance / Hessenberg / Francis double-shift QR chain (EISPACK balance/elmhes/hqr lineage, written fresh, no LAPACK dependency), unit-tested standalone against companion matrices, complex conjugate pairs, a 50x50 tridiagonal with known spectrum, and a badly-scaled balance case. Roots are emitted as the same PZtrial list the Muller driver produces, so PZpost and all downstream output are unchanged. Option plumbing follows the standard task-option path (`OPT_PZEIG -> TSKpzEig -> CKTpzEig` flag -> dispatch in `pzan.c`). RESULTS: RLC bandpass -- same roots as Muller with NO iteration-limit warning; 10-section RC ladder -- all ten poles matching the analytic tridiagonal-eigenvalue formula $s_k = -(2-2\cos((2k-1)\pi/21))/RC$ and Muller root-for-root; twin-T -- all 6 roots under both solvers; RLC bandstop -- purely imaginary zeros $+j1e6$ EXACT; balanced/differential output works; purely resistive circuits yield no roots and no crash. DEFAULT REMAINS MULLER (`pzeig` is opt-in). Dense $O(n^2)$ memory / $O(n^3)$ QR, capped at 2000 unknowns with a clear error above. Accuracy is dense-QR class (absolute error $\sim \epsilon s^* \text{ dominant root magnitude}$): an exact origin-zero can print as $\sim 1e-7$ when poles sit at $1e6$ rad/s; Muller refines locally and can resolve wider per-root dynamic range, while `pzeig` never misses or invents a root -- the two methods agree to ≥ 6 digits on every battery circuit. Verified (`verify_pzeig.py`, 13 checks: both solvers x both methods + analytic anchors + default-method check). Full example regression 136/136.

E-174 `src/frontend/commands.c, examples/helpcmd_examples/*`

help all crash fix (help string used as printf format) -- running `help all` (or `help montecarlo`) in the interactive shell crashed the shipped ngspice with 'Error: tvprintf failed / ERROR: fatal error in ngspice, exit(-1)'. Root cause: ngspice's help printer uses each command's help TEXT as a printf FORMAT string -- `out_printf(ccc[i]->co_help, cp_program)` in `com_help.c` (and the parallel `com_ahelp.c`) -- passing only one argument (`cp_program`, the program name, intended for a single `%s` as in the legitimate 'Report a %s bug.'). Any other '%' in a help string is an invalid/incomplete conversion specifier, so `tvprintf` fails and ngspice calls a fatal `exit(-1)`. The `montecarlo` command's help (added in E-151) read '... reports the yield with a Wilson 95% CI ...'; the '%' (percent-space) is the invalid specifier. `help all` walks the command table alphabetically and died at `montecarlo`; `help montecarlo` crashed directly too. Present in BOTH command tables (`spcp_coms` for ngspice, `nutcp_coms` for nutmeg). Not X11/terminal/solver specific -- plain format-string handling, so it reproduced on every build; it was invisible in headless CI only because `help all` is an interactive command not exercised by the batch example decks. Fix: escape the literal percent as `%%` (printf convention) in both tables, so the line renders 'Wilson 95% CI' and `help all` completes. New regression guard `examples/helpcmd_examples/verify_helpcmd.py` with two layers: RUNTIME -- drives an interactive ngspice on a pseudo-terminal and asserts `help`, `help all`, and `help montecarlo` each run to completion (no crash, no 'tvprintf failed') and that the `montecarlo` line renders a literal '95% CI'; and a STATIC class-guard that scans `commands.c` and fails if any command help string carries a format hazard (a '%' that is not '%s' or '%%', or more than one '%s' since only one argument is passed), catching a future unescaped '%' even when that command is never exercised at runtime. Full example regression 137/137.

E-175 [spicelib/analysis/dcpss.c, examples/rfconv_examples/*](#)

RF-suite audit -- the dropped parametric term in every periodic small-signal analysis -- a full audit of the RF suite (.sp/Touchstone, PSS/PAC/pnoise/pxf, PSP, HB, QPSS + the two-tone family, hbosc/phasenoise) driven by analytic anchors and cross-analysis invariants. CLEAN: .sp + Touchstone (17 checks: series-R/shunt-C/3-port star S-parameters exact to 10 digits, the asymmetric-2-port S21/S12 column-order trap, reciprocity, RI/MA/DB + Y/Z (v1 normalization) + GHz round-trips through wrsnp/rdsnp, N-port 4-pairs-per-line layout); HB (analytic cubic: fundamental $a1A+(3/4)a3A^3$ and $H3=(1/4)a3A^3$ exact to 7 digits, even harmonics at machine zero, Sparse/KLU identical); QPSS-HB/QPSS-tran (two-tone fundamentals, IM3 $(3/4)a3A^3$ exact, commensurate aliasing guard fires correctly). THE BUG: every LPTV small-signal conversion matrix scaled its reactive coupling by the COLUMN (input-sideband) frequency, $H_{nm} = G_{\{n-m\}} + j\omega_m C_{\{n-m\}}$. For a small signal riding a FROZEN periodic capacitance $C(t)=dQ/dv|_{\text{orbit}}$ the exact linearization is $di = d/dt[C(t)dv] = C\dot{v} + C\dot{v}_{\text{dot}}$, and the product rule collapses to the ROW frequency: $H_{nm} = G_{\{n-m\}} + j\omega_m nC_{\{n-m\}}$. Using ω_m silently DROPS THE PARAMETRIC-PUMPING TERM $C\dot{v}$ -- the very term that makes a pumped varactor convert. Every conversion sideband through a time-varying capacitance came out scaled by exactly ω_{in}/ω_{out} : on the audit circuit (1k -> OSDI varactor $C(v)=c0(1+\alpha v)$ pumped 1V@1MHz, 1mV probe at 250kHz) the $+1/+2$ sidebands were 3x/5x/7x/9x too small, with measured pre-fix/truth ratios matching the ω ratios to FIVE DIGITS (0.3333/0.2000/0.14286/0.11111) -- diagnosis confirmed beyond doubt. AFFECTED: pac, psp, pnoise, pxf, qpac, qpnoise, qpxf, and the phasenoise adjoint, for any circuit with pumped/nonlinear capacitance (varactors, junction and MOS capacitances -- every real mixer and oscillator). UNAFFECTED: LTI circuits (no off-diagonal C harmonics; $\omega_n == \omega_m$ on the diagonal) -- exactly what every prior regression check used; the E-171 accidental-correctness pattern again. THE SUBTLETY: the harmonic-balance residual uses the SAME builders and must KEEP the column frequency -- its reactive current is the exact chain rule $d/dt Q(v(t)) = C(v(t))\dot{v}(t)$ whose n -th harmonic is $\sum_m C_{\{n-m\}}(j\omega_m V_m)$. That is why HB and QPSS-HB always matched transient ground truth while the small-signal analyses did not, and why a blanket fix would have broken HB. FIX: a smallsig mode flag on pac_build_matrix and qp_build_matrix plus the two inline copies (the PAC adjoint and the phasenoise adjoint): smallsig=1 -> row frequency (restores the parametric term); smallsig=0 (HB/HBOSC/QPSS-HB residual + Jacobian) -> chain-rule column frequency, byte-for-byte unchanged. Regression safety proven directly: the linear-circuit PAC control produces IDENTICAL output pre/post fix, and HB/QPSS-HB outputs are byte-identical. Ground truth = plain transient + one-beat Fourier projection (independent of all conversion-matrix machinery); the fixed qpac matches all five sidebands within 2% (integrator error class). Also noted: cktspnoise.c is a dead fossil (#ifdef RFSPICE_ never compiled). Verified (examples/rfconv_examples/verify_rfconv.py, 5 checks x both solvers: truth vs QPSS-HB unchanged, truth vs fixed qpac, pre-fix signature absent, HB + QPSS-HB analytic anchors). Full example regression 138/138.

E-176 [spicelib/analysis/dcpss.c, examples/_setup.py, examples/pssdriven_examples/*](#)

Driven-mode PSS shooting -- no frequency hunt, no breakpoint flood -- diagnosis of the E-175 audit's flagged residual ('PSS shooting robustness: the varcap PSS ran 17+ minutes without converging'). Instrumentation showed 9,627,161 accepted timepoints by $t=2\mu s$ (~0.2 ps per step, zero LTE rejections): ngspice's PSS (the Lannutti shooting) is OSCILLATOR-oriented -- it hunts the fundamental frequency through mid-period forced breakpoints whose spacing is proportional to steady_coeff. On driven circuits (everything the E-117..126 periodic small-signal stack runs on) this is doubly pathological: (1) at the standard decks' steady_coeff=5e-6 the grid spacing is ~0.5 ps -- millions of steps per shooting cycle, and the convergence-tolerance knob doubles as grid density so tighter tolerance quadratically explodes runtime (even the default 1e-3 forces thousands of points/period -- why a linear RC .pss took ~4 minutes); (2) the frequency estimate can never settle EXACTLY on the source frequency, the cycle endpoint phase-drifts against the source, and

the period residual FLOORS ($\sim 1e-4$) -- shooting never converges and burns the full `sc_iter` budget at flood speed. FIX: a DRIVEN mode, auto-detected when the circuit contains a time-varying independent source (SIN/PULSE/... V or I source, `funcTGiven`): the period is pinned to the exact source period (no estimation, no estimator grid; each cycle is one plain-transient period; varactor residual $2e-1 \rightarrow 8e-9$ in 17 cycles / 0.3 s); the shooting-phase max step is clamped to `T/pspoints` so the orbit is integrated on the SAME discretization as the retained samples (validation caught this: without the clamp LTE grows steps toward `T/2` and shooting converges -- to $7e-9$! -- onto the fixed point of that coarse discretization, several percent off the true orbit: retained swing 0.1651 vs analytic 0.15718; the old flood had been masking the effect by brute force); the post-FFT 'relaunch at the strongest spectral line' is disabled when driven (a rectifier-like circuit with a dominant 2nd harmonic must not retain the wrong period); the AUTONOMOUS oscillator path is untouched (byte-identical behavior verified on oscillator decks; DC-only-supply circuits never trigger detection). IMPACT: `rc_pss` ~ 4 min $\rightarrow 0.05$ s with the retained fundamental EXACTLY 1000000 Hz (the old hunt returned 999999.8976); psp example 406 s $\rightarrow 0.9$ s; the rfpss battery (5 verifies, 61 checks) minutes-each $\rightarrow <0.5$ s each; rfanalyses 'several minutes even Sparse-only' $\rightarrow 0.56$ s; the KLU PSS pass 'prohibitively slow' $\rightarrow 0.14$ s. HARNESS: rfpss and rfanalyses removed from both REGRESSION_EXCLUDE and SPARSE_ONLY -- the whole RF periodic small-signal suite now runs in EVERY regression sweep under BOTH solvers; the direct .pac-on-pumped-varactor check became affordable, closing the E-175 loop on the direct PAC path (sidebands match transient ground truth within 1%). Verified (examples/psdriven_examples/verify_pssdriven.py, 6 checks x both solvers: driven detection + exact retained frequency, retained swing == $|H|$ within 0.1%, retained-period self-consistency, direct .pac vs transient truth, autonomous non-trigger). Full example regression 145/145.

E-177 `spicelib/analysis/dcpss.c, examples/pnoise/pnoise_examples/*`

pnoise noise-folding referee + the folded-flicker frequency fix -- closes the E-175 audit's last honesty-ledger gap: periodic noise analysis (pnoise/qnoise/phasenoise) was 'correct by construction, not correct by measurement' -- the folding of device noise through LPTV conversion had never been checked against anything independent. THE REFEREE (three independent layers, on a 1-node LPTV-conductance circuit $G(t) = g_0 + g_1 \sin(\omega_0 t)$ chosen so conversion transfers are $O(1)$): (1) a from-scratch Python conversion matrix, $\text{onoise}(f) = \sum_k |Z_k(f)|^2 * S(|f + kf_0|)$ -- same physics, zero shared code; (2) a TRNOISE transient Monte-Carlo (no shared code OR theory; 198-segment Welch PSD, with the pumped/unpumped RATIO cancelling the TRNOISE amplitude convention); (3) the LTI (pnoise == .noise) and white limits. VERDICT ON THE WHITE PATH: measured-correct -- pnoise matched the referee to 6 digits at every frequency and the MC arbiter confirmed the $\sim 1.86x$ folding ratio within statistical error; the adjoint, sideband sum and thermal densities are right. THE BUG THE REFEREE CAUGHT: the stationary sideband loops set the noise-evaluation frequency ONCE, to the output frequency (`data.freq = freq` outside the `k` loop). But the sideband-`k` adjoint carries noise ORIGINATING at $|f + kf_0|$ -- a frequency-dependent source PSD (flicker $1/f$, noise_table E-109) must be evaluated at the source-side frequency. The proof was digit-exact in both directions: pre-fix pnoise reproduced the referee's deliberately-wrong evaluate-at-f model DIGIT-FOR-DIGIT ($1.114323e-14$ vs $1.114323e-14$ at 1 kHz -- a 21% overestimate on this circuit, unbounded as $f \ll f_0$); post-fix it reproduces the correct model digit-for-digit (9.175472e-15). Why three generations of checks missed it: white noise is frequency-flat and LTI circuits have no $k \neq 0$ transfer -- the THIRD occurrence of the accidental-correctness pattern (E-171 determinant, E-175 parametric term). FIXED in all three stationary loops: pnoise (`data.freq = |freq + kf_0|` per sideband), qnoise ($|f_{in} + k_1 f_1 + k_2 f_2|$ per harmonic), and phasenoise -- where it matters most: folded far-sideband blocks near a carrier were evaluated at the tiny OFFSET frequency, inflating the flicker ($1/f^3$) region. CYCLO SCOPE: the cyclo mode's time-domain identity $\text{onoise} = (1/P) \sum_s S(t_s) |A_s|^2$ treats the source PSD as frequency-flat across the folding span -- exact for modulated-white noise and for flicker on conversion-free circuits (the shipped E-125/126 use cases), not for flicker folded through $k \neq 0$ conversion (cannot be collapsed into that product with the aggregate device-noise API); documented at both cyclo sites, with the (now exact) stationary

mode as the right tool when folded flicker matters. Verified (examples/pnoise_examples/verify_pnoise.py, 5 checks x both solvers, ~2 s thanks to E-176): white folding vs referee ($\leq 0.1\%$), flicker folding vs referee ($\leq 0.5\%$, sample-0 bias read from the retained orbit), pre-fix signature absent, pump/no-pump ratio == referee ratio, LTI limit == .noise. The rfpss/qpss/rcyclo suites pass unchanged (conversion-free or white -- provably unaffected). Full example regression 146/146.

E-178 *spicelib/analysis/dcpss.c, spicelib/devices/dio/dionoise.c, examples/cyclofold_examples/, examples/rfpss_examples/, examples/qpss_examples/**

exact separable cyclostationary folding + the HB DC-source fix -- removes the cyclo mode's frequency-flat limitation E-177 documented. For the separable model $S(t,f) = m(t)^2 g(f)$, $onoise(f) = \sum_g \sum_q |B_q(g)|^2 * g_g(|f + qf_0|)$ with $B_q(g) = (1/P) \sum_s m_g(t_s) dA_s(g) e^{+jq \theta_{t_s}}$ and $A_s = \sum_k \Psi_k e^{-jk \theta_{t_s}}$. *Per-generator amplitudes recovered with NO device-API change: the noise-summary machinery (prtSummary/outpVector, one density per generator; family totals excluded by the name-prefix rule) plus load POLARIZATION against a fixed reference load $R = \Psi_0$ -- five DEVnoise sweeps per orbit sample ($A+-R, A+-jR, R$) give $c_{g,s} = \sqrt{S_g(t_s)} dA_s(g)$ up to a constant phase that cancels in $|B_q|^2$. Spectral shapes are MEASURED pointwise at the folded frequencies (no $1/f^{EF}$ assumption -- noise_table works) at several probe biases; flat generators keep the original time-domain identity (exact for ANY quadratic form incl. NevalSrc2-correlated pairs); stationary limit == the E-177 sum, flat limit == the old identity (Parseval). Ported 2-D to qpnoise ... cyclo ($B_{\{q1,q2\}}$ at $|f + q1 f_1 + q2 f_2|$). THE PHYSICS: 'flicker sees $\wedge^2$, white sees $\langle m^2 \rangle$ ' -- only the envelope's DC feeds the $1/f$ band; the old identity was 23% high ($8/\pi^2$) on $|\sin|$ -modulated flicker and 34% high on the conversion circuit; the new sweep lands on a from-scratch referee to $\leq 3e-4$ (the old binary reproduced the deliberately-flat model to $6e-5$ -- the error was the model, not the code). BONUS BUG (4th accidental-correctness occurrence): the cross-machinery check (same circuit through the QPSS-HB orbit) came out exactly 4x -- BOTH HB Newtons (hb E-134, qpss ... hb E-136) double-subtracted the DC sources (the settle-mode rhs folded into I_R already carries them; -lambdaIs subtracts them again): every DC bias voltage converged to exactly 2x, silently corrupting all bias-dependent noise and conversion (flicker $\sim I^{AF}$ off by 2^{AF} , shot 2x) while AC content and every AC-driven validation stayed exact. Fixed by adding the unscaled DC source block back (net DC drive exactly -lambdaIs_DC); qpnoise cyclo now agrees with pnoise cyclo DIGIT-FOR-DIGIT across the two orbit machineries. rcyclo's flicker law updated to the exact $\langle |I| \rangle^2 = (8/\pi^2) \langle I^2 \rangle$ (pinned to $1e-4$). 6 checks x both solvers; regression 147/147. FOLLOW-UP (same enhancement): the E-139 hard-pumped-diode deck exploded the cyclo path ($\sim 6.5e1 \text{ V}^2/\text{Hz}$) -- dionoise.c computes its three sidewall generators only when DIOresistSWGIVEN but the prtSummary block copies ALL DIONSRCS slots, so without a sidewall model the summary vector carried uninitialized stack garbage (latent in stock ngspice for decades; E-178 is the first consumer of per-generator densities). Fixed at the device (unwritten slots zeroed, dio/dionoise.c) and hardened in both cyclo consumers (finite non-negative PSD guards; garbage probe ratios fall back to the flat path). verify_qpnoise's 'cyclo > 2x stationary' diode expectation was an artifact of the garbage slots + doubled HB DC bias; replaced with a closed-form torus-average referee (resistive network => memoryless LPTV, instantaneous adjoints) matched to ~1%.*

E-179 *spicelib/analysis/tfanal.c, spicelib/analysis/cktsens.c, frontend/com_measure2.c, examples/stdaudit_examples/**

standard-analyses audit: referee battery + three fixes. The 'Standard analyses (analog)' gap-table rows are 1990s SPICE3 code whose VALUES had never been physics-checked. FIX 1 (tfanal.c, inherited from Berkeley SPICE3 1988): .tf i(vm) vin output impedance was ALWAYS $1e20$ -- the output-impedance solve forces 1V on the sense branch and inverts the branch current as $1/\text{MAX}(1e-20, \text{rhs})$, but that current is NEGATIVE for every passive network (the input-impedance path right above divides by -rhs); fixed with the same sign + fabs guard; feedback-amp referee exact ($1000.990 == \text{RL} + \text{node Thevenin}$). FIX 2 (cktsens.c): KLU AC sensitivity

silently truncated to ONE frequency point -- the E-114 KLU complex-conversion block inside the frequency loop reused *i* (the OUTER loop variable) for its DEVmaxnum walk, so after the first point *i*=DEVmaxnum ended the sweep; the surviving point was correct, which is why E-62's single-point check passed. Fixed with a local index (dead second block hardened too); full sweep == analytic $dV/dC = -j\omega R/(1+j\omega RC)^2$ to 6 digits, both solvers. FIX 3 (com_measure2.c): . meas DERIV{ATIVE} was parsed but evaluation fell into an explicit 'currently not supported' stub (empty #if 0 measure_deriv() placeholder) -- implemented: 3-point Lagrange-quadratic derivative on the nonuniform time grid, AT= and WHEN forms sharing FIND's parse/locate flow, verified vs the analytic sine slope to 5 digits; DERIVATIVE/INTEGRAL long spellings accepted (only DERIV/INTEG matched); pre-existing -Wmissing-prototypes warning silenced (str_has_arith_char2 static). MEASURED-CORRECT: .disto HD2/HD3 == analytic diode Volterra kernels $\leq 0.06\%$ incl. a frequency-dependent load (harmonics at $Z(2\omega)/Z(3\omega)$ -- no E-177-style frequency bug), SIM2 intermods (f_1+f_2 , f_1-f_2 , $2f_1-f_2$) incl. cascades, exact DISTOF1/DISTOF2 scaling, junction-cap harmonics == Harmonic Balance to ~6 digits (two independent engines); .noise onoise_total^2 == band analytic (kT/C) 0.06% + flicker log-integral 6 digits; .sens DC == central finite differences at a nonlinear OP (5-6 digits, model params incl.); nonlinear-OP .pz works via the driving-point form (the E-62 'quirk' was a bias source shorting the injection node -- ngspice's refusal is correct); tran/op/meas battery exact. 8 checks x both solvers; regression 148/148.

E-180 *frontend/com_checkpoint.c, examples/checkpoint_examples/**

checkpoint/restart under KLU (the last Sparse-only feature falls). E-131 guarded savestate/loadstate off under KLU, recording the reason as 'KLU factorization objects absent on the restore path' -- a mis-diagnosis. REAL cause: .option klu lives in the task (TSKkluMODE) and reaches the circuit only in CKTdoJob at analysis dispatch, but loadstate calls CKTsetup DIRECTLY first, so the matrix was built Sparse (SMPkluMatrix NULL); the resume dispatch then set *ckt->CKTkluMODE*=1 without re-setup (by design -- it must preserve the restored state), and the first NIiter dereferenced the NULL SMPkluMatrix (KLUloadDiagGmin store; reproduced EXC_BAD_ACCESS +0x90). savestate was never broken (the checkpoint file holds only solver-agnostic vectors). FIX: com_loadstate copies the task's solver selection and the E-152 KLU knobs (ordering/scale/BTF/memgrow) into the circuit BEFORE CKTsetup, exactly CKTdoJob's block; both guards removed. BONUS: cross-solver restore -- all four save x load combinations (fresh processes) continue exactly on the uninterrupted trajectory ($\leq 1e-9$ at $t=3.5\text{ms}$ of a 4ms diode+RC run). verify_checkpoint's 'KLU rejected' check replaced by the 4-combo matrix (22/22); solver-notes and gap-analysis tables updated. Nothing in the simulator is Sparse-only any more (E-172 closed the last analysis, E-180 the last feature). Regression 148/148.

E-181 *spicelib/analysis/dctran.c, spicelib/analysis/cktsopt.c, spicelib/analysis/ckntask.c, spicelib/analysis/cktdojob.c, include/ngspice/cktdefs.h, include/ngspice/tskdefs.h, include/ngspice/optdefs.h, examples/corenum_examples/**

core-numerics audit: the integrator certified exactly + **.options ordfix**. Nobody had ever verified that the Gear (BDF) corrector coefficients deliver their nominal order (orders 3-6 were dead code for ~30 years until E-128 selected them; E-128 verified order SELECTION, not coefficient CORRECTNESS). NEW OPTION **ordfix=K** (off by default): a fixed-order verification mode -- every converged step is accepted (the step is ruled by *tmax*, not the LTE), the first step is shrunk 1000x (an order-1 starter proportional to the pinned step imposes an $O(h^2)$ floor), the order ramps to *K* while the step is tiny, then the step grows gently ($\times 1.15$; large ratios destabilize high-order variable-step BDF). Each guard defeats a measurement trap that initially presented as a bug but is real numerics (incl. the LTE step preference $h_k=(c_k \cdot \text{tol})^{1/(k+1)}$ INVERTING at loose tolerance -- the E-128 controller's refusal to raise there is correct). CERTIFIED: accepted trajectories satisfy the exact variable-step BDF-k formula (Lagrange differentiation on the actual nodes) to $\leq 1.3e-13$ at every order 1-6 (~90 uniform stencils each); measured global slopes confirm orders 1-3; trap-vs-Gear dissipation on a lossless LC matches theory ($|R(i\omega)|=1$ vs order-dependent damping).

Also referee'd with NO defects: solver precision == conditioning theory under both solvers (1e18-conductance-spread asymmetric mesh == exact-rational MNA to 1e-15); all four DC convergence-aid paths (default/noopiter/gminsteps=0/srcsteps=0) agree on a 40-diode hard-DC chain (spread 2.2e-6, KLU==Sparse 7e-13); xmu=0.5 == trap bit-identical, xmu=0.45 damps; lvltim=1 works. Drive-by: DCtran_step_quit prototype added to cktdefs.h (pre-existing -Wmissing-prototypes). 8 checks x both solvers; regression 149/149.

E-182 *frontend/plotting/plotit.c, frontend/plotting/pyplot.c, examples/pyplot_examples/**

pyplot autoscale by default -- the matplotlib bridge (E-94/95/98/99) pinned every generated axis with set_xlim/set_ylim taken from ngspice's grid-rounded internal plot ranges. User request: rely on matplotlib's autoscaling + fig.tight_layout() instead. plotit() always fills xlims[]/ylim[] (user args OR data-derived), and the statics that distinguish the user case (xlim/ylim from getlims) are freed BEFORE the device dispatch -- plotit now captures user_xlim/user_ylim flags at the free point and passes NULL limits to ft_pyplot unless the user gave xlimit/ylim explicitly (ft_pyplot already handled NULL; gnuplot/asciiplot back-ends untouched). verify_pyplot +2 checks (no pins + tight_layout on auto plots; exact pins for explicit limits). Regression 149/149.

E-183 *frontend/com_pyplot.c, frontend/plotting/pyplot.c, examples/pyplot_examples/, features/*

pyplot usability: distinct default names, deck-folder output, linewidth & backend. (1) REAL BUG: in window (interactive) mode pyplot launches the Python viewer in the BACKGROUND (python3 pyplot.py &); two plots that both omit the file name shared pyplot.py/pyplot.data, and the 2nd call overwrote them before the 1st viewer read them -- both windows showed the 2nd plot (its title and data). Invisible in PNG mode (synchronous). Fix (com_pyplot.c): a per-session counter names successive no-name plots pyplot, pyplot-2, ... (first unchanged). (2) The .py/.data/.png are written next to the CIRCUIT FILE via ngdirname(ci_filename) rather than ngspice's cwd; the script's loadtxt/savefig carry the deck-dir path so Python finds the data from any cwd; a bare in-folder deck name still lands in the cwd (unchanged). (3) set pyplot_linewidth=<w> -> linewidth in the plot()/step() calls. (4) set pyplot_backend=<name> -> matplotlib.use() ahead of import pyplot, overriding the auto backend (incl. the Agg forced by the png/svg/pdf terminals). All four opt-in; default output byte-unchanged. filename buffers in pyplot.c grown 128->1024 for full paths. verify_pyplot grows to 13 checks x both solvers; packaged as a standalone feature bundle under features/. Regression 149/149.

E-184 *frontend/outitf.c, examples/progressbar_examples/**

sweep progress bar reaches 100%. Fix to the E-129 progress bar: it sometimes froze below 100% at end of run (e.g. '[====] 82%' then 'No. of Data Rows'). CAUSE: the 'Reference value' status line is throttled -- redrawn only when >0.25 s has elapsed since the last update -- so the FINAL sweep point's print is almost always skipped (it lands within 0.25 s of the previous tick), freezing the bar at the last throttled value; reaching 100% was never reliable. FIX: outp_print_reference records the latest refval + a bar-shown flag every call; a new outp_finish_reference, called from fileEnd/plotEnd just before the 'No. of Data Rows' line, reprints the bar full at 100% in place using the sweep's TRUE endpoint (TRAN CKTfinalTime, AC ACstopFreq, NOISE NstopFreq, DC last source value). One-shot (reset at analysis start and after firing), silent for non-swept analyses (op/tf), same ft_optimizing/ft_noreprint/cp_background guards -- quiet/background/optimizer runs and the rawfile unaffected. verify_progressbar's end-of-run check tightened from >=90% to ==100% (all four swept analyses); 22/22. Regression 149/149.

E-188 *spicelib/analysis/dcop.c, include/ngspice/cktdefs.h, frontend/com_sweep.c, examples/warmstart_examples/**

Monte Carlo warm-start (**montecarlo ... -warm**). Each MC sample (E-151) re-sources the deck and COLD-solves its DC bias point, running the full gmin/source-stepping homotopy from scratch even though consecutive samples move the operating point only slightly. Profiling: the solve, not

the deck re-source, dominates for non-trivial circuits, and every op cold-starts -- two IDENTICAL ops back-to-back each burn the full homotopy (~52 Newton iters on a diode ladder; a warm .nodeset at the solution cuts it to ~5). CKTop tries a direct Newton from the caller's firstmode before any homotopy, but DCop always passes MODEINITJCT (cold junction voltages), discarding the previous solution. FIX: montecarlo -warm wraps the loop in CKTsetWarmStart(1/0); DCop keeps a buffer (outside the CKTcircuit that reset recreates; indexed by equation number, stable across resets of an identical-topology deck) of the last converged CKTrhsOld, preloads it and calls CKTop with MODEINITFLOAT (use existing node voltages) when a valid same-size guess exists, and snapshots each converged solution. A poor guess just fails the first NIiter and falls through to the normal cold homotopy, so the converged point is identical -- only a cheap failed attempt is wasted. First sample cold, rest warm; when -warm is absent dco_p_warm_enable=0 and the path is byte-identical (one extra SMPmatSize read). CORRECTNESS: warm changes only Newton's starting point, so it converges to the same OP to within convergence tolerance -- warm==cold yield EXACTLY at reltol=1e-6, agreeing to within a couple boundary samples at default reltol ($v(3) \sim 3.8V$, where $reltol * |v| \sim$ the narrow spec band; a tolerance effect, removed by tightening reltol). Iteration count ~52->~4 (~13x); wall-clock benefit scales with per-iteration cost (large/hard-converging designs win most; small circuits are dominated by deck re-source + command dispatch). Shared DC-op code, so Sparse+KLU; composes with -lhs. 5 checks (exact at tight reltol, close at default, iteration cut, +lhs, +klu). Regression 152/152.

E-189 *frontend/com_sweep.c, examples/sweepwave_examples/**

Sweep waveform overlay (**sweep ... -overlay**). A usability follow-up to the E-146 sweep command: sweep steps any knob, runs an inner analysis per point, and records each output's LAST value into a summary sweep transfer curve -- but overlaying the full per-point WAVEFORMS to compare shapes meant setplot-ing each retained run by hand. -overlay (alias -ov) collects every point's whole output waveform, resamples them onto a common independent-variable grid, and builds one sweepwave plot with a <output>_<value> vector per (output, knob value), left current so plot <out>_<val> ... works at once -- the HSPICE .step overlay in one step. Three static helpers in com_sweep.c: sw_eval_vec (evaluate the output expr and copy its FULL waveform + scale into caller buffers; complex/AC stored as magnitude), sw_interp (binary-search piecewise-linear resample, flat outside range -- each run lands on its own adaptive time/freq grid, so a shared grid is required before a common scale vector), sw_wavename (clean nutmeg name, non-alnum -> _). The sweep loop captures each waveform alongside the unchanged last-value recording; after the loop a common grid over [min,max] of all runs at the finest resolution is built, every waveform resampled onto it. GRACEFUL: a scalar inner analysis (op, no waveform) prints "-overlay ignored (analysis '...' has no waveform to overlay)" and yields the ordinary summary -- never an error. Without the flag the path is byte-identical to E-146. Front-end only, solver-independent. Verify: overlaid RC step responses match $1 - \exp(-t/RC)$ ($RC = 1/2/4\mu s$) to worst $< 1e-4$, distinct per-value vectors, graceful op, summary curve intact. 5 checks. Regression 153/153.

E-190 *frontend/com_sweep.c, examples/nestedsweep_examples/**

Nested multi-knob sweep (**sweep ... -vs ...**). Second usability follow-up to the E-146 sweep command (after E-189 -overlay). sweep stepped ONE knob; -vs <knob> <spec> (alias -family) adds OUTER knobs whose CARTESIAN PRODUCT forms a curve family: the inner (positional) knob stays the x-axis of the summary sweep plot, and each combination of the outer knobs yields one curve per output, named <output>_<outerknob>_<value>... -- the parametric family of .dc ... SWEEP ... / HSPICE nested .dc. Up to SW_MAXKNOB=4 knobs, cartesian points capped at SW_MAXPTS. The three-form spec parser (lin/dec/oct, list, start/stop/step) was factored into sw_parse_spec, now shared by the inner knob and each -vs knob so all knobs accept the same forms. SET ORDERING (the subtle part): each knob is applied with its E-146 mechanism -- a .param needs alterparam+reset (re-source), an instance/model knob is in-place alter/altermod -- and a reset drops every in-place alter, so the loop iterates inner-fastest and re-resources ONCE

only when a .param knob's value actually changed (outer-loop boundaries; an inner .param still resets every point, a pure in-place sweep never resets), staging all .params via alterparam before the single reset, then RE-APPLYING every in-place knob after the reset (so an inner alter survives an outer .param re-source). By construction a single knob reduces EXACTLY to E-146: one outer combination -> curve named <output>, E-146 messages, identical data indexing (existing sweep/sweepwave examples pass untouched). -overlay composes: each cartesian point's full waveform becomes a sweepwave vector named with EVERY knob value (inner first), so a single knob is <output>_<value> exactly as E-189. New static helpers: sw_parse_spec, sw_set_deck/sw_set_inplace (split from the old sw_set), sw_familyname (summary curve name over outer knobs), sw_pointname (overlay name over all knobs; replaces sw_wavename). Verify: RC low-pass step, R inner x C outer, each family curve == $1 - \exp(-T/(R \cdot C))$ vs R to worst <5e-6, incl. a .param-OUTER knob (reset/alter ordering); cartesian run/curve counts; -overlay family names carry both knob values; single knob still names its curve vo. 5 checks. Regression 154/154.

E-191 *spicelib/analysis/acan.c, spicelib/analysis/span.c, examples/aclin2_examples/**

.ac lin 2 / .sp lin 2 off-by-one fix. Found while auditing the data-output path (wldata/wrsnp/save/rawfile write): a ragged-length wldata test wrote an AC vector one point short, which traced not to the writer but to the AC LINEAR frequency sweep. The AC (acan.c) and S-parameter (span.c) analyses computed the linear step behind if (numberSteps - 1 > 1) -- true only for N>=3 -- so exactly lin 2 fell through to the single-point patch (freqDelta = 0). With a zero step the sweep loop (freq += freqDelta; if (freqDelta == 0) goto endsweep;) stops after ONE point, silently dropping the fstop point: a two-point linear AC/SP sweep produced one row. dec/oct (different step formula) and lin 1/3/4/... were all correct; only N=2. The single-point patch (Richard McRoberts) was meant for lin 1; - 1 > 1 over-reached by one. FIX: guard on numberSteps > 1 in both files (the form noisean.c already used correctly), so N=1 -> one point at fstart (patch preserved), N=2 -> freqDelta = fstop-fstart -> two endpoints, N>=3 unchanged. The distortion analysis uses a different N+1-point convention (unaffected); noise was already correct. Verify: RC low-pass, both lin 2 endpoints are GENUINE solved points matching $1/(1+j\omega RC)$ to <1e-6 (not a duplicate); lin 1 stays one point; lin N gives N linearly-spaced points (N=3,5,10); .sp lin 2 gives two points with analytic series-R S-params (S11=1/3, S21=2/3). Both solvers. The wider output-path audit found the writers themselves CLEAN: Touchstone wrs2p/wrsnp/rdsnp round-trip across RI/MA/DB, S/Y/Z (Rbase normalization $Y \cdot R / Z / R$), Hz-GHz incl. complex angles; 2-port special S11 S21 S12 S22 order + N-port row-major both correct for round-trip AND spec interop; wldata complex (re,im)/singlescale/ragged padding; save filtering restricts the vector set; rawfile write/load bit-exact at 12 digits. 5 checks. Regression 155/155.

E-192 *frontend/com_checkpoint.c, frontend/com_checkpoint.h, frontend/runcoms.c, examples/autosave_examples/**

Auto-checkpoint on interrupt (**set autosave=<file>**). Follow-up to E-131 checkpoint/restart, prompted by "does savestate fire on Ctrl-C?" -- it did not. E-131 gave savestate/loadstate but a checkpoint had to be taken by hand, so an interrupted long transient lost its progress. E-192: set autosave=<file> makes an interrupted transient auto-write an E-131 checkpoint. SAFE HOOK POINT (the crux): writing a file from a signal handler is undefined behavior, so E-192 does NOT checkpoint in ft_sigintr -- that handler only sets ft_intrpt and returns (during a run ft_setflag is TRUE, so no longjmp). At the next ACCEPTED timepoint, dctran's IFpauseTest (which is OUTstopnow) sees the flag and returns 1, so dctran returns E_PAUSE; if_run maps E_PAUSE to return code 1; dosim (runcoms.c) reports "interrupted" (err equals 1). E-192 hooks exactly that interrupted branch -- on the MAIN thread, at a consistent timestep boundary, with ordinary buffered I/O. CHANGE: com_savestate's body factored into ckt_write_checkpoint(ckt, fname) (the command is now a thin wrapper); the hook calls the SAME function, so an autosave file is byte-for-byte what savestate writes. In the interrupted branch: if cp_getvar("autosave", ...) yields a filename AND the run was a transient (CKTmode and MODETRAN set,

CKTtime above 0), call `ckt_write_checkpoint`. OPT-IN (no var means byte-identical old behavior, nothing written), TRANSIENT-ONLY (the MODETRAN / CKTtime guard blocks stale state after, e.g., an interrupted AC). Interactive-mode only: the SIGINT handler is installed only under non-batch mode, so a batch-mode Ctrl-C still terminates (unchanged). A real Ctrl-C (SIGINT), a stop when ... breakpoint, and a Verilog-A stop task all funnel through the same interrupted branch (they differ only in how the flag is set, upstream of E-192). CORRECTNESS: autosave file byte-identical to a manual savestate at the same pause; loadstate resumes it exactly (E-131 round-trip). Verify (examples/autosave_examples): stop-paused transient writes a checkpoint; equals a manual savestate (byte-identical); loads/resumes; opt-in gate (no var means no file); and a REAL Ctrl-C (SIGINT) via a PTY (interactive) fires it (lenient: skip if the run finishes before the signal, the hook being covered deterministically). E-131 checkpoint example still passes (savestate refactor). 5 checks. Regression 156/156.

E-193 *spicelib/analysis/dcpss.c, examples/pnoiseunits_examples/*, examples/rfpss_examples/verify_rcpnoise.py, examples/rfpss_examples/verify_rcyclo.py, examples/cyclofold_examples/verify_cyclofold.py, examples/pnoiseexamples/verify_pnoiseexamples.py, examples/rfpss_examples/rc_pac.cir*

.pnoise honors **sqrnoise** (noise-density units fix). Found auditing the RF/PSS suite. The ngspice manual (sec. **.noise**, pp. 326-327) defines `onoise_spectrum/inoise_spectrum` as a noise spectral DENSITY in $V/\sqrt{\text{Hz}}$ (or $A/\sqrt{\text{Hz}}$) by default, switchable to the squared V^2/Hz form by the `sqrnoise` control variable. **.noise** obeys this (noisean.c reads `cp_getvar("sqrnoise")` and `sqr`'s `onoise/inoise` when unset). But **.pnoise** (periodic noise) ALWAYS emitted the squared V^2/Hz density and IGNORED `sqrnoise` -- so the same-named vectors carried DIFFERENT units across the two analyses (differing by a square), and the documented `sqrnoise` control silently did nothing for `pnoise`. Confirmed by the toggle: **.noise** flips $4.063\text{e-}9$ ($V/\sqrt{\text{Hz}}$) $\leftrightarrow$ $1.651\text{e-}17$ (V^2/Hz) with `set sqrnoise`; **.pnoise** returned $1.651\text{e-}17$ both ways. FIX: `pnoise_sweep` (`dcpss.c`) now reads `sqrnoise` via `cp_getvar` (added `#include cpextern.h`) and, when unset (the default), applies `sqr` to the emitted `onoise/inoise` densities $\rightarrow V/\sqrt{\text{Hz}}$; `set sqrnoise` $\rightarrow V^2/\text{Hz}$. BOTH emit paths get the conditional `sqr`: the E-178 cyclostationary folding branch and the standard adjoint-folding branch. Now **.pnoise** and **.noise** agree by default (match to 0.00). SCOPE: **.pnoise** only. The two-tone `qpnoise` command's single-point diagnostic already prints BOTH forms explicitly (V^2/Hz and $V/\sqrt{\text{Hz}}$) so there is no ambiguity, and its swept + single-point paths share the V^2/Hz convention; an initial `QpnoiseSweep` change was REVERTED to avoid an internal swept-vs-single-point split (`verify_qpss_sweep` [3] compares them). EXAMPLES: `rc_pnoise` now checks the default $V/\sqrt{\text{Hz}}$ form (`sqr` of the analytic) and cross-checks vs **.noise** without forcing `sqrnoise` (match to 0.00); `rc_cyclo/cyclofold/pnoiseexamples` are power-density relations (`onoise == ..., referee sums |Z|^2S`) so they set `sqrnoise` to keep V^2/Hz , and two LTI-limit checks that compared `pnoise(V^2)` to `.noise(V/\sqrt{\text{Hz}})^2` became direct equalities. WIDER AUDIT (all correct): PAC driving-point + the +-1 CONVERSION SIDEBANDS (via the `.pac` card's trailing `maxsideband` arg) match a clean transient DFT to ~0.1%; Harmonic Balance matches a transient DFT to 4-5 digits (an apparent 12x `.disto` HD3 and 2.4x PAC-sideband discrepancy both dissolved into `fourier`-command interpolation error at high harmonics, resolved by a proper DFT). Also documented the previously-undocumented `.pac` trailing `maxsideband` arg (`dcpss.c` comment + `rc_pac.cir`). Verify (examples/pnoiseunits): default `pnoise == sqrt(4kTR/(1+(wRC)^2))` ($V/\sqrt{\text{Hz}}$); `set sqrnoise ==` the density (V^2/Hz); `default^2 ==` squared (exact `sqr` relation); default `pnoise ==` default `.noise` (0.00). 4 checks. Regression 157/157.

E-194 *frontend/com_optimize.c, examples/pssopt_examples/**

Particle swarm optimization (**optimize -method pso**). A new search method for the built-in optimizer. E-130/143 gave two LOCAL methods -- Nelder-Mead simplex (`-method nm`, scalar -- minimize) and Levenberg-Marquardt (`-method lm`, gradient least-squares over `-targets`) -- both of which descend into whichever basin the start point sits in, so on a MULTIMODAL objective they

return the local minimum nearest the start. E-194 adds a GLOBAL, population-based, derivative-free method: PSO. N particles fly through the normalized $[0,1]^n$ box; each keeps its own best-seen point (pbest) and the swarm shares a global best (gbest); each iteration updates every velocity $v = \text{chi} * (v + \text{phi} * r1 * (\text{pbest} - x) + \text{phi} * r2 * (\text{gbest} - x))$ and moves $x = \text{clamp}_{01}(x + v)$ with the standard Clerc-Kennedy constriction ($\text{chi}=0.72984$, $\text{phi}=2.05$ so the swarm converges without exploding), velocities clamped to half the box, positions to $[0,1]$. Particle 0 starts at the user init point, the rest uniform-random, so the whole box is explored. Reuses the existing `opt_eval` (returns the scalar cost for both objective kinds), so PSO works for scalar `-minimize` AND `-target` least-squares -- unlike LM, which needs targets. New options: `-swarmsize <N>` (aliases `-swarm/-npart`; default `auto = 10+4np capped at 60`), `-seed <s>` (PSO draws from a small self-contained `splitmix64` PRNG independent of `ngspice`'s global RNG, so a seed gives byte-identical results), `shared -maxiter/-tol` (`tol` is the gbest-stagnation tolerance held over several iterations). The `nm/lm` dispatch is unchanged. Verify (examples/`psoopt`): on the classic multimodal $f(p) = \sin(p) + \sin(10 * p / 3)$ over $[2.7, 7.5]$ (global $p=5.1457$ $f^*=-1.8996$, several higher local minima) started at the $p=2.7$ corner, PSO finds the global (-1.8996 , $p=5.145$) while Nelder-Mead from the same start is trapped at a local minimum (-1.19992 , $p=3.387$); a fixed seed is exactly reproducible; independent seeds all reach the global basin; PSO also solves a `-target` least-squares objective (recovers p to a $1e-12$ residual) and a 2-D separable multimodal min ($f^*=-3.799$). Existing optimize example (`nm+lm`, 20 checks) unchanged. 6 checks. Regression 158/158.

E-195 `frontend/com_optimize.c, examples/deopt_examples/*`

Differential evolution (**optimize -method de**). A second GLOBAL method for the built-in optimizer, after PSO (E-194), from the same population-based derivative-free family. DE keeps a population of NP vectors in the normalized $[0,1]^n$ box; where PSO pulls trials toward remembered bests, DE (classic DE/rand/1/bin) builds each trial from a scaled DIFFERENCE of random members and crosses it with the target: $v = a + F * (b - c)$ ($F=0.8$; a, b, c distinct and $!= i$), $u[j] = v[j]$ if $\text{rand} < CR$ or $j == j_{\text{rand}}$ else $x[i][j]$ ($CR=0.9$ binomial crossover, one dimension always taken), then greedy selection (keep u in place of $x[i]$ iff $\text{cost}(u) \leq \text{cost}(x[i])$). The difference vector $b-c$ self-scales to the population's own spread -- large while members are far apart (broad exploration), shrinking as they converge (fine refinement) -- which makes DE robust on rugged / discontinuous / poorly-scaled landscapes with no per-problem step tuning. Reuses the existing `opt_eval` (scalar cost for both objective kinds), so it works for scalar `-minimize` AND `-target` least-squares like PSO. Shares PSO's infrastructure: the same self-contained `splitmix64` PRNG (so `-seed` is byte-reproducible), the `-swarmsize` population option (default `auto 10+4np capped at 60; clamped to >=5 because DE needs 4 distinct members to form a mutant`), and `shared -maxiter/-tol` (gbest-stagnation). The `nm/lm/psa` dispatch is unchanged; DE is only selected by `-method de` (aliases `difevol, differentialevolution`). Verify (examples/`deopt`): on the multimodal $f(p) = \sin(p) + \sin(10 * p / 3)$ over $[2.7, 7.5]$ (global $p=5.1457$ $f^*=-1.8996$) started at the trapping $p=2.7$ corner, DE finds the global (-1.8996 , $p=5.14574$ -- matches p^* to 5 digits) while Nelder-Mead from the same start is trapped at a local minimum (-1.19992 , $p=3.387$); a fixed seed is exactly reproducible; independent seeds all reach the global basin; DE also solves a `-target` least-squares objective (recovers p to a $4e-13$ residual) and a 2-D separable multimodal min ($f^*=-3.799$); and DE and PSO both reach the global while the local NM does not. Existing optimize (`nm+lm`, 20 checks) and `psoopt` (6 checks) examples unchanged. 7 checks. Regression 159/159.

E-196 `frontend/com_optimize.c, examples/saopt_examples/*`

Simulated annealing (**optimize -method sa**). A third GLOBAL method for the built-in optimizer, completing the set: after PSO (E-194) and DE (E-195), both population-based, SA is the classic SINGLE-WALKER global optimizer. It holds one current point in the normalized $[0,1]^n$ box; each step it proposes a random neighbour $xn = \text{clamp}_{01}(x + \text{step} * U(-1, 1))$ and applies the Metropolis rule -- accept if $\text{cost}(xn) \leq \text{cost}(x)$ (downhill), else accept anyway with probability $\exp(-(\text{cost}(xn) - \text{cost}(x)) / T)$ (uphill). While the temperature T is high uphill moves are readily accepted so the walker climbs OUT of local minima; T is then cooled geometrically

($T \leftarrow \alpha$, α set so T drops ~ 4 decades over the `-maxiter` cooling levels, with $L=8+4np$ moves per level), so uphill moves become rare and it settles into a minimum. The best point ever visited is returned. ONE candidate per step (no population) $\rightarrow$ the lightest-weight global method, the natural choice when each analysis is expensive. Everything auto-scales to the problem: $T_0 = \text{mean } |D_{\text{cost}}|$ of a dozen random probes (so an uphill move of that size is accepted about half the time when hot), and the step size $= 0.30 \cdot \sqrt{T/T_0} + 0.02$ shrinks as T cools (broad exploration hot, fine refinement cold) -- nothing to tune. Reuses `opt_eval` (scalar cost both kinds) $\rightarrow$ works for scalar `-minimize` AND `-target` least-squares. Shares the self-contained `splitmix64` PRNG (`-seed` byte-reproducible); `-maxiter` is the number of cooling levels. The `nm/lm/pso/de` dispatch is unchanged; SA is only `-method sa` (aliases `anneal`, `simulatedannealing`). A single walker refines more loosely than a population, so SA reliably reaches the global BASIN rather than machine precision (a short `nm/lm` polish from its result sharpens it when that matters). Verify (examples/saopt): on the multimodal $f(p) = \sin(p) + \sin(10 \cdot p/3)$ over $[2.7, 7.5]$ (global $p=5.1457$ $f^*=-1.8996$) from the trapping $p=2.7$ corner, SA finds the global basin ($f < -1.85$, $|p-p^*| < 0.05$) while Nelder-Mead is trapped at a local minimum (-1.19992 , $p=3.387$); a fixed seed is reproducible; independent seeds all reach the global basin; SA solves `-target` least-squares and a 2-D multimodal min; and all THREE global methods (`sa/pso/de`) reach the global while the local `nm` does not. Existing `optimize` (20) + `pssopt` (6) + `deopt` (7) examples unchanged. 7 checks. Regression 160/160.

E-197 `frontend/com_optimize.c`, `examples/opt100_examples/*`

A 100-parameter circuit optimization (raised **optimize** limits). The built-in optimizer sized its per-run arrays from two compile-time macros; E-197 raises them from `OPT_MAXP` 16 to 128 (parameters) and `OPT_MAXT` 64 to 128 (least-squares targets), and lifts the auto-population cap for the PSO/DE population methods from 60 to 256, so a genuinely large problem fits in a single call. Exceeding a cap is reported cleanly (`optimize: too many -param (max 128)`), not crashed. Testbed: 100 independent 1 mA sources, each into an unknown resistor R_i to ground (so $v(n_i) = 1e-3 \cdot R_i$), fitted to 100 targets $t_i = 1e-3 \cdot (100 + 9800 \cdot i/99)$ ohm -- a linear voltage ramp 0.1 .. 9.9 V, a well-posed separable 100-D least-squares. Levenberg-Marquardt is the right tool for a well-posed high-D fit and solves it to machine precision (sum-sq residual $2.5e-29$, rms $5e-16$, 619 evaluations, ~ 2 s) with all 100 parameters and 100 targets honored. Two fixes make the derivative-free GLOBAL methods FUNCTION at that scale (they are intrinsically far slower than LM in high dimension, but were failing outright): (1) adaptive DE crossover -- classic DE/rand/1/bin uses $CR=0.9$, mutating almost every coordinate of a trial; in high dimension a trial that perturbs $\sim n$ coordinates at once is essentially always worse than its parent and rejected by greedy selection, so DE freezes at its starting guess. E-197 sets $CR = (n \leq 16) ? 0.9 : 15.0/n$, capping the expected mutated-coordinate count at ~ 15 for large n (exactly the classic 0.9 for $n \leq 16$, so low-D DE is unchanged), and DE descends again -- from ~ 1240 to ~ 882 over 35 generations at 100-D. (2) Dimension-scaled convergence patience -- high-D population runs plateau for several generations between improvements, so the gbest-stagnation counter that stops PSO/DE would cut a still-descending run short; the threshold now grows as $8 + n/4$ (exactly 8 for $n \leq 3$, so small- n tests are unchanged), giving 33 generations of patience at $n=100$. DE at 100-D remains a genuinely hard global search that does not fully converge in a short budget -- intrinsic to global optimization in high dimension, not a defect; the message is the division of labor between LM (well-posed high-D fits) and the global methods (multimodal, now usable as dimension grows). Verify (examples/opt100): LM fits all 100 params to 100 targets at machine precision (both raised caps proven at once); DE descends substantially at 100-D (self-calibrated against its own starting cost, $\geq 15\%$ reduction); a fixed `-seed` is byte-reproducible even at 100 dimensions; and a 129-parameter over-cap run is reported cleanly. Existing `optimize` (39) + `pssopt` (6) + `deopt` (7) + `saopt` (7) examples unchanged. 4 checks. Regression 161/161.

E-198 `frontend/com_stb.c`, `frontend/commands.c`, `frontend/com_commands.h`, `frontend/Makefile.am`, `examples/stb_examples/*`

stb stability / loop-gain analysis. A new front-end command that measures a feedback loop's

small-signal loop gain $T(f)$ and reports its phase and gain margin -- the everyday stability check for op-amps, regulators, PLLs -- which ngspice lacked (you had to hand-roll a `.control` script). It is the one-command equivalent of Spectre's `stb` (Middlebrook-Tian double injection). Usage `stb <Vprobe> <Iprobe> <ac-sweep>`: the user marks the loop break with a bias-transparent probe pair between the driving node A and the loaded node B -- a series `Vstb A B dc 0 ac 0` (+node = driver A, -node = load B) and a shunt `Istb 0 B dc 0 ac 0` (ground to load B). Breaking a loop and injecting a single test signal mis-reads the loop gain because of the LOADING at the break (a series voltage injection is exact only from a zero-impedance source into an infinite-impedance load); the double injection cancels that error. `stb` runs two AC sweeps by altering the probe AC magnitudes: voltage injection (`Vstb ac=1`) gives $T_v = -v(A)/v(B)$; current injection (`Istb ac=1`) gives the current loop gain from the probe branch current $a = i(Vstb)$ as $T_i = -a/(a+1)$; combined $T = (T_v T_i - 1)/(T_v + T_i + 2)$. *The combined T is INDEPENDENT of where the probe sits in the loop -- the property a single injection lacks -- verified directly (a loaded break gives the same margins as a clean break in the same loop). Numerically, the common clean-break case has a high load impedance so $a \rightarrow -1$ (all injected current returns through the shorted voltage probe) and $T_i \rightarrow \text{inf}$, which makes $T_i = -a/(a+1)$ divide by zero; `stb` evaluates the algebraically identical singularity-free form $T = -(T_v a + a + 1)/(T_v a + T_v + a + 2)$, exact at $a = -1$ where it reduces to $T = T_v$.* The complex loop gain is stored as vector `loopgain` vs frequency in a new `stb` plot (`plot db(loopgain) / plot 180/PI*cph(loopgain)`); phase margin = $180 + \text{angle}(T)$ at the first $\text{abs}(T)=1$ crossing, gain margin = $-\text{abs}(T)_{\text{dB}}$ at the first $\text{angle}(T)=-180$ deg crossing (both log-interpolated on an unwrapped phase). Implementation (`com_stb.c`, registered in `commands.c`, declared in `com_commands.h`): recovers the probe's two node names via `CKTinst2Node`, dispatches `alter/ac` through the command table (as `com_optimize/com_sweep` do), and reads the complex vectors $v(A)/v(B)/Vstb\#branch/frequency$ with `ft_evaluate`. Reuses the AC analysis, so both the Sparse and KLU solvers. Verify (examples/stb): PM/GM match the closed-form loop gain of an ideal-output op-amp loop; a loaded break (finite op-amp output impedance) gives the SAME margins as a clean high-impedance break (proving the current injection corrects the loading -- a single voltage injection is ~10% off and diverges past crossover); a 10x-gain design's reduced phase margin is tracked and matches analytics; the complex `loopgain` is stored; and an unknown probe is reported cleanly. 5 checks. Regression 162/162.

E-199 *examples/nport_examples/snp2va.py, examples/nport_examples/**

N-port Touchstone device (`snp2va.py`). Instantiate a measured or simulated S-parameter block (a filter, cable, connector, package, or amplifier stored in a Touchstone `.sNp` file) as a circuit element that works in AC AND transient. ngspice had no native n-port element: its `.sp` "ports" are tagged voltage sources used to MEASURE a circuit's S-parameters, and `rdsnp/wrsnp` (E-64/E-72) are reader/writer commands, not a device. Rather than a native convolution device (with its own history buffers and stability headaches), `snp2va.py` converts the file into a Verilog-A n-port realized with `laplace_nd`, so OpenVAF's OSDI laplace machinery (E-31 complex poles) supplies BOTH the AC response and the transient (recursive-state) response with no convolution code. Pipeline: parse Touchstone (any port count; S/Y/Z data; MA/DB/RI formats; Hz..GHz; arbitrary reference impedance) $\rightarrow S(f) \rightarrow Y(f) = (1/z_0)(I-S)(I+S)^{-1} \rightarrow$ common-pole vector fit (Gustavsen) $\rightarrow$ emit `I(p_i) <+ sum_j [ laplace_nd(V(p_j), num_ij, den) + e_ijddt(V(p_j)) ]`. *Two numerical details are essential. (1) The vector fit MUST be done in NORMALIZED frequency -- without it the se column is $\sim 1e9$ while the $1/(s-p)$ columns are $\sim 1e-7$, so the least-squares is hopelessly ill-conditioned and the poles never relocate (the pole relocation is done as polynomial roots of the sigma numerator via Durand-Kerner, not a matrix eigensolve). (2) Y is IMPROPER whenever a network has shunt capacitance ($Y \sim sC$ at high frequency), which `laplace_nd` -- a ratio of polynomials -- cannot represent (it gave a 63 percent AC error at the band edge); the fix is to give `laplace_nd` only the strictly-proper (pole) part and emit the s^*e term as an explicit `ddt` (a capacitance), dropping the AC error to $\sim 1e-6$. HARDENING: automatic order selection CLIMBS the pole count to capture high-order and delay-dominated responses (a pure delay is uniformly poor at low orders, so the selection must not quit early on a small between-order improvement), but stops at the KNEE once the*

error has already come down near a floor -- past that knee extra poles just fit measurement noise into an over-fitted, ill-conditioned, unstable model, so stopping there fits clean data to machine precision and noisy data at its noise floor. It breaks when a fit turns unstable (polynomial-root finding degrades at high degree) and returns the best STABLE fit. Stability is enforced (right-half-plane poles reflected into the left half plane) so the emitted model is always BIBO-stable even for a noisy non-passive fit; passivity is checked and reported (enforcement by residue perturbation is a documented non-goal for now). Pure Python standard library, NO numpy: least-squares via Householder QR, complex matrix inverse via Gauss-Jordan, and polynomial roots via Durand-Kerner are all reimplemented (each validated against numpy to machine precision during development). Verify (examples/nport): each check generates a Touchstone file from a network whose response is known exactly, runs snp2va.py + OpenVAF, and confirms the device matches the ORIGINAL network in ngspice -- an R-L-C resonator to 5e-6 in AC (including the transmission peak) and 5e-3 in transient (one model, both analyses); a 5-pole LC ladder (order selection scales past 2 poles); a 3-port star network to 4e-9 on both coupled outputs; and a noisy measurement fit at its noise floor, reported non-passive, staying bounded in transient. Stress-tested to near machine precision on resonators, notches, high-Q blocks, and multi-pole ladders; a delay-dominated transmission line degrades gracefully (AC accurate over the fitted band, transient pulse edges ring -- inherent to rational fitting of a pure delay, for which a T/LTRA line is the right model). 6 checks. Regression 163/163.

E-200 `src/frontend/snp2va.c, src/frontend/snp2va.h, src/frontend/com_presnp.c, src/frontend/commands.c, src/frontend/inp.c, src/frontend/inpcom.c, src/frontend/Makefile.am, examples/presnp_examples/*`

Built-in `pre_snp` command. Folds the E-199 Touchstone-to-Verilog-A converter into ngspice itself as a C command, so the workflow no longer needs an external Python script: inside a `.control` block, `pre_snp <file.sNp> [module]` parses the `.sNp`, common-pole vector-fits it, writes the `.va`, and invokes `openvaf-r` to compile the `.osdi` -- a one-line analog of `pre_osdi`. `snp2va.c` is a faithful C port of `snp2va.py` with identical numerics: it uses C99 double *Complex* (`typedef'd cplx`, so it does not collide with ngspice's own `complex.h` macros) and reimplements the three vector-fitting primitives -- least-squares via Householder QR, complex matrix inverse via Gauss-Jordan, polynomial roots via Durand-Kerner -- plus the same parse $\rightarrow S(f) \rightarrow Y(f) = (1/z_0)(I-S)(I+S)^{-1} \rightarrow$ vector fit $\rightarrow$ emit $I(p_i) \leftarrow + \sum_j [\text{laplace_nd}(V(p_j), \text{num}, \text{den}) + e_{ij} \text{ddt}(V(p_j))]$ pipeline, the same normalized-frequency fit, the improper *es* (*shunt-C*) term split out as an explicit *ddt*, and the same automatic order selection returning the best STABLE fit (RHP poles reflected $\rightarrow$ BIBO-stable). It was checked numerically identical to the Python reference (resonator RMS 4e-6, ladder 7e-7, 3-port 3e-8). The public entry point `snp2va_convert(snpfile, vafile, module, msg, msglen)` is called by the command wrapper `com_pre_snp` (`com_presnp.c`), which then shells out to `openvaf-r`. ORDERING: the command is registered under the name `snp`, so writing `pre_snp` in a deck triggers ngspice's existing *pre* mechanism (`inp.c` strips the `pre_` prefix and runs the command before the circuit is parsed). That alone is not enough -- `pre_snp` must run before the `pre_osdi` that loads its output, and control lines otherwise execute in deck order -- so the collected pre-commands are now executed in TWO PASSES: pass 0 runs every `snp` command (in deck order), pass 1 runs the rest. The `.osdi` a `pre_snp` generates is therefore guaranteed to exist by the time any `pre_osdi` loads it, EVEN IF the deck lists the `pre_osdi` line first; the ordering is a guarantee, not a convention the user must remember. `snp/pre_snp` are also added to `inpcom.c`'s case-preservation list (beside `osdi/pre_osdi`) so the file-path argument keeps its original case. The compiler is located by `find_openvaf()`: the `openvaf` ngspice variable (set `openvaf=...` in `spinit` or on the command line -- a set inside `.control` runs too late, after the pre-commands), then the `OPENVAF` environment variable, then `$SPICE_LIB_DIR/openvaf-r`, then `PATH`; if none resolve, `pre_snp` reports the failing `openvaf-r` invocation and lists all four options. Verify (examples/presnp): each check builds a Touchstone file from a network whose response is known exactly, lets `pre_snp` do the whole `convert+compile+load` inside ngspice, and confirms the device matches the ORIGINAL network -- `pre_snp` writes the `.va` and compiles the `.osdi`; the device matches an R-L-C resonator to 5e-6 in AC (including the transmission peak) and to 5e-3 in

transient (one compiled block, both analyses); a 3-port star network compiles and matches on both coupled outputs to 4e-9; and with pre_osdi written BEFORE pre_snp the model still loads, proving the ordering guarantee. 5 checks. Regression 164/164. HARDENING (from an N-port stress sweep -- generate an N-port RLC network, extract S-params with .sp, round-trip through pre_snp, compare DC/AC/transient vs the original subcircuit): two defects the 2-pole checks never reached, both fixed in snp2va.c. (1) Order-selection double-free: the climb that keeps the best STABLE fit let its best and prev buffers alias and then freed the same allocation twice, so pre_snp aborted (SIGABRT) on any network whose fit needed ≥ 3 pole orders (the 2-pole resonator/star escaped by breaking at tolerance first); fixed by not freeing a buffer the other pointer still holds. (2) Transient divergence from a non-passive realization: each Y_{ij} is fit independently, so the model could blow up in transient ($v \rightarrow \sim 1e284$) even though every pole is in the LHP and AC is exact -- a single degree-Np rational per element has coefficients spanning $\sim \text{abs}(p)^{Np}$ ($\sim 1e79$ for 8 poles at $1e10$ rad/s, an ill-conditioned companion form), and the improper e^s terms form a capacitance matrix with no passivity constraint (an indefinite "negative-capacitance" matrix $\rightarrow$ unstable DAE); fixed by realizing each element as a PARALLEL bank of first/second-order laplace_nd sections (every coefficient $\leq O(\text{abs}(p)^2)$) and projecting the e-matrix onto the symmetric PSD cone (symmetrize $\rightarrow$ Jacobi eigendecompose $\rightarrow$ clamp negative eigenvalues to 0). The presnp example gained a 4-port coupled-ladder check (fit climbs past two orders; transient must stay bounded and match) that fails on the pre-fix converter. Regression 164/164.

E-201 *src/frontend/snp2va.c, examples/presnp_examples/**

pre_snp scalability (fast vector fitting + shared-pole realization). The E-200 converter struggled from about 8 ports up (an 8-port took $\sim 190s$ and larger port counts were infeasible). Two things scaled badly, both in snp2va.c: the pole identification stacked all NN *elements into one dense least-squares of size* $(2NsNN) \times (NN(Np+2)+Np)$ -- $O(N^4)$ memory, $O(N^6)$ compute, re-solved dozens of times over the order climb ($\sim 8TB$ of RAM at $N=100$); and the emitted model gave every one of the NN *elements its own laplace_nd bank* -- $O(N^2Np)$ filter sections and OSDI state. Four fixes rebuild the scalability. (1) FAST VECTOR FITTING (Deschrijver/Gustavsen): the stacked pole-ID matrix has block-arrow structure -- each element's residue columns are independent and couple only through the shared common-pole (sigma) columns -- so instead of forming that matrix, each element block is Householder-reduced on its own (a small $2Ns \times (Np+2)$ QR) and only its residual rows (orthogonal to that element's own columns) are stacked into one tall $((2Ns-Np-2)Ne) \times Np$ system solved for sigma; memory $O(N^4) \rightarrow O(N^2)$, compute $O(N^6) \rightarrow O(N^2NsNp^2)$. (2) RECI-PROCITY: a passive network has symmetric Y, so when detected only the upper triangle $(N(N+1)/2)$ elements is fit and residues are mirrored ($\sim 2x$ fewer solves). (3) ITERATION CAP: pole relocation stops once the poles settle (relative move $< 1e-4$) instead of a fixed 12 sweeps (typically 3-5). (4) SHARED-POLE REALIZATION: all NN elements share the same poles, so the pole-filters are computed ONCE per input port -- a real pole gives one filter $V(pj)/(s-p)$, a conjugate pair two real basis filters $(s-\sigma)/D$ and ω/D , with the residue entering each output as a real scalar weight -- and each output current is a cheap weighted sum; this is $O(NNp)$ laplace_nd and OSDI state instead of $O(N^2Np)$, the difference between a model that compiles/simulates at large N and one that does not. Three latent bugs the stress test also surfaced, all fixed: (a) an order-selection double-free -- at higher pole counts a relocated pole set could lose its adjacent-conjugate-pair structure and build_basis/ctil_to_cres (which walk $i+=2$ on a complex pole) wrote one slot past the residue array; fixed by a canon_poles step that snaps near-real poles to real and rebuilds exact adjacent conjugate pairs. (b) a 256-entry stack array -- the Touchstone parser assembled each frequency's matrix in a fixed cplx pv[1616], hard-capping the port count at 16 ($N=32$ smashed the stack); moved to the heap, auto-detect limit lifted to 512. (c) an OpenVAF crash -- a laplace_nd coefficient array assigned to a variable (as the shared realization does) crashes OpenVAF when a coefficient is an integer literal like '{1}' (which %.12g prints for 1.0); the emitter now forces a real literal '{1.0}'. Results: fit at $N=8$ 220s $\rightarrow 3.6s$; $N=64$ fits in 47s (was $\sim 860GB$ of RAM); $N=100$ fits in 2.75min at $9.6e-4$ error; model laplace_nd count now NNp (65 at $N=8$, 449 at $N=32$) not $\sim N^2Np/2$ (256, 7168). Correctness preserved: DC/AC/transient still overlay the original subcircuit ($N=8$ unchanged; $N=16$,

previously infeasible, matches to 1.3%). The remaining bottleneck is OpenVAF's compile time for the (now far smaller) shared model (~26s at N=16, ~43s at N=20), so the practical end-to-end ceiling is ~N=20-24 while the fit itself scales to N=100. Verify (examples/presnp): an 8-port coupled ladder converts+compiles in a few seconds with the compact shared model (65 laplace_nd for 8 ports) and matches the original network in AC; the E-200 order/realization guards and resonator/3-port checks still pass. 9 checks. Regression 164/164.

E-202 *src/maths/dense/dense.c, examples/spyscale_examples/**

.sp S-parameter matrix inverse is $O(N^3)$, not $O(N!)$. The pre_snp stress test (E-201) noted a separate wall: producing the input Touchstone file with ngspice's own .sp analysis was itself pathologically slow -- an 8-port took ~13s and a 10-port ~18 minutes, growing by roughly a factor of N per added port. That is in the RFSPICE analysis, not the converter. Profiling put ~99 percent of an 8-port .sp run inside CKTspCalcSMatrix (cktspdum.c), which builds the port S-matrix at every frequency by inverting $N \times N$ complex matrices -- several per frequency, for the S, Y and Z matrices. The complex inverse cinverse (maths/dense/dense.c) was computed by the adjugate/determinant method (Cramer's rule): $\text{cinverse} = \text{cadjoint}(A) * (1/\text{cdet}(A))$. cdet is a RECURSIVE COFACTOR EXPANSION -- $O(N!)$ -- and cadjoint computes N^2 cofactors (each an $(N-1) \times (N-1)$ determinant), so cinverse was $O(N^3N!)$; with several inverses per frequency across a whole sweep the cost explodes by a factor of N per port (6! -> 7! -> 8! is exactly the measured x7, x8, x9 per added port). FIX: cinverse and cinversedest now invert by GAUSS-JORDAN ELIMINATION WITH PARTIAL PIVOTING on the augmented $[A \ I]$, in $O(N^3)$ -- a drop-in replacement with the same signature and the same result (the true inverse). Behavior preserved: the old cinverse never returned NULL (the adjugate always produced a matrix, garbage for a singular input), so the new one returns NULL only on a true allocation failure and zero-fills a singular matrix; the singular case does not arise for a well-posed .sp network, and the two cktspdum.c call sites that do not null-check are safe. cdet/cadjoint remain for any other callers. RESULTS: .sp at N=8 13s -> 0.3s, N=10 ~18min -> 0.02s (~50000x), N=16 0.03s, N=32 0.09s; the extracted S-parameters are unchanged, verified against the closed-form network to ~2e-7. The same inverse is used by the periodic S-parameter path (.psp/PSS in dcpss.c), which benefits too; full example regression (touchstone, rfanalyses, psp, ...) unchanged. Together with E-201 the whole extract-or-import -> convert -> simulate workflow now scales to many ports. Verify (examples/spyscale): a 12-port RLC ladder -- a port count that took minutes before -- runs through .sp + wrsnp, completes in a fraction of a second, and every entry of the extracted 12x12 S-matrix matches the closed-form network across the sweep to ~2e-7. 2 checks. Regression 165/165.

E-203 *src/frontend/com_measure2.c, src/misc/string.c, examples/acmargin_examples/**

.meas ac gain/phase margin (+ batch vdb/vp auto-save fix). An audit against a 'richer TRIG/TARG, FIND...WHEN, integral/RMS/derivative, AC margins' wish-list found that ngspice's .measure already covers TRIG/TARG (with VAL, RISE, FALL, CROSS, LAST, TD, AT), FIND...WHEN, AVG/RMS/INTEG/DERIV, MIN/MAX/MIN_AT/MAX_AT/PP, ERR, and param='...', across tran/ac/dc/sp -- all present and working. The one genuine gap was AC margins. The textbook way to get them from the existing primitives, .meas ac gm FIND vdb(out) WHEN vp(out)=-180, CANNOT work: vp() returns the phase WRAPPED to $(-180, 180]$, so on any loop whose true phase sweeps past -180 the wrapped phase jumps from about -179 straight to about +179 and never equals -180, and the crossover is never found. FIX: two first-class .meas ac functions computed on the UNWRAPPED phase -- phase_margin (PM = 180 + the unwrapped phase at the 0 dB gain crossover) and gain_margin (GM = -gain(dB) at the -180 deg phase crossover). The phase is unwrapped by undoing the +-360 deg atan2 jumps, the 0 dB and -180 deg crossings are found by linear interpolation, and the crossover frequency is interpolated in log space. A loop with no finite -180 crossover (a one- or two-pole loop, infinite gain margin) is reported as such rather than as a bogus number. Implemented as measure_margin() dispatched from get_measure2() in com_measure2.c. One subtlety: ngspice's phase (vp, and the measure 'p' vector type) honors a runtime cx_degrees flag that defaults to RADIANS, so radtodeg() is a no-op by default; the margin code computes phase in

degrees explicitly. TWO BUGS the work surfaced, both fixed: (1) `gettok_iv` (`misc/string.c`), used only by the batch measure auto-save pass, copied the leading `v/i` then broke after the FIRST following character on `n_paren==0` -- so `vdb(out)` became `vd` (and `vm/vp/...` were all truncated); the stray `vd` then failed to save ('can't parse `vd`') and left the node unsaved, so a batch `.meas ac ... vdb(out)` died with 'no data saved for AC'. Fixed to stop only once the parenthesised part has actually closed. (2) a companion `strip_ac_vec` reduces `vdb(out)` to `v(out)` so the base node is what gets saved. Verify (examples/acmargin): 5 checks against closed-form loop gains (buffered one-pole sections, so poles sit exactly where placed) -- a stable 2-pole loop's phase margin matches analytic to under 1 deg (84.89 vs 84.89) and its gain margin is correctly reported infinite; a 3-pole loop's phase margin (-42.75 deg) and gain margin (-15.92 dB) both match the closed form exactly; and a batch `.meas ac ... vdb(out)` dot-card auto-saves the node and returns the right value. 5 checks. Regression 166/166.

E-204 `src/include/ngspice/optdefs.h, src/include/ngspice/tskdefs.h, src/include/ngspice/cktdefs.h, src/spicelib/analysis/cktsopt.c, src/spicelib/analysis/ckntask.c, src/spicelib/analysis/cktdojob.c, src/spicelib/analysis/cktop.c, examples/convhelp_examples/*`

auto-triggering DC convergence aids (`.option convhelp`). ngspice's CKTop operating-point cascade already auto-escalated through `gmin` stepping and source stepping, but the two strongest aids -- the globalized damped-Newton line search (E-111) and pseudo-transient continuation (E-127) -- only fired when the user hand-set `.option linesearch` or `.option ptcont`, and nothing reported which aid rescued a hard operating point. The ngspice-vs-Spectre gap analysis flagged exactly this: what remained was auto-triggering heuristics (reaching for these aids without the user asking) and folding them into the robustness presets. `.option convhelp` makes the whole cascade one automatic ladder -- Newton (with line search), then `gmin` step, then source step, then pseudo-transient, then `optran` -- and prints a one-line note naming the aid that produced the point. Under `convhelp` the op-point Newton, and its fallback homotopies (which call `NIiter` internally), run with line search enabled; E-111's line search reduces to a plain full Newton step whenever the full step already reduces the weighted residual norm, so points that converge easily take the same iterates and are unaffected. The E-127 pseudo-transient gate was widened so it also fires when `convhelp` is set, with no explicit `.option ptcont` needed. CKTop was refactored to a single done: exit that saves the caller's line-search flag on entry and restores it on every return (so the setting never leaks into the transient analysis that follows) and emits the aid note there. OFF BY DEFAULT: the default (`convhelp` unset) path is byte-for-byte the historical behavior with no new output. `errpreset=conservative` enables it through the E-110 `TSKtolGiven` given-flag mechanism (a new `ERRP_CONVHELP` bit), so an explicit `.option convhelp=0` still overrides the preset regardless of `.options` line order. Plumbing mirrors E-127: `OPT_CONVHELP` enum and `ERRP_CONVHELP` bit (`optdefs.h`), `TSKconvhelp` bitfield (`tskdefs.h`), `CKTconvhelp` bitfield (`cktdefs.h`), the `OPT_CONVHELP` setter plus `convhelp` in the `errpreset` value table plus the `convhelp OPTtbl` keyword (`cktsopt.c`), copy and default-off (`ckntask.c`), `CKTconvhelp = TSKconvhelp` (`cktdojob.c`), and the auto-escalation plus reporting in CKTop (`cktop.c`). Verify (examples/convhelp): the option is accepted; on a battery of normal circuits (diode, BJT, two-diode, resistor network) `convhelp` is result-neutral (node voltages identical to a plain run) and silent (no aid note, since plain Newton already converges); on a deliberately stiff behavioral-exponential deck with no junction limiting -- root of $1e-14 * (\exp(V/0.026) - 1) = (100 - V)/100$, i.e. $V(1) = 0.837922$ V -- where plain Newton overshoots to a spurious 70.5 V root (and even the transient-op fallback lands there), `convhelp` reaches the correct 0.837922 V via auto-fired pseudo-transient continuation, reports the aid, and changes the outcome versus plain Newton; `errpreset=conservative` enables the same ladder; and an explicit `convhelp=0` overrides the preset. Both linear solvers (KLU and Sparse 1.3). 35 checks. Regression 167/167.

E-205 `src/frontend/snp2va.c, examples/lowrank_examples/*`

`pre_snp` low-rank residue factorization. The E-200/201 shared-pole realization emits an $O(N^2)$ -term, $O(NNp)$ -filter Verilog-A model, so OpenVAF's compile time grows super-linearly with the port count and caps the practical N at ~20-24 (an $N=32$ block compiles in ~80s -- the vector fit itself scales

to $N=100$, but the emitted model's compile does not). When an N -port block's ports couple through only a few shared modes -- the common case for multi-port filters, cavities, and packages with a shared plane -- each pole's $N \times N$ residue matrix has rank r much less than N . E-205 exploits that: per channel (the conductance d , the capacitance e , and each pole section) the emitter builds the $N \times N$ real weight matrix and, via a new in-C one-sided Jacobi SVD, picks the cheaper of a dense emit (a term per significant entry, one `laplace_nd` per input port) or a low-rank emit $W = UV^T$ (rank kept above `tol_rank = 1e-7`, so a clean singular-value gap compresses fully while a slowly-decaying channel stays dense). The key is INPUT COMBINING: because `laplace` is linear, $\sum_j V[j][m] \text{laplace}(V(p_j)) = \text{laplace}(\sum_j V[j][m]V(p_j))$, so the low-rank form filters the r COMBINED inputs $u_m = \sum_j V[j][m]V(p_j)$ ONCE (r filters, not N) and distributes them to the outputs via U -- collapsing filters from $O(NNp)$ to $O(rNp)$ (the piece that dominates the compile) and coupling terms from $O(N^2)$ to $O(Nr)$. A full-rank block has no such structure, so every channel keeps the dense form and is emitted bit-identically (the auto-detect leaves it alone); the E-201 PSD projection of the capacitance e is preserved, so the low-rank transient stays stable. The env var `PRE_SNP_DENSE` forces the old dense realization (escape hatch / A-B test). Measured on a 3-shared-mode block (openvaf-r compile of the emitted .va): $N=16$ 16.9s to 1.5s (96 to 8 filters), $N=24$ 42s to 5.7s (144 to 10), $N=32$ 81s to 11.6s (192 to 8) -- super-linear to roughly linear, about 7-14x, with the response exact (AC $\sim 4e-7$, transient $\sim 3e-5$ versus a forced-dense build of the same fit). New code: a `jacobi_svd` and an `emit_filter` helper plus the per-channel chooser and input-combined emit, all in `snp2va.c`. Verify (examples/lowrank): a 12-port/3-mode block emits far fewer filters low-rank than dense (7 vs 73); its AC ($3.9e-8$) and transient ($1.2e-4$, bounded) equal the forced-dense build of the same fit; a full-rank 8-port ladder gets no compression with bit-identical AC. 5 checks. Regression 168/168.

E-206 `src/frontend/com_optimize.c`, `examples/dcenter_examples/*`

design centering / yield optimization (`optimize -center`). The capstone that fuses two whole sub-systems: the optimizer (Nelder-Mead / LM / PSO / DE / SA, E-130-197) and the Monte-Carlo yield suite (agauss/mccorr sampling + montecarlo specs, E-149-151). Design centering optimizes the NOMINAL design point to maximize parametric yield / Cpk under process variation. The outer optimizer searches the design knobs (`-dparam/-param/-mparam`, which feed the deck's agauss centres); the objective at each candidate is an INNER Monte Carlo of `-samples` N draws -- each reset re-samples the deck's process variation (agauss/.param, and any mccorr correlations) around the current centre -- scored against `-spec [-max hi] [-min lo]` limits (the same spec syntax as `montecarlo`) and reduced to the worst-case Cpk and the pass-fraction yield. The fusion is small and clean in `opt_eval`: the in-place-knob application was factored into `opt_apply_inplace` so the MC loop can re-apply the centre after each reset re-sources the deck, and a new `opt_eval_center` runs the inner MC and returns `-min(Cpk)`; the method drivers (`nm/ps/de/sa`) are unchanged -- they call `opt_eval`, which now returns the centering cost. Cpk is the objective (yield is reported alongside) because raw yield is a granular step function that stalls the simplex, whereas $\text{Cpk} = \min(\text{gap to each active limit}) / (3 \sigma)$, worst-cased across specs, is continuous and monotone in yield -- so maximizing it centres the distribution in its spec window. A fixed inner seed gives every candidate the SAME process draws (common random numbers), so the objective is deterministic and smooth; with `-lhs` the stratified sample-mean is about zero, so the optimum lands on the analytic centre. Reuses the E-149/151 sampling (`mc_lhs_config`, `mc_sss_off`, `setseed`); `lm` is rejected for `-center` (Cpk is not a least-squares residual) and the default method is Nelder-Mead; results are published as `dcenter_yield` and `dcenter_cpk` vectors. `com_optimize.c` only. Verify (examples/dcenter): a synthetic problem output $\sim N(xc, 0.5)$ with spec [4,6] (agauss(m,v,3) makes the actual $\sigma = v/3 = 0.5$, so the analytic centre is the midpoint 5 and the analytic Cpk is $1/(3*0.5) = 0.667$) is centred from an off-centre start $xc=4.0$ to $xc = 5.00$ with the expected Cpk 0.667 and a high yield; a montecarlo baseline at $xc=4.0$ sits near 50% while the centred design reaches $\sim 96\%$; and a two-knob problem with a lower and an upper spec on two outputs centres both params independently ($xa \rightarrow 5$ on [4,6], $xb \rightarrow 10$ on [9,11]). 5 checks. Regression 169/169.

E-207 `src/frontend/com_eye.c` (new), `src/frontend/com_eye.h` (new), `src/frontend/commands.c`, `src/`

*frontend/com_commands.h, src/frontend/Makefile.am, examples/eye_examples/**

eye diagram / jitter analysis (eye command). A new application domain: the core signal-integrity measurement for serial links / SerDes / high-speed I/O, which ngspice lacked. `eye -ui [-tstart t0] [-threshold vth] [-window frac]` is a self-contained front-end post-processing command (in the `fft / linearize / meas` family, registered in `commands.c/com_commands.h/Makefile.am`), solver-independent. It reads the transient waveform and its time scale via `ft_evaluate`; auto-detects the two logic rails from the 20th/80th percentiles of the samples (`eye_level0`, `eye_level1`, `eye_amplitude`) and the decision threshold (default the midpoint, or `-threshold`); finds every threshold crossing by linear interpolation; estimates the UI phase as the circular mean of the crossings modulo UI; computes each crossing's time-interval error (TIE, relative to the ideal UI grid) and reduces it to `eye_jitter_rms` (standard deviation of TIE) and `eye_jitter_pp` (peak-to-peak); measures `eye_height` (the vertical opening at the sampling instant, one half-UI from the transitions -- min of the samples above threshold minus max of the samples below, taken in a $\pm$ `windowUI slice`) and `eye_width` (UI minus the peak-to-peak jitter, the horizontal opening at the threshold), plus `eye_width_ber12 = UI - 14.069eye_jitter_rms` (the eye width at BER 1e-12 from the ± 7.035 sigma Gaussian random-jitter tail); and folds the waveform modulo $2*UI$ into the `eye_wave` vs `eye_t` vectors (built with `dvec_alloc/vec_new`, as `stb` and `sweep` build their result plots), whose scatter plot IS the eye diagram (plot `eye_wave` vs `eye_t`). No numerical-core changes. Publishes 11 permanent result vectors. Verify (examples/eye): an ideal 0/1 clock (perfectly periodic edges) recovers rails [0,1] and amplitude 1 with jitter about 1e-24 s (machine zero) and a full-UI, fully open eye; a clock whose edges carry a KNOWN injected Gaussian timing jitter ($\sigma = 20$ ps) is recovered to within a few percent (measured about 18.6 ps rms, 99 ps pp) and `eye_width` equals UI minus the measured peak-to-peak jitter; and the folded eye vectors are published and `wrdata eye_wave` writes the folded eye without crashing (the `eye_wave/eye_t` vectors live in their own fresh eye plot -- as `stb` does -- so `wrdata/plot` pair equal-length vectors; dumping the length-m eye vectors into the transient plot, whose scale is the full much-longer time vector, segfaults). 7 checks. Regression 170/170.

E-208 *src/frontend/plotting/pyplot.c, src/frontend/plotting/pyplot.h, src/frontend/com_pyplot.c, examples/pyplot_examples/**

`pyplot -eye`: one-line matplotlib eye diagrams. A convenience bridge between the E-207 eye analysis and the E-94 matplotlib back-end: `pyplot [name] -eye -ui [-tstart t0] [-threshold vth] [-window frac]` runs the eye analysis AND renders the folded eye in one command. `plot eye_wave vs eye_t` draws a single connected line through the folded samples in file order (a scribble, not an eye); `pyplot -eye` instead draws the samples as a 2-D histogram (`hist2d`, log-scaled `LogNorm`, `turbo colormap`, empty cells masked with `cmin=1`) so overlapping traces accumulate into a persistence eye -- exactly how a sampling scope paints one -- annotated with the decision threshold (dashed), the eye-centre sampling instant at 1 UI (dotted), a vertical eye-height double-arrow, a horizontal eye-width double-arrow along the threshold, and a title carrying UI / eye height / eye width (% UI) / RMS jitter. Implementation: `com_pyplot()` detects the `-eye` marker (a single bare token before it is taken as the output base name; default `eye`), runs `com_eye()` on the trailing tokens -- which folds the waveform and leaves `eye_wave/eye_t` plus the scalar metrics in a fresh current eye plot -- then calls the new `ft_pyplot_eye()` in `plotting/pyplot.c`, which reads those back from the current plot, writes `.data` (the folded sample pairs) and a `.py` matplotlib script, and launches Python with the SAME system(`python ...`) mechanism as `ft_pyplot()`, honouring every existing `pyplot_*` setting: `pyplot_terminal=png/svg/pdf` renders headless (Agg) and runs synchronously, otherwise an interactive window is launched in the background; `pyplot_python` picks the interpreter, `pyplot_backend` forces a matplotlib backend, `pyplot_figsize` sets the figure size, and `pyplot_style` applies a style sheet (a dark sheet flips the annotation colours to stay legible on a dark ground). No numerical-core changes; the eye math is entirely E-207's, and `ft_pyplot()` and the `gnuplot/asciiplot` back-ends are untouched. Verify (examples/pyplot): a self-contained data eye -- a pseudo-random NRZ bit stream (PWL) through a bandwidth-limiting RC channel ($\tau \sim 0.5$ UI) -- is rendered with

pyplot eyefig -eye v(rx) -ui 0.5n: the eye analysis runs and reports its metrics, eyefig.png is a valid PNG, the generated script is a hist2d persistence eye referencing the metrics, and the no-name form defaults the base to eye.png. 4 checks (17 total in the pyplot suite). Regression 170/170 (extends the pyplot suite, no new folder).

E-209 *src/spicelib/analysis/dcpss.c, src/frontend/com_hb.c, src/include/ngspice/cktdefs.h, examples/hb_examples/verify_hb.py*

hb publishes its spectrum as nutmeg vectors. The E-134 harmonic-balance command (hb) used to only PRINT a spectrum table: HBAalyze() freed the converged two-sided Fourier solution Vr/Vi as soon as the table was written, so the result could not be plot/print/wrdata-ed without scraping stdout (the amp2n2222 HB study had to parse the table in Python for every figure). Now, after a converged hb, a fresh current 'hb' plot holds a real scale hbfrequency (0, f0, 2f0, ..., Kf0; K+1 points) and one COMPLEX vector per node (named by the node, e.g. out, vcc, q1#branch) carrying the single-sided amplitude, so mag(out)/vp(out) reproduce the printed

E-210 *src/frontend/inp.c, src/spicelib/analysis/dcpss.c, examples/rfpss_examples/rc_pss.cir, examples/rfpss_examples/verify_rcpss.py*

PSS polish: .pss dot-card runs in batch + complex spectrum vectors. Two usability fixes to the E-117 shooting-method periodic steady state (pss), in the spirit of E-209 (hb). (1) The pss COMMAND (inside .control) worked, but a top-level .pss card by itself silently did nothing in batch mode -- ngspice -b deck.cir with just .pss ... reported "no simulations run", unlike .tran/.ac/.hb which auto-run; a deck had to wrap it in .control ... runendc. .pss is now dispatched to the pss command the same way E-146 (.sweep) and E-162 (.hb) handle their cards: during deck loading (frontend/inp.c) a top-level .pss ... line is stripped of its leading '.' and appended to the post-parse control list, so it executes as pss ... once the circuit is built. A boundary check (next char is whitespace or end-of-line) keeps it from matching a longer name, and it never matches .psp (periodic S-parameters, a distinct card). Decks that already wrap .pss in .control ... run keep working (one PSS run, unchanged output). (2) The pss analysis already published two plots -- pss1 (time-domain periodic waveform) and pss2 (frequency-domain spectrum) -- but pss2 stored MAGNITUDE ONLY (IF_REAL): the DFT computed each harmonic's phase (pssphases) and then discarded it, so vp(out) was unavailable and print out gave a bare magnitude, inconsistent with AC analysis (whose node vectors are complex) and with the hb command's E-209 vectors. pss2 now publishes COMPLEX node vectors (IF_COMPLEX): each harmonic is stored as mag * e^{j*phase} (the DFT returns the phase in degrees; the DC term is kept real), built into IFcomplex rows and emitted with OUTpData (as PAC/PXF/PSP already do) instead of the real-only CKTdumpp. mag(out) reproduces the old magnitude spectrum exactly and vp(out) now recovers the phase; the max-frequency-verification logic still reads the retained magnitude array (pssResults), so it is unchanged, and the time-domain pss1 plot is unchanged (a real waveform). BACKWARD COMPATIBILITY: because pss2 node vectors are now complex, print out / wrdata out emit re,im columns (exactly as AC analysis does) rather than a bare magnitude -- decks that parsed the magnitude directly should use mag(out); the bundled rc_pss.cir was migrated from print b to print mag(b). This is the same magnitude-vs-complex convention the rest of ngspice's frequency-domain analyses use. The shooting-PSS numerical core, the retained periodic operating point, and every analysis built on it (PAC/Pnoise/PXF/PSP) are untouched. Verify (examples/rfpss): a top-level .pss card (no .control) auto-runs in batch and reaches convergence; the frequency-domain spectrum resolves as complex -- both vm(out) (magnitude) and vp(out) (phase) return, with a driven diode rectifier's harmonics showing the expected non-trivial phase; the existing PSS/PAC/PXF/solver-parity checks pass unchanged. 2 checks (27 total in the rcpss suite). Regression 170/170 (extends the rfpss suite, no new folder).

E-211 *src/spicelib/analysis/dcop.c, src/frontend/inpcom.c, examples/codeanalysis_examples/**

Code-analysis bug fixes: XSPICE DC-op else + LTSPICE table free. A static-analysis pass over the ngspice tree (clang's analyzer run on every project-modified file, each finding then confirmed

by reading the code, false positives -- e.g. symbolic-size array indexing in static math helpers -- discarded) surfaced two real defects. BUG 1 (dcp.c): DCop() chooses between the event-driven XSPICE solver (EVTOP) and the analog solver (CKTOP) with a BRACELESS else whose body was originally converged = CKTOP(...), so CKTOP ran only when there were no event-driven instances. Enhancement-188 (DC warm-start) inserted its preload code between the else and CKTOP, so the else then covered only its first statement (wsiz = SMPmatSize(...)) and CKTOP ran UNCONDITIONALLY: for a deck with event-driven code-model instances (A-devices) BOTH EVTOP and CKTOP ran, CKTOP re-solving the analog part and overwriting the event-driven DC result, and wsiz was read UNINITIALISED at the warm-start guard (line ~110) and snapshot (line ~129) on the EVTOP path (undefined behaviour, masked today only because DC warm-start is off by default). FIX: hoist wsiz = SMPmatSize(ckt->CKTmatrix) above the branch (always initialised; the warm-start snapshot reads it on both paths) and add braces so the else covers the warm-start preload + CKTOP, restoring the original exactly-one-solve semantics. The #ifdef XSPICE braces vanish together when XSPICE is not compiled, so the non-XSPICE build is byte-unchanged. BUG 2 (inpcom.c): inp_compat() converts an LTSPICE E ... table=(...) controlled source using an UNINITIALISED char *ckt_array[100] stack array; the LTSPICE branch sets only ckt_array[1], and a single-pair table (table=(x0,y0), ipairs==1) then tfree(ckt_array[2]) -- a free of a garbage stack pointer (ckt_array[2] is written only on the other, non-LTSPICE, branch). FIX: zero-initialise the array (char *ckt_array[100] = { NULL });; ngspice's tfree(NULL) is a safe no-op (txfree guards if (ptr) free(...)), so a tfree of any slot not written on a given branch does nothing instead of freeing garbage. No functional change to correct analog decks -- the fixes remove undefined behaviour and a redundant analog solve on the XSPICE and LTSPICE-import paths respectively. Verify (examples/codeanalysis, 5 checks x both solvers): a mixed-signal DC op (an analog divider + an adc_bridge, an event-driven instance) solves the analog node to v(a)=0.5 without error; a single-pair LTSPICE E ... table=(0.5, 3) imports and yields the constant 3 V; a plain analog DC op is unchanged (a would-fail before/after test is impractical -- the redundant CKTOP leaves a simple circuit's DC point unchanged, and freeing a garbage pointer is UB that may or may not crash). 5 checks. Regression 171/171 (new example folder).

E-212 *src/frontend/breakp.c, src/frontend/device.c, src/frontend/vectors.c, src/frontend/dotcards.c, src/frontend/inpcom.c, src/spicelib/analysis/cktpzset.c, examples/crashfix_examples/**

Crash hardening: seven user-triggerable crashes. A static-analysis pass (clang analyzer) plus argument fuzzing of every frontend command and of malformed/degenerate netlists surfaced seven reproducible, user-triggerable crashes (SIGSEGV/SIGABRT) in STOCK ngspice, all in user-input handling (command parsers, the .op output path, netlist recursion) and none in the numerical core (device models, OSDI loader, Sparse 1.3, KLU, and the analysis drivers were audited and found clean). Each is a missing edge-guard. (1) iplot (breakp.c, com_iplot): iplot -w/-d with the flag as the trailing token consumes its value word (wl=wl->wl_next), guards it with if (wl) (proving wl can be NULL), then the loop-bottom wl=wl->wl_next dereferences NULL. FIX: guard the advance. (2) altermod (device.c, com_altermod->com_alter_common): com_alter guards if (!wl) but com_altermod does not, so a bare altermod reaches com_alter_common with an empty list -> wlin=wl_head=NULL -> wl_nthelem(100,NULL)=NULL -> eq(wlin->wl_word, "]") derefs NULL. FIX: guard the empty wl_head with a clean error. (3) measure (vectors.c, vec_get): a malformed measure statement (meas tran x FIND, ... WHEN, FIND v(1) with no WHEN/AT, across tran/dc/ac/sp) leaves meas->m_vec NULL, which flows into com_measure_when -> vec_get(NULL) -> copy(NULL)=NULL -> strchr(NULL, '.'). FIX: make vec_get NULL-safe (if (!vec_name) return NULL) -- the common sink; every caller already handles a NULL return, so each trigger reports "no such vector" instead of crashing. (4) empty .op (dotcards.c, ft_cktcoms): a circuit with no non-ground nodes -- e.g. a deck whose every component is commented out -- produces an op plot with no data vectors, so pl_dvecs is NULL; the former assert(pl_dvecs != NULL) aborted (SIGABRT; ngspice's autotools build defines no NDEBDEBUG so asserts are live) and with -DNDEBDEBUG the next line dereferenced NULL. FIX: guard if (plot_cur && plot_cur->pl_dvecs) and skip the empty op printout (the solver had handled the empty matrix correctly; this

was purely output formatting). (5) KLU pole-zero (cktpzset.c): a bsearch miss for the drive pointer was reported with a printf and then dereferenced anyway (job->PZdrive_pptr = matched->CSC_Complex); latent (SMPmakeElt creates the element just above, so it is always in the bind table and the miss cannot occur today) but the guard was objectively broken. FIX: add the missing else/NULL path, matching the correct idiom in cktsetup.c. (6) .include recursion (inpcom.c, inp_read): inp_read tracked call_depth but never limited it, so a file that .includes itself -- directly, via an A->B->A cycle, or via a .lib section that .includes its own file -- recursed until the C stack overflowed (SIGSEGV); .lib files are de-duplicated by find_lib and read once, so the crash is always via .include. FIX: cap at INP_MAX_INCLUDE_DEPTH=50; a cycle now reports an error (a 49-level chain still reads fine). (7) .func recursion (inpcom.c, inp_expand_macro_in_str): the expander re-expands a substituted function body with no depth guard, so .func f(x)={f(x)} (or a mutual f<->g) expanded forever and overflowed the stack (each frame also holds a params[FCN_PARAMS] array ~8KB, so it fails near ~1000 deep). FIX: a macro_depth counter capped at INP_MAX_MACRO_DEPTH=100, reset on every error path. The netlist-recursion caps mirror the include/expression-depth hardening OpenVAF-r received in Enhancement-148. No functional change to correct decks -- the fixes turn crashes on malformed or degenerate input into graceful errors. Verify (examples/crashfix, 17 checks x both solvers): every repro now exits gracefully (no signal -- crashes are detected from the child's SIGSEGV/SIGABRT exit code, which the pre-fix binary returned) while every valid form still works (iplot -w 3 v(1); altermod nm vto=0.7; meas tran vmax MAX v(1) -> 1.0; a normal .op; KLU .pz; a legitimate nested .include; .func sq(x)={x*x} -> sq(4)=16). 17 checks. Regression 172/172 (new example folder).

E-215 *src/main.c, src/osdi/osdiload.c, src/include/ngspice/osdiitf.h, examples/plusargs_examples/**

command-line plusargs (+name[=value]) served to Verilog-A models. Lets ngspice -b deck.cir +corner=ff pick a corner or toggle a feature flag at run time without editing the deck -- the simulator half of E-215 (the compiler half is in the openvaf report). main.c treats a +-prefixed command-line argument as a plusarg rather than an input file (both the batch and interactive file loops) and calls the new ngspice_register_plusarg (osdiitf.h). get_simparams (osdiload.c) exposes each plusarg on the OsdSimParas channel a compiled model already reads for \$simparam\$str (E-25), as namespaced simparams: numeric \$test\$plusargs\$name=1 (present), \$valset\$plusargs\$name=1 (given as name=value), \$valnum\$plusargs\$name (value via strtod), and string \$value\$plusargs\$name (value text). The base sim_params arrays are spliced onto the extended arrays only the FIRST time get_simparams sees any plusarg present, so a run with NO plusargs allocates nothing and is byte-for-byte unchanged; the extended arrays are built once (plusargs are constant for the run) and only the base numeric values + analysis_name are refreshed per call. No OSDI ABI change (the string/number channel already existed). Verify (examples/plusargs, 12 checks both solvers): a conductance set from +boost/+gain=25/+scale=2.5/+corner=ff|ss plus unmatched/unrelated/no-value edge cases. Regression 175/175 (new example folder).

E-216 *src/frontend/com_optimize.c, examples/pareto_examples/**

NSGA-II multi-objective / Pareto optimization. Adds a -method nsga2 to the optimize command that trades TWO OR MORE competing objectives and returns a Pareto FRONT of non-dominated designs, rather than the single optimum the scalar methods return (E-130 Nelder-Mead, E-194 PSO, E-195 DE, E-196 SA) -- the multi-objective generalization of E-206 design centering. Objectives are given as repeated -minimize / -maximize ; the existing optimize infrastructure is reused wholesale (the -param/-dparam/-mparam knobs, the -analysis-driven expression objectives via opt_eval_expr, the normalized [0,1] search space and the opt_rand RNG). New: a vector-valued objective evaluator opt_eval_objs (maximized objectives negated to a common minimization convention) and nsga2() implementing standard NSGA-II (Deb 2002) -- fast non-dominated sorting (Pareto rank via domination counts), crowding distance (per-front normalized neighbour gaps, boundaries at infinity), binary-tournament selection by the crowded-comparison operator, real-coded SBX crossover (eta=15) + polynomial mutation (eta=20, rate 1/n), and elitist

parent+offspring (2N->N) survivor selection so a non-dominated design is never lost. The command prints the final front (each design's objectives + parameters, sorted by the first objective) and publishes each objective column as a permanent vector `pareto1/pareto2/...` so the front plots directly (plot `pareto2` vs `pareto1`). `nsga2` needs ≥ 2 objectives and a single -analysis stage; population defaults to $20+4*np$ (override `-swarmsize`), generations `-maxiter`, `-seed` for reproducibility. The scalar methods and their single-objective/least-squares/centering paths are unchanged (a lone `-minimize` still seeds the scalar objective for `nm/pso/de/sa`). Verify (examples/pareto): the Schaffer-1 benchmark -- minimize $f_1=px^2$ and $f_2=(px-2)^2$ over px in $[-1,3]$, whose Pareto set is exactly px in $[0,2]$ tracing $f_2=(\sqrt{f_1}-2)^2$, with the knob range wider than the front so NSGA-II must discover $[0,2]$. The returned front is a reasonable size, every point is genuinely non-dominated, it spans the trade-off (px reaches ~ 0 and ~ 2), and the published vectors lie on the analytic curve to machine precision ($8.9e-16$). 6 checks. Regression 176/176 (new example folder).

E-217 *src/frontend/com_pyplot.c, src/frontend/plotting/pyplot.c, src/frontend/plotting/pyplot.h, src/frontend/plotting/plotit.c, examples/pyplohist_examples/**

`pyplot -hist`: histogram plots. `pyplot -hist <sig> ...` renders each listed signal's VALUE distribution as a matplotlib histogram (`plt.hist`) instead of a trace vs time/frequency -- for the amplitude spread of a transient, a Monte-Carlo measurement, or a noise sample. Unlike `-eye` (E-208, a distinct analysis) `-hist` is a render MODE, so it is small and reuses the whole `pyplot` pipeline: `com_pyplot` detects the `-hist` marker anywhere in the arg list, unlinks it, and dispatches the remaining [name] signal-list to `plotit` with a distinct device string `pyplohist`; `plotit` resolves the signal expressions exactly as for a normal plot (`v(out)`, `db(...)`, vector arithmetic, node names) and passes the resolved vectors to `ft_pyplot` with a `hist` flag; `ft_pyplot` writes the same .data table and emits `ax.hist(value_column, bins, alpha, label)` per signal instead of `ax.plot(x,y)`. All existing `pyplot` facilities carry over unchanged (set `pyplot_terminal=png/svg/pdf` headless render, `pyplot_subplots` panels, style sheets, `figsize`, `backend`); two histogram-specific options are added -- `pyplot_hist_bins` (bin count, default matplotlib 'auto') and `pyplot_hist_density` (normalize to a density). Correctness details: the x-axis is the signal value and the y-axis the count (labels set accordingly); histogram panels do NOT share an x-axis (`sharex=False` -- different signals have unrelated value ranges, vs a line plot's shared time/freq axis); the FULL value length is histogrammed -- the data-table loop normally walks the shared scale vector, but a histogram of a raw `let` vector (e.g. a Monte-Carlo result) can have a scale whose length differs from the values', so for `hist` it walks the longest VALUE vector and pads shorter ones with NaN (the generated script filters via `~np.isnan`), fixing a silent truncation to the scale length; overlaid histograms get `alpha=0.6` transparency and a legend while a single one is drawn opaque. No change to the ordinary `pyplot` trace path or to `pyplot -eye`. Verify (examples/pyplohist): two signals with closed-form distributions -- `ramp=i/(N-1)` uniform on $[0,1]$ (flat histogram), `sine=sin(2pi i/100)` arcsine on $[-1,1]$ (U-shaped, a sinusoid dwells near its peaks so the edges tower) -- checked analytically from the generated .data (uniform bins within 15% of the mean; sine edge bins » middle), plus the `-hist` path taken (`plt.hist` not `plt.plot`, non-shared x), the full 20000-sample length not truncated, and a valid non-trivial PNG. 6 checks. Regression 177/177 (new example folder).

E-218 *src/frontend/com_pyplot.c, src/frontend/plotting/pyplot.c, src/frontend/plotting/pyplot.h, src/frontend/plotting/plotit.c, examples/pyplotcontour_examples/**

`pyplot -contour`: 2-D contour maps for 2-D parameter sweeps. `pyplot -contour <z> <x> <y>` renders a filled matplotlib contour map of a quantity `z` over the `(x,y)` plane -- the natural view of a 2-D parameter sweep (a device current over a `(Vgs,Vds)` grid, a gain over an `(R,C)` grid, a noise figure over a `(freq,bias)` grid). Like `-hist` (E-217) and unlike `-eye` (E-208, a distinct analysis), `-contour` is a render MODE dispatched over the ordinary signal list: `com_pyplot` detects the `-contour` marker anywhere in the arg list, unlinks it, and dispatches the remaining [name] `z x y` to `plotit` with a distinct device string `pyplotcontour`; `plotit` resolves the three expressions exactly as for a normal plot (`v(out)`, `db(...)`, vector arithmetic, node names) into a three-vector list (`z` first, then `x`, then

y); the new `ft_pyplot_contour` writes an (x,y,z) data table and a matplotlib script that TRIANGULATES the (x,y) points (`ax.tricontourf`). Because the plane is triangulated rather than reshaped into a rectangular mesh, gridded OR scattered sweep data plots identically -- no grid-dimension metadata (nx,ny) is needed, and a sweep with a ragged tail still works; x/y/z come from any 2-D sweep leaving three equal-length vectors (nested .dc, the sweep command family E-146/E-190, a .step grid, or .control vectors). All existing pyplot facilities carry over (set `pyplot_terminal=png/svg/pdf` headless render, style sheets, figsize, backend); three contour-specific options are added -- `pyplot_contour_levels` (level count, default matplotlib-auto), `pyplot_contour_lines` (overlay labelled black contour lines on the filled map), `pyplot_contour_cmap` (colormap, default viridis). The colorbar is labelled with z's name and the axes with x's and y's. Robustness: -contour needs exactly three vectors, else it is rejected with a clear message (pyplot -contour needs exactly three vectors) rather than a confused plot; a degenerate 1-D (collinear) sweep cannot be triangulated, so the generated script catches the qhull/matplotlib failure and exits with a helpful line (a contour needs a genuine 2-D sweep -- the points must not be collinear) instead of a bare traceback. No change to the ordinary pyplot trace path, pyplot -hist, or pyplot -eye. Also removed a stale unused-variable warning (ft_pyplot's scale) left by E-217. Verify (examples/pyplotcontour): a closed-form surface $z=x^2+y^2$ over x,y in [-2,2] (a paraboloid whose contours are concentric circles, z=0 at the centre and z=8 at the corners) -- checked analytically from the generated .data: the -contour path is taken (tricontourf not plot/hist, a colorbar labelled z, axes x/y), the data table has three columns (x,y,z) for all 1681 rows, the column mapping is correct (z reconstructs x^2+y^2 to 9e-16), the sweep is genuinely 2-D (x and y each span the full range, z runs from ~0 at the centre to ~8 at a corner), and a valid PNG is rendered. A second demo (bridge_dc_demo.cir) exercises the feature on genuine simulation output: a diode-OR bridge whose output V(c) follows whichever of its two inputs is higher, swept with a nested .dc (v1 drives node a, v2 drives node b, so V(a)/V(b) are the two swept values at each point); pyplot -contour v(c) v(a) v(b) maps the output over the (V(a),V(b)) plane -- a max-like corner surface. It checks the nested .dc yields the flattened 51x51 = 2601-row 3-column grid, the axes are the two real swept sources (each spanning [-1,1]), and the output is ~0 when both inputs are low and rises when either input is high (the diode-OR signature, which confirms the columns map onto a real circuit result). 19 checks. Regression 178/178 (new example folder).

E-221 *src/frontend/inpcom.c, examples/busnodes_examples/**

array/bus node ranges in the netlist. ngspice has no bus/array node syntax: a node is an opaque string, and a range like `a[0:1]` is read as a single literal node name (so `R1 a[0:1] r=2k` gives R1 only one node token and misparses `r=2k`). E-221 adds a netlist preprocessing pass that expands a range token `base[lo:hi]` into the scalar node sequence `base[lo] base[lo±1] ... base[hi]`, so buses can be written compactly: `R1 a[0:1] r=2k -> R1 a[0] a[1] r=2k`; `X1 n[0:3] mysub -> X1 n[0] n[1] n[2] n[3] mysub`; `.subckt mysub p[3:0] -> .subckt mysub p[3] p[2] p[1] p[0]`. The range supplies node tokens POSITIONALLY -- the 'two terminals' reading: `R1 a[0:1]` is a single resistor from `a[0]` to `a[1]`, not an array of resistors; multi-terminal instances, subcircuit calls and `.subckt` port lists expand the same way. Semantics: `base[lo:hi] -> base[lo] base[lo±1] ... base[hi]`, DESCENDING when `lo>hi` (Verilog convention `d[3:0] -> d[3] d[2] d[1] d[0]`); `a[0:0]` is a single element; the scalar names use the same bracket form so a bus `a[0:1]` and an explicit `a[0]` denote the SAME node (the two forms interoperate); only a WHOLE whitespace-delimited token that is exactly `base[int:int]` is expanded, so already-scalar names (`a[0]`), non-integer ranges (`a[x:y]`), malformed ones (`a[1:2:3]`), device values, model names and `XSPICE %vd[...]` port groups are left untouched; `.control ... .endc` blocks are skipped and only element lines and `.subckt` lines are processed (`.model/.param/other` dot cards ignored); a range wider than 8192 elements is left literal (the device parser then reports it) so a typo like `a[0:1000000000]` cannot exhaust memory. Implementation: one pass `inp_expand_buses` in `frontend/inpcom.c`, run in `inp_readall` right after whitespace normalization and BEFORE subcircuit expansion and device parsing -- so every downstream consumer (`numparam`, `.subckt` expansion, the device parser, OSDI) sees the already-scalar node list and needs no changes; continuation lines are stitched earlier in `inp_read`, so each logical line is processed whole; the token

test parses base (a plain node name) then `[:]` filling the rest of the token and emits the element list with `DSTRING`, a non-match copying the token verbatim. Verify (examples/busnodes): 6 checks, both solvers, all from node voltages/branch currents (solver-independent) -- `R1 a[0:1] r=2k` connects `a[0]-a[1]` so $I(R1)=(1-3)/2k=-1mA$; the scalar `R2 a[0]` addresses the same node as the bus element ($I=+0.25mA$, only correct if `a[0]` is one shared node); a subcircuit with a descending port bus `p[1:0]`, called with `X1 n[0:1]`, connects positionally ($I(Rs)=-5mA$); and non-integer (`b[x:y]`), malformed (`c[1:2:3]`) and scalar tokens are left literal while a `.control let w=a[0:1]` is not rewritten. One file; no device/solver/OSDI change; a deck with no `base[lo:hi]` token is untouched (a fast path skips any line without `[`). Regression 181/181 (new example folder).

E-222 *src/frontend/inpcom.c, src/frontend/subckt.c, src/misc/string.c, examples/parserfuzz_examples/**

ngspice netlist-parser hardening (fuzzing): seven crashes/hangs -> clean errors. Fuzzing the netlist parser -- mutating real decks (byte flips, truncation, line duplication, delimiter/bracket/keyword/directive/number injection, and the new E-221 bus-range tokens) and running each under `ngspice -b` -- found seven ways to CRASH (SIGSEGV/SIGABRT, real memory-safety bugs since ngspice is C) or HANG on malformed input; a 12000-iteration campaign dropped the crash+hang rate from 26 to 1 (96%, all hangs eliminated). ROOT CAUSES, all fixed: (1) `misc/string.c` `gettok_instance()` returns an empty token WITHOUT advancing `*s` when it sits on `'(` or `)'`, and `frontend/inpcom.c` `get_number_terminals()`'s OSDI (`'n'`) case was the only multi-token case with no iteration cap -- a line starting with `'n'` containing `'(` (e.g. `'nan.func f(x)=...'`, `nan` lexing as a number) spun forever `tmalloc`-ing each pass; `gettok_instance()` now always advances (consumes the bracket as a one-char token) and the `'n'` loop gained the same cap as its siblings. (2) `inp_expand_macro_in_str()`: a truncated `'.model m'` (no model type) ran `nexttok` off the end, leaving the scan pointer NULL, then `strchr(NULL)`; NULL guards added on the `.model` skip and the scan loop. (3) `frontend/subckt.c` `doit()`: `MAXNEST` bounds subcircuit DEPTH, but a subckt that instantiates itself with branching factor ≥ 2 blows up to 2^{MAXNEST} instances (an effective hang) before the depth cap trips -- a total-instantiation cap now catches recursive subcircuits. (4) `doit()`: a bare `'X'` invocation (nothing after the `refdes`) makes the `find-last-token` walk run off the FRONT of the line buffer -- an out-of-bounds read ending in `strcmp(NULL)`; such a line is now skipped. (5) `inp_modify_exp()`: two copies into a fixed `buf[512]` with no bound -- an unterminated `'v'` (no closing `)'`) and a long identifier token (a run of `'[`, which is an accepted char) overran the STACK buffer (stack smashing); both loops are now bounded by the buffer size. (6) `doit()`: a malformed `'.subckt'` can register a NULL name; the invoked-name match `eq(su_name,...)=strcmp()` then `deref NULL`; guarded. (7) `translate()`: `gettok_noparens()` returns NULL at the end of a malformed controlled-source line, then `strcmp(NULL,'POLY')`; guarded. Method: the fuzzer classifies OK/clean-error/CRASH/HANG over self-contained seed decks exercising the parser surface (devices, `.subckt`, `.param`, `.model`, `.control`, expressions, continuations, buses); each crash was triaged to its faulting instruction with `lldb` (conditional `strchr/strcmp`-with-NULL breakpoints), sample for the hangs, and Guard Malloc (`libgmalloc`) to pin the stack overrun, then read at the source; fixes applied one at a time, verified against the recorded crash corpus, and a fresh 12000-iteration run confirmed convergence. No valid deck changes behaviour. Verify (examples/parserfuzz): 10 checks -- a minimal deck per root cause (`nan.func`, `.model` in `.control`, a self-recursive subckt, a bare `X` line, an unterminated `v`, a 4000-`'[` identifier, 4000 `'{'` in a subckt E-source) now yields a clean/bounded outcome (no signal/abort, no hang), and three valid decks (subckt hierarchy with an E-source and a parameter, a `v(x)*v(x)` B-source expression, a POLY source in a subckt) compute `V` to the analytic value. The recorded 27-input crash corpus is all clean. RESIDUAL: one fuzz input in 12000 still crashes in XSPICE (`xspice/mif/mifgetmod.c` `MIFgetMod`) -- an `'a'`-device whose model name matches a non-code-model (e.g. a diode `.model`) is processed as a code model, reading device->modelParms as the wrong type; a code-model type-confusion in the XSPICE subsystem, out of scope for this netlist-parser pass and flagged as a follow-up. Three files (`frontend/inpcom.c`, `frontend/subckt.c`, `misc/string.c`); no device/solver/OSDI change. Regression 181/181 (new example folder).

E-223 *src/xspice/mif/mifgetmod.c, examples/xspicemodel_examples/**

XSPICE a-device model-type validation (MIFgetMod): fixes the one residual crash left by the E-222 netlist-parser fuzzing pass. An a device is an XSPICE code-model instance; the last token on its card is a model name, which MIF_INP2A (xspice/mif/mif_inp2.c) hands to MIFgetMod (xspice/mif/mifgetmod.c). MIFgetMod looks the model up in the pass-1 table and, if it has not been instantiated yet, processes its .model card AS a code model -- it casts INPmodfast to MIFmodel, reads the device's XSPICE DEVpublic fields (param, num_param), and matches the .model parameters through the MIF parameter path -- with NOTHING checking that the resolved model is actually a code model. If an a-device's model name happens to match a NON-code-model .model (a diode, a BJT, a resistor model), MIFgetMod walked all of that machinery over a structure that is not a MIFmodel and over DEVpublic fields that a built-in SPICE device leaves unset, reading unrelated memory as the wrong type -> SIGSEGV. The crash bit hardest when the non-code-model's parameter keywords COLLIDE with the a-device's: a diode .model has is/n, so the parameter loop found a keyword match and dereferenced the type-confused setModelParm path. The fuzz-found trigger (the E-222 residual) was exactly this, reached through an E-221 bus-expanded a-device inside a subcircuit (a[0:0] D1 a b dm with .model dm d(is=1e-14 n=1)). FIX: a code model is exactly a device that carries a code-model evaluation function -- cmpp emits .cm_func = <fn> for every code model (xspice/cmpp/writ_ifs.c) and it is the pointer XSPICE calls to evaluate the model (mifload.c/evtload.c), while every built-in SPICE device sets .cm_func = NULL. MIFgetMod now validates DEVICES[INPmodType]->DEVpublic.cm_func != NULL (and INPmodType < DEVmaxnum) BEFORE the MIFmodel cast and the parameter loop, returning a clean 'model X is not a code model; an a device requires an XSPICE code-model .model' error otherwise. The guard sits before the cast, so it catches the model whether or not an ordinary device of the same name already instantiated it. A NULL guard on the inner strcmp was rejected as whack-a-mole (the crash merely shifts to the next type-confused access); the model TYPE is the thing to validate. Legitimate code models are unaffected (cm_func != NULL). Scope: XSPICE only, one file (xspice/mif/mifgetmod.c); no netlist-parser, device, solver, or OSDI change. Verify (examples/xspicemodel): 8 checks, both solvers -- six pathological decks (a-device -> diode with and without colliding params, BJT, resistor model, the bus-expanded-in-a-subckt fuzz repro, and an undefined model) now yield a clean bounded error (no signal/abort, no hang), and a legitimate adc_bridge code-model a-device still binds and simulates (analog divider node v(a)=0.5). Regression 182/182 (new example folder).

[E-224](#) *src/frontend/parse.c, src/frontend/evaluate.c, examples/arraynodeprint_examples/**

Reference array/bus node voltages in **print/plot/let**. Enhancement-221 added array/bus node ranges to the netlist (R1 a[0:1] -> R1 a[0] a[1]), naming the nodes with LITERAL brackets so the node voltage vector is stored as a[0]. But the frontend expression parser uses [...] as its vector-INDEX operator, so in print/plot/let the token a[0] was parsed as 'element 0 of a vector named a'. There is no vector a, so print v(a[0]) and print a[0] silently failed (with a 'vector a is not available' warning) even though the node existed and solved correctly (it shows up in the .op node table and display). FIX: a literal-node fallback -- when the base name of an index expression is an UNRESOLVED name (a zero-length placeholder dvec) and the index is a constant and a node/vector named base[index] exists, use that node instead of indexing. It only fires when the base does not resolve, so ordinary vector indexing (realvec[3], where the base IS a real vector with non-zero length) is unaffected. Three frontend spots, because a[0] reaches the machinery three ways: (1) frontend/parse.c checkvalid() -- the pre-evaluation validity check rejected the whole term because base a is a zero-length placeholder; for an INDX node over an unresolved base it now accepts the term when the literal node base[index] exists. (2) frontend/evaluate.c op_ind() -- evaluates the bare a[0] form; it now resolves the literal node name BEFORE evaluating the bare base, so no spurious 'no such vector a' is emitted. (3) frontend/evaluate.c apply_func() -- the v() node-access special case needed arg->pn_value, which an INDX op node lacks (it errored 'bad v() syntax'); it now reconstructs the literal node name from the INDX argument and resolves it. RESULT: print a[0], print v(a[0]), array nodes inside expressions (v(a[0])-v(a[1])), and plot/let all work; the recovered vector is a normal waveform (indexable, aggregatable) in

transient. Branch current through the bus resistor stays the device parameter `@r1[i]` (ngspice has no current-between-two-nodes syntax; `i()` takes a source name). Scope: frontend only, two files (`parse.c`, `evaluate.c`); no device/solver/OSDI change. Verify (examples/arraynodeprint): 8 checks, both solvers -- `print a[0]/v(a[0])/v(a[1])`, the expression `v(a[0])-v(a[1])`, and `@r1[i]` read the correct DC values; ordinary indexing `unitvec(4)[2]` is unchanged; a genuinely-missing node `v(a[9])` stays a clean miss (no value, no crash). Regression 183/183 (new example folder).

E-225 *src/frontend/evaluate.c, src/frontend/com_measure2.c, src/frontend/fourier.c, src/math/cmaths/cmath4.c, examples/cmdfuzz_examples/**

Command/expression-evaluator crash hardening (fuzzing). With the netlist parser hardened (E-212/E-222/E-223), the fuzzer was turned on a previously-unfuzzed surface: the interactive `.control` command interpreter and the vector-expression evaluator (`print/let/plot/meas/fft/fourier` and the operators `?:`, indexing `[ ]`, ranges `[ : ]`, arithmetic). Mutating command lines/expressions over a running circuit and running each under ngspice `-b` found memory-safety crashes (SIGSEGV / SIGABRT, real since ngspice is C); 62 crashes in the first 2500 iterations, driven to 0 by FIVE fixes in three groups. GROUP A -- degenerate-input heap overflows in the math transforms (`fft/deriv/fourier` each crashed on a too-short (< 2-point) or synthetic vector like `fft(1)`, `deriv(vecmin(v(1)))` (a plot-derived scalar), `fourier 1k deriv(vecmin(v(1)))`; the malloc abort fired on a LATER allocation, so they looked non-deterministic): (1) `maths/cmaths/cmath4.c cx_fft()` -- the Green's FFT (`rffts`) derefs bit-reversal tables `BRLowArray[(M-1)/2]` that `fftInit()` only allocates for `M>2`, so an input of `<=4` points read unallocated memory; and `time/xscale` were sized by the data length while the scale loops fill `pl_scale->v_length` entries and the scale vector uses `fpts` points, so a vector shorter than the plot scale overran both -- pad to `N>=8`, size buffers for the largest fill, reject `length<2/scale-too-short`. (2) `maths/cmaths/cmath4.c cx_deriv()` -- the polynomial-fit derivative reads `degree+1` points per fit; a `length-1` input ran the fit/edge loops over too few points and overran the heap -- for `length<2` return a zero vector (derivative of a single point is undefined). (3) `frontend/fourier.c` -- a degenerate input gives a zero time span and overran the interpolation grid -- require `>=2` data points. GROUP B -- `frontend/evaluate.c ft_ternary()`: the `?:` operator evaluated its condition and selected branch with no NULL check, so when either failed to evaluate (`1?0[3]:9`, where `0[3]` indexing-a-scalar returns NULL) it hit `vec_copy(NULL) / cond->v_link2` and crashed -- the dominant crash; fixed with NULL guards (an audit confirmed `op_ind`, `op_range`, `binary-op` and `apply_func` already guard NULL). GROUP C -- `frontend/com_measure2.c`: measure error messages were formatted into a shared `char errbuf[100]` (passed by pointer to the `measure_parse_*` helpers); a failed measure writes its ENTIRE expression string as the vector name (`sprintf(errbuf, "no such vector as '%s'", meas->m_vec)`), so a measure expression longer than ~80 chars (trivially `meas tran m MAX (v(1)+v(1)+...)`) overran it -> SIGABRT (content irrelevant, only length, which is why it resisted reduction) -- fixed with a file-scope `MEAS_ERRBUF_SIZE` and a bounded `snprintf` at every write (28 sites). Note: `dowhile 1 ... end` (true constant condition, no break) is an intentional infinite loop, flagged as a HANG but correct. Scope: frontend/math only, four files (`evaluate.c`, `com_measure2.c`, `fourier.c`, `cmath4.c`); no device/solver/OSDI change. Verify (examples/cmdfuzz): 13 checks -- a minimal repro per root cause now yields a clean bounded outcome, and regressions confirm `fft(v(1))/deriv(v(1))`, a valid `?:`, and ordinary indexing still work. Regression 184/184 (new example folder).

E-226 *src/frontend/rawfile.c, examples/rawfuzz_examples/**

Rawfile-load crash hardening (fuzzing). Continuing the fuzzing campaign onto another untrusted-input surface: the nutmeg rawfile reader (`raw_read` in `frontend/rawfile.c`, reached by the `load` command). A `.raw` file is external data (often written by another tool), so its parser must survive malformed input. Mutating a valid `.raw` (produced by `write`) -- tweaking the header counts, corrupting the `Flags:/Variables:` lines, adding/removing value rows, truncating, flipping bytes -- and `load-ing` each found a NULL-dereference crash (SIGSEGV, `memmove` into

0x0). ROOT CAUSE: raw_read resets flags = VF_PERMANENT at the start of each plot and sets VF_REAL / VF_COMPLEX only when it parses a Flags: line; each vector is then allocated with dvec_alloc(NULL, SV_NOTYPE, flags, npoints, NULL), which allocates v_realdata OR v_compdata according to those bits. A rawfile with NO valid Flags: line -- the line missing entirely, or carrying only unknown flags (Flags: xyz) -- leaves flags with neither VF_REAL nor VF_COMPLEX, so dvec_alloc allocates NO data array at all, and the value-reading loop does fread(&v->v_realdata[i], ..) / writes v->v_compdata[i] into a NULL buffer -> memmove into 0x0 -> SIGSEGV. (A bad No. Points: / No. Variables: count -- huge, negative, non-numeric -- was already handled cleanly; only the missing TYPE was fatal.) FIX: when the header is complete and vector allocation begins (at the Variables: line), if flags carries neither VF_REAL nor VF_COMPLEX, default to real (the common case) with a warning instead of allocating a typeless vector; the VF_PERMANENT reset is per-plot so the guard runs for every plot in a multi-plot file, and a valid file is unaffected. Also fuzzed in the same pass: the interactive .control command interpreter and shell layer (variable substitution \svar, alter/altermod/alterparam, analysis reruns, structured control if/while/foreach/repeat/dowhile) -- clean over 7500 iterations (already hardened by E-212/E-222/E-225). Scope: frontend only, one file (rawfile.c); no device/solver/OSDI change. Verify (examples/rawfuzz): 4 checks -- a valid .raw is written, then loaded with its Flags: line removed and with an unknown Flags: xyz, each now a clean bounded outcome (was SIGSEGV); a regression confirms a valid .raw still round-trips write -> load with the right length. Regression 185/185 (new example folder).

E-227 *src/frontend/snp2va.c, examples/snpfuzz_examples/**

Touchstone .snp parser crash hardening (fuzzing). Continuing the fuzzing campaign onto the Touchstone S-parameter parsers -- external network-data files (VNAs, EM solvers, other tools), so their parsers must survive malformed input. Two exist: rdsnp (the E-72 Touchstone-v1 reader) and pre_snp (the E-200/E-201 convert-to-Verilog-A path in frontend/snp2va.c that vector-fits the data and compiles it through openvaf-r). Mutating valid .s1p/.s2p/.s3p files -- corrupting the # option line, data values (NaN/Inf/overflow tokens, added/removed columns), duplicating/deleting rows, truncating, flipping bytes -- AND mismatching the port count encoded in the .snp filename extension, then running each through both readers: rdsnp was clean (4,500 iters); pre_snp hit a heap-corruption crash (SIGSEGV) 103 times, ALL on a filename with an absurd port count (e.g. .s2147483647p). ROOT CAUSE: snp2va.c infers the port count from the extension (N = atoi(dot+2)) with no upper bound; a huge N slips past the parse itself because the record stride rec = 1 + 2*N*N and the per-record inner-loop bound N*N use the same wrapped int (N*N overflows to a small/zero value, so the token layout stays self-consistent and no error is raised), but N is stored in out->N and used DOWNSTREAM to size the vector-fit allocations (malloc(nf*N*N*sizeof(cplx)), malloc(N*N*sizeof(cplx))) -- the overflowing/mismatched allocation and the writes that follow corrupt the heap (the program aborts later on an unrelated malloc, hence flaky). The brute-force fallback (used when the filename gives no usable count) already caps inference at 512 ports, so 512 is the parser's own sanity bound -- the filename path just wasn't held to it. FIX: hold the filename-derived count to the same limit -- if (N > 512) N = 0; falls back to inferring N from the data. A valid file is unaffected (its extension is far below the cap) and a genuinely mis-extended file now recovers N from the data layout instead of trusting an absurd filename. Scope: ngspice frontend only, one file (frontend/snp2va.c); no device/solver/OSDI change. Verify (examples/snpfuzz): 2 checks -- a valid .s2p copied to a .s2147483647p name loads via pre_snp without crashing (run repeatedly, was SIGSEGV), and a valid .s2p still converts through the full pipeline (pre_snp -> .va -> openvaf-r -> .osdi). Regression 186/186 (new example folder).

E-228 *src/osdi/osdiregistry.c, examples/osdifuzz_examples/**

OSDI .osdi loader descriptor-count hardening (fuzzing). A .osdi model is an external, compiled shared library that ngspice dlopens and reads descriptor metadata from (load_object_file, reached by the pre_osdi command). The loader strided/indexed the descriptor arrays by the counts the file declares -- OSDI_NUM_DESCRIPTORs (outer loop stride by *OSDI_DESCRIPTOR_

SIZE) and each descriptor's num_params (inner loop indexing param_opvar[param_id] and dereferencing its name[] pointers), num_opvars/num_terminals/num_noise_src/... (offset arithmetic) -- with NO bounds. A .osdi that still dlopen but declares a count larger than its real array (a truncated file, a byte-corrupted one, or one built by a compiler whose OSDI descriptor ABI differs -- the drift the loader's own comments warn about) read past the array -> SIGSEGV during pre_osdi (same unbounded-count class as E-227, here on a binary rather than a text surface). FIX: two generous ceilings in osdiregistry.c (#define OSDI_MAX_DESCRIPTOR 4096u, #define OSDI_MAX_DESC_COUNT (1u<<20)) far above any real compact model (the biggest ship a few dozen modules and ~1k params); after reading OSDI_NUM_DESCRIPTOR, reject the file (clean diagnostic, INVALID_OBJECT) if it exceeds the first, before it sizes the registry or strides the array; inside the descriptor loop, reject if any of a descriptor's count fields (num_nodes/num_terminals/num_params/num_instance_params/num_opvars/num_noise_src/num_jacobian_entries/num_collapsible/num_states) exceeds the second, before those fields index/size anything -- a single load-time choke point protecting every downstream consumer (setup, eval), since a rejected file never registers its devices. HONEST SCOPE: this targets the egregious/out-of-range count (the signature of a truncated/byte-corrupted/ABI-drifted object, where corruption reliably lands -- an ABI mismatch reads a pointer's bits as a count -> huge value); it is NOT a claim that the loader is proof against every malformed .osdi -- a count that lies within the plausible range is an internally inconsistent binary indistinguishable from a valid one, and a .osdi is native code whose mere loading executes it, so a hostile object is out of scope by construction. On macOS byte-corrupting a signed .osdi also breaks its code signature so dlopen rejects it before the loader runs; the residual exposure is chiefly Linux (no signing gate) and toolchain/ABI drift. Scope: ngspice only, one file (osdi/osdiregistry.c); no device/solver/compiler change. Verify (examples/osdifuzz): 3 checks -- two malformed .osdi built from tiny C stubs (one declaring a huge OSDI_NUM_DESCRIPTOR, one with a valid count but a huge per-descriptor num_params built from the real OsdiDescriptor struct in osdi.h) now load-reject cleanly (both a SIGSEGV rc=139 on the pre-fix binary), and a positive control confirms a real openvaf-r-built .osdi still loads and simulates (v(a)=1). Regression 187/187 (new example folder).

E-229 *src/frontend/com_dl.c, src/frontend/commands.c, src/spicelib/devices/dev.c, src/spicelib/devices/dev.h, examples/osdireload_examples/**

pre_osdi -f: reload a recompiled OSDI model without restarting. A Verilog-A developer's inner loop is edit .va -> recompile .osdi -> re-run, but interactively ngspice loaded each .osdi once: load_osdi (dev.c) is guarded by a path dedup, a dlopen-handle dedup (osdiregistry.c), and a device-name dedup (a type already in the global DEVices[] is ignored because the model-card lookup scans front-to-back and the first/stale entry wins -- see E-223); underneath, dlopen caches by path (a re-open returns the cached handle, never re-reading the recompiled file). So a recompiled model never took effect without restarting. FIX -- an opt-in -f/-force flag: com_osdi (com_dl.c) parses -f anywhere in the argument list and passes a force flag to load_osdi(path, force); on a forced reload of an already-loaded path it STAGES a byte-for-byte copy of the (recompiled) file under a fresh unique path and loads that (a new path is always re-read, sidestepping dlopen's cache; this avoids dlclose, which on macOS does not reliably unmap and where overwriting a still-mapped file faults the process), then removes the staged copy once it is mapped (the mapping keeps it alive on POSIX), and osdi_add_device(..., replace) swaps the registered device to the freshly loaded descriptor IN PLACE at its existing table index so no other device type's index shifts. The previous SPICEdev and the library mapping it points into are intentionally left resident: a circuit still built against the old model keeps a valid device, and freeing descriptor-owned memory would risk a double free (documented caveat: don't re-run a pre-existing circuit built against the prior model version after reloading it). Behaviour is unchanged without -f (plain re-load still skipped, now with a hint pointing at -f); the osdi command help gains the flag. Scope: frontend + device registry, four files (frontend/com_dl.c, frontend/commands.c help text, spicelib/devices/dev.c, spicelib/devices/dev.h); no solver/analysis/OSDI-ABI/compiler change. Verify (examples/osdireload): compiles a resistor model at 1k then 2k, drives ngspice in pipe mode to load the 1k ver-

sion ($i(v1)=-1\text{mA}$), recompiles the same file to 2k mid-session, confirms a plain re-load is skipped (still 1k), then `osdi -f` reloads it and the operating point changes to the 2k result ($i(v1)=-0.5\text{mA}$) -- one process, no restart. Regression 188/188 (new example folder).

E-231 *src/frontend/com_gnuplot.c, src/frontend/plotting/gnuplot.c, examples/csv_examples/**

CSV output for `wrdata (set wr_csv + wrdata -csv)`. `wrdata` only ever wrote whitespace-separated columns (each vector prefixed by its own scale column); the knobs `wr_singlescale/wr_vecnames/wr_onespace/numdgt` tune spacing and precision but never the separator, so a genuine comma-separated export needed a post-processing pass. The only `-csv` in the tree was `show -csv` (device-parameter tables, not data). ADDED a first-class CSV mode. (1) `set wr_csv --` a boolean read in `ft_writesimple` (`plotting/gnuplot.c`) alongside the existing `wr_*` vars; when set it takes a dedicated, self-contained branch that writes a single shared scale column once (implies `wr_singlescale`, so the existing equal-length sanity check applies), a header row of vector names (implies `wr_vecnames`), and comma-separated values with `%.*e` (no sign-padding, no trailing separator) honouring `numdgt`; complex `.ac` vectors become two columns (real, imag). The branch writes and returns before the space-formatted writer, so DEFAULT `wrdata` output is byte-for-byte unchanged. (2) `wrdata -csv <file> <vec...> --` a per-call alias accepted in ANY argument position: `com_write_simple` (`com_gnuplot.c`) scans the wordlist for a `-csv` token, enables `wr_csv` for just that write, and restores the prior global state afterward (no leak). Because `plotit` COPIES its wordlist (never mutates/frees the caller's), the `-csv` node is spliced out for the call and relinked afterward, leaving the caller's list intact to free; the head case is detected as `csvnode == wl` (the incoming head's `wl_prev` is not necessarily null), and no-vector forms fall through the `if (wl)` guard writing nothing, node still relinked. Scope: ngspice frontend only, two files (`frontend/com_gnuplot.c`, `frontend/plotting/gnuplot.c`); no solver/analysis/device/OSDI/compiler change; default output preserved exactly. Verify (`examples/csv_examples`, 9 checks): default output unchanged (no commas); `set wr_csv` writes a header + comma-separated values with the right numbers; `wrdata -csv` in first/mid/last position is byte-identical to the option form; CSV numbers equal the default-format numbers bit-for-bit; the flag does not leak into the next plain `wrdata`; `.ac` emits real/imag columns under a frequency scale; `.tran` emits a time scale with uniform csv-parseable rows. Regression 190/190 (new example folder).

E-232 *src/maths/KLU/klusmp.c, examples/solverfix_examples/**

KLU solver-glue correctness hardening (source audit). A deep read of the active SMP dispatch layer -- `klusmp.c`, the ONLY sparse-matrix-package layer compiled when KLU is enabled (`sparse/Makefile.am` builds `spsmp.c` only if `!KLU_WANTED`; `klusmp.c` routes BOTH solvers via its `CKTkluMODE` branches) -- found four latent defects, all fixed, all behaviour-preserving for any circuit that actually solves (prior functional KLU bugs E-113/114/115/171/172 were already fixed and are confirmed still in place). (A) `SMPluFac` and `SMPsolve` dereferenced `KLUmatrixCommon->status` BEFORE their `== NULL` guard; the complex twin `SMPclufac` checks `NULL` first with an early return -- reordered both to match (`SMPsolve`, being void, now reads status only inside the non-`NULL` else). `Common` lives for the matrix lifetime so it is never actually `NULL` here; the fix aligns the paths and satisfies a static analyzer. (B, the substantive one) the complex solves `SMPcSolve/SMPcaSolve` copied the RHS with a plain identity map (`RHS[i+1]`), while the real `SMPsolve` routes RHS through the node-collapse permutation `KLUmatrixNodeCollapsingNewToOld`; if a structural-zero column ever collapsed a node (`map != identity`) the AC/noise/pz RHS would be silently mis-ordered relative to the DC/`tran` solve. Dormant -- the map is the identity for any circuit that factors (an empty MNA column is structurally singular and does not solve), which is why KLU AC/noise/pz have always matched Sparse -- but the two paths were inconsistent; both complex solves now apply the identical gather/solve/scatter so a future change to collapse cannot quietly break complex analyses. (C) `SMPcZeroCol` indexed `KLUmatrixAp[Col-1]` with no ground guard (`Ap[-1]` if `Col==0`); added the `Col >= 1` guard its sibling `SMPfindEl` already has (pole-zero passes indices `>= 1`, so it never fired). (D) `SMPmultiply`'s real branch ended with a dead `iSolution = iRHS`; (reassigns a by-value local; a real matrix has no imaginary product)

-- removed. Scope: ngspice only, one file (maths/KLU/klusmp.c); no analysis/device/OSDI/compiler change and no behavioural change for any solvable circuit. Verify (examples/solverfix, 4 checks): on an ASYMMETRIC net (a VCVS makes the MNA matrix non-symmetric, so the noise ADJOINT solve genuinely differs from the forward solve) KLU still agrees with SPARSE 1.3 to the bit on AC, on the noise spectrum (exercising the most-changed SMPcaSolve), and on pole-zero roots (SMPcaSolve + SMPcZeroCol + SMPcAddCol), plus a clean DC-op + transient run under KLU. Regression 191/191 (new example folder).

E-233 *src/math/KLU/klusmp.c*

KLU glue deeper audit: finish the null-check reorder + correct the collapse map. Follow-up to E-232 from a deeper read of the same file; two more latent defects, both behaviour-preserving (one path never triggers, the other is confirmed-dead code). FIRST two scarier candidates were investigated and CLEARED empirically, not just by inspection: (i) a suspected per-analysis memory leak -- .pz/.sens reuse `ckt->CKTmatrix` and re-call `SMPconvertCOOtoCSC + SMPpreOrder` without freeing the prior arrays/Symbolic, which LOOKS like a leak -- was REFUTED by instrumenting the lifecycle over `op->pz->sens`: every convert/preorder sees a FRESH matrix (`Ap==NULL`, `Symbolic==NULL`) and `SMPnewMatrix/SMPdestroy` are balanced (5 new / 4 destroy / 0 overwrite-without-destroy); ngspice tears down and rebuilds the matrix per analysis, so nothing leaks (no fix made -- no bug). (ii) node-collapse is DEAD code: instrumenting the collapse branch and sweeping the whole 191-example regression (every OSDI model, pz, sens, noise) plus pathological topologies fired it 0 times, because OSDI collapses nodes at allocation time (`osdisetup.c`: a collapsed node gets no matrix column) so the COO never has a gap and `Gmin` fills every diagonal. FIXES: (3) E-232 fixed the dereference-before-NULL-check in `SMPluFac/SMPsolve` but the SAME pattern survived in `SMPcReorder` and `SMPpreorder` (they read `KLUmatrixCommon->status` before the `==NULL` guard) -- reordered both to check NULL first with an early return, matching `SMPcLUfac`, and corrected two mislabeled diagnostics (the Common-NULL and Symbolic-NULL branches each printed a stray 'KLUnumeric object is NULL'). Common is never actually NULL, so behaviour is unchanged. (2) in `SMPconvertCOOtoCSC` the structural-zero-column compaction recorded `NewToOld[reduced] = MatrixC00[i].col`, but with TWO OR MORE gaps the RHS is a partially-reduced index from an earlier pass, not the true original, so the map would resolve to the wrong node and the real `SMPsolve` gather/scatter would mis-address the RHS -- fixed by chaining through the existing map (`NewToOld[reduced] = NewToOld[MatrixC00[i].col]`, identity on the first pass so single-gap behaviour is unchanged); defensive since the path is confirmed dead. Scope: ngspice only, one file (maths/KLU/klusmp.c); no analysis/device/OSDI/compiler change and no behavioural change for any circuit. Verify: behaviour-preserving, so no new example -- the existing examples/solverfix_examples (KLU vs SPARSE bit-agreement on AC/noise/pz over an asymmetric net) plus the full 191/191 regression remain unchanged; fix 2's dead path cannot be exercised by construction (correctness is by the chaining argument). Regression 191/191.

E-234 *src/frontend/com_loadpull.c (new), src/frontend/com_loadpull.h (new), src/frontend/commands.c, src/frontend/com_commands.h, src/frontend/Makefile.am, examples/loadpull_examples/**

loadpull: power-amplifier load-/source-pull analysis. Load-pull (sweep the load impedance a PA output sees over the Smith chart and contour Pout/gain/PAE/efficiency to find the optimum load) is a PA-design staple ngspice lacked, and open-source load-pull is scarce. NEW `loadpull` command rides the existing `.tran` engine + alter (as optimize/sweep do) + `pyplot -contour` (E-218). `loadpull -load <R> <L> <C> -out <node> -drive <Vsrc> -f <f0> [-supply <Vsrc>]` sweeps Gamma over ‘

E-235 *src/frontend/com_stb.c, examples/stbfix_examples/**

stb: fix a latent use-after-free in the probe lookup. The defensive follow-up flagged in E-234. `com_stb` resolved its voltage probe with `INPretrieve(&name, symtab)`, which replaces the pointer with the interned symbol-table string -- the same memory the voltage source's own name field points at -- and does NOT free the old copy; the following `tfree(name)` therefore double-

freed the source's live name (and the symbol-table entry). It never bit in practice because stb runs once with no re-setup, but it is real memory corruption and the identical defect that bit loadpull in E-234 (there the per-point re-setups re-created the drive's branch node from that freed, allocator-reused memory). Fixed by dropping INPretrieve (top-level device names need no subckt translation -- findInstance does its own name match) and lowercasing a private copy instead (ngspice stores instance names lowercased), so only our own copy is freed; as a bonus the probe lookup becomes CASE-INSENSITIVE -- stb Vprobe Iprobe used to fail "no such probe source 'Vprobe'" and now resolves. The now-unused complex helper stbsub was removed to keep the build warning-free. Scope: ngspice frontend only, one file (com_stb.c); no solver/analysis/device/OSDI/compiler change; loop-gain results unchanged. Verify (examples/stbfix, 3 checks): a mixed-case probe name resolves and reports a loop gain (failed before); the loop gain is still correct (a gain-1e5 buffer loop -> ~100 dB DC loop gain, measured 99.91 dB); 60 repeated stb runs stay stable. Regression 193/193 (new example folder).

E-236 *src/frontend/com_measure2.c, src/frontend/com_measure2.h, src/frontend/measure.c, src/tclspice.c, examples/measovf_examples/**

.meas: fix a stack-buffer overflow on long measurement names. Found in a memory-safety deep dive; reproduces on the stock binary. get_measure2() formats each measurement result line with sprintf(out_line, "%-20s= %.*e ...", mName, ...) into the caller's fixed stack buffer char out_line[1000] (measure.c). mName is the measurement NAME token, taken verbatim from the .meas <analysis> <name> <func> ... card via cp_unquote -- an unbounded user string with no length check. A .meas name longer than ~1000 chars overruns out_line; on macOS the stack canary fires and the process aborts (SIGABRT, exit 134), elsewhere it is plain stack corruption. All ten result-formatting sites share the flaw. E-225 had already hardened the sibling errbuf[100] in this same file to snprintf (after a measure-syntax fuzz campaign), but the fuzzing never exercised a long measurement NAME, so out_line was missed. Fix: thread the destination buffer size through get_measure2(..., size_t max_out_line, ...) and convert all ten sprintf(out_line, ...) to snprintf(out_line, max_out_line, ...). The three callers pass the real size -- measure.c batch .meas passes sizeof out_line, the autostop check and the tclspice.c interactive path pass NULL/0 (all ten sites are already guarded by if (out_line), so they never reach snprintf). An over-long name is safely truncated to 999 chars; the measurement value itself is computed and stored exactly as before. Scope: ngspice frontend only; no solver/analysis/device/compiler change; measurement results unchanged. Verify (examples/measovf, 2 checks): a .meas with a ~4000-char name no longer crashes (exit 0, was 134); normal measurements still correct (MAX 1.0, AVG 0.6, 0.1->0.9 rise time 0.8 ns). Regression 194/194 (new example folder).

E-237 *src/frontend/dotcards.c, src/frontend/vectors.c, examples/nameovf_examples/**

.print/.plot/.four: fix stack-buffer overflows on long vector/node names. Sibling of E-236, same deep dive. The SPICE2-compat rewrites for output tokens, and the vector-name helper they feed, each copied an unbounded user name into a fixed BSIZE_SP[512] stack buffer: fixem() (dotcards.c) rewrites v(a,b)->v(a)-v(b) (+ vm/vp/vi/vr/vdb variants) via sprintf(buf, "v(%s)-v(%s)", a, b) (19 writes); gettoks() (dotcards.c) rewrites i(x)->x#branch via sprintf(buf[513], "%s#branch", x); vec_basename() (vectors.c) does strcpy(buf, v->v_name), reached by .print/fft/spec/linearize. A .print/.plot/.four output token whose node/branch name(s) exceed the buffer overran the stack -- macOS aborts (SIGABRT/SIGTRAP, shell exit 133), elsewhere plain corruption. Fixing fixem alone was NOT enough: with valid long node names the rewrite succeeds and the long v(a)-v(b) name then flows into vec_basename and crashes there, so all three needed hardening. Fix: size each scratch buffer to its input -- fixem allocates strlen(string)+32 and every write is a bounded snprintf(buf, bufsz, ...); gettoks builds the string with tprintf; vec_basename replaces the fixed buffer + strcpy with copy() of the source (allocates to fit), preserving branch semantics. No truncation, so a valid long differential still computes the correct value. The vectors.c recompile also surfaced 3 pre-existing

-Wconversion warnings (v->v_type, enum simulation_types, passed to dvec_alloc's int param), silenced with explicit (int) casts (build stays warning-free). Scope: ngspice frontend only; no solver/analysis/device/compiler change; output values unchanged. Verify (examples/nameovf, 4 checks): long v(a,b) on nonexistent nodes no longer crashes; a VALID long differential v(A,B) with v(A)=2,v(B)=1 runs and prints v(A)-v(B)=1.0 exactly (no truncation); long i(x) via .four (gettoks) no longer crashes; short v(1,2) still rewrites to v(1)-v(2)=2.0. Regression 195/195 (new example folder).

E-238 *src/frontend/dotcards.c, examples/malftoken_examples/**

gettoks: fix a NULL-deref crash on a malformed differential token. Third find of the same deep dive (after E-236/E-237), from fuzzing degenerate output tokens; reproduces on the stock binary. gettoks() parses the output tokens of .save/.print/.plot/.four (via ft_savedotargs) and .measure vars (via com_measure2). For a differential v(a,b) it finds the comma c and close paren r (r = strchr(t, ')'), c = strchr(t, ',') and splits the second operand at r: if (c != r) { *r = '\0'; ...}. A MALFORMED token like v(1, has a comma but NO ')', so r is NULL while c is not -- c != r is then true and *r = '\0' dereferences NULL -> SIGSEGV (EXC_BAD_ACCESS at 0x0; lldb: gettoks <- ft_savedotargs). A one-character typo (missing paren) in a .print/.save/.plot/.four/.measure token crashes the simulator. Fix: require r non-NULL before splitting -- if (r && c != r). When there is a comma but no ')', the split is skipped and the token degrades to a harmless parse; every well-formed case is unchanged (v(a,b) both non-NULL c!=r -> split; v(a) c==r -> skip; i(x) sets c=r=NULL -> skip). Scope: ngspice frontend only, one guard in dotcards.c; no solver/analysis/device/compiler change; well-formed parsing unchanged. Verify (examples/malftoken, 3 checks): .print tran v(1, no longer crashes; the malformed shape across .save/.plot/.four/.meas + i()/vdb()/vm() no longer crashes; well-formed v(1,2) still rewrites to v(1)-v(2)=1.5. A broad malformed-token sweep (prefixesxbodiesxcards, hundreds) is crash-free after the fix. Regression 196/196 (new example folder).

E-239 *src/spicelib/parser/inpptree.c, examples/funcarity_examples/**

Expression parser: fix a NULL-deref on a one-argument **min/max/pow/pwr**. Fifth find of the same deep dive (after E-235-E-238), from fuzzing the behavioral-source expression parser; reproduces on the stock binary. PTEval() (ifeval.c) evaluates the two-argument functions pow/pwr/min/max by reading their operands as tree->left->left and tree->left->right -- it assumes the argument is a PT_COMMA pair. The function-node builder PT_mkfnnode() (inpptree.c) looks the name up but NEVER validates argument count, so a one-argument call like min(1) builds a node whose left is a scalar PT_CONSTANT (its ->left is NULL); at load time PTEval does tree->left->left -> NULL deref -> SIGSEGV (PTEval <- IFeval <- ASRCLoad <- CKTload). Reachable from any B-source, E/G source, or .param expression -- a one-line typo crashes the simulator during setup. Crash is specific to these 4 functions and the scalar (non-comma) one-arg form: min() is caught by the existing if (!arg) guard, min(1,2,3) yields a PT_COMMA (no deref), and single-arg forms of other functions (hypot(1), atan2(1)) fall through the single-arg path and error cleanly. Fix: reject the wrong arity at parse time in PT_mkfnnode -- 'if((POW

E-240 *src/xspice/icm/analog/s_xfer/cfunc.mod, bin//codemodels/analog.cm (rebuilt), examples/sxfer_examples/*

XSPICE **s_xfer**: fix an out-of-bounds crash on a static-gain transfer function. Sixth find of the same deep dive (after E-235-E-239), from fuzzing XSPICE a-device instantiation; reproduces on the stock binary. First find in the XSPICE code-model library rather than the core simulator. The s_xfer analog code model (a Laplace transfer function H(s)=num/den in controller-canonical form) allocates its integrator state array to den_size elements, then in the DC/transient branch unconditionally read pout_pin = *(integrator[1]) (the d(out)/d(in) partial) -- assuming a first- or higher-order denominator. A 0-order denominator (den_size == 1, i.e. a STATIC-GAIN transfer function like s_xfer(num_coeff=[k] den_coeff=[1]) = constant k) allocates integrator with a single element, so integrator[1] is one past the end -> read garbage -> SIGSEGV at load

(cm_s_xfer <- MIFload <- CKTload). Depends only on den_size==1, NOT the coefficient values: den_coeff=[1] (valid unity denominator), [5], [0] all crash; [1 1]/[1 2 1] fine -- so it bit ordinary static-gain input, not just degenerate decks. AC branch already safe (loops bounded by den_size/num_size). Fix: guard the read -- if (den_size > 1) pout_pin = *(integrator[1]); else pout_pin = (num_size>0) ? *gain * *(num_coefficient[0]) : 0.0 (for den_size==1, out = gainnum[0]/in, so the partial is gainnum[0]); the output value was already computed correctly. Zero denominator (H=inf) now gives clean non-convergence instead of a crash. Fix is in a .cm code model, so analog.cm was regenerated via cmpp and redeployed under bin//codemodels/; the ngspice binary is unchanged. Scope: XSPICE analog code-model library only (cfunc.mod + rebuilt analog.cm); no core simulator/solver/analysis/device/compiler change; dynamic-filter results unchanged. Verify (examples/sxfer, 4 checks): static-gain s_xfer(num=[3] den=[1]) no crash & exact (H=3 -> 12 V); another static gain H=0.4 -> 1.6 V; zero denominator no crash; dynamic H=1/(s+1) still correct (

E-241 *src/frontend/com_fft.c, examples/fftnorm_examples/**

fft/spec: fix the amplitude normalization for non-power-of-2 records. First numerical-correctness (wrong-number, not crash) find of a correctness-audit campaign that checked ngspice analyses against analytic + numpy oracles. com_fft.c has two paths: FFTW3 (exact length-point transform, scale = length/2, correct) and a Green's radix-2 FFT (no FFTW3) that zero-pads the length input samples up to the next power of two N. The non-FFTW path normalized the single-sided amplitude by the PADDED size (scale = N/2) instead of the true sample count (length/2), so every FFT whose sample count is not a power of two reported amplitudes too small by length/N -- up to 2x. A .tran essentially never yields exactly 2^k points, so this bit the common case silently. The DC bin is the clean discriminator (always exactly bin 0, no window/scalloping ambiguity): a signal with DC offset D read D*length/N, e.g. a 1025-sample record (padded to 2048) reported a DC of 2.0 as 1.0009. spec carried the same error in its power normalization (intres = N*N instead of length*length; the FFTW path already used length*length). Fix: scale = length/2 and intres = length*length in the non-FFTW path, matching FFTW. The frequency axis was already padding-aware (freq[i] = i/span * length/N), so only the magnitude scale needed the change; power-of-2 records (length == N) and FFTW-linked builds are unchanged. Validated against numpy as an independent oracle: after the fix ngspice's fft matches numpy.fft.rfft of the same zero-padded record (normalized by length) bin-by-bin to ~1e-9. Scope: ngspice frontend only, two normalization constants in com_fft.c; fft/spec amplitudes now change for non-power-of-2 records (they become correct). Verify (examples/fftnorm, numpy-free, 4 checks): DC offset 2.0 reads 2.0 on a non-power-of-2 record; same on a maximally-padded record (1025->2048, where the bug halved it to ~1.0); DC reading independent of sample count across four record lengths; spec PSD tone peak length-independent. Regression 199/199 (new example folder).

E-242 *src/spicelib/devices/nport/* (new device dir), src/frontend/snp2va.c, src/frontend/snp2va.h, src/frontend/com_presnp.c, src/spicelib/parser/inp2n.c, src/spicelib/devices/dev.c, configure.ac, src/spicelib/devices/Makefile.am, src/Makefile.am, examples/nport_native_examples/, bin//ngspice (rebuilt)*

Native C n-port device + **pre_snp -native:** a pole/residue Y-parameter block stamped straight into the matrix, no OpenVAF compile. pre_snp (E-200/201/205) converts Touchstone -> Verilog-A -> openvaf-r -> .osdi; the fit is cheap but the compile is the bottleneck and walls at ~24-32 ports (0(N*Np) laplace_nd sections, 0(N^2) terms for full-rank coupling). This adds a built-in nport device (new src/spicelib/devices/nport/) that evaluates the same model $Y_{ij}(s) = d_{ij} + s * e_{ij} + \sum_k \text{res}_{ijk} / (s - p_k)$ directly. It stamps as a multi-terminal admittance (four-corner conductance per (i,j); no per-port branch unknowns -> N equations not 2N): DC stamps $Y(0)$, AC the complex $Y(j\omega)$ into the (real,imag) slots, and transient a trapezoidal companion -- d a conductance, s*e a capacitor, each $\text{res}/(s-p)$ a first-order state $dx/dt = p*x + u$ whose shared pole states x_{jk} are advanced once per load and parked in CKTstate0 (the stamp phase

recovers the history $B=x_{n+1}-a*u_j$, so no per-instance scratch). Complex poles come as conjugate pairs so the summed current is real. It rides the generic N dispatcher (inp2n.c broadened from OSDI-only to also accept the nport model): `N1 p1..pN ref mymodel / .model mymodel nport(file="m.nport"), ref` an explicit reference terminal (tie to 0 for a ground-referenced Touchstone block; the extra terminal makes floating/differential blocks expressible). KLU supported via nportbindCSC.c (real/complex/complex->real bindings; a ground row/col keeps its Sparse trash pointer, like the RLC devices). snp2va.c was refactored so the shared parse+vector-fit lives in snp_fit(), feeding two emitters: `pre_snp -osdi` (default, unchanged VA->OSDI route) and new `pre_snp -native`, which writes the compact .nport for this device -- same fit, no compiler, scaling to hundreds of ports. Verify (examples/nport_native, 6 checks x 2 solvers, all vs closed-form oracles, no openvaf-r needed): RC one-port e-term (AC $1/Y$ + transient exponential); RLC one-port complex conjugate pole pair through resonance; Pi two-port off-diagonal d/e cross-coupling vs analytic MNA; `pre_snp -native` fits a known 2-port .s2p and reproduces its Y; a 20-port 6-pole full-rank model vs an analytic linear solve. Cross-checks (dev notes): native AC == OSDI backend to 1e-11 on 137mhz_bpf.s2p; a 100-port block solves an AC point in ~90 ms (4.85e-10 vs analytic); KLU == Sparse bit-for-bit. Scope: ngspice only, additive; default -osdi path unchanged so existing decks are unaffected; no solver/analysis/compiler change. Regression 200/200 (new device dir + example folder).

E-243 *src/frontend/snp2va.c, examples/presnp_examples/verify_presnp.py, bin/*/ngspice (rebuilt)*

pre_snp -osdi: emit an explicit **ref** terminal so the instance line is identical for both **-osdi** and **-native**. E-242's native nport device carries an explicit reference terminal (`N1 p1..pN ref model`), but the original `pre_snp -osdi` emitted a Verilog-A module with only the N signal ports referenced to implicit global ground (`N1 p1..pN model`) -- so the two backends were not drop-in interchangeable. snp2va.c now emits module `m(p1..pN, ref)` and takes every branch relative to ref: `I(pi,ref) <+ ... V(pj,ref) ...` (and `laplace_nd(V(pj,ref), ...)`, `ddt(V(pj,ref))`). ref is a PLAIN electrical port -- NOT a ground-declared node: that keeps it connectable (tie to 0 for a ground-referenced block, or any node for floating/differential), and it appears in every `I(pi,ref)` stamp so it is well-posed (declaring ground ref would make it an internal 0V node -- no longer a terminal -- and risks the E-116 all-zero-row/structurally-singular-KLU failure a ground node in no Jacobian entry caused). The instance line is now identical for both backends. Compatibility: the -osdi module goes from N to N+1 terminals, so existing -osdi decks must add the ref node (tie to 0 to reproduce the old behavior); the presnp example decks were updated. Verify: the identical line `N1 p1 p2 0 model` through both backends matches to 1e-11 on 137mhz_bpf.s2p (ref=0 reproduces the prior result exactly); examples/presnp (9 checks: 2-/3-/4-/8-port pre_snp cases, AC+tran+ordering) passes with the ref terminal. Scope: ngspice only, the VA emitter in snp2va.c; -native / .nport device / vector fit / .osdi compile path all unchanged. Regression 200/200.

E-244 *src/spicelib/devices/nport/nportsetup.c, src/frontend/com_pyplot.c, examples/crashfix2_examples/, bin/ngspice (rebuilt)*

Crash-hardening round 2: two user-triggerable crashes in recently-added code (both reproduce on the shipped binary). Found by argument-fuzzing the newer devices/commands. (1) nport device (E-242): NPORTsetup never checked that the instance line connected all N+1 nodes the .model's port count claims. A .nport with more ports than the instance wires (`N1 a 0` mod for a 2-port), or nports beyond NPORT_MAXTERMS, made setup stamp an UNBOUND (-1) terminal -> the generic N dispatcher leaves unconnected terminals as -1, and passing a negative row/col trips the sparse builder assert `Row>=0 && Col>=0` (spbuild.c:267 -> SIGABRT) / out-of-bounds. Fix (nportsetup.c): reject a .nport whose port count leaves no room for the reference (`N+1 > NPORT_MAXTERMS`) and verify every node `[0..N] >= 0` before stamping -> clean E_BADPARAM with a diagnostic naming the shortfall. (2) pyplot (E-217/E-218): the -hist/-contour render-mode markers were stripped by UNLINKING AND FREEING a node of the command's OWN argument wordlist -- but that list is owned/freed by the command loop. When the marker was the FIRST word, com_pyplot freed

the list HEAD, which the command loop then freed again -> use-after-free/double-free (SIGSEGV) on `pyplot -hist v(out) / pyplot -contour ... (pyplot v(out) -hist, marker not first, was fine -- so the E-217/E-218 examples, which never put it first, always passed.)` Fix (`com_pyplot.c`): detect the marker by scanning (no mutation) and build a filtered COPY without it, freed by us at the end; the caller's list is untouched. Bonus: `-hist` as first arg now renders the histogram it silently degraded to a line plot before. Verify (examples/crashfix2, 9 checks x2 solvers): under-bound + over-max nport decks exit gracefully with the diagnostic; a correctly-bound 2-port still simulates; `pyplot -hist/-contour first-arg` and `-hist-alone` no longer crash; `-hist-first` now emits `plt.hist()`; trailing-marker form still works. Scope: ngspice only, device input validation + command-arg ownership; no numerical change, valid usage unaffected. Regression 201/201.

E-245 *src/frontend/com_measure2.c, src/frontend/spiceif.c, examples/crashfix3_examples/, bin//ngspice (rebuilt)*

Crash-hardening round 3: two NULL-deref crashes in CORE ngspice command parsers (stock code, both reproduce on the shipped binary). Found by argument-fuzzing `meas` and `altermod`. (1) `meas`: `measure_parse_stdParams (com_measure2.c)` splits each trailing param token on '=' with `strtok`; a LONE '=' token (a stray one, e.g. `meas tran m find v(x) when v(x)=0.5 = y`) makes `strtok` return NULL for the name, which was then passed to `strcascmp()` -> NULL deref SIGSEGV. The `pValue==NULL` branch was guarded but `pName` was dereferenced first. Only the interactive `meas` command hits it (the `.meas` dot-card pre-tokenizes differently); `key=val/=val/val=` are all fine, only a standalone '=' produces the NULL. Fix: treat a NULL `pName` as a clean syntax error before `strcascmp`. (2) `altermod`: `altermod nm c` where the 2nd token is a bare device-type letter (c/e/i/...) or a digit makes `com_altermod` treat it as another MODEL to alter with no `param=value`, so `paramlookup (spiceif.c)` is called with `param==NULL`; its model-parameter loop passed that straight to `eq()/strcmp()` (the instance-parameter loop just above already guarded `!param`) -> NULL deref SIGSEGV. Input-dependent: only tokens resolving to a real device type crash (c/e/i/0/1/2), others are dropped earlier. Fix: guard the model loop against a NULL `param` (and defensively a NULL keyword), matching the instance loop. Verify (examples/crashfix3, 12 checks x2 solvers): the stray '=' `meas` forms and `altermod nm <c`

E-246 *src/xspice/icm/analog/pwl/cfunc.mod, src/xspice/icm/analog/pwlts/cfunc.mod, bin//codemodels/analog.cm (rebuilt), examples/pwlfix_examples/*

XSPICE **pwl/pwlts**: fix an out-of-bounds heap read on `x_array/y_array` length mismatch. Second find in the XSPICE code-model library (after E-240's `s_xfer`), from fuzzing code-model parameters; confirmed with AddressSanitizer. `pwl` and `pwlts` take two INDEPENDENT vector parameters `x_array` (breakpoints) and `y_array` (ordinates); the XSPICE framework allocates each to its OWN length. Both models size their working table from `x_array` only (`size = PARAM_SIZE(x_array)+2`) then copy BOTH arrays with that single index (`y[i]=PARAM(y_array[i-1])`), so a `y_array` SHORTER than `x_array` reads past the end of the `y_array` parameter -> heap-buffer-overflow READ (ASan: `cm_pwl cfunc.c:321 / cm_pwlts cfunc.c:202`), pulling adjacent/uninitialised heap into the interpolation table (undefined behaviour; can crash outright once the mismatch runs past the mapped heap). Reachable from any deck with such an a-device (`.model m pwl(x_array=[0 .25 .5 .75 1 1.5 2 2.5] y_array=[0 1])`). The sibling models already guard this -- `oneshot` checks `cntl_size != pw_size`, `multi_input_pwl` checks `PARAM_SIZE(x) != PARAM_SIZE(y)` -- but `pwl/pwlts`, the oldest PWL models, predate the check. A sweep of the other analog models with paired arrays found no further cases (sine/square/triangle copy into their own `cntl_size`-sized buffers, so a short array zero-pads rather than overrunning; `summer/mult` are validated by the framework against the connection size). FIX: add the same equal-length guard -- `if (PARAM_SIZE(x_array) != PARAM_SIZE(y_array)) {cm_message_send(size_error); return;}` -- before the alloc/copy in both models, aborting setup cleanly with `x_array` and `y_array` must have the same length instead of reading out of bounds. Fix is in a .cm code model, so `analog.cm` was regenerated via `cmpp` and redeployed under `bin/*/codemodels/`; the ngspice binary is unchanged. Scope: XSPICE analog code-model library only;

no core simulator/solver/analysis/device change; valid equal-length pwl/pwlts usage unaffected. Verify (examples/pwlfix, 5 checks x2 solvers): a valid pwl interpolates exactly (in 0.5 -> 1.0 V); pwl with x longer than y, and the other way, both report the size error without crashing; a valid pwlts still runs; a mismatched pwlts reports the size error; self-skips if codemodels unavailable. The OOB was reproduced and shown fixed under an AddressSanitizer build of the code models. Regression 203/203 (new example folder).

E-247 *src/xspice/icm/table/table2D/cfunc.mod, src/xspice/icm/table/table3D/cfunc.mod, bin/codemodels/table.cm* (rebuilt), examples/tablefix_examples/

XSPICE **table2d/table3d**: fix out-of-bounds read + interpolation UB on degenerate or too-small tables. Third area of the XSPICE code-model deep dive (after E-240 s_xfer, E-246 pwl/pwlts), found by fuzzing table dimensions; confirmed with ASan/UBSan. Both models read a lookup table from a file whose first lines declare the axis point counts (ix, iy [, iz]); the loader (init_local_data) never validated them, so a small table broke two ways. (1) OUT-OF-RANGE RAMP (heap OOB read): for an input past the table edge the model ramps the derivative to zero using xcol[1]-xcol[0] (low side) or xcol[ix-1]-xcol[ix-2] (high side); a single-point axis (ix==1) has no xcol[1]/xcol[-1] -> heap-buffer-overflow READ (ASan: cm_table2D cfunc.c:298); table3d has the same for all three axes. (2) ENO INTERPOLATION (OOB/UB): the Madagascar ENO of order p reads a 2*(p-1)-point stencil per axis (order 2 needs 2 pts, 3 needs 4, 4 needs 6, 5 needs 8); the order param DEFAULTS to 3 (needs >=4 pts/axis) but the loader only clamped it UP to a minimum of 2, never DOWN to what the table supports -> the default order on any table with a <4-point axis (a 3x3 table, or a 2-plane z axis like the shipped test-3d-1.table with iz=2) ran off the stencil -> UB/OOB in eno2.c/en03.c. FIX (both models, right after the order>=2 floor): compute mindim = min(ix,iy[,iz]); reject mindim<2 with a clean error (a 1-point axis can be neither ramped nor interpolated); and clamp interporder to mindim/2+1 (the inverse of the 2*(p-1) stencil rule) so a small table interpolates at a reduced but VALID order instead of reading out of bounds. Tables large enough for the requested order are byte-identical (every shipped example has mindim/2+1 >= order; the smallest are 7x7). Fix is in .cm code models, so table.cm was regenerated via cmpp and redeployed under bin/*/codemodels/; the ngspice binary is unchanged. No regression coverage impact: the regression table_model example uses OSDI Verilog-A table models, not these XSPICE code models. Scope: XSPICE analog code-model library only; no core simulator/solver/analysis change; tables large enough for the requested order unaffected. Verify (examples/tablefix, 6 checks x2 solvers): a valid 8x8 order-3 table2d interpolates exactly (out=x+10y at (2.5,3.0) -> 32.5); a single-x-point table2d and a degenerate table3d (iz=1) are rejected with a clean error; small 3x3 / 3x3x3 order-3 tables run (order clamped) without crashing. The OOB read and the ENO UB were reproduced and shown fixed across a full (dimension x order) grid under ASan/UBSan builds. Regression 204/204 (new example folder).

E-248 *src/spicelib/devices/cpl/cplsetup.c, bin/ngspice* (rebuilt), examples/cplfix_examples/

CPL coupled-line device: fix two out-of-bounds accesses (conductor count + R/L/C/G matrix sizes). First find in a CORE ngspice device (the XSPICE dive E-240/246/247 was the code-model library); found by fuzzing the p-device conductor count and matrix sizes, confirmed with ASan/UBSan; reproduces on the shipped binary. A CPL line has noL conductors (the instance dimension, set from the node count on the p card); each symmetric R/L/C/G matrix is its upper triangle, noL*(noL+1)/2 values. CPLsetup/ReadCpL (spicelib/devices/cpl/cplsetup.c) validated NEITHER count. (1) UNDER-SPECIFIED MATRIX (heap OOB read): ReadCpL fills the matrices with f = CPLmodPtr(herer)->Rm[counter] (and Lm/Cm/Gm) where counter runs 0..noL*(noL+1)/2-1, but Rm/Lm/Cm/Gm are allocated to the count the user actually gave (Rm_counter, ...); fewer values than the triangle needs -> read past the end (ASan: heap-buffer-overflow READ at cplsetup.c:474). (2) TOO MANY CONDUCTORS (fixed-array overflow): ReadCpL uses RLINE*lines[MAX_CP_TX_LINES] and CPLLine.in_node[MAX_CP_TX_LINES] with MAX_CP_TX_LINES==8 (include/ngspice/swec.h), indexed by noL; a p card with >8 conductors writes past them (UBSan: index 8 out of bounds for 'NODE *[8]' at cplsetup.c:430). Both reachable from a

valid-syntax netlist; on the release binary the reads pull adjacent heap -> heap-layout-dependent garbage (here a downstream fatal-error exit, but could equally be wrong results or a crash). FIX: validate both up front in CPLsetup, right after noL = here->dimension and before any allocation -- reject noL<1

E-249 *src/spicelib/devices/urc/urcsetup.c, src/spicelib/devices/ltra/ltraset.c, bin//ngspice (rebuilt), examples/tlinefix_examples/*

URC + LTRA transmission-line input validation: reject a degenerate/unbounded lump count and negative line parameters. Two CORE transmission-line devices accepted out-of-range parameters that gave a silently wrong result or a resource-exhaustion hang instead of a clean error (same class as E-248 CPL). (1) URC (urcsetup.c): the uniform-RC line is expanded at setup into a ladder of n (the n= lumps) resistor+capacitor(/diode) sections -- one node group per lump -- and n is also an exponent of pow(k,n); the user's n was NEVER bounds-checked. n<=0 ran the loop zero times -> no lumps -> output node left silently unconnected (v(out)=0, a wrong answer, no diagnostic); a large n (typo n=100000000, or 2e9 near INT_MAX) exhausted memory / hung building the ladder (even ~20000 lumps takes tens of seconds; the source even carried a may want to limit lumps to <= 100 comment). Fix: require 1 <= n <= URC_MAX_LUMPS (1000; auto-computed count is ~3..30), clean E_BADPARAM otherwise, validated after the auto-lump default so it covers both given and computed counts. (2) LTRA (ltraset.c): the lossy line picks its RLGC special-case with != 0 tests, so a NEGATIVE L or C passed and fell through to sqrt(L/C)/sqrt(L*C) in LTRAtemp -> NaN characteristic impedance/delay and a degenerate all-zero run (a zero/too-few-param line was already the clean "at least two ... must be ... nonzero" error). Fix: reject any negative R/L/G/C up front (per-unit-length params must be physical). Scope: core ngspice, two device setups; binary rebuilt; no solver/analysis/numerical change; a validly-specified URC/LTRA line is unaffected. Verify (examples/tlinefix, 8 checks x2 solvers): a valid 5-lump URC divider and a valid LTRA lossless line still simulate; URC n=0/-3/100000000 each rejected cleanly (the huge case INSTANTLY, no hang); URC n=1000 (max) still runs; LTRA negative C rejected cleanly while C=0 keeps the pre-existing nonzero error. Regression 206/206 (new example folder).

E-250 *src/xspice/icm/digital/d_lut/cfunc.mod, src/xspice/icm/digital/d_genlut/cfunc.mod, bin//codemodels/digital.cm (rebuilt), examples/dlutfix_examples/*

XSPICE d_lut/d_genlut: fix an undefined-behaviour 1<n shift on the input-port count. Found with UBSan while sweeping the XSPICE code-model library (E-240/246/247 dive). Both digital lookup-table models size their table as 2^(#inputs) via a left shift: d_lut size=PORT_SIZE(in); tablelen=1<<size (cfunc.c:143-144), d_genlut isize=PORT_SIZE(in); entrylen=(1<<isize) (cfunc.c:154/177). The input-port count was NEVER bounded: a left shift of a 32-bit int by >=32 is UNDEFINED BEHAVIOUR (UBSan: "shift exponent 32 is too large for 32-bit type 'int'" at cfunc.c:146/:177), and a count in the high twenties asks for a multi-gigabyte table alloc. Reachable from a valid-syntax netlist -- an a-device with ~30+ digital inputs fanned into the LUT (e.g. all tied to one pulled-up net). On arm64 the release UB is benign (shift amount taken mod 32 -> 1<<32==1, no crash) so the table size is silently WRONG rather than loud -- which is exactly why it's worth fixing. Fix: cap the input-port count at D_LUT_MAX_INPUTS / D_GENLUT_MAX_INPUTS (24; a real LUT has only a handful of inputs, and the table_values string is 2^n chars) BEFORE the shift, cm_message_send clean error + return otherwise. Fix is in .cm code models, so digital.cm was regenerated via cmpp and redeployed under bin/*/codemodels/; the ngspice binary is unchanged. Scope: XSPICE digital code-model library only; no core simulator/solver/analysis change; a normal LUT is unaffected. Verify (examples/dlutfix, 4 checks x2 solvers): valid 2-input d_lut and d_genlut still simulate; a 32-input d_lut and d_genlut are each rejected with a clean "out of range" error, no crash. The 1<<32 UB was reproduced under a UBSan build and shown fixed. Regression 207/207 (new example folder).

E-251 *examples/hb_examples/verify_hb.py (verification only -- no ngspice source change)*

Harmonic-balance tight nonlinear correctness verification. A rigorous tightening of the HB (E-

134/135) correctness cross-check. The existing checks compare the HB spectrum against ngspice's own transient+fourier steady state, but at K=8 harmonics and a 1.5-3% tolerance. This proves -- and enforces in the regression -- that HB converges to the EXACT periodic steady state, not merely within a few percent. AUDIT: HB and tran+fourier are independent code paths, so their agreement is a real oracle once the transient reference is accurate. Using a purely RESISTIVE rectifier (cjo=0/tt=0 -> steady state reached instantly, so the transient's only error is its $O(dt^2)$ timestep, which is Richardson-extrapolated to $dt \rightarrow 0$) and sweeping K=8..48, the HB harmonics converge MONOTONICALLY to the $dt \rightarrow 0$ transient; at K=48 the match is $\sim 1e-7$ (h1/h2 exact, h3 -1.9e-7, h4 -3.1e-6, h5 +8.9e-7). The 13% apparent 5th-harmonic 'error' at K=8 is entirely HB truncation (aliasing of the discarded harmonics), NOT a modelling error -- it vanishes as K grows. TIGHTENED CHECK (verify_hb.py [10]): on the resistive rectifier, HB at K=24 matches a fine-dt (period/2000) transient fourier to < 0.5% for DC..4th (6x tighter than the loose checks, observed worst-case $\sim 2e-5$), and raising K from 8 to 24 strictly shrinks the 5th-harmonic mismatch (encoding 'converges as K grows' so a future regression that broke HB's harmonic content would be caught). Scope: verification only (verify_hb.py); no ngspice binary / HB-engine / analysis change; documents that HB is EXACT, not merely within tolerance. Regression 207/207.

E-252 *src/xspice/icm/analog/xfer/cfunc.mod, src/xspice/icm/analog/file_source/cfunc.mod, bin/codemodels/analog.cm* (rebuilt), *examples/filefix_examples/*

XSPICE xfer/file_source: fix two heap OUT-OF-BOUNDS WRITES in the file parsers. Found while sweeping the file-parser code models (d_state/d_source/xfer/file_source); confirmed with ASan. Unlike the earlier code-model finds (E-246/247/250, reads or bounded errors) these are OOB WRITES -- silent heap corruption. (1) xfer (read_file): reads a transfer-function file (Touchstone-style # option line + data), sscanf's up to 9 values/line and a state machine stores EVERY value (freq/real/imag triples), so a line with more than one record stores >3; but the alloc check reserved only 3 (if(i+3>size), ALLOC=1024) -> a multi-record line writes past the buffer at the 1024-double boundary (ASan: heap-buffer-overflow WRITE, cm_xfer cfunc.c:232). Fix: reserve the sscanf max of 9 (if(i+9>size)); the loop appends at most count<=9/line since offset>=1 && span>=offset+2 are enforced). (2) file_source: stores one record/line = a timepoint + size channels = stepsize(=size+1) doubles, but reserved only size (vecallocated - size), one short; at the realloc boundary (count lands exactly at vecallocated-size once vecallocated grows out of stepsize alignment) the last channel writes one double past the end (ASan: heap-buffer-overflow WRITE, cm_filesources cfunc.c:326). Fix: reserve a full record (-stepsize). Both are heap OOB WRITES reachable from a valid-syntax netlist + crafted data file; release build corrupts adjacent heap silently rather than always crashing (UB either way). Siblings checked: d_source validates per-line token count vs declared width; d_state's fixed line buffer is fgets-bounded -- no analogous overrun. Fix is in .cm code models, so analog.cm was regenerated via cmpp and redeployed under bin/*/codemodels/; the ngspice binary is unchanged. Scope: XSPICE analog code-model library only; no core/solver/analysis change; a well-formed data file is unaffected. Verify (examples/filefix, 4 checks x2 solvers): valid transfer-function + file_source files simulate; an xfer multi-record file and a boundary-crossing file_source file run without overrunning. Both reproduced under ASan (cm_xfer/cm_filesources) and shown fixed. Regression 208/208 (new example folder).

E-253 *src/frontend/com_rfstab.c, src/frontend/com_rfstab.h, src/frontend/commands.c, src/frontend/com_commands.h, src/frontend/Makefile.am, bin/ngspice* (rebuilt), *examples/rfstab_examples/*

rfstab: two-port stability & gain report -- the first RF design-aid on top of the S-parameter machinery. ngspice can compute S-parameters (.sp / n-port / Touchstone) but not turn them into the stability/gain figures of merit an RF designer works with. After a .sp analysis publishes S_1_1..S_2_2 vs frequency, the new rfstab command post-processes them into the standard linear-two-port metrics per frequency: determinant $D=S_{11}S_{22}-S_{12}S_{21}$; Rollett $K=(1-$

E-254 *src/frontend/plotting/pyplot.c, src/frontend/plotting/pyplot.h, src/frontend/plotting/plotit.c, src/*

frontend/com_pyplot.c, bin//ngspice (rebuilt), examples/pyplotsmith_examples/

pyplot -smith: Smith-chart plotting mode -- the second RF design-aid (after rfstab, E-253). A Smith chart is the working surface for two-port RF design (reflection coefficients, matching trajectories, stability/gain circles), but ngspice's modern matplotlib output (pyplot) had no Smith mode -- only the legacy in-terminal plotter did. `pyplot [name] -smith <complex vectors>` draws each listed complex vector (S_{11} , S_{22} , a load reflection Γ , a stability/gain-circle trace) as a curve in the reflection-coefficient plane ($\text{Re} \rightarrow x$, $\text{Im} \rightarrow y$) over the standard Smith grid: the unit circle

E-255 *src/spicelib/analysis/cktdisto.c, bin//ngspice (rebuilt), examples/distoexact_examples/*

.disto proven machine-exact vs Harmonic Balance / QPSS-HB + behavioral-source (B) distortion warning. RF/distortion correctness frontier, in the E-251 HB-proof mold. (1) PROOF: .disto (the 1990 Volterra small-signal distortion analysis) reports the pure 2nd/3rd-order kernels scaled by the DISTOF1/DISTOF2 magnitudes; HB (hb) and two-tone QPSS-HB (qpss) are INDEPENDENT large-signal engines including all orders, whose k -th harmonic / mixing amplitudes converge onto the Volterra result with $\sim A^2$ leakage. So HB(A) $\rightarrow$.disto as $A \rightarrow 0$ -- the HB-proof structure applied to distortion; because HB and .disto share ngspice's identical device model, model-constant ambiguity (VT, DC-op) cancels, so the agreement measures the .disto engine itself. Measured (diode + series R, both solvers): single-tone HD2 vs HB 2f tightens rel err $4.0e-5 \rightarrow 2.6e-7$ as $A: 1e-3 \rightarrow 1e-4$ (HB $\rightarrow$ disto), HD3 $\sim 4e-6$; two-tone IM3 (2f1-f2) vs QPSS-HB (2,-1) $\sim 1e-5$; .disto output scales EXACTLY as A^2/A^3 (DISTOF1/2 applied exactly, kernels amplitude-independent). The residual floors are the reference engine's own readout/

E-256 *src/math/ni/niiter.c, src/spicelib/analysis/cktop.c, src/include/ngspice/cktdefs.h, src/spicelib/parser/ptfuncs.c, bin//ngspice (rebuilt), examples/bsrcconv_examples/*

Fix SILENT spurious DC operating point for singular-derivative behavioral sources. A behavioral (B) source whose small-signal conductance is INFINITE at the $v=0$ initial guess -- $B I = \sqrt{v(n)}$ ($dI/dv = 0.5/\sqrt{v} \rightarrow \infty$), $I = 0.1/v(n)$, $I = \ln(v(n))$, $v^{0.5}$ -- pins the node at $v \sim 0$ so Newton takes a vanishing step; the convergence test (iterate-to-iterate VOLTAGE change only) sees no change and declares "converged" at a point that grossly VIOLATES KCL (0.42A imbalance on a 0.6V/1ohm divider). The false convergence also PRE-EMPTED the gmin/source-stepping fallbacks (they never ran). Silent + data-dependent ($i = 1m\sqrt{v}$ converges fine; a .nodeset finds the true root). The B-source AC derivative itself is correct (matches analytic $dI/dv \sim 1e-10$ across the function library) -- the defect is purely the DC convergence ACCEPTING a KCL-violating point. FIX (3 coordinated parts): (1) *niiter.c* false-convergence guard -- after the voltage test declares convergence, also check the KCL residual $F = Gx - b$ (the node-current imbalance, already computed for the E-111 line search); if the worst node merit $> 100x$ the current tolerance the point is spurious $\rightarrow$ decline convergence so CKTop falls through to gmin/source stepping, which regularizes the singular node and finds the TRUE point. Separation is decisive: false convergence sits at merit $\sim 1/\text{reltol} \sim 1000$, every legitimate circuit < 1 (measured $1e-7 \dots 0.35$). (2) First-attempt isolation -- a new one-bit CKTdcFirstTry flag (*cktdefs.h*) set by CKTop (*cktop.c*) ONLY around the initial plain-Newton attempt, so the guard never fires inside a convergence-aid sub-solve (gmin/source stepping, pseudo-transient/optran, convhelp); an earlier unconditioned version broke convhelp/ptcont/corenum/tempphys by firing inside their homotopies. (3) *ptfuncs.c* `pwr(0,negative)` guard $\rightarrow$ HUGE instead of raw $+\infty$ (matches the existing `/0` and `sqrt(neg)` guards) so a singular derivative stays finite rather than NaN-poisoning the Jacobian. Result-neutral for every circuit that already converged. Verify (examples/bsrcconv, both solvers): [1] $I = \sqrt{v}$ reaches analytic $v_0 = 0.178045$ with KCL satisfied ($i_{B1} = i_{R1}$), not the spurious $v \sim 0$; [2] $I = 0.1/v$ reaches upper-branch $v_0 = 0.887298$; [3] result-neutral: $v^2/v^3/\exp/\tanh$ unchanged. Full regression 212/212 incl. the convergence-aid suites convhelp/ptcont/corenum/tempphys. Scope: DC Newton loop + a one-bit CKT flag + a parser math guard; binary rebuilt; confined to the first operating-point attempt.

E-257 *src/math/ni/niiter.c, bin//ngspice (rebuilt), examples/bsrcconv_examples/ (verify [4])*

Extend the E-256 false-convergence guard to the TRANSIENT operating point. E-256 fixed the silent spurious DC operating point for singular-derivative behavioral sources but gated the guard to MODEDCOP only; the transient operating point (MODETRANOP -- the bias a .t ran starts from without uic) was left uncovered and still false-converged. So a .t ran of a biased $B\ I=\sqrt{v(n)}$ started from the spurious $v\sim 0$ and showed a FAKE startup transient ($t=0\ v(n)=9.3e-17$ ramping $0\rightarrow 0.178$ over $\sim 20ms$) when the biased circuit should sit at 0.178045 from $t=0$. E-256 gated to MODEDCOP conservatively because the pseudo-transient/optran aids run in transient modes -- but the guard already fires only when CKTdcFirstTry is set (the initial plain-Newton attempt), and optran runs as a FALLBACK with CKTdcFirstTry=0, so it is safe to broaden the mode condition to MODEDCOP

E-258 *src/maths/ni/niiter.c, src/spicelib/analysis/dctrcurv.c, bin/ngspice (rebuilt), examples/bsrconv_examples/ (verify [5])*

Extend the false-convergence guard to the .dc sweep cold-start point. E-256/257 fixed the silent spurious operating point for singular-derivative behavioral sources for .op (MODEDCOP) and the transient op (MODETRANOP), but a .dc sweep still false-converged on its FIRST point: .dc V1.. of a $\sqrt{v}$ source gave $v(n)=9.3e-17$ at the first sweep value (true root 0.178045) while every later point (warm-started from the previous solution, away from the $v=0$ singularity) was correct. Why the guard missed it: it fires only when CKTdcFirstTry is set, which happens only inside CKTop -- but in the default (non-HS) path the .dc sweep (dctrcurv.c) solves each point with a DIRECT NIiter call, bypassing CKTop, and only falls back to CKTop if that solve FAILS; the false convergence "succeeds", so CKTop (and the guard) never ran. FIX (2 parts): (1) niiter.c -- the E-256/257 guard, which enumerated the op modes (MODEDCOP

E-259 *examples/integaccuracy_examples/* (verification only -- no ngspice source change)*

Transient integration accuracy PROOF (oracle-based, both solvers). A permanent regression guard for the transient integrator, in the E-251 (HB) / E-255 (.disto) mold: rather than checking one waveform, it proves the integrator MATH PROPERTIES vs closed-form analytics + Richardson self-convergence. The audit that produced it swept the whole transient path and found it CORRECT. [1] ORDER of accuracy on RC decay $\exp(-t/RC)$: global error $\sim dt^p$ with the theoretical p -- TRAP $p=1.94$, Gear2 $p=1.93$, Backward-Euler $p=0.94$ (order test freezes the LTE step controller $trtol=1e11$ so steps are uniform and the convergence is purely truncation order). [2] ENERGY signature on an LC $\cos(\omega t)$ over 30 periods: TRAP amplitude ratio 0.9999 (marginally stable/energy-conserving, AND the trapezoidal-ringing BE-switch is NOT firing spuriously), Backward-Euler 0.057 (dissipative) -- the fingerprint of a correct trapezoidal vs BDF. [3] BREAKPOINT: a PULSE edge into RC matches the piecewise-analytic charge (max err $6e-7$, pre-edge pinned 0). [4] RLC damped sinusoid vs $\exp(-at)(\cos \omega_d t + (a/\omega_d) \sin \omega_d t)$ (max err $4.3e-6$, correct ω_d /envelope). [5] NONLINEAR charge: a diode+CJO rectifier converges under Richardson (TRAP $dt\rightarrow dt/2$) at order $p=2.00$, so the nonlinear device charge integration is 2nd order too. Both solvers give identical results (integrator is solver-independent). Separately confirmed (needs .options dynorder + set ngdebug, not asserted): higher-order Gear engages -- maxord=6 reaches order 6 and rejects $\sim 10x$ fewer steps than order 2 (14 vs 143 on the RLC), so the high-order BDF coefficients are functional not a no-op. Scope: verification only; no source/binary change. Regression 213/213 (new example folder).

E-260 *examples/integaccuracy_examples/* (verify [6]; verification only -- no source change)*

LTE step-controller accuracy PROOF on a stiff circuit. Extends E-259 (which verified the integration METHODS with a FIXED step, freezing the LTE controller to isolate truncation order) with the complementary property: the ADAPTIVE local-truncation-error step controller actually DELIVERS the accuracy requested via reltol -- on a stiff circuit, the hard case. Two independent decays $v_f=\exp(-t/1\mu s)$, $v_s=\exp(-t/1ms)$ (stiffness ratio 1000) integrated with one adaptive step (the controller must take tiny steps to resolve the fast decay then coarsen for the slow tail). Swept over reltol the delivered error vs the closed form shrinks MONOTONICALLY with no plateau: reltol

1e-3->1e-9 gives err(fast) 9.0e-3->2.0e-6 while internal steps grow 1032->1884. The error scales $\sim \text{reitol}^{0.6}$ -- the theoretical global-error rate for a 2nd-order method under LTE control (reitol is a LOCAL per-step tolerance; $\text{step} \sim \text{reitol}^{(1/3)} \Rightarrow \text{global err} \sim \text{reitol}^{(2/3)}$). Separately confirmed (not asserted): extreme stiffness ratio 1e6 ($\tau = 1\text{ns}/1\text{ms}$) still tracks reitol, and a re-excited fast mode (1MHz square wave) is handled by refining $\sim 6\times$ more steps at each edge. Verify (examples/integaccuracy [6], both solvers): on the stiff circuit err at reitol 1e-3/1e-5/1e-7 is strictly decreasing and improves $\geq 10\times$ (measured $\sim 247\times$). Completes the transient-engine proof (E-259 methods + E-260 step controller). Scope: verification only; no source/binary change. Regression 213/213.

E-266 *src/spicelib/analysis/cktsetup.c, src/spicelib/analysis/cktpzset.c, src/include/ngspice/cktdefs.h, examples/solverannounce_examples/* (binaries rebuilt by CI)*

Announce the direct linear solver once, not on every analysis. CKTsetup() (and CKTpzSetup()) printed Using SPARSE 1.3 as Direct Linear Solver (or the KLU line) unconditionally at the top of setup, and CKTsetup runs once per analysis -- so the sweep command, which sets the knob and re-runs the -analysis command for every point (alter/altermod in place, or alterparam+reset which re-sources the deck), reprinted the line on EVERY point (a 5-point sweep, 5 times; a 100-point sweep, 100). Monte Carlo, optimize, and the pso/de/sa optimizers -- all of which re-run the analysis per point -- were equally noisy. FIX: an announce-on-change helper CKTannounceSolver(int klu) (defined in cktsetup.c, prototyped in cktdefs.h, called by both cktsetup.c and cktpzset.c) prints the line only when the active solver differs from the last announcement. The last-announced solver is tracked PROCESS-WIDE (a function-static), not per-circuit, because a .param sweep re-sources the deck and rebuilds the circuit each point -- a per-circuit flag would still repeat. Consequences: a multi-point sweep/Monte/optimize run announces the solver ONCE; a single analysis still announces once (unchanged); a genuine solver switch (.option klu/.option sparse) has a different mode and re-announces; a fresh ngspice process starts un-announced, so batch runs and the dual-solver test harness (which re-execs once per solver) still print it once each, and the HB (verify_hb) and benchmark suites that detect KLU by grepping for the line still see it. Result-neutral: the line is informational and part of no numerical comparison; full dual-solver regression unchanged. Verify (examples/solverannounce, 4 checks): a 5-point sweep announces SPARSE once (was 5x); a single op still announces once; a sparse->klu switch re-announces both; a KLU analysis still prints the KLU line. Scope: 3 source files (one small helper + two call-site swaps + one prototype); no analysis/result change.

E-267 *src/frontend/com_sweep.c, examples/sweepbus_examples/* (binaries rebuilt by CI)*

sweep records array/bus nodes under their natural names (**ph[0]**, not **ph_0_**). Sweeping a circuit whose outputs are array/bus nodes -- a node named **ph[0]** (E-221) -- stored the result vectors under MANGLED names: **ph[0]** became **ph_0_**, **ph[1]** **ph_1_**, so the sweep plot listed **ph_0_**, **ph_1_**, ... The sweep command (E-146/-189/-190) builds each result-vector name from the output plus, for a -vs family or overlay waveform, an appended **_<knob>_<value>** segment; that segment carries a float (**_rt_1.5k**, with **./-/e** illegal in a nutmeg vector name), so the name builders **sw_familyname/sw_pointname** mapped every non-alphanumeric char (except **_**) to **_** -- but over the WHOLE string, INCLUDING the user's base output name, destroying a bus node's brackets. FIX (com_sweep.c): (1) a helper **sw_append_sanitized(base, suffix)** sanitizes ONLY the appended **_<knob>_<value>** segment, leaving the base name byte-for-byte intact, so **ph[0]** stays **ph[0]** (and a -vs family curve is **ph[0]_rt_2k**); ordinary names are unchanged (a plain node has no special chars; an explicit **-output name=expr** still uses the clean given name). (2) A bare **-output** token that is a bus range **ph[0:3]** now expands to one output per index (**ph[0]** **ph[1]** **ph[2]** **ph[3]**), matching the netlist bus expansion (E-221), so **sweep rt .. -output ph[0:3]** records the four taps directly. Values are unaffected -- only the NAME the result is stored under changes (and range-token expansion); each output is still evaluated exactly as before. Verify (examples/sweepbus, 4 checks, a 4-tap bus divider): **-output ph[0:3]** records four vectors named **ph[0]..ph[3]** (range expanded, natural names); the mangled **ph_0_..** names are gone; the recorded values are the correct divider ratios; a plain **-output vo=v(out)** is unaffected. Full dual-

solver regression unchanged (the existing sweep suites use the `-output name=expr` form, whose clean names are untouched). Scope: one source file; no analysis/result change.

E-268 *src/frontend/spiceif.c, src/frontend/device.c, src/frontend/com_sweep.c, src/include/ngspice/fteext.h, examples/sweepwild_examples/* (binaries rebuilt by CI)*

Wildcard model-parameter knob **@*[param]**. Multiple `.model` cards can share a model-parameter NAME (two Verilog-A models both with a `wavelength` parameter). Nothing let you sweep that shared parameter across all such models in ONE knob: `sweep @dev1[wavelength]` (altermod) targets one model, and the `.param` idiom (`.param wl, wavelength={wl}` in each card, `sweep wl`) works but its `.param` knob commits with `reset -- RE-SOURCING THE ENTIRE DECK` at every point (slow) and needing the cards to reference the `param`. NEW: **@*[param]** sets `param` on every loaded model that has it, IN PLACE (altermod, no reset), so `sweep @[wavelength]` ... co-varies all of them fast; it also works standalone (altermod `@*[wavelength]=1.55u`). IMPL: (1) `spiceif.c if_setparam_wildcard(ckt, param, val)` -- model params are a device-TYPE property, so one `paramlookup` per device type decides whether that type's models carry `param`; every model of a matching type (walking `ckt->CKThead[typecode]->GENnextModel`) is set with `doset` (the same path a concrete altermod `@model[param]=uses`); returns the count set and runs `CKTtemp` to push into live instances when called mid-run; models/types without the `param` are skipped. (2) `device.c com_alter_common` -- a wildcard device name `@*` routes to `if_setparam_wildcard`; a concrete `@model[param]` or device is unchanged; if no loaded model has the `param`, one clear warning. (3) `com_sweep.c sw_kind` -- **@*[param]** classified as an in-place altermod knob (like `@model[param]`), so the sweep applies it without reset; prototype in `fteext.h`. Additive -- `@*` was not previously a valid model name, so no existing alter/altermod/sweep behaviour changes. Verify (examples/sweepwild, 4 checks, both solvers): two `wlmodel` cards (`R=wavelength1k`) + *one unrelated model without wavelength*; `sweep @[wavelength]` co-varies BOTH `wlmodel` devices (equal currents tracking $1/(wavelength1k)$); the unrelated model is untouched; a concrete `@dev1[wavelength]` still targets only `dev1`; `@*<absent>` warns and changes nothing. Full dual-solver regression passes. Four source files.

E-269 *src/frontend/spiceif.c, src/frontend/device.c, src/frontend/com_sweep.c, src/include/ngspice/fteext.h, examples/sweepwild_examples/* (binaries rebuilt by CI)*

Instance wildcard **@#[param]** (alias **@*[[param]]**). E-268's **@*[param]** sets a MODEL parameter on every model that has it; this adds the INSTANCE counterpart -- set an instance parameter on EVERY device instance that has it, in place (alter/altermod, no re-source), so `sweep @#[scale]` ... co-varies an instance parameter across all instances. Two spellings: **@#[param]** (canonical; `#` marks instance) and **@*[[param]]** (alias -- the outer `@dev[param]` parse leaves `param` as `[param]`, whose leading `[` is stripped). IMPL: (1) `spiceif.c if_setparam_wildcard_instance(ckt,param,val)` -- instance params are a device-TYPE property, so one `paramlookup(...,do_model=0,...)` per type gates membership; every instance (walking each model's `GENinstances->GENnextInstance`) is set with `doset(dev=inst,mod=NULL) -> setInstanceParm` (the path a plain alter `@dev[param]=uses`). (2) `device.c alter_set` -- the device-name token picks the wildcard: **@#*** (token `#*`) or **@*[[param]]** (token `* + param` starting `[`) -> instance wildcard; **@*** -> model wildcard (E-268); else ordinary `if_setparam`; no match -> one warning. (3) `com_sweep.c sw_kind` -- **@#*/@*[[...]]** classified as in-place altermod-style knobs (no reset); prototype in `fteext.h`. **@#=models**, **@#*/@*[[...]]=instances**; concrete `@dev[param]` unchanged. Additive -- these were not previously valid names. Verify (examples/sweepwild, 7 checks, both solvers) on a `wlmodel` with a MODEL param `wavelength` and an INSTANCE param `scale` (`R=wavelengthscale1k`): model wildcard co-varies both model cards AND reaches `.model` cards inside SUBCIRCUITS (they flatten to per-instantiation model copies); instance wildcard **@#[scale]** and alias **@*[[scale]]** co-vary every instance; concrete `@dev1[wavelength]/@NX2[scale]` target just one; absent-param wildcards warn. Full dual-solver regression passes. Four source files.

E-270 *src/frontend/com_sweep.c, examples/sweepbounds_examples/* (binaries rebuilt by CI)*

sweep validates its numeric bounds (ASan/UBSan fuzz find). An ASan/UBSan instrumented build fuzzing the newest command parsers (sweep/alter/altermod/loadpull/rfstab/stb/pyplot) surfaced a bug in sweep's spec parser `sw_parse_spec`. `sweep` reads each bound of its `<start>` `<stop>` `<step>` (and `'lin`

E-271 *src/frontend/com_let.c, examples/letoob_examples/* (binaries rebuilt by CI)*

let no longer reads one byte before its buffer on an empty LHS (ASan/UBSan fuzz find). The same command-parser fuzz that produced E-270 flagged a one-byte out-of-bounds READ in `com_let`. `com_let` flattens everything after `let` into a heap string `p`, splits the RHS off at the first `=`, then NUL-terminates the destination vector name at the first `[`. It next trims trailing whitespace from the name: `for (q=p+strlen(p)-1; *q<=' ' && p<=q; q--); *++q='\0';`. When the LHS is a bare bracket -- `let [[ = ...`, where the leading `[` has just been turned into `'\0'`, leaving `p` an EMPTY string -- or is entirely whitespace, `strlen(p)==0` so `q` starts at `p-1`; the loop test evaluates `*q` BEFORE the `p<=q` guard can short-circuit, dereferencing `p[-1]` -- a read one byte before the allocation (ASan: heap-buffer-overflow READ of size 1 at `com_let.c`, from `let [[ = @#[r] v1 -o i(v1)`). The stray byte does not fault deterministically on the shipped build (almost always valid heap metadata), but it is a genuine OOB read on malformed input and could fault if `p` sat at the start of a page. FIX (`src/frontend/com_let.c`): test the bound FIRST -- `for (q=p+strlen(p)-1; p<=q && *q<=' '; q--)` -- so `&&` short-circuits and `q` is never read while it points below `p`. An empty or all-whitespace name then falls through unchanged to the existing "bad variable name" check, which rejects it cleanly. No valid `let` is affected (a one-line reordering of a loop condition). Verify (`examples/letoob`, 5 checks): the fuzz-found `let [[ = ..`, a single-bracket `let [ = 5`, and an all-whitespace LHS each error quickly with `bad variable name` and no crash/hang; a plain `let a = 5` and an indexed `let a[1] = 7` still assign correctly; the original input under an ASan build no longer reports the overflow. Full dual-solver regression passes. One source file.

E-272 *src/frontend/device.c, examples/alternull_examples/* (binaries rebuilt by CI)*

alter/sweep no longer deref a NULL parameter for an m-named device (ASan/UBSan fuzz find). The same command-parser fuzz crashed on `sweep mag(v(b)) 0 1k 5k 1k`. `com_alter_common` (`src/frontend/device.c`) supports altering a device's PRINCIPAL value with no named parameter -- `alter r1 = 2k` -- in which case the parsed param is NULL. Just before applying the value it runs a binned-MOSFET guard: `'if ((dev[0]=='m') && (eq(param,"w"))`

E-273 *src/maths/cmaths/cmath2.c, examples/mathcast_examples/* (binaries rebuilt by CI)*

`cmaths` integer operators guard their (**int**) casts against out-of-range doubles (ASan/UBSan fuzz find). Continuing the expression-layer fuzz flagged `cmath2.c:765` on `1e30 % 5`; it is one of a cluster where a routine casts a double operand to `int` with a plain cast (UB when the value is non-finite or outside `int` range). (1) `cx_mod` (the `%` operator) computed `(int) floor(fabs(op))` per operand -- `1e30 % 5` is UB (UBSan) and on the SHIPPED build returned a garbage 2 rather than an error (a silently wrong result). (2) `cx_vector/cx_cvector/cx_unitvec` set the result length with `(int)fabs(arg)` -- `vector(1e30)/unitvec(1e30)` produced a saturated (INT_MAX-ish) length and then allocated and filled it: a multi-gigabyte `calloc` (macOS over-commit lets it succeed) plus a billion-iteration fill loop, i.e. an apparent HANG on the shipped build (`rc=142` under a watchdog), not merely a sanitizer report. FIX (`cmath2.c`): range-check before every (`int`) cast. `cx_mod` floors each operand into a double, then `rchecks` it lies in `[0-or-1, INT_MAX]` before the cast; a non-finite/too-large operand takes the existing argument out of range for `mod` path (`rcheck -> EXITPOINT -> NULL`). A shared helper `cx_vec_len()` validates the length for the three vector builders (`fabs/cmag` are `>= 0`, so one `<= INT_MAX` check also rejects `NaN/inf`); an out-of-range length prints `vector length ... is out of range` and returns NULL, which `apply_func` already handles (it checks `if (!data)` return NULL). Valid expressions unchanged -- `17 %`

5 == 2, vector(4) == [0,1,2,3], unitvec(3) == [1,1,1], and any length within int range behaves exactly as before. Pre-existing (cmaths, not touched by the recent command work). Verify (examples/mathcast, 5 checks): 1e30 % 5, vector(1e30), unitvec(1e30) each error cleanly and fast (no UB, no hang, no garbage); valid 17 % 5 and vector(4) still evaluate correctly. Full dual-solver regression passes. One source file.

E-274 *src/frontend/evaluate.c, examples/idxcast_examples/* (binaries rebuilt by CI)*

An out-of-range vector index no longer invokes undefined behaviour (ASan/UBSan fuzz find). A fuzz of the expression layer flagged evaluate.c:779 on v(a) [1e308]. op_ind (vector indexing v[expr]) rounds the index with (int)floor(value+0.5); casting a double that is non-finite or outside int range to int is UB. The code ALREADY clamps the resulting index to [0, majsize-1] with a warning -- the only defect was the cast, which runs before that clamp. FIX (evaluate.c): a helper idx_floor() clamps the value to int range (NaN->0) before the cast, applied to all three casts (the real index and the real/imag parts of a complex [lo,hi] range index); v(a) [1e308] then resolves to the last element with the pre-existing warning, exactly as a large in-range index already did. Verify (examples/idxcast, 5 checks): vx [1e308], vx [1e30], a NaN index each resolve cleanly (no UB); valid vx [0] and range vx [0:2] still return the right elements. One source file; no change to any valid index.

E-275 *src/maths/cmaths/cmath4.c, examples/iffreal_examples/* (binaries rebuilt by CI)*

ifft() on a real vector no longer reads past its buffer (ASan/UBSan fuzz find). cx_ifft unconditionally did ngcomplex_t *indata = (ngcomplex_t *) data;. For a VF_REAL input, data is a plain double[length] (length8 bytes), so reading indata[i] for i<length walks length COMPLEX elements (length16 bytes) -- twice past the buffer (ASan heap-buffer-overflow READ: a 528-byte/33-element region read as 66 complex). cx_fft (Enhancement-225) already distinguishes real and complex input; cx_ifft did not. FIX (cmath4.c): for a VF_REAL input build a proper complex array (imag=0) sized length, use it, and free it before returning; a VF_COMPLEX input still points straight at data; a length>=2 guard (matching cx_fft) rejects a degenerate input. Verify (examples/iffreal, 3 checks): ifft(vx) on a real vector runs with no overflow; the fft->ifft round-trip keeps the length; complex ifft still works. One source file; no change to a valid complex ifft.

E-276 *src/maths/cmaths/cmath2.c, examples/rndcast_examples/* (binaries rebuilt by CI)*

rnd() guards its (int) cast against an out-of-range operand (ASan/UBSan fuzz find). cx_rnd (the rnd() random-integer function) turned each operand into a modulus with j = (int)floor(dd[i]); rand() % j. Casting a double outside int range to int is UB, so rnd(1e30) (or inf/NaN) tripped UBSan (cmath2.c:162). This is the same class as Enhancement-273, which hardened cx_mod and the cx_vector builders but not cx_rnd (at the time rnd(1e30) did not reproduce; the expression fuzz found the path). FIX (cmath2.c): a helper cx_rnd_i() clamps the value to int range (mapping NaN to 0) before the cast, applied to all three casts (the real operand and the real/imag parts of a complex operand); the result is unchanged for any in-range operand. Verify (examples/rndcast, 3 checks): rnd(1e30) and rnd(inf) resolve cleanly (no UB); a valid rnd(5) still yields an integer in [0,5). One source file; completes the E-273 (int)-cast hardening for cx_rnd.

E-277 *src/maths/cmaths/cmath4.c, examples/derivcx_examples/* (binaries rebuilt by CI)*

deriv() of a complex vector: fix a heap overflow AND a wrong result (ASan/UBSan fuzz find). The fuzz flagged cmath4.c:314 on deriv((0.5, time)) (complex input). cx_deriv fits a degree-d polynomial over a sliding window; its COMPLEX branch diverged from its (correct) REAL branch in three places, each an out-of-bounds access on the last block AND a numerical error: (1) it read the data window c_indata[j+i+base] while fitting against the scale window scale + i - degree + base -- offset by degree, so it fit mismatched (x,y) pairs and read degree points past the end (heap-buffer-overflow READ); the real branch reads indata + i - degree + base. (2) The real-part output loop used j <= i + degree/2 where the imag-part loop and the whole real branch use j <= i - degree/2, overrunning scale[]/c_outdata[]. (3) The tail indexed scale[j+base]/c_

outdata[j+base]; the real branch's tail uses j (its FIXME notes j+base crashed), so a grouped (base>0) complex derivative overran both arrays. FIX (cmath4.c, complex branch only): align to the real branch -- read c_indata[j + i - degree + base], bound the real-part output loop with i - degree/2, index the tail with j. Every overflow is gone AND the complex derivative is now numerically correct. Verify (examples/derivcx, 3 checks): deriv(0.5, time) runs with no overflow; deriv(t + 2t*i) now returns (1, 2i) -- the correct derivative, mid-vector; a real deriv(2t) still returns 2; a focused deriv stress over complex/real/grouped inputs is clean under ASan. One source file; real derivatives unchanged, complex derivatives now memory-safe and correct.

E-278 *src/maths/cmaths/cmath4.c, examples/scaleguard_examples/* (binaries rebuilt by CI)*

Transform functions integ/deriv/fft guard a length/scale mismatch (ASan/UBSan fuzz find). Continuing the expression-layer fuzz, integ/deriv/fft overran the heap on a vector whose length differs from its plot scale (cmath4.c:500 on integ(unitvec(200)), :1065 on fft(vector(5)), :407 on deriv(unitvec(200))). cx_integ/cx_deriv/cx_fft index both the data (by length) and its plot scale; a SYNTHETIC vector (vector(n)/unitvec(n)/an expression result) carries the current plot's scale, whose length need not equal n. Two directions: (1) data LONGER than the scale (integ/deriv of unitvec(200) on a 66-pt plot) -- the loops read pl_scale->v_realddata[i] for i<length, past the end of the scale; (2) data much SHORTER than the scale (fft(vector(5)) on a 66-pt plot) -- the Green's inverse-FFT sizes datax from N=pow2(length) but the output loop writes tpts(=scale length) points from datax[2*i], so N<tpts overran datax. Only cx_fft had been guarded (Enhancement-225); the siblings had not. FIX (cmath4.c): cx_integ/cx_deriv reject a data vector LONGER than its scale (the direction that reads OOB) with a clean error -- a SHORTER data vector, as produced by fft (freq-domain, where group-delay deriv(vp(3)) is legitimate), stays valid; cx_fft grows its transform size N (and M) to cover the output length tpts before allocating datax, a no-op for a well-formed transform where pow2(length)>=tpts already (the fft->fft roundtrip is unchanged). Verify (examples/scaleguard, 5 checks): integ(unitvec(200)), deriv(unitvec(200)), fft(vector(5)) resolve cleanly (no overflow); the fft->fft roundtrip still returns a full-length vector; valid integ(vx)/deriv(vx) still work; a transform stress over mixed lengths/types is clean under ASan. One source file; completes for integ/deriv/fft the length/scale guarding E-225 gave fft.

E-279 *src/frontend/com_let.c, src/frontend/options.c, src/frontend/com_measure2.c, src/maths/cmaths/cmath4.c, examples/castguard_examples/* (binaries rebuilt by CI)*

The remaining unguarded (**int**) casts and the last unguarded scale-dependent transform (systematic audit, not fuzz). After the fuzz rounds a systematic AUDIT -- grep every (int) floor/fabs(...) of a runtime double, plus every cmath routine taking a struct plot * and whether it guards its length against the scale -- found the tail of two already-fixed classes at sites the fuzzer never reached. (1) Unguarded (int)floor(x+0.5) of a user-supplied double (UB outside int range, the E-273/-274/-276 class): com_let.c (an index expression in an indexed let), options.c (set numdgt / rawfileprec / measureprec -- set numdgt=1e30 tripped UBSan at options.c:359, set rawfileprec=1e30 at :340), com_measure2.c (meas ... rise/fall/cross=<n> at :1580/:1584/:1588). (2) An unguarded scale-dependent transform (the E-278 class): cx_mtimeavg walks the scale data dsc[j] for j < length-1, so a vector longer than its plot scale -- mtimeavg(unitvec(200)) on a shorter plot -- read past the scale (ASan heap-buffer-overflow @ cmath4.c:1188). FIX: a small clamping helper per file (let_idx_int, opt_int, meas_int) converts to int only after clamping to int range (NaN -> 0), so an out-of-range option/count saturates instead of invoking UB; values already in range are unchanged. cx_mtimeavg gets the same length-vs-scale guard cx_integ/cx_deriv received in E-278. Verify (examples/castguard, 7 checks): set numdgt=1e30, set rawfileprec=1e30, set numdgt=-1e30, meas ... rise=1e308, mtimeavg(unitvec(200)) all clean where they previously tripped UBSan/ASan; a valid set numdgt=6 and mtimeavg(vx) still work. Four source files, guard-only.

E-280 *src/frontend/com_let.c, examples/letidxoob_examples/* (binaries rebuilt by CI)*

An out-of-range `let v[i] = x` wrote past the end of the vector (heap-buffer-overflow WRITE). Hardening the index cast (E-279) exposed a far more serious defect underneath. `get_index_values` (`com_let.c`) parses either one index (`v[i]`) or a range (`v[lo:hi]`). It validates `low > high` and `high >= n_elem_this_dim` -- but those checks lived INSIDE THE RANGE BRANCH ONLY. The single-index path set `high = low` and returned completely unchecked, so `let vx[100] = 1` on a 66-element vector walked straight into the byte-offset arithmetic and performed a heap-buffer-overflow WRITE (ASan: WRITE of size 8) -- memory corruption from an ordinary typo, no exotic input required; a very large index additionally overflowed the `index * n_byte_elem` product in `int` (`com_let.c:771`). The range form `vx[0:999]` was correctly rejected all along. FIX: move both checks out of the range branch so they validate a single index too; an out-of-range assignment now reports `index/high range (N) exceeds the maximum value (M) and changes nothing`. Reads are unaffected -- `op_ind` clamps an out-of-range READ index with a warning (E-274) -- and every valid assignment, including the last element and the range form, behaves exactly as before. Verify (examples/letidxoob, 6 checks): `let w[10]=99` on a 10-element vector, `let w[999]=99`, `let w[1e308]=99` each rejected cleanly where the first two corrupted the heap; the last valid index (`let w[9]=42`) and a mid-vector assignment still assign; the range form still rejected. One source file, moving two existing checks.

E-281 *src/maths/cmats/cmath4.c, examples/derivgroup_examples/* (binaries rebuilt by CI)*

`deriv()` over-read on a partial block (grouping != length) (ASan fuzz find). `cx_deriv` walks its input in blocks: `for (base = 0; base < length; base += grouping) for (i = degree; i < grouping; i += 1) ... fitting a window around i + base`. `grouping` is the vector's `v_dims[0]`; for an ordinary vector that equals `v_length`, so there is a single block and the window always fits. But a vector whose declared dimension differs from its length -- as produced by a binary op on operands of UNEQUAL length, e.g. `min(v(b), ac.v(b))` where a 66-point real and a 5-point complex vector yield length 66 with `dims[0] = 5` -- leaves a PARTIAL LAST BLOCK: `base` climbs to `length - 1` while the window still spans `base + grouping - 1`, reading past the end of the input (ASan heap-buffer-overflow READ @ `cmath4.c:326`, the same line E-277 had already realigned). FIX (`cmath4.c`): bound the inner loop with `i + base < length` in BOTH the real and complex branches, so the fit window can never reach past the input; a no-op whenever `grouping == length` (every ordinary vector), and a partial trailing block too short to hold a full window is simply skipped. Verify (examples/derivgroup, 4 checks): `deriv(min(v(b), ac.v(b)))` and `deriv(max(v(b), ac.v(b)))` clean where they previously over-read; an ordinary real `deriv(2t)` still returns 2; the complex derivative from E-277 is still correct (`deriv(t + 2t*i) = 1 + 2i`). One source file, a loop bound in each branch.

E-282 *src/frontend/plotting/agraf.c, examples/plotlabel_examples/* (binaries rebuilt by CI)*

`asciiplot` read past its axis-label buffer on a 3-digit exponent (audit lead). Following up the plotting (`int`)`floor(...)` casts noted during the E-279 audit turned up a different, real defect one layer over. `ft_agraf` allocates its two axis-label lines as `maxy + margin + FUDGE + 1` bytes (`FUDGE=7`) and NUL-terminates them at the last byte, budgeting the exponent width with `sprintf(buf, "%1.1e", 0.0); shift = strlen(buf) - 7;`. Formatting 0.0 ALWAYS yields a 2-digit exponent, so `shift` is 0 and the whole layout silently assumes two digits -- but real data can need three: denormal or very large values produce labels like `1.00e-320` (9 chars vs 8). The last label's `memcpy(&line2[i + margin - ((j<0)?2:1) - shift], buf, strlen(buf))` then runs one byte too far and overwrites `line2`'s terminating `'\0'`, so the following `out_printf("%s\n%s\n", line2, line1)` walks off the end of the heap allocation -- ASan reports heap-buffer-overflow READ of size 82 on an 81-byte region inside `vsnprintf`. On the shipped build it does not necessarily fault; whether it prints heap garbage or crashes depends on what follows the buffer. FIX (`agraf.c`): remember the allocation bound in `line_end` (`maxy` is REASSIGNED further down as `maxy = spacing*nsp`, so the bound cannot be recomputed at the label loop), clamp the label `memcpy` so it can never reach the terminator slot (handling a negative offset too), and re-assert `line2[line_end] = '\0'` after the loop. Verify (examples/plotlabel, 5 checks):

asciipLOT of denormal 1e-320, 1e300, 1e-300 data and of both extremes together are clean where they previously over-read; an ordinary plot still renders its legend, axis rule and points; rendering confirmed BYTE-IDENTICAL to the pre-fix binary. One source file; no change to any rendered plot.

E-283 *src/frontend/plotting/agraf.c, src/frontend/points.c, src/frontend/display.c, examples/plotcoord_examples/* (binaries rebuilt by CI)*

Plot coordinate math no longer casts a non-finite double to **int** (extreme-data fuzz). The same output/formatting fuzz that produced E-282 found three more sites where a plot coordinate goes through `mylog10()` -- which returns +/-inf for zero, denormal or overflowed values -- or through a division by a range that can be zero, and is then cast to **int** (undefined behaviour). All are reached by ORDINARY plot commands on pathological data. (1) `agraf.c`: the decade `mag = (int)floor(mylog10(...))` -- data that overflows ($-1e308 * 6 \rightarrow -\text{inf}$, so `ylims[0]` is `+inf`) or is zero/denormal makes the argument non-finite; the limits `lmt/hmt = (int)floor(ylims[0] / tenpowmag)` are equally exposed since `tenpowmag = pow(10, mag)` can itself be 0 or inf. (2) `points.c` `ft_findpoint(): (int)((mylog10(pt) - tl) / (th - tl)) * (maxp - minp) + minp` is 0/0 for a degenerate range (`lims[0]==lims[1]`, or log endpoints sharing a decade). (3) `display.c`: the four screen-coordinate casts (log and linear, x and y). FIX: per-site clamping chosen so the result stays MEANINGFUL, not merely defined -- `agraf_decade()` bounds the exponent by `DBL_MAX_10_EXP` (the widest decade a double can represent) and maps NaN to 0, `agraf_int()` clamps the `lmt/hmt` ratios into **int** range; `ft_findpoint` computes the fraction in double, maps NaN (degenerate range) to 0 and clamps to [0,1] so the result always lands in `[minp, maxp]`; `disp_clamp_coord()` clamps each screen coordinate into the viewport. Verify (examples/plotcoord, 5 checks): asciipLOT of overflowing ($-1e308$), denormal and all-zero data, and a min/max over denormal data, are clean where they previously invoked UB; an ordinary asciipLOT `v(b)` still renders. Rendering additionally diffed against the pre-fix binary and confirmed BYTE-IDENTICAL for ordinary data. Three source files, guard-only.

E-284 *src/frontend/spiceif.c, src/frontend/device.c, src/frontend/com_sweep.c, src/include/ngspice/ftext.h, examples/wildparam_examples/* (binaries rebuilt by CI)*

Wildcard parameter diagnostics, and the **sweep** knob label (reported from real use). A Verilog-A parameter `L_um` could not be swept as `sweep @*[L_um]`, and ngspice's message named '`L_um`' -- which reads like the mixed-case name had been mangled by ngspice's case-insensitive netlist tokens. The case theory is a RED HERRING: the OSDI parameter lookup IS case-insensitive (`@*[L_UM]` resolves `L_um` correctly, verified). The real cause is a LEVEL mismatch that the diagnostics `hid: @*[param]` (and `@#*[param]`) is the INSTANCE wildcard, while a plain parameter `real L_um` in Verilog-A is a MODEL parameter -- an instance parameter needs `(* type = "instance" *)`. So the instance wildcard correctly matched nothing and reported no loaded instance has parameter '`L_um`': lower-cased, with no hint that the parameter exists at the MODEL level, and nothing pointing at the fix. Separately, sweep labelled every knob it applies in place as (model param) -- including the instance wildcards -- because `sw_kind()` returns `SW_MODEL` as a DISPATCH flag ("apply via altermod, no re-source") and the banner reused that flag as a description. FIX: (a) `spiceif.c` gains a probe `if_hasparam_wildcard(ckt, param, do_model)` reporting whether any loaded model (or instance) carries a settable parameter of that name -- it only probes, nothing is changed; (b) `device.c` uses it so a wildcard that matched nothing checks the OTHER level and names the form that would work, in both directions (... but a loaded model does -- use the model wildcard '`@*[L_um]`' (an instance parameter needs `(* type = "instance" *)`) in the Verilog-A), and the mirror suggesting `@#*[...]`, while a parameter that genuinely does not exist still gets the plain message with no misleading suggestion; (c) `com_sweep.c` gains `sw_knobdesc()`, which classifies the knob from the NAME TOKEN, so the banner reads (instance param, wildcard) for `@#*[...]/@*[...]` and (model param, wildcard) for `@*[...]`. Verify (examples/wildparam, 8 checks) on a model with a mixed-case MODEL parameter `WaveLength` and a mixed-case INSTANCE parameter `L_um`: both cross-level hints appear

and name the right form; a truly absent parameter gets no bogus hint; sweep labels instance and model wildcards correctly; the matching wildcard still sweeps to the right values; and ALL-CAPS `@*[ [L_UM] ]` still resolves the mixed-case `L_um`, confirming the lookup was never case-sensitive. Four files; diagnostics and one banner string, no change to which parameters any wildcard actually sets.

E-285 `src/frontend/plotting/plotit.c, src/frontend/plotting/agraf.c, src/frontend/plotting/gnuplot.c, src/frontend/com_measure2.c, examples/veclenmix_examples/*` (binaries rebuilt by CI)

Output paths indexed one vector by another's length (extreme-data fuzz). The remaining findings from the extreme-data output fuzz all proved to be ONE class in four places: a vector's OWN length need not equal its plot SCALE's length -- any synthetic vector carries the current plot's scale, so let `y = vector(8)` on a 66-point transient plot has 8 points and a 66-point scale -- and a COMPLEX vector holds its data in `v_compdata`, leaving `v_realdata` NULL. (1) `plotit.c`: the transient resampling passed `v->v_realdata` together with `v->v_scale->v_length` to `ft_interpolate()`, which indexes the DATA by that length, reading ~58 elements past a shorter vector; for a complex vector it passed NULL outright -- a HARD SEGV on the shipped build (`asciiplot sqrt(-1*vector(10))` returns `rc=139`). (2) `agraf.c`: the bracketing indices lower/upper are computed against the X scale and bounded by `xscale->v_length`, but were then used to index each plotted VECTOR. (3) `gnuplot.c` (`wrdata`): the "no more data" guard tested only `i >= scale->v_length` while the branch it protects indexes `v->v_realdata[i]`. (4) `com_measure2.c`: the plain tran/dc branch read `d->v_realdata[i] / dScale->v_realdata[i]` with no NULL check, though the ac/sp branches beside it already had one -- measuring a complex vector dereferenced NULL (present twice, in `measure_at()` and `measure_deriv_at()`). FIX: clamp each index to the vector it actually addresses -- `plotit.c` passes the length the data and scale SHARE and skips the transient resampling entirely when any vector is not real (`all_vecs_real()`), since that block replaces the shared scale and must convert all vectors or none (complex vectors keep plotting through the normal path); `agraf.c` clamps lower/upper into `[0, v->v_length-1]` per vector; `gnuplot.c`'s guard also requires `i < v->v_length`; `com_measure2.c` takes the real part when `v_realdata` is NULL, exactly as its ac/sp branches do. Verify (examples/veclenmix, 7 checks): a vector shorter than its scale, a vector longer than its scale, a complex vector (previously `rc=139`), `wrdata` of a short vector, and a measure over a complex vector are all clean and still render/measure; an ordinary `asciiplot v(b)` and `meas tran ... max v(b)` are unaffected. Ordinary plot output AND `wrdata` output diffed against the pre-fix binary and confirmed BYTE-IDENTICAL; all eight decks from the output fuzz are clean. Four source files, guard/clamp only.

(E-25's `simparam` exposure lives in the OSDI callback table populated at load time; its diff rides inside the `osdiload.c/callbacks` changes.)

Build-warning cleanup (post-E-127): a full rebuild under the repo's `-Wconversion` flags surfaced latent warnings. `configure.ac` had `-s` (a strip flag) in the compile CFLAGS, so clang warned "argument unused during compilation: '-s'" on every source file (~1526 warnings) — removed (binaries are stripped separately by the fold and CI). `-Wconversion` also flagged a real latent bug in `maths/ni/nipzmeth.c` (`b = 0.0` was an assignment where `b == 0.0` was intended) — fixed — and benign `double→bool` conversions in `maths/cmaths/cmath4.c` — made explicit. Full-build warnings dropped 1903 → 394; the remainder are `-Wconversion` warnings in third-party SuiteSparse (KLU) and CMC device models (HiSIM), which are left to their upstream sources.

Every entry above is guarded by a verify script under `examples/` and was regression-checked against the full example arsenal, the compiler's integration suite, and the VA_TEST industry-model corpus at the time it landed. | **E-296** | `src/frontend/plotting/pyplot.c` | `pyplot` appearance controls. Seven additive set variables: `pyplot_grid` (on/off/x/y, overriding the grid-follows-axis default), `pyplot_legend` (off, or a matplotlib location -- underscore form since set keeps only the first word), `pyplot_markers` (a cycling marker `o s ^ D v * P X` ON the line, distinct from `pointstyle=markers` which drops the line), `pyplot_axhline/pyplot_axvline` (SI-aware reference lines via `ft_numparse`), `pyplot_dpi`,

`pyplot_transparent`. With none set the emitted script is byte-identical to before. Verify (examples/pyplotmore): each control appears in the script AND the script executes; the default path is unchanged. One file. || [E-297](#) | `src/frontend/plotting/pyplot.c`, `plotit.c`, `com_pyplot.c` | `pyplot -fft` magnitude spectrum. `ft_pyplot`'s `bool hist -> int mode (LINE/HIST/FFT)`. `-fft <sig>` reuses the time-domain (t,y) table; the emitted Python RESAMPLES onto a uniform grid (`np.interp`) before `np.fft.rfft`, because transient data is adaptively sampled and a raw FFT over non-uniform samples is wrong. One-sided, scaled by $2/\text{sum}(\text{window})$ so a pure tone reads back its amplitude. Options `pyplot_fft_window` (hann/hamming/blackman/rect), `pyplot_fft_db`, `pyplot_fft_points`, `pyplot_fft_logf` (log-f in Python, dropping the DC bin -- NOT the command's `xlog`, which would validate the `t=0` TIME scale and abort). Verify (examples/pyplotmore): a 2.0@1kHz + 0.5@3kHz tone reads 2.0 and 0.5 (amplitude oracle, $\text{rel} < 2\%$). Three files. || [E-298](#) | `src/frontend/plotting/pyplot.c`, `plotit.c`, `com_pyplot.c` | `pyplot` complex-aware AC views **-bode/-nyquist/-polar**. An ordinary `pyplot v(out)` on AC data silently keeps only the REAL part (0.5 not 0.7071 at -3 dB). New `ft_pyplot_ac()` emits a `<vec-index> <freq> <re> <im>` table -- the imaginary part is CARRIED -- grouped by the first column so multiple/ragged vectors stay separate; `-bode = 20log10|H| dB / unwrapped phase deg vs log-f stacked`, `-nyquist = imag vs real (equal aspect)`, `-polar = |H| at angle(H) on a polar projection`. Reuses the shared file/launch scaffolding and `pyplot_*` settings; auto-names `bode/bode-2/...` via the shared counter. Verify (examples/pyplotmore): RC low-pass at `fc` gives -3.010 dB and -45.00 deg (exact first-order values, from the preserved complex data). Three files. || [E-299](#) | `src/frontend/plotting/pyplot.c` | `pyplot` overlay robustness + hover cursor + export note. Cross-run overlay via ngspice's `plotname.vector` refs (`pyplot tran1.v(out) tran2.v(out)`) already worked, but the data table was sized by the FIRST vector's scale, so a longer second run was truncated; it now walks the LONGEST scale across the plotted vectors (each still uses its own scale for x, shorter ones pad with NaN which matplotlib skips; a single run is unchanged since all share one scale). `set pyplot_cursor` adds a hover crosshair via matplotlib's built-in Cursor widget (no extra package), interactive-only (not emitted for a hardcopy). Pan/zoom/save (matplotlib toolbar) and the `<name>.data` export were already provided. Verify (examples/pyplotmore): a finer second run keeps every sample; cursor present in a window script, absent from a hardcopy. One file. || [E-300](#) | `src/frontend/plotting/pyplot.c` | `pyplot pyplot_mplcursors` cursor backend. On top of E-299's `pyplot_cursor` (built-in `matplotlib.widgets.Cursor` crosshair), `set pyplot_mplcursors` selects the `mplcursors` package -- data cursors that snap to a trace and show its (x,y) on hover -- AND turns the cursor on by itself. Unset, the built-in crosshair remains the default. The emitted script imports `mplcursors` inside a `try` and FALLS BACK to the built-in `Cursor` on `ImportError`, so a deck written where `mplcursors` exists still runs where it does not. Window-only (not emitted for a hardcopy), like `pyplot_cursor`; `mplcursors` must be importable by the interpreter named by `pyplot_python`. Verify (examples/pyplotmore): the `mplcursors` branch + fallback are emitted on `pyplot_mplcursors` alone, `pyplot_cursor` alone still emits only the built-in `Cursor`, neither appears in a hardcopy; the `mplcursors` branch was executed under an interpreter that has the package (calls `mplcursors.cursor(hover=True)`) and the fallback under one that does not. One file. || [E-301](#) | `src/frontend/plotting/pyplot.c` | `pyplot pyplot_cursor` is the single cursor master switch. E-299 (built-in `Cursor`) and E-300 (`mplcursors`) had drifted: `pyplot_mplcursors` turned the cursor ON by itself, so there was no one "off" knob. The enable condition went from `pyplot_cursor OR pyplot_mplcursors` to `pyplot_cursor` alone -- off by default, the ONLY thing that enables any cursor; `pyplot_mplcursors` now only SELECTS the backend when the cursor is on and does nothing on its own. Predictable, and matching every other optional `pyplot` behaviour. Verify (examples/pyplotmore): the full gating truth table -- default off, cursor-alone -> built-in, `mplcursors-ALONE` -> off, both -> `mplcursors` (with fallback), off in a hardcopy. One condition, one file. || [E-302](#) | `src/frontend/com_measure2.c` | **.meas avg** did not clip its window to [from,to]. Found by a closed-form oracle (a same-binary comparison is structurally blind to a uniformly present error). `measure_minMaxAvg()` trapezoid-summed only the samples INSIDE the window and divided by their span, without interpolating either boundary -- so the first and last partial intervals were dropped. `rms/integ (measure_rms_integral)` already clip exactly, so ngspice's own two answers for one window disagreed: `avg=1.27114` vs `integ/(to-from)=1.27324`, closed form 1.2732395; and the echoed `to=` was the first sample OUTSIDE the window (2.50280e-04 for a requested 2.5e-4). $O(dt/\text{window})$ error: 1.6e-3 -> 4e-7 here. Fix interpolates both ends and reports `sprev` as the window end, all guarded to `AT_AVG`

so min/max/min_at/max_at keep whole-sample semantics. The dc branch is deliberately EXCLUDED -- a dc sweep may descend and enters its window at the high end, where a clip extrapolates (an early version of the fix turned a descending-sweep avg of 0.270250 into -3.79e+11); every dc measurement is byte-identical. Known remaining: `.meas dc avg` still truncates (0.270250 vs 0.2708333), and on a descending sweep `integ` returns 0.0 with `from= nan`. Verify (examples/measwindow): 28 checks vs closed-form integrals, windows landing BETWEEN samples; fails 14/28 on the pre-fix binary. | | [E-303](#) | `src/frontend/com_measure2.c` | `.meas dc avg` window truncation, both sweep directions. Completes E-302, which fixed tran/ac/sp but excluded dc because a sweep may DESCEND (dc v1 2 0 -0.001) and enters its window at the HIGH end, where an unguarded clip extrapolates outside the bracketing samples. `.meas dc avg` averaged over [first sweep point inside, last inside] instead of [from,to]: 0.270250 against a closed-form mean-of-x^2 of 0.2708333, echoing `to=7.49000e-01` for a requested 0.75, while `.meas dc integ` over the same window was already correct. Fixed direction-agnostically in `measure_minMaxAvg()`: the clip works from the ACTUAL crossing between the previous raw sample and the current one -- open the window at the boundary crossed on the way in, close it on the way out -- so it fires at `from->to` ascending, `to->from` descending, and never on a sweep that does not cross. Guarded to AT_AVG, so min/max/min_at/max_at keep whole-sample semantics. The echoed window reports the upper end of the range actually covered, so a descending sweep no longer prints the nonsensical `from= 0.25 to= 0.25`. STILL OPEN (different function, `measure_rms_integral`): on a descending sweep `integ/rms` return 0.0 with `from= nan` -- the loop meets the first sample already above `to` and interpolates with `index i-1 == -1`, an OUT-OF-BOUNDS read, then breaks with an empty array. Verify (examples/measwindow): 38 checks, the same window run in BOTH sweep directions vs closed form; 32/38 on the pre-303 binary. | | [E-304](#) | `src/frontend/com_measure2.c` | `.meas dc integ/rms` heap-buffer-overflow on a descending sweep. Completes E-303. On dc v1 2 0 -0.001 these returned 0.0 with `from= nan` (rms printed nothing). `measure_rms_integral()` builds its clipped window assuming the scale ASCENDS -- it clips at `from` going in and `to` going out, and the widths `x[i+1]-x[i]` must be positive for the Simpson/trapezoid sums. A descending sweep starts at the top, so the FIRST sample is already above `to`: the end-clip fired at `i==0` and interpolated against `index i-1 == -1`, reading `v_realdatal[-1]`. AddressSanitizer: heap-buffer-overflow ... `measure_interpolate com_measure2.c:195 / measure_rms_integral com_measure2.c:1415`. It then broke with a one-element array, so while (`i < xy_size-1`) ran ZERO times (sums 0) and `first` was never set (`m_measured_at` unwritten = the nan). Fix: walk the samples in order of INCREASING SCALE rather than index order -- the identity for ascending data, so tran/ac/ascending-dc are byte-identical -- and make the has-a-previous-sample guard the traversal position `n > 0` instead of `i > 0`, correct in both directions. `integ 0.0 -> 0.135416` (closed form 0.135416667); `rms -> 0.307458` (0.307459347). Also removed two dead reads in the `toVal < 0.0` tail that indexed `d->v_realdatal[]` with the INTEGRATION loop counter (past the end of the data) into variables nothing reads again, and took the window end from `xScale[v_length-1]` -- the smallest scale value when the sweep descends, not the last traversed. Verify (examples/measwindow): 44 checks, both sweep directions for avg/integ/rms vs closed form; 40/44 on the pre-304 binary; and the OOB is PROVEN gone -- ASan reports it before the fix and is clean after. | | [E-305](#) | `src/math/misc/randnumb.c`, `src/include/ngspice/randnumb.h`, `src/frontend/com_sweep.c`, `src/frontend/commands.c` | **wcd** -- worst-case distance / MPFP high-sigma. The named remainder of the statistical suite, and the complement E-150's own write-up says it does not implement. `highsigma` uses direction-free scaled-sigma IS; `wcd` works in STANDARDISED NORMAL space and asks a geometric question -- which point of the failure region is most probable? With the margin as `g(u)>0` for pass, that is the boundary point closest to the origin; its distance `beta` is the worst-case distance and, the density being spherical, `P_fail = Phi(-beta)` -- the sigma number a designer quotes. Hasofer-Lind/Rackwitz-Fiessler iteration with forward-difference gradients: cost BOUNDED at a few iterations of (1+ndim) sims, versus the 1e6-1e9 samples plain MC needs merely to SEE a 4.5-6 sigma event. --is N refines FORM (exact for a linear margin, approximate when it curves) with mean-shift IS centred on the MPFP, carrying the exact ratio `phi(z)/phi(z-u*)`: 986 failures seen in 2000 shifted samples for a 3.4e-6 event. Implemented as two modes behind the EXISTING single Gaussian-draw funnel `mc_sample_gauss()`, so it composes with LHS/SSS rather than duplicating them: `MC_MODE_WCD` makes draws deterministic (dimension `d` returns a chosen `u[d]`, so the deck is a plain function `g(u)`) and `MC_MODE_SHIFT` does the shifted sampling. Dimen-

sionality is DISCOVERED (a deck's Gaussian-draw count is unknown until it has been evaluated once) via a running maximum cleared per evaluation -- a single margin evaluation triggers several deck-copy passes and a later pass that draws nothing would otherwise wipe the count. Verify (examples/wcd): 19 checks vs the ANALYTIC GAUSSIAN TAIL, never a previous build -- FORM is exact for a linear margin so beta is exact to machine precision at 3/4/5/6 sigma and P matches Phi(-beta) to $\sim 1e-8$; a 2-D case puts the MPFP on the symmetric point ($4/\sqrt{2}$, $4/\sqrt{2}$), which is what shows the search finds the right DIRECTION and not a 1-D degenerate answer; a nominal-already-failing case gives signed beta=-3. The four existing statistical suites pass unchanged. | | [E-306](#) | src/math/cmaths/cmath4.c | The E-241 twin: **fft()** expression function normalised by the padded N. E-241 fixed an amplitude normalisation that divided by the ZERO-PADDED transform size instead of the input length, in the **fft** COMMAND (frontend/com_fft.c). The identical mistake survived in math/cmaths/cmath4.c **cx_fft** -- the vector-expression function reached by **let F = fft(v)**, a separate implementation of the same computation. Found by continuing the oracle campaign that produced E-241, and located with E-241's own discriminator: bin 0 is unambiguous (no window/scalloping) and $X[0] = Dlength$ for a DC offset D, so dividing by N reads back $Dlength/N$. On one signal (4001 samples padded to 4096, DC 2.0) the COMMAND read 2.000000 and the EXPRESSION read $1.953613 = 2.04001/4096$. NOT a convention question: **cx_fft** holds TWO complete implementations (complex-input and real-input) and in BOTH the FFTW branch already used the input length $((double)length)/2.0$, fpts) while Green's radix-2 used the padded size $((double)N)/2$, (double) N) -- the correct version sitting a few lines from the wrong one, an internal contradiction like avg vs integ in E-302. HAVE_LIBFFTW3 is undefined in this build so Green's is live. Fixed both Green branches to the input length. INDEPENDENT confirmation: nothing here touches **ifft**, yet **ifft(fft(x))** went $2.3e-02 \rightarrow 1.1e-16$ (the pair only inverts when the forward normalisation is right), so **cx_ifft** was audited and deliberately left alone. Full audit of every Green-kernel caller: **com_fft.c** **fft x2** and **spec/PSD x2** correct (E-241); **cx_ifft** correct; **trannoise/1-f-code.c** cannot pad (n_pts grown to 2^n_exp by construction); **fft/ifft** are the only transform functions in the expression table so there is no spec twin. Verify (examples/fftexpr): 6 checks vs closed form (DC bin, round trip padded and unpadded, complex-input bin 0 = $|mean(H)|$); 3/6 on the pre-fix binary. | | [E-311](#) | src/frontend/measure.c | **param/expr** measurements now work in a **.control** block. Found oracle-checking the less-common **.meas** modes. The **param/expr** measure types are handled only by **do_measure()**'s second pass (via **nupa_eval**); the interactive/**.control** **meas** command (**com_meas**) skips **do_measure** and calls **get_measure2()** directly, which has no **param/expr** case -- so **.meas tran pp param='a1-a2'** (dot-card) worked but **meas tran pp param='a1-a2'** (**.control**) failed with 'no such function as ...', down to a bare **param=a1**. Fix: in **com_meas** the prior results (**a1,a2**) are already single-valued ngspice VECTORS, so a **param/expr** measurement is **let <name> = (<expr>)** via the ordinary vector evaluator -- detected BEFORE **com_meas**'s single-valued-vector substitution loop (which would mangle the expression), reassembled from the tokens after **param=/expr=**, delimiter single quotes stripped, and re-lexed with **cp_lexer** so a spaced expression tokenises as the shell would. Works: **param=a1-a2** (unquoted), **param='a1-a2'** (quoted no-space), **param={sqrt(a1*a1 + a2*a2)}** (braces, spaces OK), **expr=**. Normal **meas** types (max/min/avg/when/...) route through **get_measure2** unchanged (byte-identical vs shipped). KNOWN LIMITATION left in place: **param='sqrt(a1*a1 + a2*a2)'** (single quotes + internal spaces) is mangled by the **.control** shell's own quote pre-expansion UPSTREAM of **meas** -- use braces; the dot-card handles it (no shell tokenizer). Verify (examples/measparam): 8 checks vs closed-form triangle values; 7/8 fail on the pre-fix binary. | | [E-312](#) | src/xspice/cm/cm.c, src/xspice/icm/analog/s_xfer/cfunc.mod | XSPICE integrating code models are second-order accurate. Found oracle-checking the **s_xfer** Laplace transfer function against a closed-form transient: its error fell only as $O(h)$, not the $O(h^2)$ every native ngspice storage element delivers. Two independent causes. (a) SHARED across all integrating code models (**s_xfer**, **int**, **d_dt**, ...): **cm_static_integrate()**'s TRAPEZOIDAL order-2 arm was a self-flagged stand-in -- **cur = -0.5*ag[0]*intgr[1]** with a compensating **geq *= 0.5**, carrying its own /* WARNING - This code needs to be redone */ -- which works out algebraically to backward Euler ($y(n)=y(n-1)+hu(n)$), a first-order method. Fixed to the true rule $y(n)=y(n-1)+(h/2)(u(n)+u(n-1))$; the "one previous value" the WARNING said was missing is $u(n-1)$, stored in the SPARE state double **cm_analog_alloc** already reserves per integrator (**doubles_needed = bytes/sizeof(double)+1**), rotated through the **CKT-states[]** history for free -- the standard SPICE companion-model idiom of keeping a state and its rate in adjacent slots. **geq** is now **ag[0]** like the order-1 and Gear arms; a guarded legacy fallback preserves

old numerics only where no spare exists (no real code model). (b) SPECIFIC to s_xfer: its controller-canonical feedback read the fed-back integrator states from the PREVIOUS timestep (cm_analog_get_ptr(i,1)), making the loop explicit/lagged one step, which caps H(s) at O(h) on its own even after (a). Now reads current-iteration states (offset 0) so the loop is solved implicitly within each Newton step (the DC/init pass already used offset 0). Both TRAP and Gear now converge O(h^2). Convergence verified vs an exact 1/(1+tau s) sine oracle (error ratio ~2 per h-halving before, ~4 after) and, for stability of the now-implicit feedback, vs SciPy signal.lsim across five transfer functions -- 1st-order, a zeta=0.05 near-oscillatory resonator, a 3rd-order Butterworth, a stiff pair with poles three decades apart, and one with a numerator zero -- all converge and match to <2.2e-4. Verify (examples/sxferorder): 6 checks under both solvers, including a self-contained convergence-ORDER test (ratios ~4, fitted p~2.1) that fails on the pre-fix binary with no reference build required. 246-example regression green under both solvers. || [E-315](#) | src/spicelib/analysis/tfanal.c, src/spicelib/analysis/cktic.c, src/spicelib/analysis/distoan.c | Three analysis crashes on adversarial-but-valid input -> clean errors (command/netlist fuzz, E-222..228 / E-270..285 family). [6] .tf on a SINGULAR circuit -- a dangling inductor (l1 2 3 1, nodes 2,3 floating) makes the operating point singular; tfanal.c IGNORED CKTop's return, so the matrix was never factored and the following SMPsolve() asserted IS_VALID(Matrix) && IS_FACTORED(Matrix) (spsolve.c:137, SIGABRT). Fix: propagate the CKTop error before solving. [7] a second .pz over a URC device -- CKTic zeroes the RHS vectors with a loop the compiler vectorises into memset(CKTrhs,...), but here ckt->CKTrhs is NULL, so it wrote through a NULL pointer (SIGSEGV in platform_memset). Fix: CKTic returns cleanly when the RHS vectors/matrix are unallocated. [8] .disto with NO distortion sources (a plain resistor) -- distoan.c's output section left OUTpBeginPlot's result unchecked, so on failure acPlot stayed NULL and OUTattributes(acPlot,...) dereferenced it (SIGSEGV at OUTattributes+268, addr 0x28); the first plot checked its result but the five output-section plots did not. Fix: capture and check OUTpBeginPlot at each output plot, bail on NULL acPlot. Legitimate .tf/.pz/.disto unaffected -- guards fire only on the failure paths (valid .tf still gives 0.5 on a 1k:1k divider). Verify (examples/ngcrashanalysis): 4 checks. || [E-316](#) | src/frontend/com_measure2.c | .meas AVG ended one timestep short of to. Found continuing to oracle-check .meas against its own contract (E-302/303/304 family): AVG is the same quantity as INTEG/(to-from) over the identical window. measure_minMaxAvg()'s final window-clip (added by E-302, which made AVG integrate a trapezoid over [from,to]) guarded the WHOLE accumulation with !AlmostEqualUlp(svalue, meas->m_to, 100), so when the first out-of-window sample fell within 100 ULPs of to the entire final trapezoid [sprev,to] -- a full timestep -- was skipped. AVG's window then ended one timestep short of to: for a 2.5e-7 timestep the AVG window echoed to=6.45990e-04 while INTEG (correct) echoed to=6.46240e-04, exactly 2.5e-7 apart, so avg != integ/(to-from) by ~1.6%. measure_rms_integral() (INTEG/RMS) does NOT have this bug -- it always adds the final point, interpolating to to only when the sample overshoots, else using the ~to sample as-is. Fix: move the !AlmostEqualUlp guard so it gates only the INTERPOLATION, not the accumulation, matching INTEG -- the final trapezoid is now always added, and AVG covers the full [from,to] window (AVG==INTEG/(to-from), rel 1.6e-6). MIN/MAX keep whole-sample semantics (branch is AT_AVG-only). Verify (examples/measavgwin): AVG==INTEG/(to-from) + AVG window reaches to, both fail on the pre-fix binary; the E-302/303/304 measwindow suite (44 checks) and E-311 measparam still pass. | [E-318](#) | src/spicelib/devices/vsrc/vsrcload.c | SFFM/AM voltage sources dropped the DC offset before the delay. Found in the correctness campaign, oracle-checking transient waveform sources. vsrcload.c evaluated SFFM and AM with if (time <= 0) value = 0; (after subtracting TD), returning 0 instead of the quiescent value at time<=0 -- so an SFFM/AM source's OPERATING POINT (common TD=0) was 0 not VO, and over any pre-delay window it read 0, injecting a spurious startup transient (an SFFM(2,0.001,...) ~ 2V source into an RC started at 0V and charged to 0.785V, vs a DC-2 source flat at 2V). Three oracles show VO is correct: the SIN case in the SAME function holds VO+VAsin(phase) at time<=0; ngspice's OWN current-source SFFM (isrload.c) has NO such zeroing and evaluates the formula directly (the two SFFM implementations in one binary disagreed); SIN/PULSE/EXP/PWL all preserve their offset. Fix: hold the waveform's time=0 value -- SFFM value = VO + VAsin(phasec + MDIsin(phasem)); AM value = VO + (VMO + VMAsin(phasem))sin(phasec); (both reduce to VO for zero phases). No example/regression circuit used SFFM/AM, so nothing else changes; the current-source path was already correct and is untouched. Verify (examples/sffmoffset): SFFM->1.5, AM->2.0 over a delayed window (SIN control->1.5), all failing on the pre-fix binary. || [E-319](#) | src/frontend/com_qpss.c | Transient-form QPSS

leaked the fundamental into every mixing bin. Found in the correctness campaign via the RF steady-state oracle (a QPSS steady state must equal the HB-form and a long transient + DFT). The transient-form qpss <expr> <f1> <f2> [periods] [maxorder] (as opposed to ... hb K1 K2) runs a two-tone transient then computes each 2-D harmonic $k_1f_1+k_2f_2$ by a *TRAPEZOIDAL integral of $v(t)\exp(-j2\pi f t)$ over the "last period" $[tt[i0], tt[end]]$ of the RAW transient grid. That window is not exactly the beat period $T=1/\gcd(f_1,f_2)$ ($tt[i0] \geq wstart$ but not $=$, short by up to one tstep), is non-uniform, and has non-periodic endpoints, so the large DC/fundamental leaks $\sim tstep/T$ into EVERY mixing bin. On a LINEAR two-tone RC (every product with $|k_1|+|k_2| \geq 2$ is analytically exactly 0) every bin read $\sim 5.8e-4|dominant\ line|$ (~ -45 dB), scaling *LINEARLY with the DC line (the leakage signature)*; confirmed against the HB-form ($\sim 1e-16$) and a plain .tran + uniform DFT ($\sim 1e-10$). Fix: resample the last beat period onto a UNIFORM grid over exactly $[wend-T, wend)$ (linear interpolation) and Fourier-project with the rectangular rule -- for commensurate tones every harmonic $k_1f_1+k_2f_2=mfb$ completes an integer m cycles in T, so a uniform DFT over exactly T is exact. The spurious floor drops ~ 4 decades to $\sim 6.6e-8$ (~ -122 dB). Real products are unchanged: the qpss_examples strong IM3 ($3.75e-4$) still matches the built-in to 4 digits, and a single-tone diode's (2,0) matches the HB-form to 5. Honest limit: the residual floor is grid-limited (identical across 4/16/64 settling periods), so the very weakest products (below ~ -120 dB) remain floor-limited -- for those the exact HB-form is the tool; but the old -45 dB floor masked far more. The ... hb form was already exact and is untouched. Verify (examples/qpssleak): linear-circuit mixing products $< 1e-5$ of the fundamental (pre-fix $\sim 7e-3$); qpss_examples (12 checks) still pass. || [E-320](#) | src/frontend/com_sweep.c, src/frontend/numparam/{spicenum.c,xpressn.c,numparam.h,numpaif.h} | Sweeping a .param no longer forces a full circuit reset at every point. A .param has no live binding into the built circuit -- numparam folds each {expr} on a device line to a fixed numeric literal at parse time -- so the only way to change one is alterparam (rewrite deck text) + reset (re-parse + subckt-expand + CKTsetup + matrix REORDER), per point, even though topology never changes. E-320 adds a fast path to the sweep command: when every swept .param feeds ONLY addressable top-level device/model VALUES, com_sweep per point overrides the swept symbol(s) in the retained numparam dico (nupa_add_param), refreshes the derived-param closure (new nupa_recompute_params replays the dependent .param lines), re-evaluates each captured device value against the dico (new nupa_eval_expr -> the numparam formula() evaluator), and pushes it into the live instance with a resolved direct set (ftsim->setInstanceParm = CKTparam -> DEVparam), refreshing each touched device type once via DEVtemperature -- NO reset, so the matrix ordering/factorization pattern is preserved. The template is the temper machinery (a temperature-dependent device value re-evaluated + alter'd in place every .dc step); the difference is that a parse tree cannot re-read an arbitrary .param cell (only time/temper/hertz bind to live ckt storage), so re-eval goes through numparam not IFeval. Two optimizations: each @dev[param] target is resolved to (instance, param-id) ONCE at arm time (no per-point lex or device-list walk), and identical value expressions are grouped so each unique expression is evaluated once per point. Conservative classifier: arms only when a scan of the ORIGINAL (pre-expansion) deck finds every swept-param occurrence to be a top-level device/model value slot (positional principal, or key={expr} named); DISARMS to the exact reset path on any subckt-body use (expression not retained past flattening), structural use (param in a node position, an instance/model/subckt name, a subckt call arg, .if/.elseif, .temp, an analysis card, .option, .ic, .nodeset, .global), or a derived .param (a non-swept param defined from a swept one). Because a disarm reproduces today's behaviour exactly, a sweep can only get faster, never change result. Measured on a 2500-stage ladder, 120-point sweep: param feeding ONE device in a large circuit $2.91s \rightarrow 0.26s$ (11.1x), param feeding ALL 5000 devices $3.42s \rightarrow 1.33s$ (2.6x); fast matches reset to $\sim 7e-9$ (numparam value-string formatting only), subckt fallback bit-identical. Scope: the sweep command (instance values fast, model values via textual altermod fallback); subckt-internal params and the optimize/Monte-Carlo .param reset sites stay on the reset path. Verify (examples/paramfastsweep): top-level param arms + matches closed-form $R2/(rval+R2)$; subckt param falls back yet identical. || [E-321](#) | src/frontend/com_sweep.c, src/frontend/numparam/{spicenum.c,numpaif.h} | .param fast-sweep, sub-circuit tier -- subckt-INTERNAL device values now take the fast path. E-320 was limited to top-level values: a swept param feeding a device inside a subckt disarmed to reset, because numparam freezes each {expr} to a literal during flattening. The expression is recoverable: a flattened instance card (r.x1.r1 in out 1e3) carries card->linenum back to its subckt-DEFINITION body line (R1 a b {rval}),*

a top-level template whose original text numparam retains in dicoS->dynrefptr[linenum] (two instances of a subckt share the linenum). So capture now walks the FLATTENED deck (ci_deck) and, per device card, recovers the value expression via the new nupa_get_dynref(card->linenum) while the card's own first token gives the full hierarchical instance name -- top-level and subckt-internal captured by one mechanism, no name reconstruction. (Empirically the key is card->linenum, NOT card->linenum_orig, which uses a different counter -- inpcom.c line_number vs line_number_inc -- and is off by one.) SAFE across subckt scope, since a subckt can shadow/derive-from/pass-through a swept global: (1) a Pass-1 scan of the original deck INCLUDING subckt bodies disarms on a swept param in a subckt header default, a subckt-call (X) argument, a subckt-body .param that shadows (swept LHS inside a subckt) or derives from (swept RHS) the param, or any structural slot (node/name, .if, .temp, analysis card, .option, .ic, .nodeset); (2) an arm-time SELF-CHECK re-evaluates each captured expression against the GLOBAL dico at nominal params and requires it to reproduce the value numparam already baked into the flattened card (read straight off the card, no ask-API), so any shadow/pass-through that slips past (1) disarms the whole path; (3) a disarm falls back to the exact reset path. Validated by a 10-deck battery -- 5 that must arm (top-level, direct subckt use, multi-instance, nested subckt, scaled expr) and 5 that must disarm (local .param shadow, formal-param shadow, formal pass-through, subckt-local derived, top-level derived) -- where EVERY deck's fast output equals reset bit-for-bit and every arm/disarm decision is as expected; the subckt fast path matches closed-form $R2/(rval+R2)$ to $9e-10$ and reset to $1.9e-9$. Speedup follows fan-out: a param feeding one device inside a subckt in a large fixed circuit $2.92s \rightarrow 0.24s$ (12.0x); a param feeding every device (400 subckt instances $\times 8 R$) $2.50s \rightarrow 1.64s$ (1.5x). Scope: the sweep command (subckt-internal instance values fast, model values via altermod fallback); optimize/Monte-Carlo .param reset sites stay on the reset path. Verify (examples/paramfastsweep): top-level arm + subckt-internal arm + local-shadow fallback, all vs closed form (4 checks). | | [E-322](#) | src/frontend/com_optimize.c, src/frontend/com_sweep.{c,h} | .param fast-sweep, optimizer tier -- the **optimize** command shares the fast-path engine. Every -dparam (symbolic .param) objective evaluation used to do alterparam + reset (re-parse + rebuild the whole circuit), and an optimization runs tens to hundreds of evaluations. E-322 exports the E-320/321 capture-classify-selfcheck-push engine (sw_fp_build / sw_fp_apply / sw_fp_free, de-static'd + declared in com_sweep.h) and wires it into the optimizer: opt_fp_arm() collects the OPT_DECKPARAM knob names and calls sw_fp_build ONCE before the search; opt_eval (Nelder-Mead / Levenberg-Marquardt / PSO / DE / SA) and opt_eval_objs (NSGA-II) push the candidate's deck-param values in place with sw_fp_apply instead of alterparam+reset; sw_fp_free at cleanup. The conservative classifier and arm-time self-check are unchanged, so a deck-param feeding a subckt shadow / structural slot / derived param disarms and the optimizer runs exactly as before. Two opt-outs: (1) -center design-centering (E-206) keeps reset -- its inner Monte-Carlo RE-SAMPLES process variation on each reset, a re-draw not a deterministic value push; (2) a device-count guard keeps the reset path on SMALL circuits -- the in-place apply has a fixed per-eval cost (numparam re-eval + dico ops) roughly independent of circuit size while a reset's cost grows with the deck, crossing at ~80 device instances, and because the in-place values differ from reset in the last few digits (numparam value-string formatting) an extremely tight -tol on a tiny circuit could otherwise send the two paths to different iteration counts; so opt_fp_arm counts flattened device instances and engages only at ≥ 80 . Correctness: across NM/PSO/DE the fast-path optimum is IDENTICAL to reset (122-device ladder fit: $r_{top}=361.914$ under NM and DE, 382.872 under PSO), and NSGA-II arms and returns the same Pareto front; all 41 existing optimize verify checks still pass. Speedup: a 4000-device fixed ladder, single -dparam knob, Nelder-Mead to convergence $0.94s \rightarrow 0.17s$ (5.6x); grows with circuit size, guarded to the sub-30ms reset on tiny circuits. Scope: optimize -dparam knobs, all methods except -center (-param/-mparam were already in-place). Verify (examples/optimize): a large-circuit -dparam optimization arms the fast path and still converges (2 new checks). | | [E-323](#) | src/frontend/com_optimize.c, src/osdi/osdiinit.c, src/include/ngspice/osdiitf.h | .param fast-sweep, OSDI-aware optimizer guard. E-322 gated the optimizer's fast path behind a device-count guard (arm at ≥ 80 device instances) because on a small circuit a reset re-parses faster than the fast path's fixed per-eval overhead. That threshold was calibrated on RESISTORS and is wrong for OSDI (compiled Verilog-A) devices: a resistor reset just re-parses an element line, but an OSDI reset RE-INITIALIZES each compiled model instance (re-runs its setup/temperature callbacks), ~30x costlier per device. Measured (fast path on vs off, same binary, OSDI-resistor ladder): the fast path wins at 3

OSDI instances (2.1x), scaling to 3.2x at 161 -- where resistors needed ~80 to break even. So the E-322 count guard would keep a SMALL OSDI -dparam optimization (common in device model extraction/fitting) on the slower reset path, forgoing 1.5-3x (a 300-instance OSDI fit measured 3.0x fast vs reset, same optimum). Fix: the crossover is about reset COST not device COUNT, so `opt_fp_arm` now walks the built circuit's instances and WEIGHTS each by device kind -- `primitive=1`, `OSDI=OPT_FP_OSDI_WEIGHT=30` (measured ~3 OSDI ~ = the reset cost of ~80 primitives) -- and compares the weighted total to the same 80 primitive-equivalent threshold. A handful of OSDI devices arms the fast path; a small all-resistor circuit still keeps the cheap reset. OSDI types are identified by a stable marker: every OSDI SPICEdev is built by `osdi_create_spicedev` and shares the `OSDIparam` instance-parameter setter (no built-in device uses it), exposed via a one-line `osdi_devtype_is_osdi(type)` helper in `osdiinit.c`; the optimizer call is `#ifdef OSDI`-guarded so a build without OSDI is unaffected (every instance weighs 1). Correctness is unchanged from E-322 -- the guard only decides WHETHER to arm, not what the fast path computes -- so a small OSDI -dparam optimization now arms and converges to the SAME optimum as the reset path. Verify (examples/optimize): a small OSDI -dparam optimization (compiled `optres.va`) arms the fast path and converges (2 new checks); all 43 optimize checks pass. || [E-338](#) | `ngspice-46/src/frontend/inpcom.c` | A full 64-bit bus range HUNG `ngspice` and grew without bound -- 3.6 GB after 4 s, 7.6 GB after 9 s, from ONE netlist line, in the released binary. `R1 a[-9223372036854775808:9223372036854775807] r=2k`. Found by fuzzing the array/bus node feature. ROOT CAUSE -- THE GUARD COULD NOT SEE THE REAL WIDTH: E-221 expands `base[lo:hi]` into scalar node names behind a deliberate `BUS_MAX_WIDTH` (8192) ceiling, computed as `long width = (hi >= lo) ? (hi - lo + 1) : (lo - hi + 1)`. `strtol` saturates at `LONG_MIN/LONG_MAX`, so the full span makes `hi - lo + 1` OVERFLOW a signed long -- undefined behaviour, which in practice wrapped to a small value, so the guard saw a tiny width and let it through. The loop then stepped from `LONG_MIN` toward `LONG_MAX`, about $1.8e19$ iterations, appending to a string each time. The guard was not missing: it was computing its input with the very overflow it existed to prevent. ONLY THE FULL SPAN TRIGGERS IT -- the one-sided cases `a[-9223372036854775808:0]` and `a[0:9223372036854775807]` have width ~ $9.2e18$, which fits and is correctly rejected. That is exactly why a corpus using only 32-BIT extremes (`a[-2147483648:2147483647]`, width $4.29e9$) missed it: that width is representable, so the guard worked. FIX: compute the span in UNSIGNED arithmetic (exact for every pair of longs) and compare the SPAN rather than the width, so the `+ 1` cannot overflow either; `strtol` setting `ERANGE` now also rejects an endpoint too large to represent rather than letting it be silently clamped to `LONG_MIN/LONG_MAX`, values that then look like an ordinary acceptable range. A netlist is untrusted input, so this is a resource-exhaustion defect reachable from a single line with no diagnostic and no bound. FUZZ: 432 decks, 0 crashes/hangs/runaway -- 54 range tokens (32- and 64-bit numeric extremes, malformed brackets, missing parts, nesting, control characters, XSPICE `%vd[...]` groups, a 4000-character base name, 40-deep bracket nesting) crossed with 8 placements (device nodes, subckt call, .subckt port list, value position, model name, print `v()`, bare print, let). CORRECTNESS: 13/13 numeric checks -- ascending AND DESCENDING ranges map ports in the right order, verified through an ASYMMETRIC subcircuit (1k on one port, 10k on the other) so a reversed mapping changes the measured current instead of being invisible; `.subckt bus4 d[3:0]` expands descending and maps positionally across four distinct currents; `R1 a[0:1] 2k` is ONE resistor (series $2k+4k = 0.5$ mA); a bus element and an explicit scalar are the SAME node; `v(a[0])` and bare `a[0]` resolve correctly (E-224) while ordinary vector indexing still works; non-integer/malformed/already-scalar tokens are untouched; the guard boundary expands at 8192 and stays literal at 8193. Verify (examples/busoverflow): 5 checks. || [E-339](#) | `ngspice-46/src/frontend/parse.c` | `v()` with THREE OR MORE node names SIGSEGV'd `ngspice` -- found by fuzzing the `pyplot` command family, but the defect is in the SHARED EXPRESSION PARSER. `print v(in,out,0)`, `let q = v(in,out,0)` and `pyplot v(in,out,0)` all died; `plot` SURVIVED, which is why an obvious spot-check passes and this was easy to miss. ROOT CAUSE -- TWO PATHS WITH DIFFERENT OWNERSHIP: `v()` takes at most two node names (`v(a)` or the differential `v(a,b)`), and a third was never rejected. `PP_mkfnnode`'s comma branch recurses, `PP_mkbnode(MINUS, PP_mkfnnode(func, left), PP_mkfnnode(func, right))` then `free_pnode(arg)` -- that branch CONSUMES its argument, while the normal path at the end of the function BORROWS it (`p->pn_left = arg; pn_use++`). `free_pnode` recurses into children guarded by `pn_use`, so: with TWO names the children are plain nodes,

each recursive call takes the BORROW path and bumps pn_use to 2, and the outer free merely decrements -- safe; with THREE OR MORE one child is ITSELF a comma node, so the recursive call takes the CONSUME path and frees it, and the outer free_pnode(arg) then walks into that already-freed child -- a DOUBLE FREE. Node existence is irrelevant: all-present v(in,out,0), all-missing v(a,b,c) and mixed v(in,out,zz) crash alike -- it is purely the arity. FIX: reject the invalid arity where it is detected (Error: v() takes at most two node names.) rather than relying on the inconsistent ownership. Repairing the mismatch instead would mean changing PP_mkfnnode's contract for every caller in order to make a syntactically invalid expression work; the arity is not legal in the first place, so rejecting is both smaller and more honest. The ownership mismatch is recorded as the underlying hazard. FUZZ: 122 pyplot invocations, 0 crashes/hangs/unexpected writes after the fix -- the eight mode flags (-eye -hist -contour -smith -fft -bode -nyquist -polar) crossed against missing/extra/repeated/conflicting arguments, hostile signal lists (unclosed v(, empty v(, 8000-character names, control characters, 200 signals) and hostile output base names (3000 characters, spaces, quotes, shell metacharacters, unicode, leading -). NOT A DEFECT, flagged and dismissed: pyplot ../name and pyplot /abs/path write outside the working directory -- that is the documented pyplot [file] plotargs contract and matches wrdata/write, so the fuzz policy was wrong rather than the command. Verify (examples/vfuncarity): 5 checks. || [E-341](#) | ngspice-46/src/frontend/com_sweep.c | **sweep -analysis reset** and **-analysis remcirc** SIGSEGV'd -- found by fuzzing sweep/optimize for RECURSION and STATE. sweep runs its -analysis argument as a command at every point, dispatched by name through sw_run_cmd, and nothing stopped a user naming reset or remcirc -- which free and rebuild (or remove) the very circuit the sweep is iterating over. The loop carried on with its resolved knob bindings, the old CKTcircuit and the old plot. *WHY THE FIX REJECTS UP FRONT RATHER THAN RECOVERING: the obvious repair -- notice the circuit changed and break out -- was tried FIRST AND MADE THINGS WORSE. The sweep's post-loop plot finalisation touches the same freed state, so breaking out merely moved the crash, AND it turned a previously WORKING case (-analysis 'optimize ...', which resets internally but re-establishes its own state) into a NEW crash. That attempt was reverted. Rejecting a destructive analysis before the loop starts cannot regress anything, and nothing legitimate is lost: a sweep re-resources the deck itself when a knob needs it, so a user reset in the per-point analysis has no purpose. optimize is deliberately NOT rejected -- it resets internally, recovers, and works today. FUZZ (E-270 already covered numeric bounds, so this EXTENDED THE STRATEGY rather than repeating seeds): ROUND 1, argument surface -- 684 invocations, CLEAN -- 20 knob kinds (instance param, @model[param], .param, the E-268/269 wildcards @[p]/@#[p]/@[p]), malformed @-forms) x 29 sweep specifications (start/stop/step, lin|dec|oct, list, plus missing/reordered/non-numeric) x 31 option shapes, plus 51 optimize invocations across all seven methods (nm lm pso de sa nsga nsga2), all three knob kinds and every objective form at degenerate values. ROUND 2, recursion and state -- WHERE THE BUGS WERE -- 23 cases: sweep inside sweep (same knob, three deep), optimize driving a sweep and vice versa, repeated sweeps into one plot, -overlay accumulation, -warm with and without a prior run, fast-path and reset-path knobs alternating, and the analysis mutating the circuit under the loop. ALSO FOUND, CHARACTERISED, NOT FIXED: a sweep near SW_MAXPTS looks like a hang because point cost grows with point count -- 171 us/point at 1000, 761 at 5000, 2994 at 20000 (20x the points for 300x the time, roughly QUADRATIC). Profiling puts it in DCop -> OUTpBeginPlot -> cp_getvar -> cp_usrvars -> cp_enqvar -> dup_string: cp_getvar materialises the ENTIRE user-variable list (a duplicated string per entry) and that list grows with the plots a long sweep accumulates, while OUTpBeginPlot calls it once per analysis for printinfo and interp. Making the construction LAZY was tried and MEASURED TO GIVE NO IMPROVEMENT (those variables are normally unset, so the search falls through to the materialised list anyway) and was reverted rather than shipped unmeasured; a real fix means making cp_usrvars() cheap or caching it with proper invalidation -- a general ngspice change outside a crash fix. Recorded with its measurements so it need not be rediscovered. Verify (examples/sweepanalysis): 5 checks. || [E-342](#) | ngspice-46/src/frontend/options.c, variable.c, variable.h | Two ownership bugs in the synthetic user-variable list: a heap-use-after-free (SIGSEGV) and a double free (SIGABRT). cp_usrvars() synthesizes \$plots, \$curplot, \$curplottitle, \$curplotname and \$curplotdate and returns a list its CALLERS free. That only works if the list owns every node. Two places assumed it did. [A] cp_enqvar() does NOT always allocate -- it clears its tbfreed out-parameter and returns a BORROWED pointer into plot_cur->pl_env when that live list already defines the requested name. cp_*

usrvars() declared tbfreed, passed it to all five calls and NEVER READ IT, so it relinked the borrowed node (tv->va_next = v, orphaning the live list's tail) and the caller then freed a node somebody else still owned. Not hypothetical: a rawfile Option: line writes pl_env, so a loaded rawfile containing Option: plots = 1 was enough -- SIGSEGV for all five names, ASan confirming allocate/free/use across three different functions (raw_read -> cp_setparse, cp_getvar -> free_struct_variable, cp_usrvars -> cp_enqvar). Fixed by honouring the flag and copying a borrowed variable; the five near-identical blocks became one loop so the rule cannot be applied to four of five sites. New var_copy() is written as the deliberate DUAL of free_struct_variable() -- deep, because that function recurses into va_vlist and a shallow copy would only move the bug. [B] found while fuzzing the fix for [A] and INDEPENDENT of it, no file involved: cp_remvar() freed the variable it looked up even when nothing had unlinked it. Only the US_OK arm unlinks; plots returns US_READONLY and the curplot names return US_DONTRECORD, leaving the node inside uv1 to be freed a second time by the trailing free_struct_variable(uv1). unset plots aborted in malloc on a plain deck. Fixed by tracking whether the node is genuinely ours to free -- which also closes the same latent corruption for variables found in variables, pl_env and ci_vars. The ci_vars borrowed path is NOT reachable for these names today (plot_cur is always set, so all five return earlier) -- measured, not assumed -- but the fix keys on the ownership flag rather than the source list, so it covers that path anyway. Fuzz: 640 combinations (5 names x 8 Option: value shapes x 16 follow-up sequences), each run twice, plain and under ASan, all clean. No leak from the copy: 20000 iterations through cp_usrvars(), peak RSS 7.1 MB vs 7.2 MB control (LeakSanitizer unavailable on this host, so measured rather than asserted). Behaviour preserved and DIFFED against the old binary, not eyeballed. Found by tracing the E-341 performance finding through cp_enqvar() and asking what happens when tbfreed comes back 0. Verify (examples/usrvarown): 6 checks. || [E-343](#) | ngspice-46/src/frontend/options.c, variable.c, include/ngspice/cpextern.h | The $O(N^2)$ sweep cost E-341 documented but did not fix: FIXED, 26.6x at 16000 points (39.55 s -> 1.49 s). cp_getvar(name,...) looks up ONE name, but its first statement built ALL FIVE synthetic user variables (\$plots, \$curplot, \$curplottitle, \$curplotname, \$curplotdate), searched the result and freed it again. Synthesizing \$plots walks the live plot list and copies a string per plot, so every call cost $O(\text{number of plots in the session})$ regardless of what was asked for; beginPlot() does two such lookups per analysis ('printinfo', 'interp') and a sweep creates a plot per point, so point k paid $O(k)$. WHY THE FIRST ATTEMPT FAILED (the whole lesson): E-341 tried making the construction LAZY -- build the list only after the variables search misses -- and measured 1.0x. printinfo and interp are normally UNSET, so that search always misses and falls through to needing the list anyway; laziness only helps the lookups that were never the problem. The gate has to be on WHICH NAME IS REQUESTED, not on when the list is built. Neither hot lookup is one of the five, so the fix builds NOTHING at all. cp_usrvars() is kept for cp_vprint(), which genuinely prints all five. Safe because both gated callers already search pl_env and ci_vars separately, which is the only other thing cp_enqvar() could return; ownership still follows E-342. Measured: 7.3x / 13.0x / 18.9x / 24.1x / 26.6x at 1k / 2k / 4k / 8k / 16k points -- the speedup GROWS with N, which is what removing a quadratic looks like. The regression suite itself, same 274 examples, went 748 s -> 479 s, because every analysis in every example was paying two of these lookups. NOT LINEAR YET, and the remainder is characterised rather than hidden: profiling the NEW binary at 64000 points puts 89% of what is left in plot_alloc(), a different function, which scans the whole plot list with a case-insensitive compare to pick a unique name. Two fixes were considered and REJECTED HERE: an exact-membership index (correct, but plot_list is mutated in ~9 places and a missed insertion hands out a DUPLICATE plot name), and a cheaper superset index (contained, but CHANGES BEHAVIOUR -- today plot names are reused after destroy all, and a superset would skip past them). Both are changes to plot lifetime, not to this lookup path. Semantics diffed byte-for-byte against the old binary, including a deck with printinfo and interp explicitly SET across op/tran/ac. E-342's two vectors re-checked. ASan clean. Verify (examples/sweepscale): 5 checks. || [E-344](#) | ngspice-46/src/frontend/com_sweep.c | .model params now take the fast sweep's DIRECT SET -- the one tier E-320 left on the textual path. Model params now cost exactly what instance params cost (they were 2.1x-2.8x slower). The E-320..323 fast .param sweep resolves each swept device value to a slot ONCE and pushes it with ft_sim->setInstanceParm, skipping the per-point reset -- but sw_fp_resolve() bailed on if (!ckt || b->mod), so every model bind ran altermod <model> <param> = <value> through sw_run_cmd on EVERY point: snprintf + cp_lexer + com_alter_common argument re-parse + finddev() resolving the model BY NAME + a string walk of the model parameter table, all to reach the one setModelParm that mattered. None of that can change between

points. Fix: `sw_fp_resolve()` resolves model binds to (*GENmodel*, type, param-id) at arm time and `sw_fp_apply()` pushes them with `ft_sim->setModelParm`. The parameter-table search instance binds already used was factored into `sw_fp_find_parm()` and `SHARED`, so both tiers apply identical rules (`IF_SET`, `IF_REAL`, keyword match, or the positional principal -- which only instances have, a .model line has no principal slot). Anything unresolvable (model not built, or a non-real param such as an integer level) keeps `rok=0` and its textual command: the fallback is the correctness backstop, NOT dead code, and an example check exercises it deliberately. The banner now reports fallbacks -- ..., N via `alter/altermod` -- appended ONLY when something fell back, so ordinary output is unchanged and a slower textual push no longer hides behind the same banner as a direct one. THIS IS A SPEED CHANGE, NOT A CORRECTNESS FIX, and that shaped the verification: BEFORE writing any code the textual path was checked against reset on a semiconductor resistor and on a compiled OSDI model and matched EXACTLY, so 'identical to reset' became the whole acceptance criterion. Measured (same binary): 200 models x 200 points 0.095 s -> 0.034 s (2.8x); 400 R models x 400 pts 0.36 -> 0.12 s (3.1x, and fast-vs-reset 2.4x -> 7.5x); 60 OSDI models x 400 pts 0.07 -> 0.03 s (2.2x, fast-vs-reset 1.7x -> 4.0x). The real result is the model/instance cost ratio landing on exactly 1.00x -- no residual per-point resolution left. Correctness: a 12-case battery run twice (fast and reset) and compared -- resistor rsh/tc1, capacitor cj, diode is, MOSFET vto, BJT bf, OSDI model r, two models sharing one param, model+instance mixed, derived-param (disarms as before), non-real integer param (falls back) -- ALL exact; plus subckt-internal models (arm + exact) and a subckt that locally shadows the swept name (correctly disarms). optimize shares the engine (E-322) so `mparam` inherits the direct set; optimize_examples 43/43 unchanged. Verify (examples/modelparamset): 7 checks, closed form to 1e-10 (deck raises `numdgt` so the comparison is not capped by the default 7-digit print). Monte Carlo remains the only tier not on this path -- it needs a re-DRAW of process variation, a different mechanism. | | [E-345](#) | `ngspice-46/src/frontend/vectors.c`, `com_fft.c`, `linear.c`, `postcoms.c`, `spec.c`, `mw_coms.c` | The residual O(N) E-343 profiled but rejected fixing: FIXED. The sweep is now LINEAR -- 87x at 64000 points (54.97 s -> 0.632 s). `plot_alloc()` and `plot_add()` pick a plot name by counting a shared, monotone `plot_num` up until `<abbrev><plot_num>` is not the typename of any plot in `plot_list`. Plots are PREPENDED, so the colliding probe hits at the head in O(1); it is the SUCCESSFUL probe -- proving a name ABSENT -- that walks the whole list, with a tolower per character. One full walk per plot created, and a sweep creates a plot per point. Only the MEMBERSHIP TEST changed: a hash index of the typenames currently in `plot_list` answers it in O(1). The search still starts at the same shared `plot_num` and still counts up by one, which matters more than it looks -- `plot_num` is shared ACROSS abbreviations and only advances on a collision, so the sequence is `op1 tran1 op2 ac2 op3 not op1 op2 op3`; and a number freed by `destroy all` is REUSED, which is exactly why E-343 rejected a 'remember every name ever issued' superset cache. Both copies of the loop (`plot_alloc` and `plot_add` carried the same one) now share `plot_unique_typename()`. E-343's OTHER objection -- an exact index is unsafe because `plot_list` is mutated in ~9 places and a missed insertion hands out a DUPLICATE plot name -- was answered by REMOVING the nine places rather than tracking them: `plot_new()` is now the SINGLE insertion point (six callers open-coded the same two lines: `com_fft x2`, `linear x2`, `postcoms`, `spec`), `killplot()` calls the new `plot_forget()`, `com_removecirc()` (which rewrites `plot_list` wholesale) calls `plot_forget_all()`, and the index builds LAZILY from `plot_list` so the static `const plot` needs no registration. PROVEN, not asserted: `plot_name_taken()` carries a `PLOTNAME_SELFHECK` block that answers the same question the OLD way and `abort()`s on disagreement; the full suite was run with that build -- 276/276 OK, ZERO disagreements. The block stays in the source, unset, so it can be re-run against future changes. A REAL BUG was hit on the way and is worth recording: converting the open-coded prepends with a REGEX produced a DOUBLE INSERTION, because three sites ALREADY called `plot_new()` with the redundant two lines wrapped AROUND it, so the edit turned the trailing line into a second `plot_new()`. The second call ran `new->pl_next = plot_list` when `plot_list` was already `new` -- a SELF-LOOP -- and `vec_gc` spun on the circular list forever. Found by a stress case (fft) that hung, located by SAMPLING the hung process, not by guessing. A mechanical bulk edit needs each site read: the pattern being replaced was not the pattern actually there. Names verified IDENTICAL to the previous binary across 15 stress cases (mixed analyses, destroy all, destroy first/middle/last, setplot new, fft, linearize x2, spec, write/load x2, load-then-destroy, twelve-then-destroy, remcirc, sweep-then-destroy, nested destroy). Measured: per-point cost pinned at 10 us and the doubling ratio converged on

exactly 2.0x -- the first time this sweep has been linear. With E-343, a 16000-point sweep went 39.55 s -> 0.163 s, about 240x. Verify (examples/plotname): 7 checks. || [E-346](#) | ngspice-46/src/frontend/com_sweep.c | A BUG FIRST, then the last tier: the fast **.param** path was silently FREEZING random draws that the reset path re-drew -- and fixing that IS the Monte Carlo tier. The fast sweep only re-evaluates brace expressions that MENTION the swept name. But numparam INLINES a .param's expression into the device line during preprocessing, so .param rv = agauss(...) + R2 a b {rv} arrives as r2 a b {(agauss(...))} -- carrying the random call but NOT the swept name, so it was skipped and never touched again. (Same discovery: nupa_recompute_params, added in E-320 to refresh the derived-param closure, replays ZERO lines on such a deck -- there is no .param line left.) Measured over 5 points: reset gave rv = 1097.7/1015.2/952.5/988.7/969.2, fast gave 1161.1 FIVE TIMES. E-320's guarantee that results never change, only speed, had not held for any deck with a random .param since it shipped. Fix: capture brace expressions that draw from the RNG (agauss/gauss/unif/aunif/limit/mvnorm) even with no swept name present. Once that exists MONTE CARLO NEEDS NOTHING MORE -- a sample is a re-draw plus an in-place push, and montecarlo arms the same machinery with NO swept knob (sw_fp_build(NULL,0)). Bit-identical to reset required THREE things, each found by a test that failed first: (1) DECK ORDER -- random binds sort ahead of the expression-grouped ones by capture index, since only they consume the RNG, so the stream matches however the deterministic binds are grouped; (2) NO CACHING -- a random bind is never served from the by-expression-text cache, because two devices with identical text must draw INDEPENDENTLY (verified: two {agauss(1000,300,3)} resistors give 1113.1 and 1099.5, not equal); (3) the SAMPLE BOUNDARY -- a deck re-copy signals one MC sample via nupa_signal(NUPADECKCOPY) -> mc_sample_advance(), which rewinds the per-sample dimension counter and steps the Latin-Hypercube sampler; skipping the re-source skipped that signal, which is exactly why plain draws matched reset (240/400 both) while -lhs did NOT (239 vs 258) and -lhs -warm did not (252 vs 258). Raising the boundary in the apply path made all four agree. The E-321 arm-time self-check is SKIPPED for random binds -- re-evaluating one draws a NEW value so it could never reproduce the baked one, and the draw would perturb the very stream that must match; their safety rests on the structural disarm, which were extended so a random gets the same scrutiny a swept name gets (a random reaching .temp, a subckt call or any structural slot disarms). Verified against an INDEPENDENT ORACLE, not a golden number: the same experiment hand-written as a reset loop in the control language -- montecarlo 124/200, oracle 124. Battery of 12 cases run on BOTH paths with discriminating yields (60.3/57.7/27.7/84.0/89.7/90.8%): plain, -lhs, -warm, gauss, aunif, limit, two randoms, same-expression-twice, random model param, two specs, derived random (disarms), structural random (disarms) -- all identical. Speed 1.8x/2.3x/2.4x at 20/200/1000 devices -- smaller than the sweep tiers because an MC sample re-draws and re-pushes EVERY random value; the saving is the re-source, not the pushes. CHANGES EXISTING RESULTS: a sweep over a deck with a random .param now re-draws it per point instead of freezing it -- which is what makes the fast path agree with reset again. Regression 278/278. Verify (examples/mcfastpath): 7 checks. || [E-348](#) | ngspice-46/src/spicelib/parser/inp2dot.c, src/spicelib/analysis/dcpss.c | **.pss** segfaulted on a short argument list -- and, separately, on a COMPLETE one whose **harmonics** was 0. Found by a new sweep shape: take a KNOWN-GOOD invocation of each command and run every PROPER PREFIX of its argument list. Existing campaigns could not see this -- cmdfuzz mutates argument VALUES (every value it supplies is present) and the parser fuzzers feed malformed FILES; a command that reads its fifth argument after checking only that a first exists is invisible to both. Across 38 commands pss was the ONLY offender: 0 args hung, 1-4 args SIGSEGV; qpss (8 args) and optimize (12) truncate cleanly. The stripped binary gave only the shape -- ldr d0, [x19] with x19 == 0, then fdiv, then a store -- and the UBSan build named it: dcpss.c:4990, load of null pointer of type 'double', i.e. Mag[0] = Phase[0] / (double)ndata. THREE failures stacked. (1) dot_pss() reads seven required values with no end-of-line check, and INPgetValue() cannot report 'nothing left to read' -- it returns 0 -- so a short card reached the analysis with points/harmonics silently 0. (2) DCpss() never validated: CKTharms is the LENGTH of every array the DFT writes into, and TMALLOC(double, 0) is tmalloc(0), which returns NULL. (3) DFT() writes index 0 OUTSIDE any loop -- Mag[0], Phase[0], nMag[0], nPhase[0], Freq[0] -- so its numFreq bound never protected them. That third point is why a parser guard ALONE would have been insufficient: pss 1meg 1u out 1024 0 50 5u is well-formed at full arity and crashed by the same route, a second entrance that never passes

the truncation check. Fixed at all three layers: a PSS_NEED_ARG macro naming the missing argument (the idiom dot_dc() already uses a few hundred lines earlier in the same file, which dot_pss simply never adopted); DCpss() rejecting fguess <= 0, points < 1 and harms < 1 before anything is sized from them; and DFT() guarding its own unconditional index-0 stores. WHAT WAS DELIBERATELY NOT TIGHT-ENED, because it was measured rather than assumed: every parameter was tested at 0 and negative, and sc_iter and steady_coeff are harmless at both (rc=0), so they are still accepted -- narrowing the accepted input on no evidence could break a working deck. NO VALID ANSWER MOVED: 22 correct invocations (op/dc/ac/tran/tf/pz/sens/noise/disto/four/pss plain and uic/pac/pnoise/hb/sp) were run on the shipped and rebuilt binaries and every printed number compared EXACTLY -- all identical, the one apparent difference being Total analysis time, i.e. the clock. Truncation sweep 0 crashes from 5. Verify (examples/pssargs): 7 checks. | | [E-349](#) | ngspice-46/src/spicelib/parser/inp2dot.c | One transposed letter in a node name killed the process -- **op** then **tf v(ouy) v1** and the session is gone, with every loaded circuit, plot and let-defined vector. Seven commands did it: tf, pz, noise, sens, pss, pxf, psp. Found by a STATEFUL CROSS-COMMAND fuzz -- random sequences of state-mutating commands followed by a canary circuit whose answer is known exactly (1 V across a 1k/3k divider, v(b) must be 0.75). The harness was first aimed at the pre-E-342 binary, where it found 38 aborts in 200 sequences, so a clean run means the detector works rather than that it is asleep. ROOT CAUSE: INPtermInsert() is CREATE-OR-FIND. That is right for a device card, where the device DEFINES its nodes (R1 in out 1k is what brings out into existence), and wrong for an analysis card, which only REFERS to one. A mistyped name was quietly created as a new unconnected node. The milder consequence is arguably the worse bug -- before setup nothing crashes and the invented node reads back as a perfectly plausible v (bogus) = 0.000000e+00, no error, no warning. The fatal one: from the .control section CKTsetup() has already snapshotted the node list into prev_CKTlastNode and sized the matrix, so the extra node makes the tail check ending CKTunsetup() fail -- and that check calls controlled_exit(EXIT_FAILURE). The check is CORRECT and was firing on a real inconsistency; the defect is upstream of it. This never bit in a plain netlist purely because of ORDERING: a .tf card in the deck is parsed BEFORE CKTsetup() runs, so the invented node is inside the snapshot and the tail still matches -- it is reachable only interactively, which is exactly where a typo is most likely and losing the session hurts most. Fix: analysis cards resolve through INPtermSearch(), the lookup-only half of the pair that already existed alongside INPtermInsert() and was simply never used here; an unknown name is reported (no such node: ouy) instead of created. DECK ORDER STILL WORKS, and this is the constraint that shaped the fix rather than an afterthought -- a .tf card may legally sit ahead of the devices that define its nodes, so creation remains permitted while the circuit is NOT yet set up, and is refused only once CKTisSetup is true, i.e. deck parsing is over. Applied at 14 call sites across dot_noise, dot_tf, dot_sens, dot_pss, dot_pac, dot_psp, dot_pnoise, dot_pxf and dot_hb; dot_pz reached the same insert INDIRECTLY through INPgetValue(IF_NODE) and its four node arguments are now read explicitly so they take the same path. Verified: 0 of 14 commands kill the process, from 7; the same reference as a deck card still parses; a .tf card placed BEFORE its devices still returns transfer_function 0.75, output_impedance 750, input_impedance 4000; 22 valid invocations compared digit-for-digit against the shipped binary, all identical; and the stateful fuzz that found it went 58/60 to 60/60. Verify (examples/nodetypo): 6 checks. | | [E-350](#) | ngspice-46/src/frontend/com_sweep.c | A sweep never put its .param back, and after two sweeps **reset** could not either. .param rload = 3k, sweep it 1k..5k, sweep it again 2k..10k, then reset + op gave 0.909 -- rload still 10k while the deck said 3k. One sweep was fine; the second broke it. ROOT CAUSE: the two implementations disagreed about the state they LEAVE BEHIND, which is exactly what E-320 guarantees they never do ("a sweep can only get faster, never change its result") -- it held for the numbers, not for the aftermath. The reset path drives each point with alterparam, which edits the parameter PERMANENTLY, so a later reset re-sources a deck that now carries the last swept value; the fast path never touches the deck, so reset did restore. Same command, same netlist, different circuit afterwards, decided only by which path happened to arm. WHY THE SECOND SWEEP: arming self-checks each captured expression against the value numparam baked into the flattened card at nominal (a guard against a subckt-local shadow), but the first sweep left the dico holding its last swept value while the flattened card still carried the nominal one -- so the check could not pass for an honest reason and DISARMED. The second sweep of a parameter therefore fell silently back to the reset path: no error, no banner, the fast .param path armed

line simply stopped appearing -- and the reset path is the one that edits the deck permanently, which is why the damage needed exactly two sweeps to surface. Both symptoms, one root cause: the sweep never restored the parameter. FIX: capture each swept .param before the first point and restore it at cleanup -- deck text via `sw_set_deck` (what reset re-sources) AND the numparam dico via `nupa_add_param + nupa_recompute_params` (what the NEXT sweep reads and what arming self-checks against). BOTH HALVES ARE NEEDED and the first attempt shipped only the deck half: with the dico still dirty the second sweep read the first sweep's last value as its own nominal and dutifully restored THAT, giving 0.833 instead of 0.909 -- better, and still wrong. The capture is all-or-nothing (if any swept name cannot be read back nothing is restored) rather than putting some parameters back and not others. Restoring is the direction that makes the two paths agree; making the fast path permanent instead would have satisfied the letter of the E-320 invariant while leaving reset unable to undo a sweep and every repeat sweep disarmed, and it matches how ngspice already treats a swept source in .dc, which does not leave it parked at the last step. DELIBERATELY UNCHANGED: the live circuit still holds the last swept point, so a bare op afterwards answers with that value -- true of both paths before this change and unchanged by it; only the deck and dico return to nominal, so reset means something again. Verified: reset gives 0.75 after 1/2/3/mixed-spec sweeps (was 0.909 from the second on); 3 of 3 sweeps arm the fast path (was 1 of 3); derived params and subckt-internal values restore exactly; and 12 sweep batteries compared digit-for-digit against the shipped binary -- plain/list/dec, repeated sweeps, derived, subckt-internal, model params, Monte Carlo draws twice over to catch RNG-stream drift, and a two-knob family, 5-11 numbers per case, none vacuous -- ALL IDENTICAL, so nothing a sweep computes moved. Regression 282/282. The example is a proven trigger, not decoration: on the pre-fix binary 4 of its 5 checks fail, and the one that passes is precisely the invariant that must hold on both. Verify (examples/sweeprestore): 5 checks. | | [E-351](#) | ngspice-46/src/osdi/osdisetup.c, ngspice-46/src/osdi/osdidefs.h | **sens** killed the process on any OSDI model carrying an INTERNAL NODE -- which is every production compact model. BSIM, HICUM, PSP and EKV all allocate internal nodes for their terminal resistances, so .sens was effectively unusable with real Verilog-A models, and it failed by taking the session down: "Internal Error: node allocation in DEVsetup() during sensitivity analysis, this will cause serious troubles !, please report this issue !" then exit(1). FOUND BY a campaign built on the observation that the previous OSDI sweep (83/83 clean) was LOPSIDED -- 40 dc and 22 ac checks against only 2 sens, 2 disto and 4 pz. Those thin three are exactly where an OSDI device is most likely to differ from a built-in, because they do not merely solve the circuit, they consume the DERIVATIVES the model returns. Three axes were run (OSDI vs equivalent built-in on both solvers; Sparse vs KLU on the same OSDI deck; repeat-invariance): 181/184, and all three failures were this one bug, on the single model in the set with an internal node. ROOT CAUSE: cktsens.c re-invokes every model's DEVsetup() to stamp the perturbation matrix and requires that doing so allocate NO nodes -- it snapshots CKTlastNode around the call and calls controlled_exit(EXIT_FAILURE) if it moved. The requirement is reasonable and every built-in device meets it by guarding allocation on "not already allocated": the inductor does `if (here->INDbrEq == 0) { CKTmkCur(...); here->INDbrEq = tmp->number; }` and INDunsetup() zeroes INDbrEq again. OSDI had NO such guard -- its setup called CKTmkVolt/CKTmkCur for each internal node unconditionally, so the second call, the one sens makes on a circuit that is still set up, allocated a fresh set and tripped the check. The check was CORRECT and firing on a real inconsistency; the defect is upstream of it. FIX: the same idiom the built-ins use. Each instance records in `OsdiExtraInstData` the node numbers its internal nodes were given, and a later setup reuses them instead of allocating. `OSDIunsetup()` deletes those nodes so it also clears the record -- the OSDI counterpart of `INDunsetup()` zeroing `INDbrEq` -- because otherwise a genuine re-setup after teardown would hand out numbers for nodes that no longer exist. Reuse is conditional on the node COUNT matching, since collapsing depends on parameters that alter can move; if it changes, the record is dropped and nodes are allocated afresh. NOT MERELY NON-FATAL -- CORRECT, checked two independent ways because "it stopped crashing" proves nothing: an internal-node model ($r=3000 + rs=10$) and a plain one ($r=3010$) are the SAME circuit (operating points agree to 14 digits), so their sensitivity to the shared built-in $R1$ must be identical -- and is, at $-1.8718770344518e-04$ from both decks, every printed digit, against a closed form $-RL/(R1+RL)^2 = -1.87187890622571e-04$ (the $1e-6$ residue is the sensitivity method's own perturbation step, not the fix). The model's own parameter agrees too: $n1_r$ $6.21886087563079e-05$ vs closed form $6.21886679809205e-05$. Verified: campaign

184/184 (from 181/184); sens completes on both solvers; sens twice in one session gives 41 bit-identical values; a reset -> sens -> reset cycle returns the operating point to 7.50623441396507e-01, guarding that the node record never outlives the nodes; all 7 other analyses on the same model unaffected. Regression 283/283. The example is a proven trigger: on the pre-fix binary 5 of its 7 checks fail, and the two that pass are precisely the ones that should -- the model WITHOUT an internal node, and the analyses other than sens. Verify (examples/osdisens): 7 checks. | | [E-352](#) | ngspice-46/src/osdi/osdidisto.c (new), ngspice-46/src/osdi/osdi.h, osdiinit.c, osdiregistry.c, osdiext.h, ngspice-46/src/include/ngspice/osdiitf.h, ngspice-46/src/spicelib/analysis/cktdisto.c | SUPERSEDED BY E-359 (the capability stands and is still verified by the same oracles; the compiler-emitted-tensor IMPLEMENTATION below was replaced by numerical differentiation of the analytic jacobian in ngspice, removing the compile-time regression, the 30MB objects, the runtime penalty and the ABI bump). **.disto** now includes a Verilog-A model's nonlinearities; previously it printed a warning and left them out. Distortion is a VOLTERRA analysis: it needs each device's Taylor expansion of $I(v)$ to THIRD order, not the operating-point linearisation the Jacobian provides. Built-in devices hand-code those coefficients (diodset.c computes id_{x2}/id_{x3}), which is why only FOUR of ~58 built-ins implement DEVdisto at all -- diode, BJT, BSIM1, MES -- and the OSDI ABI stopped at first derivatives, so cktdisto had nothing to ask an OSDI device for. KEY ENABLER, verified before anything was built: OpenVAF's autodiff is ALREADY arbitrary-order (mir_autodiff's DerivativeIntern chains through previous_order, with a comment anticipating 8th order), and nested ddx(ddx(ddx(...))) reproduces closed form to every printed digit. So the compiler work was not implementing higher-order AD but ASKING for it: build_taylor_tensors re-runs autodiff over its own first-order results for the second, and again for the third. NO 3-VARIABLE CEILING: ngspice's framework caps at three controlling variables (Dderivs holds derivatives "w.r.t 3 variables"; the DFx helpers take up to 27 scalar arguments with no fourth-variable form, DF_n2F12 taking a struct because the list outgrew the calling convention). That cap belongs to those hard-coded helpers, not to the mathematics -- the contribution is a symmetric tensor contraction -- so osdidisto.c contracts directly over however many inputs a device has and never calls D1x/DFx. Proven with an ideal mixer $I(out)=kV(a,ref)V(b,ref)$ whose ONLY nonlinearity is the mixed partial (both diagonal 2nd derivatives identically zero, so a dropped cross term gives exactly zero): matches closed form $0.5kA^2R$ EXACTLY. TWO SILENT-ZERO BUGS found on the way, both indistinguishable from a linear model: (1) the tensors were OPTIMISED AWAY -- ensure_optbarriers protects residuals/noise/jacobian but the Taylor values feed nothing else in the MIR, so the optimiser deleted them and every coefficient arrived as zero (they also take the mfactor scaling, since m parallel devices inject m times the distortion current); (2) NOTHING STORED THEM -- eval's jacobian/residual stores are gated on CALC* flags that .disto never sets. Plus one wrong-by-a-constant bug only a reference implementation could catch: D1nF12 returns $0.5S2vF12(...)$ and $S2vF12$ ALREADY sums both index orderings, so without that 0.5 the mixed-tone products came out exactly 2x -- and because IM3 consumes the $f1-f2$ kernel its error was a non-obvious 2.246x. KNOWN LIMITATION AT THE TIME, SINCE LIFTED BY E-353: a model whose contribution goes through \$limit emitted NO tensors (the residual depends on the limited value, not the raw voltage read; build_jacobian handled that via intern.lim_state and the tensor pass did not). Limiting is standard in production diode/BJT/MOS models, so this was a real restriction; E-353 folds the limited values into the derivative chain and removes it. What remains unreachable is a nonlinearity in a GROUND-REFERENCED probe (the tensors are indexed by model input and a bare $V(a)$ is not one, having no hi/lo pair). Registering DEVdisto removed cktdisto.c's blanket warning, which would have turned such a model into a SILENT ZERO -- worse than before -- so OSDIdisto warns for itself. VERIFIED against four oracles, two involving no simulator at all: HD2/HD3 vs a closed-form polynomial (0.0 and $1.2e-16$); the A^2/A^3 scaling laws (4.0000, 8.0000); $f1+f2$, $f1-f2$ and IM3 vs the built-in diode ($1.87e-06$, the bound the 0.24 ppm \$vt difference justifies); IM3 vs transient+FFT in a DIFFERENT DOMAIN, converging 5.7% -> 1.4% -> 0.21% as the drive halves, exactly the A^2 signature of the 5th-order content a 3rd-order truncation excludes; and the multi-variable cross term (0.0). TWO HARNESS TRAPS recorded because both would have produced a confident wrong answer: the first IM3 transient was measuring SPECTRAL LEAKAGE from the fundamentals rather than IM3 (exposed by the A^3 law giving 2.93 instead of 8, not by the value looking wrong); and a first parameter choice made HD3 IDENTICALLY ZERO (the two contributions cancel exactly when $c3 = 2c2^2/Ytot$, and $2(1e-4)^2/2e-3$ is precisely the $1e-5$ chosen) -- a zero that would have passed a naive check for the wrong reason, the same

trap as the old campaign's "HD2 ~ 0 on a LINEAR network". The suite now refuses to score a both-zero comparison. Regression 284/284. The example is a proven trigger: on the pre-fix binaries 5 of its 6 checks fail, and the one that passes is the one that should. Verify (examples/osdidisto): 6 checks. | | [E-353](#) | ngspice-46/src/osdi/osdidisto.c, examples/limitdisto_examples/, examples/osdidisto_examples/ | SUPERSEDED BY E-359 (\$limit models still work -- the suite below passes unchanged -- but the chain-fold no longer exists; differencing the real jacobian makes the problem unable to arise). **.disto** now works for Verilog-A models that use \$limit -- i.e. every production compact model. E-352 gave OSDI devices distortion analysis, but only for models reading their controlling voltages DIRECTLY; a model passing one through \$limit contributed NOTHING, and limiting is how every production diode/BJT/MOS model converges, so the feature did not reach the models people actually run. ROOT CAUSE: \$limit hands the model a LIMITED copy of a branch voltage and the residual is written in terms of that copy, so differentiating the residual by the RAW voltage read yields zero -- nothing in the expression depends on it. The model is not linear, but every derivative the tensor pass asked for was. This was never a problem for AC or transient: build_jacobian has always folded the limited values back in via intern.lim_state; E-352's build_taylor_tensors simply did not perform the same fold. FIX: each model input now maps to a CHAIN of autodiff unknowns rather than a single one -- new taylor_unknown_chain returns the limited values from intern.lim_state.raw (each with the sign it contributes) followed by the input's own unknown. The 2nd-order pass sums over every PAIR drawn from the two chains and the 3rd-order pass over every TRIPLE, each term carrying the XOR of its entries' signs; an entry is negated when the model limits a REVERSED branch, \$limit(V(ref,a,...). VERIFICATION uses an exact expectation with no appeal to another simulator: \$limit is a convergence aid, so at convergence the limited value equals the actual one and a model written with it is THE SAME DEVICE as the same model written without it -- every distortion product must agree. Both spellings are generated from ONE body string, which is what guarantees they differ only by \$limit and cannot drift apart. Five shapes cover the chain combinatorics (where this can go wrong -- a model limiting a single branch never puts more than one entry on either side of a pair): two independently limited diagonals; a cross term with BOTH inputs limited (2x2 pairs); a cross term with ONE input limited (asymmetric chains, 2x1); a THIRD-order cross term whose chain has 3 entries, so the triple loop genuinely runs; and a reversed-branch limit, the only spelling producing a NEGATED chain entry. All 13 live comparisons agree, worst 7.1e-10 -- convergence noise, since the two spellings settle 8.3e-11 apart (limiting changes the Newton path, not the solution) and d2/dv2 of exp(v/vt) amplifies that by 1/vt^2. THREE WAYS THIS COULD HAVE PASSED WHILE WRONG, each caught and closed because each produces a confident green: (1) BOTH-ZERO comparisons -- 'no distortion' is precisely the pre-fix behaviour, so scoring 0 == 0 as agreement would certify the bug; the suite scores a both-zero product as a FAILURE except where the mathematics requires zero (a bilinear term kv1v2 has no third derivative, so its IM3 is correctly zero and is asserted as such). (2) COMPARING MAGNITUDES -- a dropped sign on a negated chain entry flips the result and leaves |.| untouched, so the reversed-branch shape, the ONLY one testing the sign logic, would have scored its own target bug as a pass; products are compared as COMPLEX values and the shape is pinned against the forward one (+1.2488e-01 vs -1.2488e-01, exact negatives). (3) DEGENERATE DECKS -- three shapes initially read zero everywhere for reasons unrelated to \$limit: the probed node was pinned by an ideal voltage source so no injected distortion current could move it, and each nonlinearity saw only ONE tone and so had nothing to intermodulate; fixed with series resistors and both tones on both drives. A FOURTH trap was ruled out rather than fixed: several shapes agree at exactly 0.00e+00, which would be the signature of OpenVAF having elided the \$limit altogether (identical code compared against itself); dumping the chains under OPENVAF_DAE_DEBUG settles it -- every plain build reports len 1 on every input, every \$limit build reports 2 or 3. The rest of the E-352 suite is unchanged: closed-form HD2/HD3 still 0.0 and 1.2e-16, two-tone vs the built-in diode still 1.9e-06, the closed-form mixer still exact. The OSDI ABI is UNCHANGED -- this is a compiler fix, so 0.8 descriptors from before and after are interchangeable. NGSPICE SIDE: no functional change -- osdidisto.c already consumes whatever tensors a model publishes, which is why the fix is compiler-only. What did change is the no-tensors warning: it named \$limit as the case that falls here, which is now false, so both the message and the comment above it name the remaining unreachable case instead (a nonlinearity in a GROUND-REFERENCED probe, not a model input because it has no hi/lo pair). E-352's example check [6] asserted that warning on a

\$limit model and so became a REGRESSION the moment this landed -- it now uses a ground-referenced model, which still legitimately trips the guard (3 jacobian entries, 2 nodes, no model inputs), keeping the never-silently-zero guarantee under test. Verify (examples/limitdisto): 7 checks; osdidisto suite back to 6/6. | | [E-359](#) | ngspice-46/src/osdi/osdidistonum.c (new), osdidistonum.h (new), osdidisto.c, osdi.h, osdiregistry.c, Makefile.am, src/include/ngspice/osdiitf.h, examples/osdidisto_examples/* | **.disto** for Verilog-A rebuilt on NUMERICAL differentiation of the analytic jacobian; the compiler-side tensor machinery is DELETED. E-352 obtained the 2nd/3rd-order Taylor coefficients by having the compiler emit a symbolic closed form. The capability was right but the cost was out of all proportion to what .disto does: on ASMHEMT 707k intermediate derivatives -> 1.4M MIR instructions -> a 30.2MB .osdi and a 126s compile against a 3.6s baseline (35x; hisimhv 20x), plus -- because the tensors were computed inside eval() -- a 3.3-3.9x RUNTIME penalty on every DC sweep, transient, AC and noise run, and an OSDI 0.8->0.9 ABI bump that made every already-compiled model unusable without a rebuild. ROOT MISTAKE: .disto needs the coefficients at ONE point, once per instance, once per analysis; E-352 computed a closed form valid at EVERY point in order to evaluate it once. WHAT REPLACES IT: the model already publishes its FIRST derivatives analytically -- that is the jacobian every analysis loads -- so the higher ones follow by differencing THAT at the operating point ($d2 R_r/dv_p dv_q = d/dv_q[J_{rp}]$; $d3 = d2/dv_q dv_s[J_{rp}]$). Differencing an already-analytic first derivative rather than the residual is what keeps it accurate. Two properties make it cheap: the tensors are evaluated at the DC OPERATING POINT so they are frequency-INDEPENDENT (built once at D_SETUP, reused by every mode and every frequency), and they are built in NODE coordinates where an instance has 5-20 nodes rather than the 52 'model inputs' ASMHEMT reports. Entries keep the exact (row, col...) shape the analytic path produced, so the whole verified Volterra contraction is reused UNCHANGED -- only the tensor source and the kernel differ. NODE COORDINATES ALSO LIFT A REAL LIMITATION: E-352 indexed by model input (a branch voltage with a hi/lo pair), so a nonlinearity in a GROUND-REFERENCED probe V(a) had no pair, was never recorded, and the device reported 'contributes no distortion' and returned zero; it now works and is pinned to a closed form (-4.16666665e-02 vs exact -4.16666667e-02). E-353's \$limit chain-fold disappears for the same reason -- differencing the real jacobian sees what the simulator sees, so there is no chain to fold, and E-353's seven-shape suite (written specifically to break that logic) passes unchanged against an implementation containing none of it. ACCURACY, measured against EXACT CLOSED FORMS rather than against the old implementation: cubic 4.0e-09/5.4e-09, exponential diode 2e-10/~1e-10, internal-node 3.6e-12/1.2e-08, ground-referenced 4e-09, MEXTRAM vs the analytic implementation 5e-09/1e-07 -- all below the 1.9e-06 floor the built-in-diode oracle already sits at (a 0.24 ppm \$vt difference). THE STEP SIZE IS THE WHOLE GAME: truncation is $O(h^2)$ and roundoff $O(\text{eps}/h)$ at 2nd order but $O(\text{eps}/h^2)$ at 3rd, so one step cannot serve both (a single 1e-3V step put the result 8.5e-5 out); and the right yardstick is not the operating voltage but the CURVATURE scale ($v_t = 26mV$ for a diode, volts for a polynomial), so a voltage-relative step left a cubic biased at 0V 11% wrong. The 2nd-order pass measures $|dJ/dv|/|J| = 1/V0$ directly and the 3rd-order step is derived as $h = 1.2e-4V0 \sim V0\text{eps}^{(1/4)}$. ONLY .disto IS APPROXIMATE: the jacobian is still exact symbolic autodiff and the compiler now emits nothing distortion-related, so a model's .osdi is what a pre-E-352 compiler would have produced -- op/dc/ac/tran/noise verified IDENTICAL to 15 digits against a build with the tensor pass disabled, and an op before and after a .disto agrees exactly (the probe restores the unperturbed evaluation as its last act). THREE BUGS THAT PRODUCED PLAUSIBLE WRONG NUMBERS rather than failures, all now documented in osdidistonum.c with the measurement table that settles them: one step for both orders (8.5e-5); a voltage-relative step (11% on a cubic at 0V); and indexing write_jacobian_array_resist as if parallel to jacobian_entries[] when it is a DENSE array of the RESISTIVE entries only (the two coincide only for a model with no charge storage, which is exactly why simple models agreed and MEXTRAM was out by 3-10x). A fourth belonged to the prototype and mattered most: COLLAPSED NODES -- $V(a,b) <+ 0$ shorts nodes together (MEXTRAM merges three onto one) and treating them as independent perturbs the shared node once per alias, a silent ~0.2% error that briefly looked like a defect in the SHIPPED analytic implementation. It was not: E-352/E-353 are correct on every model tested. NGSPICE SIDE: the new osdidistonum.c (~300 lines) is the whole implementation -- distinct-global-node discovery (collapse-aware), a +-h jacobian sweep for 2nd order and a mixed second difference over node pairs for 3rd, emitted as sparse (row, col...) entries. osdidisto.c keeps its Volterra modes verbatim and only changes where the entries come from (numcache, built at D_

SETUP) and the kernel (node voltage rather than an input node PAIR). The *has_taylor* gate and the OSDI version floor are gone -- all this needs is *eval*, *jacobian_entries* and *write_jacobian_array_resist*, which every OSDI ≥ 0.4 object has, so models compiled before any of this work get distortion WITHOUT being recompiled. Verify (examples/osdidisto): 6 checks, now including a ground-referenced nonlinearity pinned to closed form; examples/limitdisto 7 checks unchanged. Regression 285/285. || [E-360](#) | ngspice-46/src/osdi/osdidisto.c, examples/osdidisto_examples/ | A circuit with TWO OR MORE different Verilog-A models reported distortion for only the LAST one; every other model contributed ZERO, silently. Introduced by E-359 and shipped with it. CAUSE: DEVdisto is dispatched once per DEVICE TYPE -- cktdisto.c walks DEVICES[] and calls each type's handler in turn -- and every distinct .osdi registers as its own device type, so a circuit using two Verilog-A models calls OSDIdisto TWICE for every mode, D_SETUP included. E-359 cached the numerical tensors in ONE GLOBAL array and cleared it at D_SETUP, so model B's setup destroyed model A's tensors; in the mode passes A's instance pointers then failed the consistency check and every instance was skipped. That check is the only reason this produced NO numbers rather than WRONG ones -- without it model A would have contracted model B's tensors against A's kernels and produced confident garbage. FIX: the cache is keyed by descriptor (OsdNumModelCache{descr, ent, n, cap}; numcache_for(descr,create) finds-or-adds, numcache_reset drops only that model's entries), so each model only ever clears its own; lookup is linear in the number of MODEL TYPES, a handful, not in instances. WHY THE TESTS MISSED IT: every distortion test -- the six E-352 checks, the seven E-353 \$limit shapes, every model in the corpus campaign -- used EXACTLY ONE .osdi. The bug needs two. Single-model coverage was thorough enough to hide a whole-feature failure. The new check [7] compares the cubic branch's value WITH a diode present against the same branch ALONE, so it fails on the silent zero rather than merely asserting non-zero; against the pre-fix binary it reports '0.00000000e+00 alongside a diode vs 6.24999998e-03 alone'. Also stress-tested: 50 instances (cache growth past its initial capacity of 32; first and last match the single-instance reference) and three model types with .disto run TWICE in one session (every branch correct on both runs, so the per-model reset is idempotent). Verify (examples/osdidisto): 7 checks. Regression 285/285. || [E-361](#) | ngspice-46/src/osdi/osdidistonum.c, ngspice-46/src/spicelib/analysis/distoan.c, examples/osdidisto_examples/* | Two sanitizer findings in .disto that every ordinary build tolerates silently. (1) HEAP-BUFFER-OVERFLOW, mine from E-359: numdisto_jac_at copied the solution vector with for (i = 0; i <= n; i++) where n = CKTmaxEqNum, but CKTrhsOld holds CKTmaxEqNum entries -- ngspice's own code copies it as CKTmaxEqNum * sizeof(double) -- so the last valid index is n-1 and this read one element past the end on EVERY .disto run with a Verilog-A device. ASan: 'READ of size 8 ... 0 bytes after 32-byte region', a 4-double region faulting at index 4, matching exactly. It is a READ, so nothing was corrupted and the computed values are unchanged by the fix, but it is undefined behaviour on the main path; the same off-by-one was in the node-range guard (gp > n where the valid range is gp < n). (2) (int)NaN IN THE SWEEP POINT COUNT, pre-existing ngspice in distoan.c and reproducible with NO OSDI device in the circuit at all: a linear sweep with start == stop leaves DfreqDelta at zero so the tolerance term is 0/0, and a decade/octave sweep with ZERO steps does $\exp(\log(10)/0) = \text{inf}$ and then evaluates $0 * \text{inf}$ -- converting NaN to int is undefined, so a nonsense request quietly ran with whatever point count the hardware produced instead of being refused. FIX: zero-step decade/octave sweeps are rejected the way an unknown step type already is, and the linear division is guarded (where DfreqDelta is non-zero the term is $\text{floor}(\text{DfreqDelta} \text{reltol} / \text{DfreqDelta}) = \text{floor}(\text{reltol}) = 0$ for any sane reltol, so zero is the value consistent with every other case and normal sweeps are bit-identical). NOTE: this second one had been firing under the project's OWN test deck all along -- osdidisto check [4]'s mixer sweeps lin 1 1e6 1e6, precisely the start == stop trigger -- and nothing ever failed because an un-sanitized build tolerates it. Neither defect produced a wrong number: one is an out-of-bounds READ, the other an undefined conversion that happened to yield a usable value on this hardware. That is the class only a sanitizer finds, and the argument for running the suites under one after any change involving raw index arithmetic. Verify (examples/osdidisto): 8 checks, the new one asserting a zero-step sweep produces no output rather than inventing some (an ordinary build cannot observe the UB itself). Both suites 8/8 and 7/7 under ASan+UBSan with values unchanged. Regression 285/285. || [E-362](#) | ngspice-46/src/spicelib/analysis/{distoan,dctran,cktsens,dctrurv}.c, examples/sweepguard_examples/(new) | Fuzzing analysis-card sweep parameters against an ASan+UBSan build: SEVEN pre-existing defects, all reachable from ordinary input, three of the counts

involved being used as ALLOCATION SIZES and one reaching the allocator NEGATIVE. E-361 found one (int)NaN conversion in .disto's point count; that looked like a family rather than a one-off, so this drives every sweep-taking analysis (.disto/.ac/.noise/.sens/.sp/.dc/.tran/.four/.fft) with degenerate and extreme parameters. FINDINGS: distoan.c DECADE and OCTAVE -- the count goes +-inf (a zero or denormal start frequency makes log() diverge) or exceeds INT_MAX, and the (int) cast is undefined, feeding DstorAlloc(..., NoOfPoints+2); distoan.c LINEAR -- DnumSteps+1 is plain INT arithmetic and overflows at INT_MAX BEFORE any double-valued guard can run; distoan.c LINEAR -- the final count could be negative, reaching the allocator as ASan's 'requested allocation size 0xfffffffffffff0'; dctrans.c -- CKTtimeListSize overflows and is used both by TMALLOC and by osdiaccept.c to size a buffer; cktsens.c x2 -- the same cast pattern driving the sensitivity sweep; dctrans.c -- .dc has NO bound at all, advancing by a step and comparing against the stop value, so dc V1 0 1 1e-30 (a typo for 1e-3) runs forever with no diagnostic, and a step below the ULP of the start never advances the value at all. .tran already declines the equivalent request and a zero step is already refused, so .dc was the outlier. ONE FIX MADE THINGS WORSE FIRST: clamping CKTtimeListSize to a representable value regressed tran 1e6 1e30 1e6 from erroring to HANGING, because ngspice's existing rejection of absurd transients WAS this very overflow producing an allocation size that failed -- clamping made the allocation succeed and the run proceed forever. The accidental rejection had to become a deliberate one (E_PARMVAL); absurd transients now return rc=1 exactly as before and ordinary ones are bit-identical. TWO FINDINGS WERE NOT BUGS and were deliberately not 'fixed': an enormous but finite sweep is SLOW, not hung, and a timeout cannot tell them apart -- sens ac dec 1000000 over 300 decades measured 0.05/0.20/1.78/17.4s at 1/10/100/1000 steps (linear, so simply large), and ac oct 1000000 1e300 1e6 takes 12.8s identically on the shipped and fixed builds, exceeding the fuzzer's timeout only under ASan's ~5x slowdown. The scaling check that distinguishes them is documented in the harness. Final fuzz run 400 iterations: SANITIZER=0, CRASH=0, no real hangs. Verify (examples/sweepguard): 3 checks pinning eleven repros plus six ordinary sweeps that must be unaffected; a proven trigger -- against the pre-fix build it reports three HANGs. Its 'produces no output' check passes even pre-fix, because an ordinary build cannot observe an undefined conversion, only its consequences, which is why the sanitizer harness (fuzz_analysis.py) ships alongside it rather than being folded into the regression. Regression 285/285. | | [E-364](#) | src/osdi/osditrnoise.c (new), src/osdi/{osdidefs.h,osdiload.c,osdisetup.c,osdi.h,Makefile.am} | TRANSIENT NOISE for OSDI devices: Verilog-A noise sources are now injected into .tran. .noise linearises about an operating point, so it cannot show jitter, noise-induced switching or oscillator phase noise; ngspice has had trnoise for independent sources for years, but OSDI devices were the only silent components in an otherwise noisy circuit. NO ABI OR COMPILER CHANGE was needed -- the compiler already emits a full parametric description that NOTHING read: noise_source_type[], load_noise_params() (per-source power and exponent, i.e. $S(f)=\text{power}/f^{\text{exp}}$ at the CURRENT bias), and each source's node pair and correlation name. load_noise_params and noise_source_type were dead ABI, partly because the NOISE_TYPE_* names lived only in the compiler-side header and had never been mirrored into ngspice's osdi.h (now fixed). ACTIVATION IS AUTOMATIC with no option to set: it turns on when the circuit already has a trnoise source and adopts that source's noise timestep, so all generators share one grid; it is deliberately NOT keyed on 'the device declares noise sources', which every real compact model does and would mean always-on. AMPLITUDE LAW, derived not fitted: f_alpha is Kasdin's fractional-noise method, $H(z)=(1-z^{-1})^{(-a/2)}$ so $|H|^2 \rightarrow (2\pi f t_s)^{-a}$, and discrete white noise of deviation Q on a grid ts has one-sided density $2Q^2 t_s$, giving $Q = \sqrt{(\text{power}|(2\pi f t_s)^a/(2 t_s))}$ -- with $a=0$ collapsing to the white case so both kinds share one line. The first attempt used the white law for flicker: wrong by $\sqrt{1/(2 t_s \pi)} = 399x$ at $t_s=1\mu s$, and measurement against .noise gave 397x/396x, which is what confirmed the derivation instead of inviting a fudge factor. A FIXED noise grid is essential -- scaling by the ADAPTIVE timestep makes the injected power track the LTE controller and the spectrum an artefact of the integrator. power is re-read every timepoint (shot noise follows the current), correlated same-name sources (LRM 4.6.4) share one stream, and the signed power from E-42 is carried into the amplitude for coherent addition. Injection is RHS-only (an independent current source, no Jacobian entry) and is stamped in the SERIAL post-eval loop. Verified: thermal 1.7% vs exact analytic, shot 0.6-4.5% across a 45x current range (which proves the bias is re-read), flicker 0.9939 +/- 0.0091 vs .noise with fitted slope -0.993. noise_table is deliberately NOT injected (a tabulated spectrum needs frequency

shaping, not a scalar amplitude); it warns once and stays fully accounted for in `.noise`. || [E-365](#) | `src/include/ngspice/cktdefs.h, src/spicelib/analysis/{cktpzset.c, cktsetup.c}, src/frontend/com_hb.c` | **pz** followed by **hb** returned a SILENTLY WRONG harmonic-balance result (DC term 0.2500 instead of the correct 0.2439, 2.5% out) and, underneath it, read freed memory. `CKTpzSetup` opens with `NIdestroy(ckt); NIinit(ckt);` -- freeing `ckt->CKTmatrix` and building a DIFFERENT one -- while leaving `CKTisSetup` asserted, so every device's cached matrix-element pointer (bound by `CKTsetup` into the old matrix) dangles. `com_hb` then guarded its setup with `if (ckt->CKTmatrix == NULL || SMPmatSize(...) <= 0)`, which asks 'is there a matrix?' -- and after a `pz` there is a perfectly good non-empty one -- so `CKTsetup` was skipped and `CKTload` read the stale pointers (ASan heap-use-after-free in `VSRCload`, via `HBanalyze/com_hb`; freed by `spDestroy<-NIdestroy<-CKTpzSetup<-PZan`). FOUND BY A NEW KIND OF CAMPAIGN: every prior ngspice fuzzing round varied the INPUT (netlist, model card, command, expression, rawfile, `.snp`, OSDI loader, analysis-card parameters), and each input runs in a fresh process, so none could reach a bug living in what one analysis leaves behind for the next. This one held the netlist fixed and fuzzed the ORDER of analyses and state mutators within one session -- 500 iterations, 4 sanitizer reports, all the same cause. Same shape as E-360 (a second Verilog-A model silencing the first in `.disto`). SCOPE measured, not assumed: `pz` then `op/dc/ac/tran/noise/disto/tf/sp/pss` is clean, and `hb` after any other analysis is clean; `portnum` matters only because it lets `hb` reach the load. The bug is NOT confined to a failing `pz` -- the success path corrupts identically. FIX: a new `CKTbindStale` flag records that an analysis replaced the matrix while `CKTisSetup` stayed asserted; `CKTpzSetup` sets it, `CKTsetup` clears it (a successful setup is what makes the bindings valid), and `com_hb` honours it with a `BALANCED CKTunsetup()/CKTsetup()` pair -- a bare `CKTsetup()` returns `E_NOCHANGE` since `CKTisSetup` is still 1, and calling it without the `unsetup` would re-run `DEVsetup` on already-setup devices and double-allocate their internal nodes. Circuits that never run `pz` are untouched. The example suite needs NO sanitizer: the consequence is a number, so it pins that `hb` after `pz` equals `hb` alone (fixed 4/4, before 2/4 with max deviation 4.55e-02). || [E-366](#) | `src/frontend/com_qpss.c, src/include/ngspice/klu-binding.h` | Two more sites of the E-365 stale-binding class, found by continuing the SAME sequence-fuzzing campaign against the FIXED build. Widening the pool (RF/steady-state commands, `sweep/wcd`, and SOLVER SWITCHING option `klu/option sparse` between analyses -- added deliberately, since KLU and SPARSE keep separate per-instance pointer sets) took the yield from 500 iterations/1 signature to 700/3. (1) `com_qpss.c` carried the identical guard `com_hb.c` had -- `if (ckt->CKTmatrix == NULL || SMPmatSize(...) <= 0)` -- which asks 'is there a matrix?' when the question is 'do the device bindings point into it'; after a `pz` there IS one, so `CKTsetup` was skipped and `CKTload` read freed memory. Same fix: honour `CKTbindStale` with a `BALANCED CKTunsetup/CKTsetup`. `com_checkpoint.c` also matches the grep but only sizes a buffer, so it is unaffected. (2) `CREATE_KLU_BINDING_TABLE` in `klu-binding.h` did `if (matched == NULL) { printf("not found"); } here->ptr = matched->CSC;` -- it REPORTED the NULL and then dereferenced it, so every lookup miss was undefined behaviour and silent on an ordinary build (UBSan: member access within null pointer of type 'BindElement'). Reachable from option `klu + pz + any AC-family analysis (ac/sp/stb)`, because `pz` rebuilds `ckt->CKTmatrix` and the device's COO pointer is no longer in the new bind table; SPARSE is unaffected since binding tables are a KLU construct. The companion `CONVERT_KLU_BINDING_TABLE_TO_COMPLEX/_TO_REAL` macros dereferenced the same unresolved binding and are now guarded, and a miss sets binding to NULL rather than leaving the freed value -- leaving it stale would turn a NULL deref into a use-after-free. The macro lives in a shared header, so this covers every device using it. STILL OPEN and documented: under KLU the pole-zero block at the end of `VSRCbindCSCComplex` reads through a binding that is NOT null but STALE (`CREATE_KLU_BINDING_TABLE` is skipped for that entry by its branch `> 0` precondition), so the existing NULL test passes; that needs the bindings TORN DOWN when `pz` rebuilds the matrix, not another guard, because a guard cannot distinguish stale from live. A central rebind in `CKTdoJob` was tried and REVERTED -- tracing showed `CKTbindStale` is already clear by dispatch time, so it fixed nothing, and a guard that fixes nothing is worse than none. Campaign findings 3 -> 1. || [E-367](#) | `src/frontend/typesdef.c, src/frontend/com_sweep.c` | **sweep** produced plots nobody could name. (1) `plot_alloc()` names a plot from the `plotabs[]` table in `typesdef.c` and falls back to the literal "unknown" when there is no entry -- and EVERY plot type this project added was missing: `sweep`, `sweepwave`, `hb`, `envelope`, `eye`, `loadpull`, `rfstab`, `stb`. `sweep` therefore created `unknown4/unknown7/`

unknown10. All eight registered. ORDER IS LOAD-BEARING: `ft_plotabbrev` returns the FIRST entry whose pattern is a substring of the plot name, so "sweepwave" must precede "sweep" or a waveform plot would be abbreviated "sweep" and collide with the point plot (the same rule already explains why a plot named "spectrum" matches the earlier "sp" entry, never "spect"). The struct is {p_name, p_pattern} -- the SECOND field is matched and the FIRST is returned, which reads backwards from the table and is now documented in-source, because getting it wrong yields a table that compiles and matches nothing.

(2) The summary printed 'sweep' as a LITERAL, so the one hint the command gave for returning to an earlier sweep -- `setplot sweep` -- always failed with "no such plot named sweep"; it now prints the real name (sweep3), in BOTH summary branches (single-knob and the -vs/-family curve family), and the stale header comment in `com_sweep.c` documenting the plot as named sweep is corrected. The numbers are not 1,2,3 because `plot_unique_tname` draws from a counter SHARED by every plot type and advances it only far enough to avoid a collision -- a 3-point sweep burns op1/op2/op3 and lands on sweep3. That is pre-existing behaviour and is NOT changed; what changes is that you no longer have to guess.

FOUND ALONGSIDE, LEFT OPEN: `print alle` is dead in every build configuration -- `findvec_alle()` sits behind `#if defined(XSPICE) && defined(SIMULATOR)` with a `#ifndef SIMULATOR` stub returning NULL beside it, and `-DSIMULATOR` is set only on `ngspice_CPPFLAGS/libngspice_la_CPPFLAGS` (`main.c`, `ngspice.c`, `sharedspice.c`) while every `src/frontend` source is compiled ONCE into `libfte.la` with `AM_CPPFLAGS` alone, which both the executable and the shared library link. Confirmed by inspection: the string "DigitalData" is absent from the binary. Adding `-DSIMULATOR` to `src/frontend/Makefile.am` was tried and REVERTED -- `ngproc2mod/ngmultidec/ngmakeidx` also link `libfte.la` but not the XSPICE libraries, so enabling the block breaks them at link time under `--enable-oldapps`. The guard is unsatisfiable BY DESIGN, and a real fix means moving the event-node lookup into a simulator-only object. Two latent items behind that dead code are recorded so they are not re-discovered: `findvec_alle` overwrites the `pl_tname` `plot_alloc` just assigned with a hardcoded copy("dig1") (leaking it, leaving the registered completion keyword pointing at a name no plot has, and giving every digital plot the same name), and `get_all_type` has a misplaced paren -- `tolower(word[0]) != 'a')` instead of `tolower(word[0]) != 'a'` -- defeating its own case-insensitivity, currently harmless only because callers lowercase first (verified: `print ALLV` and `print allv` behave identically). Verified by `examples/sweepname_examples`: 6/6 fixed vs 1/5 on the pre-fix binary (got ["unknown3"], no 'into plot' message). Regression 291/291. | [E-368](#) | `src/frontend/typesdef.c` | The periodic small-signal analyses named their plots wrong; E-367's claim that the naming was fixed was too broad. E-367 found its eight missing plot types by grepping for STRING LITERALS passed to `plot_alloc()`, which is blind to the path every standard and RF analysis uses: `outtif.c` calls `plot_alloc(run->type)` with a RUNTIME string handed to `beginPlot(analName)` by each analysis (mostly `ckt->CKTcurJob->JOBname`). Checking properly found seven more, four of them COLLIDING with an unrelated analysis -- worse than "unknown", because the name looks right. `ft_plotabbrev` returns the FIRST entry whose pattern is a substring of the plot name, and each of these contains a more general pattern: `PXF Analysis` matched nothing at all -> unknown; `PAC Analysis` contains "ac" -> collided with ordinary AC; `PSP Analysis` contains "sp" -> collided with S-parameters; `PNoise Analysis`, `qnoise` and `phasenoise` all contain "noise" -> collided with ordinary noise AND with each other; `Frequency Domain Periodic Steady State Analysis (QPSS)` contains "periodic" -> collided with PSS. Each now has its own pattern placed BEFORE the general one its name contains. TWO ORDERING TRAPS: "qnoise" contains "pnoise", so the qnoise entry must precede the pnoise entry; and PSS/QPSS differ only by their leading words (Time Domain vs Frequency Domain Periodic), so the QPSS pattern must carry "frequency domain periodic" rather than the "periodic" they share. Verified by `examples/periodicnames_examples`: 9/9 fixed vs 4/9 pre-fix (`pac->ac2`, `pxf->unknown2`, `pnoise->noise2`, PSS/QPSS both pss). The FOUR REGRESSION CONTROLS -- ordinary noise, ac, tran, sp -- pass on BOTH binaries, which is what shows the change is targeted: those names are relied on elsewhere in the repo (`setplot noise1`, `print noise1.noise_spectrum` in `noisecorr/noisejw/stdaudit/tempphys/physcheck`) and none moved; no deck that runs pnoise selects a plot by name either. HARNESS TRAP RECORDED: the `setplot` listing has a blank line immediately after its header, so a non-greedy regex reaching to the first blank line captures NOTHING -- an early harness reported "no plot" for all 25 analyses and called the run clean. SCOPE STATED HONESTLY: `envelope/eye/loadpull` use literals and were registered by E-367, so they are correct by construction but

no deck here reaches those paths; sens produced no plot under every invocation tried, consistent with the already-open sens breakage and not a naming problem. Regression 292/292. || [E-369](#) | src/spicelib/devices/vsrc/vsrcbindCSC.c | Closes the E-365/366 stale-binding class. E-366 fixed two sites and left the third OPEN on purpose, because it could not be closed with another guard: the binding was NOT NULL but STALE, and no NULL test distinguishes those. THE ASYMMETRY IS THE BUG: VSRcbindCSC assigns the pole-zero binding only INSIDE if (here->VSRCibrIbrPtr), and that pointer is allocated ONLY by a pole-zero analysis -- so on any later analysis the gate is false, the assignment never happens, and the binding keeps its previous value, pointing into the BindStruct that SMPdestroy() freed when the pz matrix went away. Both consumers (VSRcbindCSCComplex, VSRcbindCSCComplexToReal) then dereference it, because THEIR guard is VSRCbranch != 0 -- a property of the device -- rather than "was this binding re-established for THIS matrix". ASan on option klu ; pz ; ac: heap-use-after-free READ of size 8 in VSRcbindCSCComplex, freed by SMPdestroy, reallocated by SMPconvertCOOtoCSC. FIX: clear the binding BEFORE the gate, so a stale value can never survive a matrix rebuild -- one line, fixing both consumers at once since both NULL-test it; the duplicated (x != 0) && (x != 0) copy-paste slip is collapsed in the same pass. NO OTHER DEVICE NEEDS THIS: vsrc is the only one that hand-writes a pz binding block (grep "Pole-Zero Analysis" across every bindCSC.c returns one file); every other binding goes through the klu-binding.h macros, which are self-consistent -- CREATE and CONVERT are gated on the SAME (a>0 && b>0) condition, so a skipped create implies a skipped convert. HONEST LIMIT, stated in the doc and the example: the PRE-FIX binary PASSES every behavioural check -- the freed BindStruct entry still held the right CSC_Complex pointer, so the UAF read plausible values and the numbers came out correct, verified deliberately on a 40-node ladder with a tran between the pz and the ac to churn the heap. This is a MEMORY-SAFETY fix, not a wrong-answer fix; the behavioural checks are a regression guard, and the check that DISCRIMINATES runs under a sanitizer (the example runs it when NGSPICE_ASAN is set and SKIPS LOUDLY otherwise, so nobody mistakes a pass for proof). UB that happens to give the right answer today is still UB -- the same read against recycled memory is how E-365 presented, 2.5% wrong and silent. examples/klubind_examples 8/8. Regression 294/294. || [E-370](#) | src/spicelib/devices/urc/urcinit.c | Every .pz re-expanded the URC subcircuit, creating nodes past the allocated RHS. FOUND BY RUNNING AN EXISTING FIXTURE UNDER A SANITIZER: a sweep of all 185 shipped decks with each solver forced produced exactly one real finding, in ngrcrashanalysis_examples/pz_urc.cir -- the E-315 fixture, passing ever since E-315 stopped it SIGSEGVing. Nothing had run it under a sanitizer; the crash was gone, the corruption underneath was not. ASan: heap-buffer-overflow READ of size 8 in RESload (CKTload->NIter->CKTop->PZan), buffer allocated by NIreinit. RESload indexes ckt->CKTrhsOld by node number, so the read past the end means a resistor's node number exceeded the allocated RHS. CAUSE: urc declared .DEVsetup = URCsetup AND .DEVpzSetup = URCsetup, but URCsetup is not a matrix-entry allocator like its neighbours -- it is a SUBCIRCUIT EXPANDER (CKTmkVolt per lump, CKTtrtElt per element) with NO idempotency guard, and CKTpzSetup calls DEVpzSetup for every device on EVERY pz.job. Each .pz therefore expanded the URC again, creating fresh internal nodes AFTER NIinit had sized the RHS. That is exactly why RESsetup/CAPsetup are safe in the same slot and URCsetup is not: they only allocate matrix entries, which is idempotent. URC was the only expander wired up this way. Even a SINGLE .pz triggered it (CKTsetup expands once, CKTpzSetup expanded again); the overflow took a second pass to cross the allocation boundary. Visible symptom on an ordinary build: doAnalyses: device already exists, existing one being used + aborted run. FIX: .DEVpzSetup = NULL. The URC stamps nothing itself -- DEVload, DEVacLoad and DEVpzLoad are all NULL -- and the RES/CAP instances it creates are ordinary elements under their own device types, for which CKTpzSetup already calls RESsetup/CAPsetup; that is what actually re-binds the matrix. NOT A SOLVER BUG despite being found in a KLU/Sparse hunt: it reproduces identically under both. BEHAVIOURAL CHANGE, not hidden: on the E-315 fixture the fixed build gets FURTHER -- instead of aborting on a spurious duplicate device it runs the analysis and reports singular matrix: check node 3, which is CORRECT (that fixture's node 3 is genuinely floating); the duplicate expansion had been masking a real topology diagnostic behind a spurious one. verify_ngrcrashanalysis.py still 4/4. examples/urcpz_examples 10/10 fixed vs 5/10 pre-fix. Regression 294/294. || [E-371](#) | src/frontend/vectors.c, src/frontend/outitf.c, src/frontend/spec.c, src/frontend/postcoms.c, src/frontend/com_fft.c, src/frontend/rawfile.c | Per-type plot numbering, and a date on every plot. Reported from user observation: a 500-point sweep produced sweep500, not sweep1, and the plot's

Date field was empty. (1) NUMBERING. `plot_unique_type_name()` counted up a SINGLE counter shared by every plot type. A sweep runs one analysis per point and KEEPS its plot (deliberate -- `-overlay` and waveform inspection rely on it), so a 500-point sweep created `op1..op500`, each collision pushing the shared counter, and by the time the sweep's own plot was named the counter stood at 500. The number looked like the point count; it was really "how many plots exist". Each abbreviation now carries its own counter, so the first sweep is `sweep1` whatever else ran. A `tran/ac/tran/ac/noise/sweep/sweep` deck goes from `sweep12 op12..op3 noise3` to `sweep2 op10..op1 noise2`; the two agree except in the pathological case where one type's plots number an unrelated type. WHAT HAD TO BE PRESERVED: E-345 exists because this path was QUADRATIC over a sweep (89% of a 64000-point sweep in `plot_alloc->cieq->tolower`), fixed by starting the search from the shared monotone counter plus a hash membership index. A per-type counter restarting at 1 each time would have reintroduced exactly that quadratic, so the counter is REMEMBERED per type (nghash on the lowercased abbreviation) and the search resumes from it -- E-345's timing check still passes at 22->21 us/point, ratio 0.93 over a 4x longer sweep. **DESTROY HAS TO WALK IT BACK DOWN:** `destroy all` must restart numbering at `tran1`, which `examples/lifecycle_examples` (E-81) pins -- and that is what CAUGHT the first version of this change. `destroy` unlinks plots one at a time through `plot_forget()` rather than `plot_forget_all()`, so `plot_forget()` now lowers that type's counter to the freed number. That also makes a single `destroy op2` recycle the number (the shared counter went on to `op4`), which is the behaviour `destroy all` always had, so it makes the two consistent rather than adding a rule. (2) DATES. Only the analysis path set one: `outtif.c` did `pl->pl_date = copy(datestring())` after `plot_alloc`, so every plot created DIRECTLY by a command -- `sweep`, `hb`, `envelope`, `eye`, `loadpull`, `stb`, `rfstab`, `qpac` -- left `pl_date` NULL and `print` rendered it as `(null)`. The stamp now happens in `plot_alloc()`, the single point where every plot is created, covering all eight command paths and any future caller. The five sites that set their own date became redundant and were removed, except `rawfile.c`, where a loaded rawfile's OWN date must win -- that one frees the stamp before replacing it, so centralising does not leak. VERIFIED by updating and extending `examples/plotname_examples` -- the E-345 example whose expectations this deliberately alters, and the file that documents this behaviour: 9/9 fixed vs 4/9 on a true pre-fix build (all six files reverted). The `tran-has-a-date` row PASSES on both, as the control showing the analysis path always had one and only the command paths changed. Regression 294/294. || [E-372](#) | `src/frontend/variable.c` | **unset plots** reported a bogus "Internal Error". Found by a PLOT-LIFECYCLE sequence fuzzer -- random analyses interleaved with plot-management commands (`destroy`, `setplot`, `unset plots`, `write/load`, `fft`, `linearize`, `remcirc`) under ASan. That area was chosen deliberately: it is where E-342 (borrowed-pointer UAF via `unset plots`) and E-345 came from, and E-371 had just added a new allocation and new pointer arithmetic to it. Fired in 18/120 cases and minimised to ONE command on the SHIPPED binary: `unset plots -> "Error: plots is read-only."` followed by `"cp_remvar: Internal Error: var 112"`. TWO DEFECTS in that one line. (1) THE MESSAGE PRINTED AN ASCII CODE, NOT A NAME: the signature is `cp_remvar(char *varname)` and the format was `"...var %d", *varname`, which dereferences the string to its FIRST CHARACTER -- 112 is 'p' for plots, and `unset curplot` reported 99, which is 'c'. The diagnostic never named the variable it complained about (verified against the ordinals). (2) IT WAS NOT AN INTERNAL ERROR: the branch is guarded by `if (*p)`, and `*p` non-NULL only means the variable WAS FOUND in one of the lists walked just above -- the normal state for plots and `curplot` -- so it fired on 100% of ordinary unsets of them. An internal error valid input reaches every time carries no signal. FIX: drop both spurious prints. For `US_READONLY` the "%s is read-only" message is already the complete answer and nothing is unlinked or freed there (correct for a read-only variable), so removing the noise changes no behaviour; for `US_DONTRECORD` the case's own comment says "Do nothing..." and its siblings `curplotname/curplottitle/curplotdate` were always silent, so `curplot` now matches them. The other Internal Error in the same function ("US val %d", i) is CORRECT and left alone -- i really is an int and that default branch is unreachable from valid input. CHECKED AS A CLASS: a `grep` for a %d conversion fed a dereferenced char across the frontend returns no other occurrence. VERIFIED by `examples/unsetvar_examples` (no sanitizer needed): 9/9 fixed vs 5/9 pre-fix. The FOUR CONTROLS -- `curplotname`, `curplottitle`, `curplotdate` and an ordinary `set/unset` round-trip -- pass on BOTH binaries, which shows the change is confined to the two affected cases and did not break `unset` itself. Re-running the fuzzer against the fixed build: 400 iterations, 0 findings, 382/400 cases holding live plots at the end (a real clean result, not a harness that stopped working); harness committed under `examples/unsetvar_examples/fuzz/` and excluded from the regression sweep. PROCESS

NOTES RECORDED: the investigation began by checking whether E-371 had leaked (it added a `dstring()` allocation per plot) -- it had not, `killplot()` frees `pl_date` and peak RSS stayed flat 7152->7168 kB across 202->6060 plots created and destroyed. That leak check was FIRST attempted with `ASAN_OPTIONS=detect_leaks=1` and came back clean MEANINGLESSLY, because macOS/arm64 answers "detect_leaks is not supported on this platform" -- a sanitizer that silently does not run is indistinguishable from a clean result unless you check. Regression 295/295. | | [E-373](#) | `src/frontend/rawfile.c` | A rawfile write/load round trip lost the x-axis column and renamed the sweep axis. FRESH AXIS: E-226 fuzzed rawfile LOADING with malformed input looking for crashes; this asks instead whether a rawfile written by ngspice and then loaded by ngspice still holds the same data. Round-trip identity is a perfect oracle (no reference implementation needed) and catches silent fidelity loss a crash fuzzer cannot see; it also lands on freshly-changed code, since E-371 had just touched `rawfile.c`. Two independent defects, NEITHER of which corrupted a value. (1) `print` LOST THE X-AXIS COLUMN for every loaded plot -- before write: Index `v-sweep v(mid)` became after load: Index `v(mid)`, so nothing indicated which x-value each row belonged to. `print` prepends the scale column only when `pl_ndims` is non-zero (`postcoms.c`); `outtf.c` initialises it to 0 and sets it to 1 when analysis data arrives, but `pl_ndims` appeared NOWHERE in `rawfile.c`, so every loaded plot carried 0 and failed the gate. PROVEN NOT INFERRED: before writing the fix, a one-line probe setting `pl_ndims=1` on load was inserted, shown to restore the column exactly, and reverted. THE op CASE IS THE CONTROL THAT MATTERS: `pl_ndims=1` could have made `print` INVENT a column for a plot with no real scale (an op plot's variable 0 is just data, and the reader assigns `pl_scale` to variable 0 unconditionally) -- it does not, because `print`'s inner `pl_scale && !vec_eq(bv, pl_scale)` test still suppresses it, and op prints the same header before and after; checked explicitly, since that is why the fix is safe. (2) THE `.dc` SWEEP AXIS WAS RENAMED `v-sweep -> v(v-sweep)` in the written file. The `v(...)` form means "the voltage AT node X" -- right for a node probe, wrong for a sweep AXIS. A `.dc` plot's scale is the synthetic, voltage-typed vector `v-sweep`, which has no `v(` prefix, so the writer wrapped it. It did NOT compound (verified over three successive write/load passes: the prefix test leaves `v(v-sweep)` alone), so the damage was a one-time rename. Fix: add `|| v == pl->pl_scale` to the write-as-is condition. Checked how wide that else-branch's reach was: node voltages are already named `v(mid)` internally and take the first branch, and let vectors are `SV_NOTYPE`, so the plot's scale was the ONLY thing it was ever reached by; time and frequency escape it by not being voltage-typed. SCOPED OUT AND RECORDED, NOT SILENTLY IGNORED: ASCII rawfiles are 1 ULP lossy. The writer emits `%. *e` with `prec = DEFPREC`, i.e. 16 significant digits, while reproducing an IEEE double exactly needs 17. Measured with `numdgt=16`: 33 of 59 rows differ, ALWAYS in the scale column and ALWAYS at the 17th digit (`4.0000000000000004e-11` reads back as `3.9999999999999998e-11`, one ULP); data values identical, worst relative deviation `4.33e-16` over 118 values. The in-source comment on `raw_prec` even claims "default 15 (max)", which is wrong. Raising the default precision changes the on-disk format for every ASCII rawfile ngspice writes -- broader than these two defects -- so it is documented and NOT made. The example's ascii value check therefore compares to `1e-15` while the binary one demands bit-exactness. VERIFIED by `examples/rawtrip_examples`: 19/19 fixed vs 6/19 pre-fix. Controls pass on BOTH binaries (the op rows and both value comparisons), which is what shows the values were never the problem. Regression 296/296. |