

OpenVAF-r Robustness Campaign

Adversarial testing, mutation fuzzing, and the crash/hang paths found and fixed

Ngspice + OpenVAF Enhancements project

July 2026

Contents

OpenVAF-r robustness campaign	2
Summary	3
Method	4
Findings	5
The critical finding: exponential-time nested ?:	6
The hardening: parser depth, include cap, array cap	7
Behaviour-preserving	8
Reproducing	9

OpenVAF-r robustness campaign

A deep robustness audit of the `openvaf-r` Verilog-A → OSDI compiler: the full production-model corpus, ~50 adversarial hand-crafted inputs, and thousands of mutation-fuzzing iterations, run through a classifier that labels every outcome OK / clean-error / panic / segfault / abort / hang. The compiler came back clean on real work, and the campaign surfaced five ways to make it crash or hang on pathological input — all now fixed.

Re-run 2026-07-17. The campaign was replayed against the then-shipped compiler (production corpus + adversarial battery + mutation fuzzer). The memory-safety core was unchanged — 0 panics, 0 segfaults, every standalone model compiled — but the fuzzer found a fifth hang path that the original guards did not cover: an unbounded preprocessor macro-argument collector, plus an unbounded diagnostic-rendering flood behind it. Both are now fixed ([Enhancement-219](#)). Corpus counts below are refreshed to the current `VAF-Models-main` collection.

Round 2 (2026-07-18). A second campaign with a stronger fuzzer — diverse compact-model seeds and new mutation strategies (keyword / attribute / bracket injection) — found that ~5 % of mutated inputs still *crashed* the compiler (exit 101), the E-213 panic class. Triage grouped them into ten distinct root causes (a parser “seems stuck” assert, an empty-span `unimplemented!()`, `TextRange start>end` in span mapping, `unwrap()` on the source-map-back of synthesized exprs, `unwrap_node/unwrap_branch` on wrong-typed builtin args, an include-quote slice, a no-overload candidates `[0]`, and a too-few-arguments args `[0]`), all fixed ([Enhancement-220](#)). A fresh 12 000-iteration fuzz on the fixed compiler is now 0 crashes, 0 hangs. All the panics were already caught by the E-213 hook, so the compiler was never memory-unsafe; the bug was the crash UX and the missing diagnostic.

This is a companion to the [OpenVAF compiler internals](#) guide: that document explains how the compiler works; this one documents how hard it was pushed and what broke.

Summary

Production models compiled	92 / 92 standalone (124 .va files incl. include-fragments) — 0 crashes / hangs / panics
Adversarial hand-crafted inputs	~50 — all rejected cleanly, 0 accepted-invalid
Mutation-fuzzing iterations	thousands — 0 panics, 0 segfaults
Crash / hang paths found	5 — all fixed (E-147 , E-148 , E-219)

The compiler never crashed on random or garbage input — the only failures were specific, structured pathologies, each turned into a clean bounded diagnostic.

Method

Every input was compiled with a per-file timeout under `RUST_BACKTRACE=1`, and the result classified from the exit signal and output:

- OK — exit 0 (a `.osdi` was produced);
- clean error — nonzero exit with a diagnostic, no panic/crash;
- panic — a Rust panicked at ... (an internal compiler error);
- segfault / abort — killed by SIGSEGV / SIGABRT (e.g. a stack overflow);
- hang — did not finish within the timeout.

Only panic / segfault / abort / hang count as robustness bugs; a clean error on malformed input is the desired outcome.

Phase A — the real production corpus

Every model in the compact-model test suites (BSIM3/4/6, BSIM-CMG/IMG/SOI/BULK, PSP 102/103/104, HICUM L0/L2, MEXTRAM, VBIC, EKV/EKV3, HiSIM, ASMHMT, r2/r3_cmc, diode_cmc, ...) was compiled. The reproducible driver is `VA_TEST/compile_all.py` over the `VA-Models-main` collection — 124 `.va` files, of which 92 are standalone models and 32 are include-fragments.

Result: 92 / 92 standalone models compiled with 0 crashes / hangs / panics. The 32 include-fragment files — function libraries and macro-body pieces meant to be ``included`, not compiled standalone — report a clean error, as expected (17 of them; the other 15 happen to compile anyway).

Phase B — adversarial hand-crafted inputs

~50 inputs targeting the parser, preprocessor, and lowering: unterminated constructs (block comment, string, module, parenthesis), malformed numbers (`1e999999`, `4'b2`, `999999999'b1`, `1.2.3`), macro bombs (``define A A`, mutual recursion, an exponential expansion chain), circular / missing includes, duplicate declarations, unicode identifiers and embedded null bytes, and extreme nesting / sizes.

Result: all rejected cleanly; 0 accepted-invalid inputs. Five structured pathologies produced a crash or hang (below); everything else was a clean error. Macro recursion is already capped; ``include` recursion was not.

Phase C — mutation fuzzing

4,000 iterations, each taking a real model and applying a random mutation — byte flips, truncations, run duplication, delimiter injection, deep-nesting injection — then compiling it.

Result: 0 panics, 0 segfaults across 4,000 iterations. The compiler never crashed on mutated input; a few mutations produced hangs, all instances of the structured pathologies below (macro-expansion / nesting blow-ups).

Findings

Severity	Trigger	Symptom (before)	Status
critical	nested ?: chain, depth $\approx$ 30	$O(2^n)$ validation $\rightarrow$ compiler hangs; reachable via macros	fixed (E-147)
low	expression nested ~4k–32k deep (unary / binary / parens / calls)	recursive-descent stack overflow (SIGABRT)	fixed (E-148)
low	file that <code>`includes</code> itself	uncapped include recursion $\rightarrow$ stack overflow	fixed (E-148)
low	absurd array range <code>real x[0:100000000]</code>	expanded element-by-element $\rightarrow$ memory exhaustion / hang	fixed (E-148)
high	<code>`name(</code> with a stray directive ( <code>`include</code> , <code>`ifdef</code> , ...) in its args	preprocessor macro-arg collector never advances $\rightarrow$ infinite loop (hang)	fixed (E-219)
low	thousands of parse errors (deeply nested garbage)	every diagnostic rendered with source context $\rightarrow$ ~40 s	fixed (E-219, render cap 128)

The low-severity findings all require input a human or real model would never write. The others are worse: ~30 nested conditionals (E-147) is easily produced by macros in a real model, and the E-219 macro-argument loop is trivially reached by injecting (near any ``include`/macro token — both *hang* the compiler, the worst failure mode. E-219 was found only by re-running the campaign against the shipped binary (see the 2026-07-17 note above); its two paths sit in the preprocessor and the diagnostic sink, which the E-147/E-148 parser/validator guards did not touch. See [Enhancement-219](#) for root cause and fix.

The critical finding: exponential-time nested ?:

Root cause

The body validator visits every expression in `validate_expr`, whose match ends by recursing into the children:

```
_ => ()  
}  
self.parent.body.exprs[expr].walk_child_exprs(|child| self.validate_expr(child))
```

Arms that fully validate their own operands return early to skip that generic walk — the `Call` and `Path` arms both do. The **Select** (ternary) arm did not: it validated `cond`, `then_val` and `else_val` via `validate_condition`, then *fell through* to `walk_child_exprs`, which validated `then_val` and `else_val` a second time. Each ternary therefore validated its branches twice; for a chain of N nested ?: this compounds to 2^n validations.

Profiling a hanging compile with a live stack sample confirmed it exactly — an unbounded `validate_expr` $\rightarrow$ `validate_condition` $\rightarrow$ `validate_expr` $\rightarrow$... recursion with a doubling call count.

The fix

Add a `return`; at the end of the `Select` arm, matching the `Call` / `Path` arms. The arm already validates all three children, so the generic walk was pure redundant work — removing it is behaviour-preserving and turns validation from $O(2^n)$ into $O(N)$:

nesting depth	before	after
40	hang (> 30 s)	0.09 s
160	2^{160} (hopeless)	0.11 s
2000	—	1.6 s (linear)

Full write-up: [Enhancement-147](#).

The hardening: parser depth, include cap, array cap

The three lower-severity findings are pathological-input crashes — a robust compiler should reject them cleanly rather than die. [Enhancement-148](#) adds bounded limits.

Parser expression-depth limit. The Pratt parser recurses through `atom_expr` for every level of nesting *and* builds a left-leaning tree for operator chains — so both an over-deep recursion and an over-deep tree (which later passes `recurse` over) could overflow the stack. A shared `expr_depth` counter — incremented on each `atom_expr` (recursion) and per operator in the `expr_bp` loop (chain length) — bounds the total expression-tree depth at 1000 and recovers to the next statement boundary.

**`include** recursion cap. An `include_depth` counter caps nesting at 64 and emits a new `IncludeRecursionLimit` diagnostic — a self-including file is reported instead of overflowing the stack, mirroring the existing macro-recursion guard.

Array element cap. An array-shaped declaration is flattened into one scalar element per index. A shared `array_elem_count` helper caps the expansion at $2^{20} \approx 1.05$ M elements and emits a new `ArrayTooLarge` diagnostic, applied at every expansion site: variable arrays, parameter arrays, net/port buses and array function returns (item-tree lowering), and instance arrays — in *both* the item-tree lowering and the elaboration pass, which re-expands `dev s[0:N]()` into rendered text.

Behaviour-preserving

Both fixes are correctness/robustness changes with no effect on valid programs, and that was checked rather than assumed:

- Corpus head-to-head. Every one of the 92 standalone production models compiles to the identical verdict before and after each fix — 0 flips — and BSIM4 (12.6 k lines) still compiles in ~2.3 s. (E-219 was checked the same way: the whole corpus is byte-for-byte unchanged across it.)
- Unit tests. The parser, preprocessor, `hir_def`, `hir_ty` and `sim_back` suites pass with no snapshot changes; the build is warning-free.
- Targeted examples. `examples/nested_cond_examples/` (7/7, the exponential fix) and `examples/robustness_examples/` (17/17, the hardening) pin the behaviour: every pathological input now errors cleanly in well under a second, and valid deep-but-reasonable input (nested ternary depth 30, parentheses depth 100, a 100-term sum, small variable and instance arrays) still compiles.

Reproducing

The campaign is a small Python harness: a classifier that runs `openvaf-r` on an input under a timeout and labels the outcome, plus three drivers — compile the real corpus (Phase A), a battery of hand-crafted adversarial files (Phase B), and a mutation fuzzer that perturbs real models (Phase C). Localizing a hang used two tools: the compiler's own `--dump-j son / --dump-unopt-mir` flags to bisect the pipeline phase, and macOS `sample <pid>` on the live process to name the exact hot function.